Elixir
IN ACTION

THIRD EDITION

Saša Jurić

MANNING

Elixir in Action, Third Edition MEAP V08

Elixir

IN ACTION

THIRD EDITION

Saša Jurić

MEAP

MANNING

MEAP VERSION 8

MANNING PUBLICATIONS

Welcome

Thank you for purchasing the MEAP for *Elixir in Action, Third Edition.*

This is an intermediate level book that teaches about Elixir, Erlang, and how they assist the development of highly available server-side systems. To fully take advantage of the material, you should be an experienced software developer. It is not expected that you know anything about Elixir, Erlang, or functional programming, but you should be skillful in at least one programming language, such as Java, C#, Ruby, JavaScript.

A lot of work has been invested into making the contents approachable, and the learning process incremental. The first part of the book discusses the functional aspect of Elixir. Going forward, in the second part you'll learn about concurrent Elixir, and how it can help you improve scalability and fault-tolerance of your systems. Finally, the third part of the book deals with the production-related topics, such as distributed systems and deployable releases.

It's been four years since the second edition of this book has been published, and Elixir has moved forward in that period. This edition discusses some new language features, updates the deprecated parts, and polishes the example code, so it's more idiomatic and aligned with the latest best practices.

As you're reading, I encourage you to take advantage of Manning's [Livebook discussion forum] and post your questions and comments there. Feedback is highly appreciated and helps me make this book better.

Finally, I hope you'll enjoy reading this book and have fun working with Elixir!

—Saša Jurić

In this book

1 First steps

This chapter covers

- Overview of Erlang
- Benefits of Elixir

This is the beginning of your journey into the world of Elixir and Erlang, two efficient and useful technologies that can significantly simplify the development of large, scalable systems. Chances are, you're reading this book to learn about Elixir. But because Elixir is built on top of Erlang and depends heavily on it, you should first learn a bit about what Erlang is and the benefits it offers. So let's take a brief, high-level look at Erlang.

1.1 About Erlang

Erlang is a development platform for building scalable and reliable systems that constantly provide service with little or no downtime. This is a bold statement, but it's exactly what Erlang was made for. Conceived in the mid-1980s by Ericsson, a Swedish telecom giant, Erlang was driven by the needs of the company's own telecom systems, where properties like reliability, responsiveness, scalability, and constant availability were imperative. A telephone network should always operate regardless of the number of simultaneous calls, unexpected bugs, or hardware and software upgrades taking place.

Despite being originally built for telecom systems, Erlang is in no way specialized for this domain. It doesn't contain explicit support for programming telephones, switches, or other telecom devices. Instead, it's a general-purpose development platform that provides special support for technical, nonfunctional challenges such as concurrency, scalability, fault-tolerance, distribution, and high availability.

In late 1980s and early 90s, when most software was desktop-based, the need

for high availability was limited to specialized systems such as telecoms. Today we face a much different situation: the focus is on the internet and the web, and most applications are driven and supported by a server system that processes requests, crunches data, and pushes relevant information to many connected clients. Today's popular systems are more about communication and collaboration; examples include social networks, content-management systems, on-demand multimedia, and multiplayer games.

All of these systems have some common nonfunctional requirements. The system must be responsive, regardless of the number of connected clients. The impact of unexpected errors must be minimal, instead of affecting the entire system. It's acceptable if an occasional request fails due to a bug, but it's a major problem when the entire system becomes completely unavailable. Ideally, the system should never crash or be taken down, not even during a software upgrade. It should always be up and running, providing service to its clients.

These goals might seem difficult to reach, but they're imperative when building systems that people depend on. Unless a system is responsive and reliable, it will eventually fail to fulfill its purpose. Therefore, when building server-side systems, it's essential to make the system constantly available.

This is the intended purpose of Erlang. High availability is explicitly supported via technical concepts such as scalability, fault-tolerance, and distribution. Unlike with most other modern development platforms, these concepts were the main motivation and driving force behind the development of Erlang. The Ericsson team, led by Joe Armstrong, spent a couple of years designing, prototyping, and experimenting before creating the development platform. Its uses may have been limited in the early 90s, but today almost any system can benefit from it.

Erlang has recently gained more attention. It powers various large systems and has been doing so for three decades, such as the WhatsApp messaging application, the Discord instant messaging platform, the RabbitMQ message queue, financial systems, and multiplayer backends. It's truly a proven technology, both in time and scale. But what is the magic behind Erlang? Let's take a look at how Erlang can help you build highly available, reliable systems.

1.1.1 High availability

Erlang was specifically created to support the development of highly available systems — systems that are always online and provide service to their clients even when faced with unexpected circumstances. On the surface, this may seem simple, but as you probably know, many things can go wrong in production. To make systems work 24/7 without any downtime, you have to tackle some technical challenges:

- *Fault-tolerance*—A system has to keep working when something unforeseen happens. Unexpected errors occur, bugs creep in, components occasionally fail, network connections drop, or the entire machine where the system is running crashes. Whatever happens, you want to localize the impact of an error as much as possible, recover from the error, and keep the system running and providing service.
- *Scalability*—A system should be able to handle any possible load. Of course, you don't buy tons of hardware just in case the entire planet's population might start using your system some day. But you should be able to respond to a load increase by adding more hardware resources without any software intervention. Ideally, this should be possible without a system restart.
- *Distribution*—To make a system that never stops, you have to run it on multiple machines. This promotes the overall stability of the system: if a machine is taken down, another one can take over. Furthermore, this gives you the means to scale horizontally — you can address load increase by adding more machines to the system, thus adding work units to support the higher demand.
- *Responsiveness* —It goes without saying that a system should always be reasonably fast and responsive. Request handling shouldn't be drastically prolonged, even if the load increases or unexpected errors happen. In particular, occasional lengthy tasks shouldn't block the rest of the system or have a significant effect on performance.
- *Live update*—In some cases, you may want to push a new version of your software without restarting any servers. For example, in a telephone system, you don't want to disconnect established calls while you upgrade the software.

If you manage to handle these challenges, the system will truly become highly available and be able to constantly provide service to users, rain or shine.

Erlang gives us tools to address these challenges — that's what it was built for. A system can gain all these properties and ultimately become highly available through the power of the Erlang concurrency model.

Let's look at how concurrency works in Erlang.

1.1.2 Erlang concurrency

Concurrency is at the heart and soul of Erlang systems. Almost every nontrivial Erlang-based production system is highly concurrent. Even the programming language is sometimes called a concurrency-oriented language. Instead of relying on heavyweight threads and OS processes, Erlang takes concurrency into its own hands, as illustrated in figure 1.1.

Figure 1.1. Concurency in the Erlang virtual machine

BEAM is a single OS process.

BEAM

Erlang process is a unit of concurrent execution.

Process

Process

Scheduler is an OS thread responsible for executing multiple processes.

Scheduler

Scheduler

BEAM uses multiple schedulers to parallelize the work over available CPU cores.

CPU

CPU
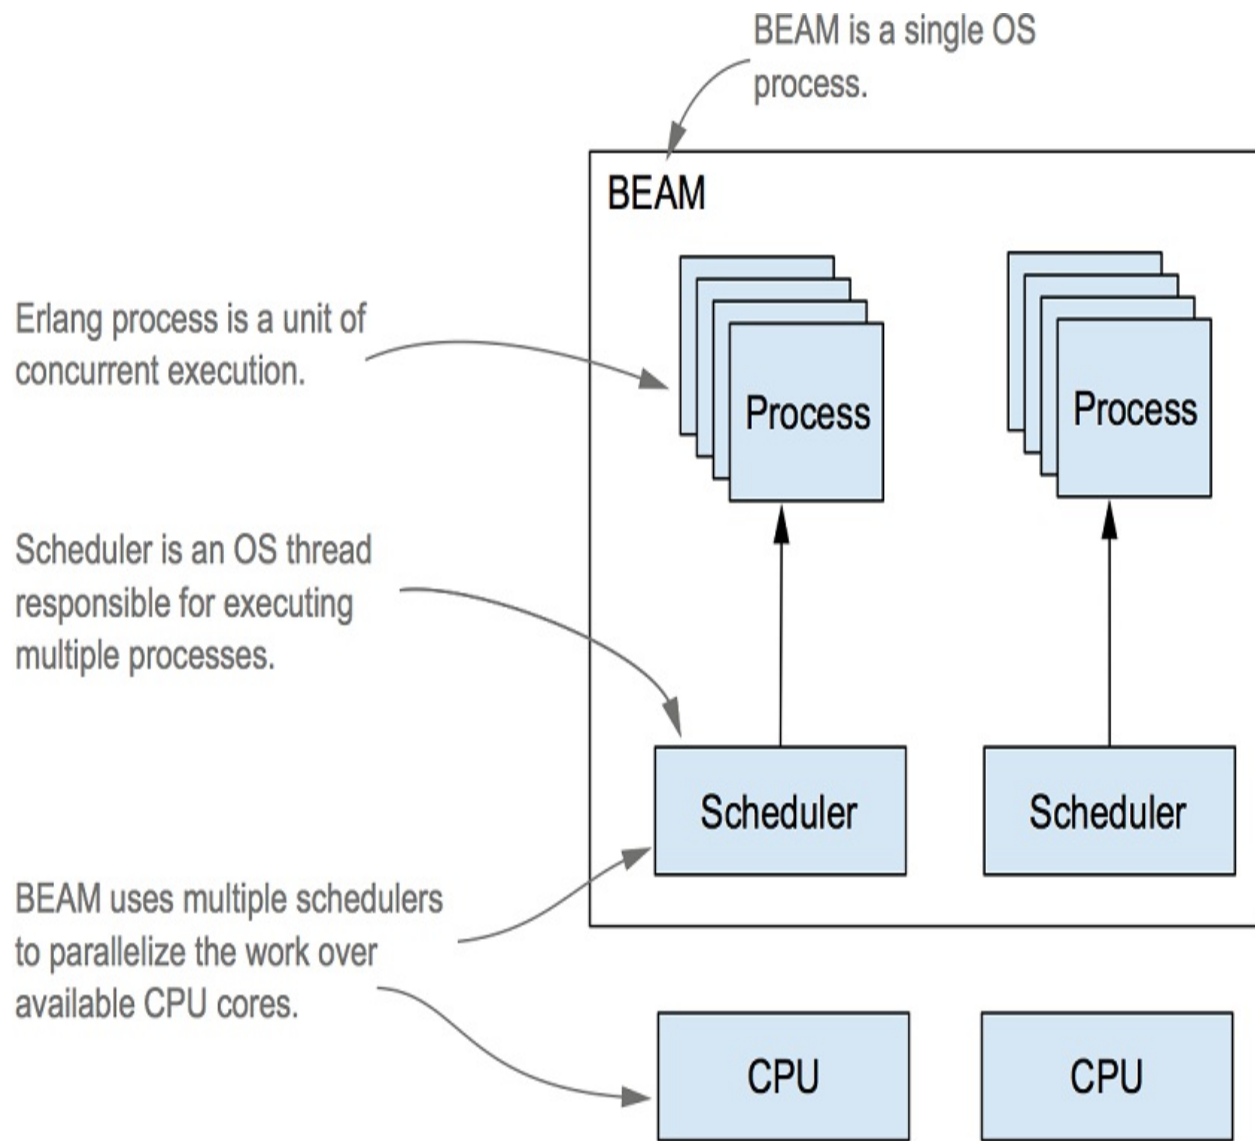
The basic concurrency primitive is called an *Erlang process* (not to be confused with OS processes or threads), and typical Erlang systems run thousands or even millions of such processes. The Erlang virtual machine, called BEAM (Bogdan/Björn's Erlang Abstract Machine), uses its own schedulers to distribute the execution of processes over the available CPU cores, thus parallelizing execution as much as possible. The way processes are implemented provides many benefits.

Fault-tolerance

Erlang processes are completely isolated from each other. They share no

memory, and a crash of one process doesn't cause a crash of other processes. This helps you isolate the effect of an unexpected error. If something bad happens, it has only a local impact. Moreover, Erlang provides you with the means to detect a process crash and do something about it; typically, you start a new process in place of the crashed one.

Scalability

Sharing no memory, processes communicate via asynchronous messages. This means there are no complex synchronization mechanisms such as locks, mutexes, or semaphores. Consequently, the interaction between concurrent entities is much simpler to develop and understand.

Typical Erlang systems are divided into a large number of concurrent processes, which cooperate together to provide the complete service. The virtual machine can efficiently parallelize the execution of processes as much as possible. This makes Erlang systems scalable because they can take advantage of all available CPU cores.

Distribution

Communication between processes works the same way regardless of whether these processes reside in the same BEAM instance or on two different instances on two separate, remote computers. Therefore, a typical highly concurrent Erlang-based system is automatically ready to be distributed over multiple machines. This in turn gives you the ability to scale out — to run a cluster of machines that share the total system load. Additionally, running on multiple machines makes the system truly resilient — if one machine crashes, others can take over.

Responsiveness

The runtime is specifically tuned to promote the overall responsiveness of the system. I've mentioned that Erlang takes the execution of multiple processes into its own hands by employing dedicated schedulers that interchangeably execute many Erlang processes. A scheduler is preemptive — it gives a small

execution window to each process and then pauses it and runs another process. Because the execution window is small, a single long-running process can't block the rest of the system. Furthermore, I/O operations are internally delegated to separate threads, or a kernel-poll service of the underlying OS is used if available. This means any process that waits for an I/O operation to finish won't block the execution of other processes.

Even garbage collection is specifically tuned to promote system responsiveness. Recall that processes are completely isolated and share no memory. This allows per-process garbage collection: instead of stopping the entire system, each process is individually collected as needed. Such collections are much quicker and don't block the entire system for long periods of time. In fact, in a multicore system, it's possible for one CPU core to run a short garbage collection while the remaining cores are doing standard processing.

As you can see, concurrency is a crucial element in Erlang, and it's related to more than just parallelism. Owing to the underlying implementation, concurrency promotes fault-tolerance, distribution, and system responsiveness. Typical Erlang systems run many concurrent tasks, using thousands or even millions of processes. This can be especially useful when you're developing server-side systems, which can often be implemented completely in Erlang.

1.1.3 Server-side systems

Erlang can be used in various applications and systems. There are examples of Erlang-based desktop applications, and it's often used in embedded environments. Its sweet spot, in my opinion, lies in server-side systems — systems that run on one or more servers and must serve many simultaneous clients. The term *server-side system* indicates that it's more than a simple server that processes requests. It's an entire system that, in addition to handling requests, must run various background jobs and manage some kind of server-wide in-memory state, as illustrated in figure 1.2.

A server-side system is often distributed on multiple machines that collaborate to produce business value. You might place different components

on different machines, and you also might deploy some components on multiple servers to achieve load balancing or support failover scenarios.

Figure 1.2. Server-side system
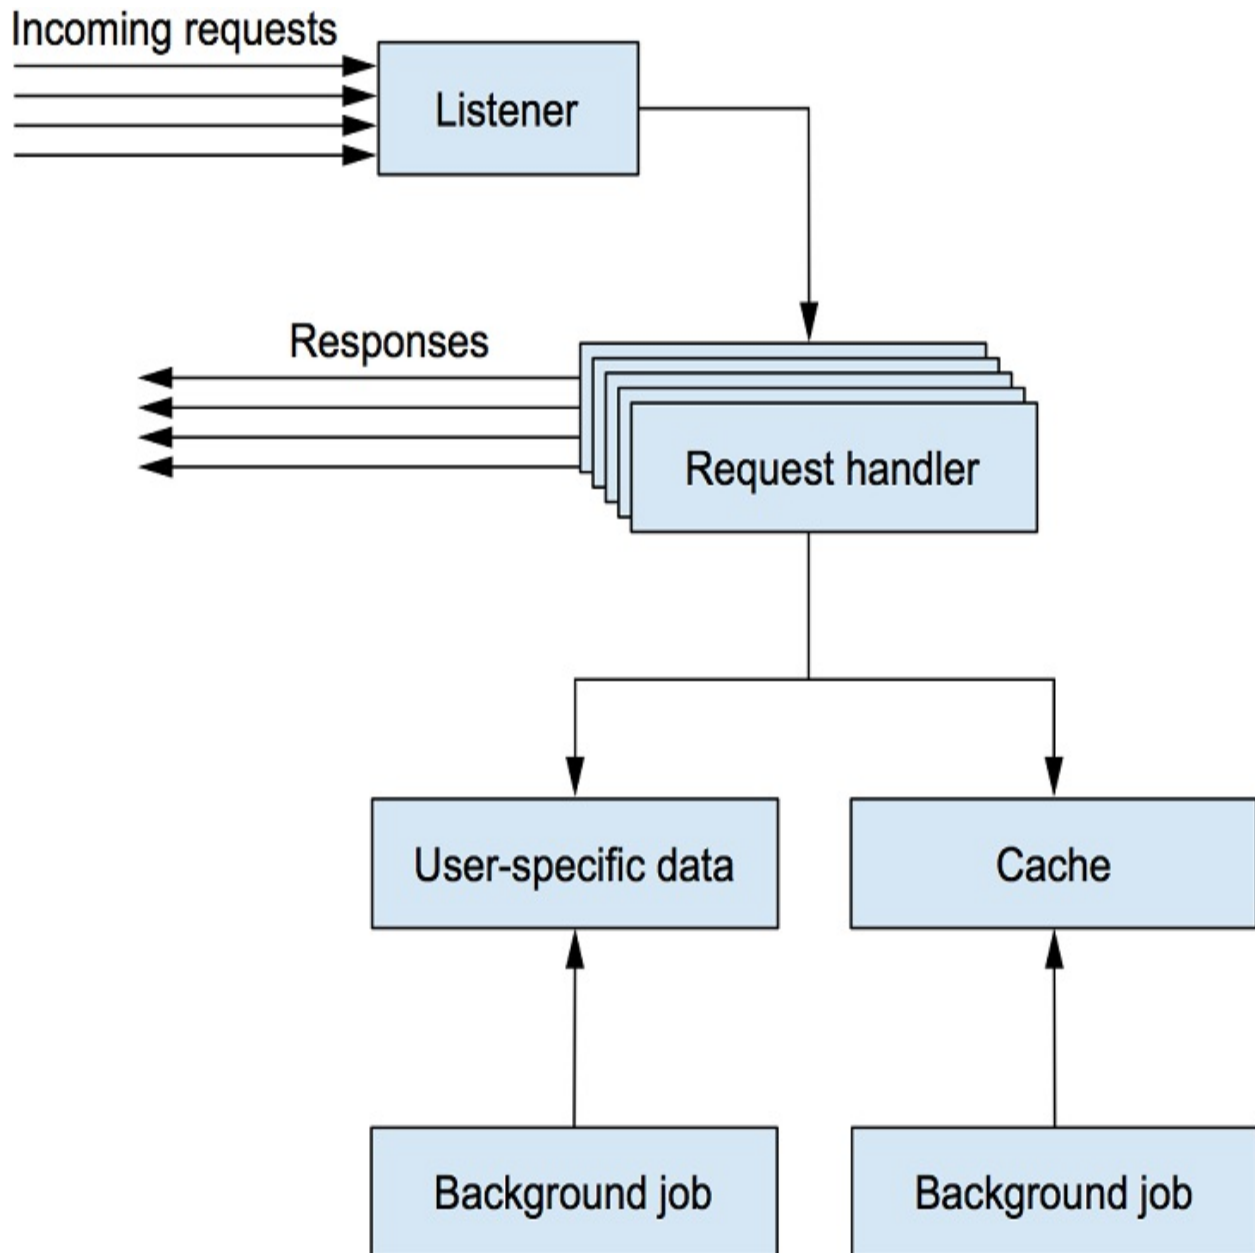
This is where Erlang can make your life significantly simpler. By giving you primitives to make your code concurrent, scalable, and distributed, it allows you to implement the entire system completely in Erlang. Every component in figure 1.2 can be implemented as an Erlang process, which makes the

system scalable, fault-tolerant, and easy to distribute. By relying on Erlang's error-detection and recovery primitives, you can further increase reliability and recover from unexpected errors.

Let's look at a real-life example. I've been involved professionally in the development of two web servers, both of which have similar technical needs: they serve a multitude of clients, handle long-running requests, manage server-wide in-memory state, persist data that must survive OS processes and machine restarts, and run background jobs. Table 1.1 lists the technologies used in each server.

Table 1.1. Comparison of technologies used in two real-life web servers

Technical requirement	Server	Server B
HTTP server	Nginx and Phusion Passenger	Erlang
Request processing	Ruby on Rails	Erlang
Long-running requests	Go	Erlang
Server-wide state	Redis	Erlang
Persistable data	Redis and MongoDB	Erlang
Background jobs	Cron, Bash scripts, and Ruby	Erlang
Service crash recovery	Upstart	Erlang

Server A is powered by various technologies, most of them known and popular in the community. There were specific reasons for using these technologies: each was introduced to resolve a shortcoming of those already present in the system. For example, Ruby on Rails handles concurrent requests in separate OS processes. We needed a way to share data between these different processes, so we introduced Redis. Similarly, MongoDB is used to manage persistent frontend data, most often user-related information. Thus there's a rationale behind every technology used in server A, but the entire solution seems complex. It's not contained in a single project, the components are deployed separately, and it isn't trivial to start the entire system on a development machine. We had to develop a tool to help us start the system locally!

In contrast, server B accomplishes the same technical requirements while relying on a single technology, using platform features created specifically for these purposes and proven in large systems. Moreover, the entire server is a single project that runs inside a single BEAM instance — in production, it runs inside only one OS process, using a handful of OS threads. Concurrency is handled completely by the Erlang scheduler, and the system is scalable, responsive, and fault-tolerant. Because it's implemented as a single project, the system is easier to manage, deploy, and run locally on the development machine.

It's important to notice that Erlang tools aren't always full-blown alternatives to mainstream solutions, such as web servers like Nginx, database servers like MongoDB, and in-memory key/value stores like Redis. But Erlang gives you options, making it possible to implement an initial solution using exclusively Erlang and resorting to alternative technologies when an Erlang solution isn't sufficient. This makes the entire system more homogeneous and therefore easier to develop and maintain.

It's also worth noting that Erlang isn't an isolated island. It can run in-process C code and can communicate with external components such as message queues, in-memory key/value stores, and external databases. Therefore, when opting for Erlang, you aren't deprived of using existing third-party technologies. Instead, you have the option of using them when they're called for rather than because your primary development platform doesn't give you

a tool to solve your problems.

Now that you know about Erlang's strengths and the areas where it excels, let's take a closer look at what Erlang is.

1.1.4 The development platform

Erlang is more than a programming language. It's a full-blown development platform consisting of four distinct parts: the language, the virtual machine, the framework, and the tools.

Erlang, the language, is the primary way of writing code that runs in the Erlang virtual machine. It's a simple, functional language with basic concurrency primitives.

Source code written in Erlang is compiled into bytecode that's then executed in the BEAM. This is where the true magic happens. The virtual machine parallelizes your concurrent Erlang programs and takes care of process isolation, distribution, and the overall responsiveness of the system.

The standard part of the release is a framework called Open Telecom Platform (OTP). Despite its somewhat unfortunate name, the framework has nothing to do with telecom systems. It's a general-purpose framework that abstracts away many typical Erlang tasks:

- Concurrency and distribution patterns
- Error detection and recovery in concurrent systems
- Packaging code into libraries
- Systems deployment
- Live code updates

OTP is battle-tested in many production systems and is such an integral part of Erlang that it's hard to draw a line between the two. Even the official distribution is called Erlang/OTP.

The tools are used for various typical tasks such as compiling Erlang code, starting a BEAM instance, creating deployable releases, running the interactive shell, connecting to the running BEAM instance, and so on. Both

BEAM and its accompanying tools are cross-platform. You can run them on most mainstream operating systems, such as Unix, Linux, and Windows. The entire Erlang distribution is open source, and you can find the source on the official site (https://www.erlang.org/) or on the Erlang GitHub repository (https://github.com/erlang/otp). Ericsson is still in charge of the development process and releases a new version on a regular basis, once a year.

That concludes the story of Erlang. But if Erlang is so great, why do you need Elixir? The next section aims to answer this question.

1.2 About Elixir

Elixir is an alternative language for the Erlang virtual machine that allows you to write cleaner, more compact code that does a better job of revealing your intentions. You write programs in Elixir and run them normally in BEAM.

Elixir is an open source project, originally started by José Valim. Unlike Erlang, Elixir is more of a collaborative effort; presently it has about 1200 contributors. New features are frequently discussed on mailing lists, the GitHub issue tracker, and the #elixir-lang freenode IRC channel. José has the last word, but the entire project is a true open source collaboration, attracting an interesting mixture of seasoned Erlang veterans and talented young developers. The source code can be found on the GitHub repository at https://github.com/elixir-lang/elixir.

Elixir targets the Erlang runtime. The result of compiling the Elixir source code is BEAM-compliant bytecode files that can run in a BEAM instance and can normally cooperate with pure Erlang code — you can use Erlang libraries from Elixir and vice versa. There's nothing you can do in Erlang that can't be done in Elixir, and usually the Elixir code is as performant as its Erlang counterpart.

Elixir is semantically close to Erlang: many of its language constructs map directly to their Erlang counterparts. But Elixir provides some additional constructs that make it possible to radically reduce boilerplate and duplication. In addition, it tidies up some important parts of the standard

libraries and provides some nice syntactic sugar and a uniform tool for creating and packaging systems. Everything you can do in Erlang is possible in Elixir, and vice versa, but in my experience the Elixir solution is usually easier to develop and maintain.

Let's take a closer look at how Elixir improves on some Erlang features. We'll start with boilerplate and noise reduction.

1.2.1 Code simplification

One of the most important benefits of Elixir is its ability to radically reduce boilerplate and eliminate noise from code, which results in simpler code that's easier to write and maintain. Let's see what this means by contrasting Erlang and Elixir code.

A frequently used building block in Erlang concurrent systems is the server process. You can think of server processes as something like concurrent objects — they embed private state and can interact with other processes via messages. Being concurrent, different processes may run in parallel. Typical Erlang systems rely heavily on processes, running thousands or even millions of them.

The following example Erlang code implements a simple server process that adds two numbers.

Listing 1.1. Erlang-based server process that adds two numbers

```erlang
-module(sum_server).
-behaviour(gen_server).

-export([
  start/0, sum/3,
  init/1, handle_call/3, handle_cast/2, handle_info/2, terminate/
  code_change/3
]).

start() -> gen_server:start(?MODULE, [], []).
sum(Server, A, B) -> gen_server:call(Server, {sum, A, B}).

init(_) -> {ok, undefined}.
```

```erlang
handle_call({sum, A, B}, _From, State) -> {reply, A + B, State}.
handle_cast(_Msg, State) -> {noreply, State}.
handle_info(_Info, State) -> {noreply, State}.
terminate(_Reason, _State) -> ok.
code_change(_OldVsn, State, _Extra) -> {ok, State}.
```

Even without any knowledge of Erlang, this seems like a lot of code for something that only adds two numbers. To be fair, the addition is concurrent, but regardless, due to the amount of code it's hard to see the forest for the trees. It's definitely not immediately obvious what the code does. Moreover, it's difficult to write such code. Even after years of production-level Erlang development, I still can't write this without consulting the documentation or copying and pasting it from previously written code.

The problem with Erlang is that this boilerplate is almost impossible to remove, even if it's identical in most places (which in my experience is the case). The language provides almost no support for eliminating this noise. In all fairness, there is a way to reduce the boilerplate using a construct called `parse transform`, but it's clumsy and complicated to use. In practice, Erlang developers write their server processes using the preceding pattern.

Because server processes are an important and frequently used tool in Erlang, it's unfortunate that Erlang developers have to constantly copy-paste this noise and work with it. Surprisingly, many people get used to it, probably due to the wonderful things BEAM does for them. It's often said that Erlang makes hard things easy and easy things hard. Still, the previous code leaves an impression that you should be able to do better.

Let's look at the Elixir version of the same server process.

Listing 1.2. Elixir-based server process that adds two numbers

```elixir
defmodule SumServer do
  use GenServer

  def start do
    GenServer.start(__MODULE__, nil)
  end

  def sum(server, a, b) do
    GenServer.call(server, {:sum, a, b})
```

```
    end

    def handle_call({:sum, a, b}, _from, state) do
      {:reply, a + b, state}
    end
end
```

This code is significantly smaller and therefore easier to read and maintain. Its intention is more clearly revealed, and it's less burdened with noise. And yet, it's as capable and flexible as the Erlang version. It behaves exactly the same at runtime and retains the complete semantics. There's nothing you can do in the Erlang version that's not possible in its Elixir counterpart.

Despite being significantly smaller, the Elixir version of a sum server process still feels somewhat noisy, given that all it does is add two numbers. The excess noise exists because Elixir retains a 1:1 semantic relation to the underlying Erlang library that's used to create server processes.

But Elixir gives you tools to further eliminate whatever you may regard as noise and duplication. For example, I've developed my own Elixir library called ExActor that makes the server process definition dense, as shown next.

Listing 1.3. Elixir-based server process

```
defmodule SumServer do
  use ExActor.GenServer

  defstart start

  defcall sum(a, b) do
    reply(a + b)
  end
end
```

The intention of this code should be obvious even to developers with no previous Elixir experience. At runtime, the code works almost exactly the same as the two previous versions. The transformation that makes this code behave like the previous ones happens at compile time. When it comes to the bytecode, all three versions are similar.

ⓘ Note

I mention the ExActor library only to illustrate how much you can abstract away in Elixir. You won't use that library in this book because it's a third-party abstraction that hides important details of how server processes work. To completely take advantage of server processes, it's important that you understand what makes them tick, which is why in this book you'll learn about lower-level abstractions. Once you understand how server processes work, you can decide for yourself whether you want to use ExActor to implement server processes.

This last implementation of the `sum` server process is powered by the Elixir macros facility. A *macro* is Elixir code that runs at compile time. Macros take an internal representation of your source code as input and can create alternative output. Elixir macros are inspired by Lisp and shouldn't be confused with C-style macros. Unlike C/C++ macros, which work with pure text, Elixir macros work on an abstract syntax tree (AST) structure, which makes it easier to perform nontrivial manipulations of the input code to obtain alternative output. Of course, Elixir provides helper constructs to simplify this transformation.

Take another look at how the sum operation is defined in the last example:

```
defcall sum(a, b) do
  reply(a + b)
end
```

Notice the `defcall` at the beginning. There's no such keyword in Elixir. This is a custom macro that translates the given definition to something like the following:

```
def sum(server, a, b) do
  GenServer.call(server, {:sum, a, b})
end

def handle_call({:sum, a, b}, _from, state) do
  {:reply, a + b, state}
end
```

Because macros are written in Elixir, they're flexible and powerful, making it possible to extend the language and introduce new constructs that look like an integral part of the language. For example, the open source Ecto project, which aims to bring LINQ-style queries to Elixir, is also powered by Elixir macro support and provides an expressive query syntax that looks deceptively like part of the language:

```
from w in Weather,
  where: w.prcp > 0 or w.prcp == nil,
  select: w
```

Due to its macro support and smart compiler architecture, most of Elixir is written in Elixir. Language constructs like `if`, and `unless` are implemented via Elixir macros. Only the smallest possible core is done in Erlang — everything else is then built on top of it in Elixir!

Elixir macros are something of a black art, but they make it possible to flush out nontrivial boilerplate at compile time and extend the language with your own DSL-like constructs.

But Elixir isn't all about macros. Another worthy improvement is some seemingly simple syntactic sugar that makes functional programming much easier.

1.2.2 Composing functions

Both Erlang and Elixir are functional languages. They rely on immutable data and functions that transform data. One of the supposed benefits of this approach is that code is divided into many small, reusable, composable functions.

Unfortunately, the composability feature works clumsily in Erlang. Let's look at an adapted example from my own work. One piece of code I'm responsible for maintains an in-memory model and receives XML messages that modify the model. When an XML message arrives, the following actions must be done:

- Apply the XML to the in-memory model.

- Process the resulting changes.
- Persist the model.

Here's an Erlang sketch of the corresponding function:

```
process_xml(Model, Xml) ->
  Model1 = update(Model, Xml),
  Model2 = process_changes(Model1),
  persist(Model2).
```

I don't know about you, but this doesn't look composable to me. Instead, it seems fairly noisy and error-prone. The temporary variables `Model1` and `Model2` are introduced here only to take the result of one function and feed it to the next.

Of course, you could eliminate the temporary variables and inline the calls:

```
process_xml(Model, Xml) ->
  persist(
    process_changes(
      update(Model, Xml)
    )
  ).
```

This style, known as *staircasing,* is admittedly free of temporary variables, but it's clumsy and hard to read. To understand what goes on here, you have to manually parse it inside-out.

Although Erlang programmers are more or less limited to such clumsy approaches, Elixir gives you an elegant way to chain multiple function calls together:

```
def process_xml(model, xml) do
  model
  |> update(xml)
  |> process_changes
  |> persist
end
```

The pipe operator `|>` takes the result of the previous expression and feeds it to the next one as the first argument. The resulting code is clean, contains no temporary variables, and reads like the prose, top to bottom, left to right.

Under the hood, this code is transformed at compile time to the staircased version. This is again possible because of Elixir's macro system.

The pipe operator highlights the power of functional programming. You treat functions as data transformations and then combine them in different ways to gain the desired effect.

1.2.3 The big picture

There are many other areas where Elixir improves the original Erlang approach. The API for standard libraries is cleaned up and follows some defined conventions. Syntactic sugar is introduced that simplifies typical idioms. A concise syntax for working with structured data is provided. String manipulation is improved, and the language has explicit support for Unicode manipulation. In the tooling department, Elixir provides the tool called mix that simplifies common tasks such as creating applications and libraries, managing dependencies, and compiling and testing code. In addition, a package manager called Hex (https://hex.pm/) is available that makes it simpler to package, distribute, and reuse dependencies.

The list goes on and on, but instead of presenting each feature, I'd like to express a personal sentiment based on my own production experience. Personally, I find it much more pleasant to code in Elixir. The resulting code seems simpler, more readable, and less burdened with boilerplate, noise, and duplication. At the same time, you retain the complete runtime characteristics of pure Erlang code. You can also use all the available libraries from the Erlang ecosystem, both standard and third-party.

1.3 Disadvantages

No technology is a silver bullet, and Erlang and Elixir are definitely not exceptions. Thus it's worth mentioning some of their shortcomings.

1.3.1 Speed

Erlang is by no means the fastest platform out there. If you look at various synthetic benchmarks on the internet, you usually won't see Erlang high on

the list. Erlang programs are run in BEAM and therefore can't achieve the speed of machine-compiled languages, such as C and C++. But this isn't accidental or poor engineering on behalf of the Erlang/OTP team.

The goal of the platform isn't to squeeze out as many requests per second as possible, but to keep performance predictable and within limits. The level of performance your Erlang system achieves on a given machine shouldn't degrade significantly, meaning there shouldn't be unexpected system hiccups due to, for example, the garbage collector kicking in. Furthermore, as explained earlier, long-running BEAM processes don't block or significantly impact the rest of the system. Finally, as the load increases, BEAM can use as many hardware resources as possible. If the hardware capacity isn't enough, you can expect graceful system degradation — requests will take longer to process, but the system won't be paralyzed. This is due to the preemptive nature of the BEAM scheduler, which performs frequent context switches that keep the system ticking and favors short-running processes. And of course, you can address higher system demand by adding more hardware.

Nevertheless, intensive CPU computations aren't as performant as, for example, their C/C++ counterparts, so you may consider implementing such tasks in some other language and then integrating the corresponding component into your Erlang system. If most of your system's logic is heavily CPU-bound, you should probably consider some other technology.

1.3.2 Ecosystem

The ecosystem built around Erlang isn't small, but it definitely isn't as big as that of some other languages. At the time of writing, a quick search on GitHub reveals about 20,000 Erlang-based repositories and about 45,000 Elixir repositories. In contrast, there are more than 1,500,000 Ruby repositories and almost 7,000,000 for JavaScript.

You should be aware that the choice of libraries won't be as abundant as you may be used to, and in turn you may end up spending extra time on something that would take minutes in other languages. If that happens, keep in mind all the benefits you get from Erlang. As I've explained, Erlang goes a long way toward making it possible to write fault-tolerant systems that can

run for a long time with hardly any downtime. This is a big challenge and a specific focus of the Erlang platform. Although it's admittedly unfortunate that the ecosystem isn't as big as it could be, my sentiment is that Erlang significantly helps with hard problems, even if simple problems can sometimes be more clumsy to solve. Of course, those difficult problems may not always be important. Perhaps you don't expect a high load, or a system doesn't need to run constantly and be extremely fault-tolerant. In such cases, you may want to consider some other technology stack with a more evolved ecosystem.

1.4 Summary

- Erlang is a technology for developing highly available systems that constantly provide service with little or no downtime. It has been battle tested in diverse large systems for three decades.
- Elixir is a modern language that makes development for the Erlang platform much more pleasant. It helps organize code more efficiently and abstracts away boilerplate, noise, and duplication.

2 Building blocks

This chapter covers

- Using the interactive shell
- Working with variables
- Organizing your code
- Understanding the type system
- Working with operators
- Understanding the runtime

It's time to start learning about Elixir. This chapter presents the basic building blocks of the language, such as modules, functions, and the type system. This will be a somewhat long, not particularly exciting tour of language features, but the material presented here is important because it prepares the stage for exploring more interesting, higher-level topics.

Before starting, make sure you've installed Elixir version 1.14.x and Erlang version 25. There are multiple ways of installing Elixir, and it's best to follow the instructions from the official Elixir site at https://elixir-lang.org/install.html.

With that out of the way, let's start our tour of Elixir. The first thing you should know about is the interactive shell.

Detailed information

This book doesn't provide a detailed reference on any of the language or platform features. That would take up too much space, and the material would quickly become outdated. Here are some other references you can check out:

- For an alternative syntax quick start, you should look at the Getting Started guide on the Elixir official site: https://elixir-lang.org/getting-started/introduction.html.

- A more detailed reference can be found in the online documentation: https://hexdocs.pm/elixir.
- For specific questions, you can turn to the Elixir forum (https://elixirforum.com/), or the Slack channel (https://elixir-slackin.herokuapp.com/).
- Finally, for many things, you'll need to look into the Erlang documentation: https://www.erlang.org/doc. If you're not familiar with Erlang syntax, you may also need to read Elixir's crash course on Erlang (https://elixir-lang.org/crash-course.html).

2.1 The interactive shell

The simplest way to experiment and learn about a language's features is through the interactive shell. You can start the Elixir interactive shell from the command line by running the command `iex`:

```
$ iex
Erlang/OTP 25 [erts-13.0] [source] [64-bit] [smp:8:8] [ds:8:8:10]

Interactive Elixir (1.14) - press Ctrl+C to exit
  (type h() ENTER for help)

iex(1)>
```

Running `iex` starts an instance of the BEAM and then starts an interactive Elixir shell inside it. Runtime information is printed, such as the Erlang and Elixir version numbers, and then the prompt is provided so you can enter Elixir expressions:

```
iex(1)> 1 + 2   #1
3               #2
```

After you type an expression, it's interpreted and executed. Its return value is then printed to the screen.

ℹ Note

Everything in Elixir is an expression that has a return value. This includes not

only function calls but also constructs like `if` and `case`.

 Tip

You'll use `iex` extensively throughout the book, especially in the initial chapters. The expression result often won't be particularly relevant, and it will be omitted to reduce noise. Regardless, keep in mind that each expression returns a result, and when you enter an expression in the shell, its result will be presented.

You can type practically anything that constitutes valid Elixir code, including more complicated multiline expressions:

```
iex(2)> 2 * (      #1
         3 + 1    #1
       ) / 4      #2
2.0
```

Notice how the shell doesn't evaluate the expression until you finish it on the last line. In Elixir, you need no special characters, such as semicolons, to indicate the end of an expression. Instead, a line break indicates the end of an expression, if the expression is complete. Otherwise, the parser waits for additional input until the expression becomes complete.

The quickest way to leave the shell is to press Ctrl-C twice. Doing so brutally kills the OS process and all background jobs that are executing. Because the shell is mostly used for experimenting and shouldn't be used to run real production systems, it's usually fine to terminate it this way. But if you want a more polite way of stopping the system, you can invoke `System.stop`.

 Note

There are multiple ways to start Elixir and the Erlang runtime, and to run your Elixir programs. You'll learn a bit about all of them by the end of this chapter. In the first part of this book you'll mostly work with the `iex` shell, because it's a simple and efficient way of experimenting with the language.

You can do many things with the shell, but most often you'll use it to enter expressions and inspect their results. You can research for yourself what else can be done in the shell. Basic help can be obtained with the `h` command:

```
iex(4)> h
```

Entering this in the shell will output an entire screen of `iex`-related instructions. You can also look for the documentation of the IEx module, which is responsible for the shell's workings:

```
iex(5)> h IEx
```

You can find the same help in the online documentation at https://hexdocs.pm/iex.

Now that you have a basic tool with which to experiment, let's research the features of the language. We'll start with variables.

2.2 Working with variables

Elixir is a dynamic programming language, which means you don't explicitly declare a variable or its type. Instead, the variable type is determined by whatever data it contains at the moment. In Elixir terms, assignment is called *binding*. When you initialize a variable with a value, the variable is bound to that value:

```
iex(1)> monthly_salary = 10000    #1
10000                    #2
```

Each expression in Elixir has a result. In the case of the = operator, the result is whatever is on the right side of the operator. After the expression is evaluated, the shell prints this result to the screen.

Now you can reference the variable:

```
iex(2)> monthly_salary     #1
10000          #2
```

The variable can, of course, be used in complex expressions:

```
iex(3)> monthly_salary * 12
120000
```

In Elixir, a variable name always starts with a lowercase alphabetic character or an underscore. After that, any combination of alphanumerics and underscores is allowed. The prevalent convention is to use only lowercase ASCII letters, digits, and underscores:

```
valid_variable_name
also_valid_1
validButNotRecommended
NotValid
```

Variable names can also end with the question mark (?) or exclamation mark (!) characters:

```
valid_name?
also_ok!
```

Variables can be rebound to a different value:

```
iex(1)> monthly_salary = 10000     #1
10000

iex(2)> monthly_salary             #2
10000

iex(3)> monthly_salary = 11000     #3
11000

iex(4)> monthly_salary             #4
11000
```

Rebinding doesn't mutate the existing memory location. It reserves new memory and reassigns the symbolic name to the new location.

Note

You should always keep in mind that data is immutable. Once a memory location is occupied with data, it can't be modified until it's released. But variables can be rebound, which makes them point to a different memory

location. Thus, variables are mutable, but the data they point to is immutable.

Elixir is a garbage-collected language, which means you don't have to manually release memory. When a variable goes out of scope, the corresponding memory is eligible for garbage collection and will be released sometime in the future, when the garbage collector cleans up the memory.

2.3 Organizing your code

Being a functional language, Elixir relies heavily on functions. Due to the immutable nature of the data, a typical Elixir program consists of many small functions. You'll witness this in chapters 3 and 4, as you start using some typical functional idioms. Multiple functions can be further organized into modules.

2.3.1 Modules

A module is a collection of functions, somewhat like a namespace. Every Elixir function must be defined inside a module.

Elixir comes with a standard library that provides many useful modules. For example, the `IO` module can be used to do various I/O operations. The `puts` function from the `IO` module can be used to print a message to the screen:

```
iex(1)> IO.puts("Hello World!")    #1
Hello World!                #2
:ok                 #3
```

As you can see in the example, to call a function of a module you use the syntax `ModuleName.function_name(args)`.

To define your own module, you use the `defmodule` expression. Inside the module, you define functions using the `def` expression. The following listing demonstrates the definition of a module.

Listing 2.1. Defining a module (geometry.ex)

```
defmodule Geometry do   #1
```

```
  def rectangle_area(a, b) do  #2
    a * b          #2
  end              #2
end                #3
```

There are two ways you can use this module. You can copy/paste this definition directly into `iex` — as mentioned, almost anything can be typed into the shell. Another way is to tell iex to interpret the file while starting:

```
$ iex geometry.ex
```

Regardless of which method you choose, the effect is the same. The code is compiled, and the resulting module is loaded into the runtime and can be used from the shell session. Let's try it:

```
$ iex geometry.ex

iex(1)> Geometry.rectangle_area(6, 7)  #1
42        #2
```

That was simple! You created a `Geometry` module, loaded it into a shell session, and used it to compute the area of a rectangle.

ⓘ Note

As you may have noticed, the filename has the .ex extension. This is a common convention for Elixir source files.

In the source code, a module must be defined in a single file. A single file may contain multiple module definitions:

```
defmodule Module1 do
  ...
end

defmodule Module2 do
  ...
end
```

A module name must follow certain rules. It starts with an uppercase letter and is usually written in CamelCase style. A module name can consist of

alphanumerics, underscores, and the dot (.) character. The latter is often used to organize modules hierarchically:

```
defmodule Geometry.Rectangle do
  ...
end

defmodule Geometry.Circle do
  ...
end
```

You can also nest module definitions:

```
defmodule Geometry do
  defmodule Rectangle do
    ...
  end
  ...
end
```

The inner module can be referenced with `Geometry.Rectangle`.

Note that there is nothing special about the dot character. It's just one of the allowed characters in a module name. The compiled version doesn't record any hierarchical relations between the modules.

This is typically used to organize the modules in some meaningful hierarchy that is easier to navigate when reading the code. In addition, this informal scoping can eliminate possible name clashes. For example, consider two libraries, one implementing a JSON encoder, and another implementing an XML encoder. If both libraries defined the module called `Encoder`, you couldn't use them both in the same project. However, if the modules are called `Json.Encoder` and `Xml.Encoder`, then the name clash is avoided. For this reason, it's customary to add some common prefix to all module names in a project. Usually the application or the library name is used for this purpose.

2.3.2 Functions

A function must always be a part of a module. Function names follow the same conventions as variables: they start with a lowercase letter or

underscore character, followed by a combination of alphanumerics and underscores.

As with variables, function names can end with the `?` and `!` characters. The `?` character is often used to indicate a function that returns either *true* or *false*. Placing the character `!` at the end of the name indicates a function that may raise a runtime error. Both of these are conventions, rather than rules, but it's best to follow them and respect the community style.

Functions can be defined using the `def` macro:

```
defmodule Geometry do
  def rectangle_area(a, b) do       #1
    ...      #2
  end
end
```

The definition starts with the `def` expression, followed by the function name, the argument list, and the body enclosed in a `do…end` block. Because you're dealing with a dynamic language, there are no type specifications for arguments.

Note

Notice that `defmodule` and `def` aren't referred to as keywords. That's because they're not! Instead, these are examples of Elixir *macros*. You don't need to worry about how this works yet; it's explained a bit later in this chapter. If it helps, you can think of `def` and `defmodule` as keywords, but be aware that this isn't exactly true.

If a function has no arguments, you can omit the parentheses:

```
defmodule Program do
  def run do
    ...
  end
end
```

What about the return value? Recall that in Elixir, everything that has a return

value is an `expression`. The return value of a function is the return value of its last expression. There's no explicit return in Elixir.

Note

Given that there's no explicit return, you might wonder how complex functions work. This will be covered in detail in chapter 3, where you'll learn about branching and conditional logic. The general rule is to keep functions short and simple, which makes it easy to compute the result and return it from the last expression.

You saw an example of returning a value in listing 2.1, but let's repeat it:

```
defmodule Geometry do
  def rectangle_area(a, b) do
    a * b           #1
  end
end
```

You can now verify this. Start the shell again, and then try the `rectangle_area` function:

```
$ iex geometry.ex

iex(1)> Geometry.rectangle_area(3, 2)     #1
6      #2
```

If a function body consists of a single expression, you can use a condensed form and define it in a single line:

```
defmodule Geometry do
  def rectangle_area(a, b), do: a * b
end
```

To call a function defined in another module, you use the module name followed by the function name:

```
iex(1)> Geometry.rectangle_area(3, 2)
6
```

Of course, you can always store the function result to a variable:

```
iex(2)> area = Geometry.rectangle_area(3, 2)  #1
6

iex(3)> area  #2
6
```

Parentheses are optional in Elixir, so you can omit them:

```
iex(4)> Geometry.rectangle_area 3, 2
6
```

Personally, I find that omitting parentheses makes the code ambiguous, so my advice is to always include them when calling a function.

Using a code formatter

Starting with version 1.6, Elixir ships with the code formatter, which you can use to format your code in a consistent style, and stop worrying about lower-level style decisions, such as layouts or parentheses usage.

For example, after formatting the following code snippet,

```
defmodule Client
do
def run do
Geometry.rectangle_area 3,2
end
end
```

you'll end up with this nice-looking code:

```
defmodule Client do
  def run do
    Geometry.rectangle_area(3, 2)
  end
end
```

You can format your code with the `mix format` task (https://hexdocs.pm/mix/Mix.Tasks.Format.html), or install a formatter extension in your editor of choice.

If a function being called resides in the same module, you can omit the

module prefix:

```
defmodule Geometry do
  def rectangle_area(a, b) do
    a * b
  end

  def square_area(a) do
    rectangle_area(a, a)     #1
  end
end
```

Given that Elixir is a functional language, you'll often need to combine functions, passing the result of one function as the argument to the next one. Elixir comes with a built-in operator, |>, called the *pipe operator,* that does exactly this:

```
iex(5)> -5 |> abs() |> Integer.to_string() |> IO.puts()
5
```

This code is transformed at compile time into the following:

```
iex(6)> IO.puts(Integer.to_string(abs(-5)))
5
```

More generally, the pipe operator places the result of the previous call as the first argument of the next call. So the following code,

```
prev(arg1, arg2) |> next(arg3, arg4)
```

is translated at compile time to this:

```
next(prev(arg1, arg2), arg3, arg4)
```

Arguably, the pipeline version is more readable because the sequence of execution is read from left to right. The pipe operator looks especially elegant in source files, where you can lay out the pipeline over multiple lines:

```
-5             #1
|> abs()       #2
|> Integer.to_string()     #3
|> IO.puts()    #4
```

If you paste the previous pipeline chain into an iex session, you'll notice that each intermediate result is printed to the console:

```
iex(1)> -5
-5                                    #1

iex(2)> |> abs()
5                                     #1

iex(3)> |> Integer.to_string()
"5"                                   #1

iex(4)> |> IO.puts()
5                                     #1
```

Recall that iex evalutes the Elixir expression as soon as it is complete and valid. In this example, each line completes a valid Elixir expression, such as `-5` or `-5 |> abs()`, and therefore each intermediate result is printed.

2.3.3 Function arity

Arity is a fancy name for the number of arguments a function receives. A function is uniquely identified by its containing module, its name, and its arity. Take a look at the following function:

```
defmodule Rectangle do
  def area(a, b) do        #1
    ...
  end
end
```

The function `Rectangle.area` receives two arguments, so it's said to be a function of arity 2. In the Elixir world, this function is often called `Rectangle.area/2`, where the `/2` part denotes the function's arity.

Why is this important? Because two functions with the same name but different arities are two different functions, as the following example demonstrates.

Listing 2.2. Functions with the same name but different arities (arity_demo.ex)

```
defmodule Rectangle do
  def area(a), do: area(a, a)  #1

  def area(a, b), do: a * b    #2
end
```

Load this module into the shell, and try the following:

```
iex(1)> Rectangle.area(5)
25

iex(2)> Rectangle.area(5,6)
30
```

As you can see, these two functions act completely differently. The name might be overloaded, but the arities differ, so we talk about them as two distinct functions, each having its own implementation.

It usually makes no sense for different functions with the same name to have completely different implementations. More commonly, a lower-arity function delegates to a higher-arity function, providing some default arguments. This is what happens in listing 2.2, where `Rectangle.area/1` delegates to `Rectangle.area/2`.

Let's look at another example.

Listing 2.3. Same-name functions, different arities, default params (arity_calc.ex)

```
defmodule Calculator do
  def add(a), do: add(a, 0)      #1
  def add(a, b), do: a + b       #2
end
```

Again, a lower-arity function is implemented in terms of a higher-arity one. This pattern is so frequent that Elixir allows you to specify defaults for arguments by using the \\ operator followed by the argument's default value:

```
defmodule Calculator do
  def add(a, b \\ 0), do: a + b    #1
end
```

This definition generates two functions exactly like in listing 2.3.

You can set the defaults for any combination of arguments:

```
defmodule MyModule do
  def fun(a, b \\ 1, c, d \\ 2) do     #1
    a + b + c + d
  end
end
```

Always keep in mind that default values generate multiple functions of the same name with different arities. The previous code generates three functions: `MyModule.fun/2`, `MyModule.fun/3`, and `MyModule.fun/4`, with the following implementations:

```
def fun(a, c), do: fun(a, 1, c, 2)
def fun(a, b, c), do: fun(a, b, c, 2)
def fun(a, b, c, d), do: a + b + c + d
```

Because arity distinguishes multiple functions of the same name, it's not possible to have a function accept a variable number of arguments. There's no counterpart of C's … or JavaScript's `arguments`.

2.3.4 Function visibility

When you define a function using the `def` macro, the function is made public: it can be called by anyone else. In Elixir terminology, it's said that the function is *exported*. You can also use the `defp` macro to make the function private. A private function can be used only inside the module it's defined in. The following example demonstrates this.

Listing 2.4. Module with a public and a private function (private_fun.ex)

```
defmodule TestPrivate do
  def double(a) do     #1
    sum(a, a)          #2
  end

  defp sum(a, b) do    #3
    a + b
  end
```

```
end
```

The module `TestPrivate` defines two functions. The function `double` is exported and can be called from outside. Internally, it relies on the private function `sum` to do its work.

Let's try this in the shell. Load the module, and do the following:

```
iex(1)> TestPrivate.double(3)
6

iex(2)> TestPrivate.sum(3, 4)
** (UndefinedFunctionError) function TestPrivate.sum/2
...
```

As you can see, the private function can't be invoked outside the module.

2.3.5 Imports and aliases

Calling functions from another module can sometimes be cumbersome because you need to reference the module name. If your module often calls functions from another module, you can import that other module into your own. Importing a module allows you to call its public functions without prefixing them with the module name:

```
defmodule MyModule do
  import IO      #1

  def my_function do
    puts "Calling imported function."   #2
  end
end
```

Of course, you can import multiple modules. In fact, the standard library's `Kernel` module is automatically imported into every module. `Kernel` contains functions that are often used, so automatic importing makes their use easier.

Note

You can see what functions are available in the Kernel module by looking in

the online documentation at https://hexdocs.pm/elixir/Kernel.html.

Another expression, `alias`, makes it possible to reference a module under a different name:

```
defmodule MyModule do
  alias IO, as: MyIO     #1

  def my_function do
    MyIO.puts("Calling imported function.")  #2
  end
end
```

Aliases can be useful if a module has a long name. For example, if your application is heavily divided into a deeper module hierarchy, it can be cumbersome to reference modules via fully qualified names. Aliases can help with this. For example, let's say you have a `Geometry.Rectangle` module. You can alias it in your client module and use a shorter name:

```
defmodule MyModule do
  alias Geometry.Rectangle, as: Rectangle     #1

  def my_function do
    Rectangle.area(...)       #2
  end
end
```

In the preceding example, the alias of `Geometry.Rectangle` is the last part in its name. This is the most common use of `alias`, so Elixir allows you to skip the `as` option in this case:

```
defmodule MyModule do
  alias Geometry.Rectangle    #1

  def my_function do
    Rectangle.area(...)     #2
  end
end
```

Aliases can help you reduce some noise, especially if you call functions from a long-named module many times.

2.3.6 Module attributes

The purpose of module attributes is twofold: they can be used as compile-time constants, and you can register any attribute, which can then be queried in runtime. Let's look at an example.

The following module provides basic functions for working with circles:

```
iex(1)> defmodule Circle do
          @pi 3.14159    #1

          def area(r), do: r*r*@pi    #2
          def circumference(r), do: 2*r*@pi
        end

iex(2)> Circle.area(1)
3.14159

iex(3)> Circle.circumference(1)
6.28318
```

Notice how you define a module directly in the shell. This is permitted and makes it possible to experiment without storing any files on disk.

The important thing about the `@pi` constant is that it exists only during the compilation of the module, when the references to it are inlined.

Moreover, an attribute can be `registered`, which means it will be stored in the generated binary and can be accessed at runtime. Elixir registers some module attributes by default. For example, the attributes `@moduledoc` and `@doc` can be used to provide documentation for modules and functions:

```
defmodule Circle do
  @moduledoc "Implements basic circle functions"
  @pi 3.14159

  @doc "Computes the area of a circle"
  def area(r), do: r*r*@pi

  @doc "Computes the circumference of a circle"
  def circumference(r), do: 2*r*@pi
end
```

To try this, however, you need to generate a compiled file. Here's a quick way to do it. Save this code to the circle.ex file somewhere, and then run `elixirc circle.ex`. This will generate the file Elixir.Circle.beam. Next, start the `iex` shell from the same folder. You can now retrieve the attribute at runtime:

```
iex(1)> Code.fetch_docs(Circle)
{:docs_v1, 2, :elixir, "text/markdown",
 %{"en" => "Implements basic circle functions"}, %{},
 [
   {{:function, :area, 1}, 5, ["area(r)"],
    %{"en" => "Computes the area of a circle"}, %{}},
   {{:function, :circumference, 1}, 8, ["circumference(r)"],
    %{"en" => "Computes the circumference of a circle"}, %{}}
 ]}
```

More interesting is that other tools from the Elixir ecosystem know how to work with these attributes. For example, you can use the help feature of `iex` to see the module's documentation:

```
iex(2)> h Circle      #1

                                Circle
Implements basic circle functions

iex(3)> h Circle.area    #2
                           def area(r)
Computes the area of a circle
```

Furthermore, you can use the `ex_doc` tool (see https://hexdocs.pm/ex_doc) to generate HTML documentation for your project. This is the way Elixir documentation is produced, and if you plan to build more complex projects, especially something that will be used by many different clients, you should consider using `@moduledoc` and `@doc`.

The underlying point is that registered attributes can be used to attach meta information to a module, which can then be used by other Elixir (and even Erlang) tools. There are many other preregistered attributes, and you can also register your own custom attributes. Take a look at the documentation for the `Module` module (https://hexdocs.pm/elixir/Module.html) for more details.

Type specifications

Type specifications (often called *typespecs*) are another important feature based on attributes. This feature allows you to provide type information for your functions, which can later be analyzed with a static analysis tool called `dialyzer` (https://www.erlang.org/doc/man/dialyzer.html).

Here's how we can extend the `Circle` module to include typespecs:

```
defmodule Circle do
  @pi 3.14159

  @spec area(number) :: number      #1
  def area(r), do: r*r*@pi

  @spec circumference(number) :: number  #2
  def circumference(r), do: 2*r*@pi
end
```

Here you use the `@spec` attribute to indicate that both functions accept and return a number.

Typespecs provide a way of compensating for the lack of a static type system. This can be useful in conjunction with the `dialyzer` tool to perform static analysis of your programs. Furthermore, typespecs allow you to better-document your functions. Remember that Elixir is a dynamic language, so function inputs and output can't be easily deduced by looking at the function's signature. Typespecs can help significantly with this, and I can personally attest that it's much easier to understand someone else's code when typespecs are provided.

For example, look at the typespec for the Elixir function:

```
List.insert_at/3:
@spec insert_at(list, integer, any) :: list
```

Even without looking at the code or reading the docs, you can reasonably guess that this function inserts a term of any type (third argument) to a list (first argument) at a given position (second argument) and returns a new list.

You won't be using typespecs in this book, mostly to keep the code as short as possible. But if you plan to build more complex systems, my advice is to seriously consider using typespecs. You can find a detailed reference in the

official docs at https://hexdocs.pm/elixir/typespecs.html.

2.3.7 Comments

Comments in Elixir start with the character #, which indicates that the rest of the line is a comment:

```
# This is a comment
a = 3.14 # so is this
```

Block comments aren't supported. If you need to comment multiple lines, prefix each one with the # character.

At this point, we're done with the basics of functions and modules. You're aware of the primary code-organization techniques. With that out of our way, it's time to look at the Elixir type system.

2.4 Understanding the type system

At its core, Elixir uses the Erlang type system. Consequently, integration with Erlang libraries is usually simple. The type system itself is reasonably simple, but if you're coming from a classical OO language, you'll find it significantly different from what you're used to. This section covers basic Elixir types and also discusses some implications of immutability. To begin, let's look at numbers.

2.4.1 Numbers

Numbers can be integers or floats, and they work mostly as you'd expect:

```
iex(1)> 3    #1
3

iex(2)> 0xFF   #2
255

iex(3)> 3.14   #3
3.14
```

```
iex(4)> 1.0e-2  #4
0.01
```

Standard arithmetic operators are supported:

```
iex(5)> 1 + 2 * 3
7
```

The division operator / works differently than you might expect. It always returns a float value:

```
iex(6)> 4/2
2.0
```

```
iex(7)> 3/2
1.5
```

To perform integer division or to calculate the remainder, you can use auto-imported Kernel functions:

```
iex(8)> div(5,2)
2
```

```
iex(9)> rem(5,2)
1
```

As added syntactic sugar, you can use the underscore character as a visual delimiter:

```
iex(10)> 1_000_000
1000000
```

There's no upper limit on an integer's size, and you can use arbitrarily large numbers:

```
iex(11)> 999999999999999999999999999999999999999999999999999999999
99999999999999999999999999999999999999999999999999999999999999
```

If you're worried about memory size, it's best to consult the official Erlang memory guide at https://www.erlang.org/doc/efficiency_guide/advanced.html. An integer takes up as much space as needed to accommodate the number, whereas a float

occupies either 32 or 64 bits, depending on the build architecture of the virtual machine. Floats are internally represented in IEEE 754-1985 (binary precision) format.

2.4.2 Atoms

Atoms are literal named constants. They're similar to symbols in Ruby or enumerations in C/C++. Atom constants start with a colon character, followed by a combination of alphanumerics and/or underscore characters:

```
:an_atom
:another_atom
```

It's possible to use spaces in the atom name with the following syntax:

```
:"an atom with spaces"
```

An atom consists of two parts: the *text* and the *value*. The atom text is whatever you put after the colon character. At runtime, this text is kept in the *atom table*. The value is the data that goes into the variable, and it's merely a reference to the atom table.

This is exactly why atoms are best used for named constants. They're efficient both memory- and performance-wise. When you say

```
variable = :some_atom
```

the variable doesn't contain the entire text, but only a reference to the atom table. Therefore, memory consumption is low, the comparisons are fast, and the code is still readable.

Aliases

There's another syntax for atom constants. You can omit the beginning colon and start with an uppercase character:

```
AnAtom
```

This is called an *alias,* and at compile time it's transformed into

:"Elixir.AnAtom". We can easily check this in the shell:

```
iex(1)> AnAtom == :"Elixir.AnAtom"
true
```

When you use an alias, the compiler implicitly adds the `Elixir.` prefix to its text and generates the atom. But if an alias already contains the `Elixir.` prefix, it's not added. Consequently, the following also works:

```
iex(2)> AnAtom == Elixir.AnAtom
true
```

You may recall from earlier that you can also use aliases to give alternate names to modules:

```
iex(3)> alias IO, as: MyIO

iex(4)> MyIO.puts("Hello!")
Hello!
```

It's no accident that the term *alias* is used for both things. When you write `alias IO, as: MyIO`, you instruct the compiler to transform `MyIO` into `IO`. Resolving this further, the final result emitted in the generated binary is `:Elixir.IO`. Therefore, with an alias set up, the following also holds:

```
iex(5)> MyIO == Elixir.IO
true
```

All of this may seem strange, but it has an important underlying purpose. Aliases support the proper resolution of modules. This will be discussed at the end of the chapter when we revisit modules and look at how they're loaded at runtime.

Atoms as booleans

It may come as a surprise that Elixir doesn't have a dedicated boolean type. Instead, the atoms `:true` and `:false` are used. As syntactic sugar, Elixir allows you to reference these atoms without the starting colon character:

```
iex(1)> :true == true
true
```

```
iex(2)> :false == false
true
```

The term *boolean* is still used in Elixir to denote an atom that has a value of either `:true` or `:false`. The standard logical operators work with boolean atoms:

```
iex(1)> true and false
false
```

```
iex(2)> false or true
true
```

```
iex(3)> not false
true
```

```
iex(4)> not :an_atom_other_than_true_or_false
** (ArgumentError) argument error
```

Always keep in mind that a boolean is just an atom that has a value of `true` or `false`.

Nil and truthy values

Another special atom is `:nil`, which works somewhat similarly to `null` from other languages. You can reference `nil` without a colon:

```
iex(1)> nil == :nil
true
```

The atom `nil` plays a role in Elixir's additional support for *truthness*, which works similarly to the way it's used in mainstream languages such as C/C++ and Ruby. The atoms `nil` and `false` are treated as *falsy* values, whereas everything else is treated as a *truthy* value.

This property can be used with Elixir's short-circuit operators `||`, `&&`, and `!`. The operator `||` returns the first expression that isn't falsy:

```
iex(1)> nil || false || 5 || true
5
```

Because both `nil` and `false` are falsy expressions, the number 5 is returned. Notice that subsequent expressions won't be evaluated at all. If all expressions evaluate to a falsy value, the result of the last expression is returned.

The operator `&&` returns the second expression, but only if the first expression is truthy. Otherwise, it returns the first expression without evaluating the second one:

```
iex(1)> true && 5
5

iex(2)> false && 5
false

iex(3)> nil && 5
nil
```

Short-circuiting can be used for elegant operation chaining. For example, if you need to fetch a value from cache, a local disk, or a remote database, you can do something like this:

```
read_cached() || read_from_disk() || read_from_database()
```

Similarly, you can use the operator `&&` to ensure that certain conditions are met:

```
database_value = connection && read_data(connection)
```

In both examples, short-circuit operators make it possible to write concise code without resorting to complicated nested conditional expressions.

2.4.3 Tuples

Tuples are something like untyped structures, or records, and they're most often used to group a fixed number of elements together. The following snippet defines a tuple consisting of a person's name and age:

```
iex(1)> person = {"Bob", 25}
{"Bob", 25}
```

To extract an element from the tuple, you can use the `Kernel.elem/2` function, which accepts a tuple and the zero-based index of the element. Recall that the `Kernel` module is auto-imported, so you can call `elem` instead of `Kernel.elem`:

```
iex(2)> age = elem(person, 1)
25
```

To modify an element of the tuple, you can use the `Kernel.put_elem/3` function, which accepts a tuple, a zero-based index, and the new value of the field in the given position:

```
iex(3)> put_elem(person, 1, 26)
{"Bob", 26}
```

The function `put_elem` doesn't modify the tuple. It returns the new version, keeping the old one intact. Recall that data in Elixir is immutable, so you can't do an in-memory modification of a value. You can verify that the previous call to `put_elem` didn't change the `person` variable:

```
iex(4)> person
{"Bob", 25}
```

So how can you use the `put_elem` function, then? You need to store its result to another variable:

```
iex(5)> older_person = put_elem(person, 1, 26)
{"Bob", 26}

iex(6)> older_person
{"Bob", 26}
```

Recall that variables can be rebound, so you can also do the following:

```
iex(7)> person = put_elem(person, 1, 26)
{"Bob", 26}
```

By doing this, you've effectively rebound the `person` variable to the new memory location. The old location isn't referenced by any other variable, so it's eligible for garbage collection.

Note

You may wonder if this approach is memory efficient. In most cases, there will be little data copying, and the two variables will share as much memory as possible. This will be explained later in this section, when we discuss immutability.

Tuples are most appropriate for grouping a small, fixed number of elements together. When you need a dynamically sized collection, you can use lists.

2.4.4 Lists

Lists in Erlang are used to manage dynamic, variable-sized collections of data. The syntax deceptively resembles arrays from other languages:

```
iex(1)> prime_numbers = [2, 3, 5, 7]
[2, 3, 5, 7]
```

Lists may look like arrays, but they work like singly linked lists. To do something with the list, you have to traverse it. Therefore, most of the operations on lists have an O(n) complexity, including the `Kernel.length/1` function, which iterates through the entire list to calculate its length:

```
iex(2)> length(prime_numbers)
4
```

List utility functions

There are many operations you can do with lists, but this section mentions only a couple of the most basic ones. For a detailed reference, see the documentation for the `List` module (https://hexdocs.pm/elixir/List.html). There are also many helpful services in the `Enum` module (https://hexdocs.pm/elixir/Enum.html).

The `Enum` module deals with many different enumerable structures and is not limited to lists. The concept of enumerables will be explained in detail in chapter 4, when we discuss protocols.

To get an element of a list, you can use the `Enum.at/2` function:

```
iex(3)> Enum.at(prime_numbers, 3)
7
```

`Enum.at` is again an O(n) operation: it iterates from the beginning of the list to the desired element. Lists are never a good fit when direct access is called for. For those purposes, tuples, maps, or a higher-level data structure is appropriate.

You can check whether a list contains a particular element with the help of the `in` operator:

```
iex(4)> 5 in prime_numbers
true

iex(5)> 4 in prime_numbers
false
```

To manipulate lists, you can use functions from the `List` module. For example, `List.replace_at/3` modifies the element at a certain position:

```
iex(6)> List.replace_at(prime_numbers, 0, 11)
[11, 3, 5, 7]
```

As was the case with tuples, the modifier doesn't mutate the variable, but returns the modified version of it, which you need to store to another variable:

```
iex(7)> new_primes = List.replace_at(prime_numbers, 0, 11)
[11, 3, 5, 7]
```

Or you can rebind to the same one:

```
iex(8)> prime_numbers = List.replace_at(prime_numbers, 0, 11)
[11, 3, 5, 7]
```

You can insert a new element at the specified position with the `List.insert_at/3` function:

```
iex(9)> List.insert_at(prime_numbers, 3, 13)    #1
[11, 3, 5, 13, 7]
```

To append to the end, you can use a negative value for the insert position:

```
iex(10)> List.insert_at(prime_numbers, -1, 13)    #1
[11, 3, 5, 7, 13]
```

Like most list operations, modifying an arbitrary element has a complexity of O(n). In particular, appending to the end is expensive because it always takes *n* steps, *n* being the length of the list.

In addition, the dedicated operator ++ is available. It concatenates two lists:

```
iex(11)> [1, 2, 3] ++ [4, 5]
[1, 2, 3, 4, 5]
```

Again, the complexity is O(n), *n* being the length of the left list (the one you're appending to). In general, you should avoid adding elements to the end of a list. Lists are most efficient when new elements are pushed to the top, or popped from it. To understand why, let's look at the recursive nature of lists.

Recursive list definition

An alternative way of looking at lists is to think of them as recursive structures. A list can be represented by a pair (*head, tail*), where *head* is the first element of the list and *tail* "points" to the (*head, tail*) pair of the remaining elements, as illustrated in figure 2.1.

Figure 2.1. Recursive structure of the list [1, 2, 3, 4]

If you're familiar with Lisp, then you know this concept as *cons cells*.

In Elixir, there's a special syntax to support recursive list definition:

```
a_list = [head | tail]
```

head can be any type of data, whereas `tail` is itself a list. If `tail` is an empty list, it indicates the end of the entire list.

Let's look at some examples:

```
iex(1)> [1 | []]
[1]

iex(2)> [1 | [2 | []]]
[1, 2]

iex(3)> [1 | [2]]
[1, 2]

iex(4)> [1 | [2, 3, 4]]
[1, 2, 3, 4]
```

This is just another syntactical way of defining lists, but it illustrates what a list is. It's a pair with two values: a head and a tail, the tail being itself a list.

The following snippet is a canonical recursive definition of a list:

```
iex(1)> [1 | [2 | [3 | [4 | []]]]]
[1, 2, 3, 4]
```

Of course, nobody wants to write expressions like this one. But it's important that you're always aware that, internally, lists are recursive structures of (*head*, *tail*) pairs.

To get the head of the list, you can use the `hd` function. The tail can be obtained by calling the `tl` function:

```
iex(1)> hd([1, 2, 3, 4])
1

iex(2)> tl([1, 2, 3, 4])
[2, 3, 4]
```

Both operations are O(1), because they amount to reading one or the other value from the (*head*, *tail*) pair.

 Note

For the sake of completeness, it should be mentioned that the tail doesn't need to be a list. It can be any type. When the tail isn't a list, it's said that the list is improper, and most of the standard list manipulations won't work. Improper lists have some special uses, but we won't deal with them in this book.

Once you know the recursive nature of the list, it's simple and efficient to push a new element to the top of the list:

```
iex(1)> a_list = [5, :value, true]
[5, :value, true]

iex(2)> new_list = [:new_element | a_list]
[:new_element, 5, :value, true]
```

Construction of the `new_list` is an O(1) operation, and no memory copying occurs — the tail of the `new_list` *is* the `a_list`. To understand how this works, let's discuss the internal details of immutability a bit.

2.4.5 Immutability

As has been mentioned before, Elixir data can't be mutated. Every function returns the new, modified version of the input data. You have to take the new version into another variable or rebind it to the same symbolic name. In any case, the result resides in another memory location. The modification of the input will result in some data copying, but generally, most of the memory will be shared between the old and the new version.

Let's take a closer look at how this works.

Modifying tuples

Let's start with tuples. A modified tuple is always a complete, shallow copy of the old version. Consider the following code:

```
a_tuple = {a, b, c}
```

```
new_tuple = put_elem(a_tuple, 1, b2)
```

Figure 2.2. Modifying a tuple creates a shallow copy of it.
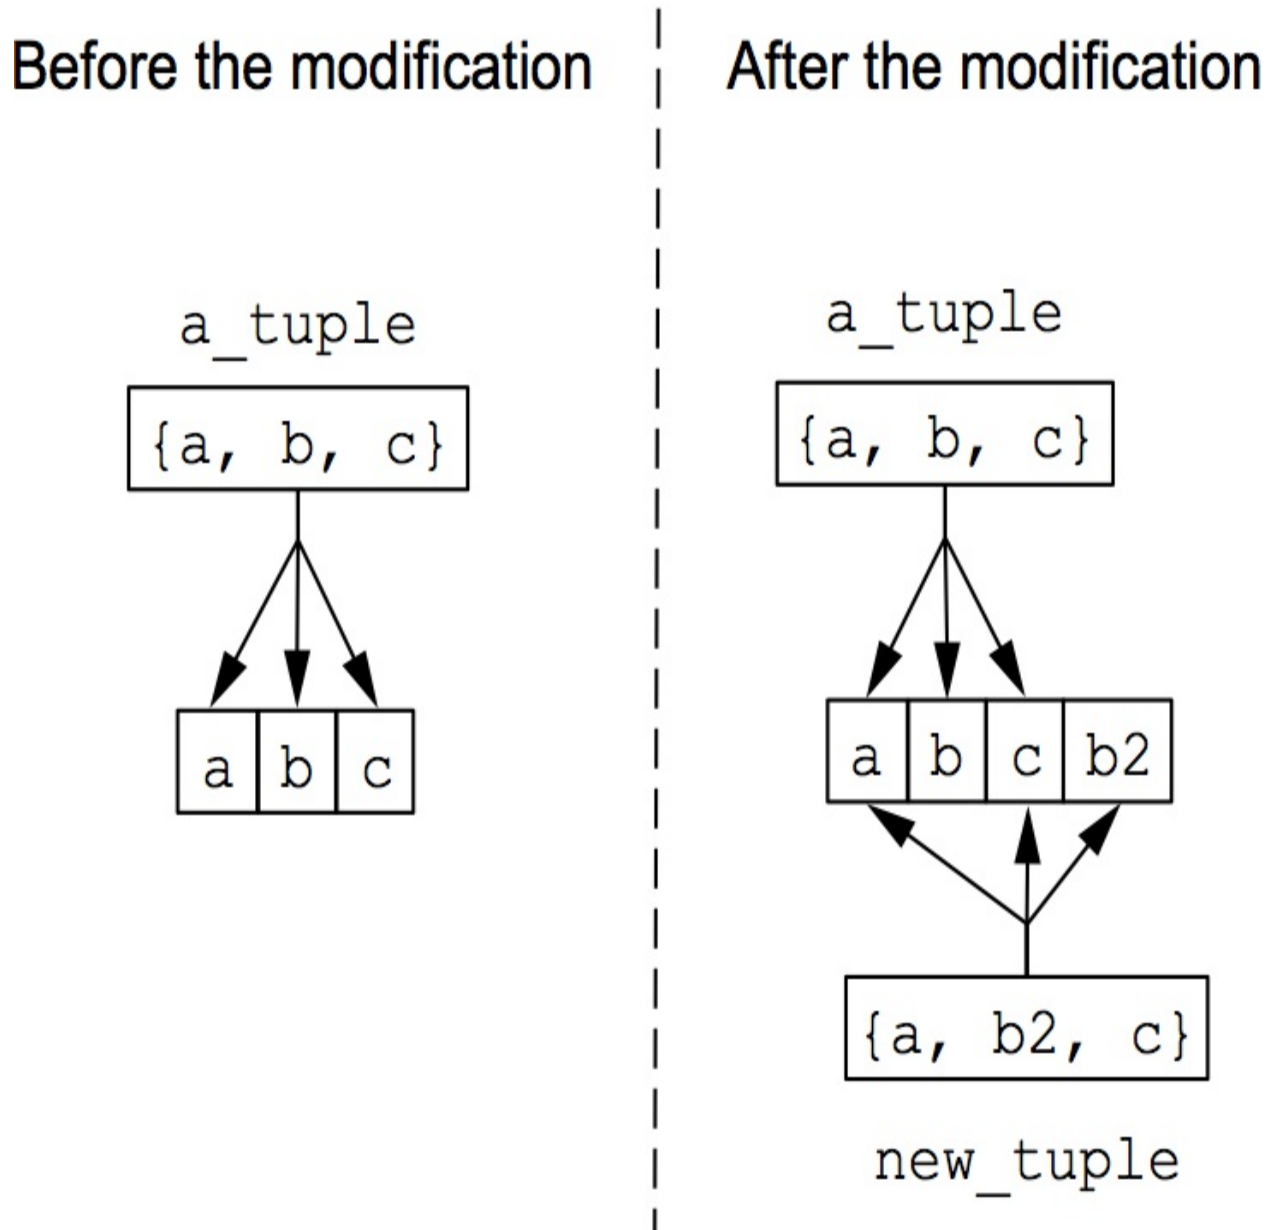
The variable `new_tuple` will contain a shallow copy of `a_tuple`, differing only in the second element, as illustrated in figure 2.2.

Both tuples reference variables a and c and whatever is in those variables is shared (and not duplicated) between both tuples. `new_tuple` is a shallow copy of the original `a_tuple`.

What happens if you rebind a variable?

In this case, after rebinding, the variable `a_tuple` references another memory location. The old location of `a_tuple` isn't accessible and is available for garbage collection. The same holds for the variable `b` referenced by the old version of the tuple, as illustrated in figure 2.3.
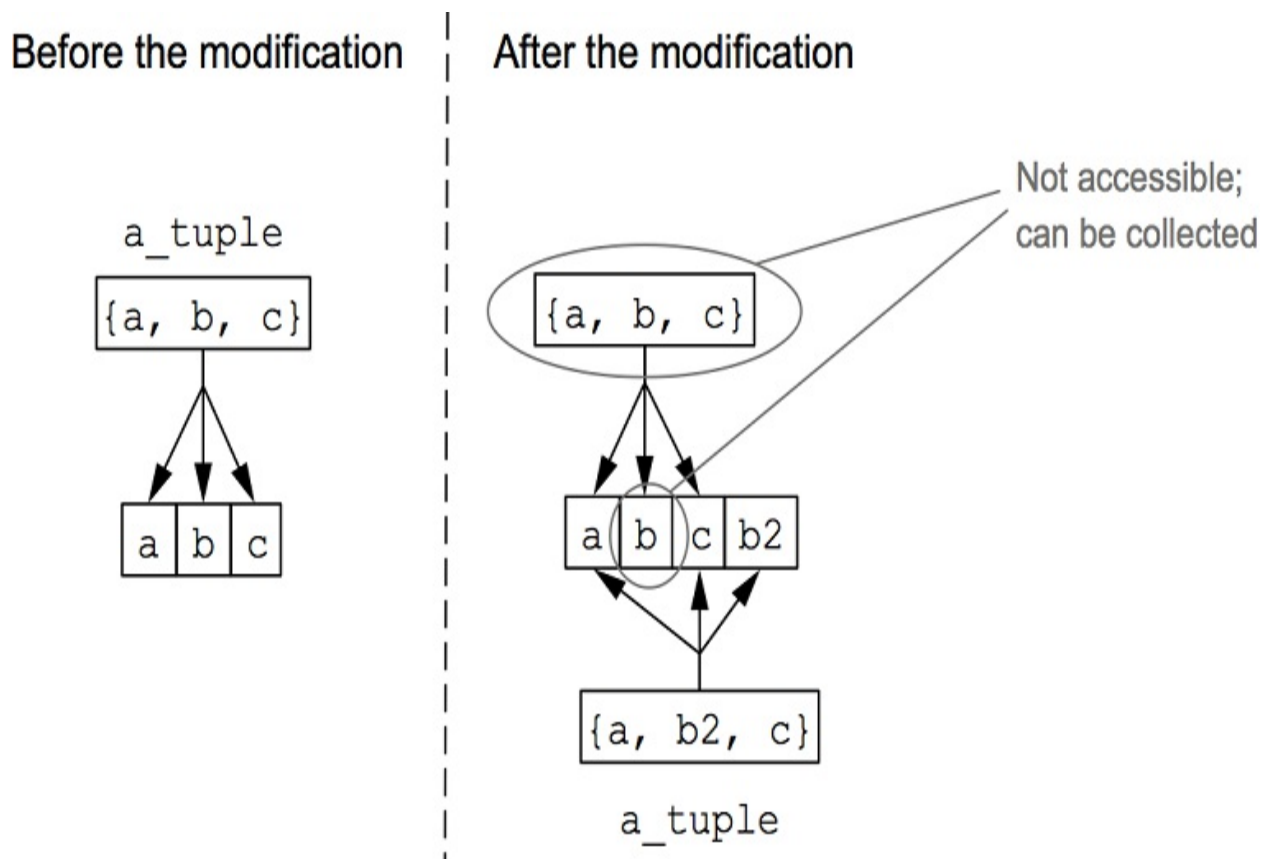
Figure 2.3. Rebinding a tuple makes the old data garbage-collectible.

Keep in mind that tuples are always copied, but the copying is shallow. Lists, however, have different properties.

Modifying lists

When you modify the nth element of a list, the new version will contain shallow copies of the first $n - 1$ elements, followed by the modified element. After that, the tails are completely shared, as illustrated in figure 2.4.

Figure 2.4. Modifying a list

This is precisely why adding elements to the end of a list is expensive. To append a new element at the tail, you have to iterate and (shallow) copy the entire list.

In contrast, pushing an element to the top of a list doesn't copy anything, which makes it the least expensive operation, as illustrated in figure 2.5.

Figure 2.5. Pushing a new element to the top of the list

Before the push

After the push

a | b | c | d | e

new_top

a | b | c | d | e

The tail of the new list is the old list.
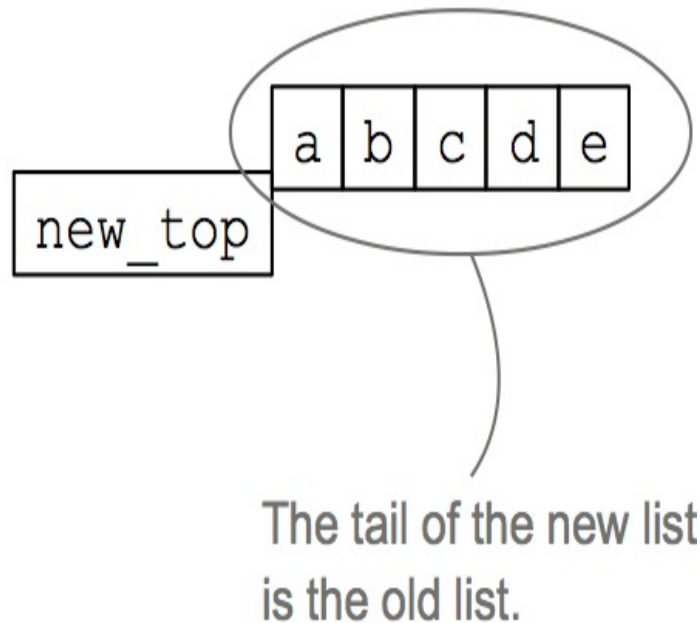
In this case, the new list's tail *is* the previous list. This is often used in Elixir programs when iteratively building lists. In such cases, it's best to push consecutive elements to the top, and then, after the list is constructed, reverse the entire list in a single pass.

Benefits

Immutability may seem strange, and you may wonder about its purpose. There are two important benefits of immutability: side-effect-free functions and data consistency.

Given that data can't be mutated, you can treat most functions as side-effect-free transformations. They take an input and return a result. More complicated programs are written by combining simpler transformations:

```
def complex_transformation(data) do
  data
  |> transformation_1(...)
  |> transformation_2(...)
  ...
  |> transformation_n(...)
end
```

This code relies on the previously mentioned pipe operator that chains two functions together, feeding the result of the previous call as the first argument of the next call.

Side-effect-free functions are easier to analyze, understand, and test. They have well-defined inputs and outputs. When you call a function, you can be sure that no variable will be implicitly changed. Whatever the function does, you must take its result and do something with it.

Note

Elixir isn't a pure functional language, so functions may still have side effects. For example, a function may write something to a file and issue a database or network call, which causes it to produce a side effect. But you can be certain that a function won't modify the value of any variable.

The implicit consequence of immutable data is the ability to hold all versions of a data structure in the program. This, in turn, makes it possible to perform atomic in-memory operations. Let's say you have a function that performs a series of transformations:

```
def complex_transformation(original_data) do
  original_data
  |> transformation_1(...)
  |> transformation_2(...)
  ...
end
```

This code starts with the original data and passes it through a series of transformations, each one returning the new, modified version of the input. If something goes wrong, the function `complex_transformation` can return `original_data`, which will effectively roll back all of the transformations

performed in the function. This is possible because none of the transformations modifies the memory occupied by `original_data`.

This concludes our look at basic immutability theory. It may still be unclear how to properly use immutable data in more complex programs. This topic will be revisited in chapter 4, where we'll deal with higher-level data structures.

2.4.6 Maps

A map is a key/value store, where keys and values can be any term. Maps have dual usage in Elixir. They're used to power dynamically sized key/value structures, but they're also used to manage simple records — a couple of well-defined named fields bundled together. Let's take a look at these cases separately.

Dynamically sized maps

An empty map can be created with the `%{}` expression:

```
iex(1)> empty_map = %{}
```

A map with some values can be created with the following syntax:

```
iex(2)> squares = %{1 => 1, 2 => 4, 3 => 9}
```

You can also prepopulate a map with the `Map.new/1` function. The function takes an enumerable where each element is a tuple of size two (a pair):

```
iex(3)> squares = Map.new([{1, 1}, {2, 4}, {3, 9}])
%{1 => 1, 2 => 4, 3 => 9}
```

To fetch a value at the given key, you can use the following approach:

```
iex(4)> squares[2]
4

iex(5)> squares[4]
nil
```

In the second expression, you get a `nil` because no value is associated with the given key.

A similar result can be obtained with `Map.get/3`. On the surface, this function behaves like `[]`. But `Map.get/3` allows you to specify the default value, which is returned if the key isn't found. If this default isn't provided, `nil` will be returned:

```
iex(6)> Map.get(squares, 2)
4

iex(7)> Map.get(squares, 4)
nil

iex(8)> Map.get(squares, 4, :not_found)
:not_found
```

Notice that in the last expression you don't precisely know whether there's no value under the given key, or if the value is `:not_found`. If you want to precisely distinguish between these cases, you can use `Map.fetch/2`:

```
iex(9)> Map.fetch(squares, 2)
{:ok, 4}

iex(10)> Map.fetch(squares, 4)
:error
```

As you can see, in the successful case you'll get a value in the shape of `{:ok, value}`. This format makes it possible to precisely detect the case when the key isn't present.

Sometimes you want to proceed only if the key is in the map, and raise an exception otherwise. This can be done with the `Map.fetch!/2` function:

```
iex(11)> Map.fetch!(squares, 2)
4

iex(12)> Map.fetch!(squares, 4)
** (KeyError) key 4 not found in: %{1 => 1, 2 => 4, 3 => 9}
    (stdlib) :maps.get(4, %{1 => 1, 2 => 4, 3 => 9})
```

To store a new element to the map, you can use `Map.put/3`:

```
iex(13)> squares = Map.put(squares, 4, 16)
%{1 => 1, 2 => 4, 3 => 9, 4 => 16}

iex(14)> squares[4]
16
```

There are a bunch of other helpful functions in the `Map` module, such as
`Map.update/4` or `Map.delete/2`. You can look into the official module
documentation at https://hexdocs.pm/elixir/Map.html. In addition, a map is
also enumerable, which means that all the functions from the `Enum` module
can work with maps.

Structured data

Maps are the go-to type for managing key/value data structures of an
arbitrary size. But they're also frequently used in Elixir to combine a couple
of fields into a single structure. This use case somewhat overlaps that of
tuples, but it provides the advantage of allowing you to access fields by
name.

Let's look at an example. In the following snippet, you'll create a map that
represents a single person:

```
iex(1)> bob = %{:name => "Bob", :age => 25, :works_at => "Initech
```

If keys are atoms, you can write this so it's slightly shorter:

```
iex(2)> bob = %{name: "Bob", age: 25, works_at: "Initech"}
```

To retrieve a field, you can use the `[]` operator:

```
iex(3)> bob[:works_at]
"Initech"

iex(4)> bob[:non_existent_field]
nil
```

Atom keys again receive special syntax treatment. The following snippet
fetches a value stored under the `:age` key:

```
iex(5)> bob.age
```

With this syntax you'll get an error if you try to fetch a nonexistent field:

```
iex(6)> bob.non_existent_field
** (KeyError) key :non_existent_field not found
```

To change a field value, you can use the following syntax:

```
iex(7)> next_years_bob = %{bob | age: 26}
%{age: 26, name: "Bob", works_at: "Initech"}
```

This syntax can be used to change multiple attributes as well:

```
iex(8)> %{bob | age: 26, works_at: "Initrode"}
%{age: 26, name: "Bob", works_at: "Initrode"}
```

But you can only modify values that already exist in the map. This makes the update syntax a perfect choice for powering maps that represent structures. If you mistype the field name, you'll get an immediate runtime error:

```
iex(9)> %{bob | works_in: "Initech"}
** (KeyError) key :works_in not found
```

Using maps to hold structured data is a frequent pattern in Elixir. The common pattern is to provide all the fields while creating the map, using atoms as keys. If the value for some field isn't available, you can set it to `nil`. Such a map, then, always has all the fields. You can modify the map with the update expression, and fetch a desired field with the `a_map.some_field` expression.

Of course, such data is still a map, so you can also use the functions from the `Map` module, such as `Map.put/3`, or `Map.fetch/2`. But these functions are usually suitable for the cases where maps are used to manage a dynamical key/value structure.

2.4.7 Binaries and bitstrings

A binary is a chunk of bytes. You can create binaries by enclosing the byte sequence between << and >> operators. The following snippet creates a 3-byte

binary:

```
iex(1)> <<1, 2, 3>>
<<1, 2, 3>>
```

Each number represents the value of the corresponding byte. If you provide a byte value bigger than 255, it's truncated to the byte size:

```
iex(2)> <<256>>
<<0>>

iex(3)> <<257>>
<<1>>

iex(4)> <<512>>
<<0>>
```

You can specify the size of each value and thus tell the compiler how many bits to use for that particular value:

```
iex(5)> <<257::16>>
<<1, 1>>
```

This expression places the number 257 into 16 bits of consecutive memory space. The output indicates that you use 2 bytes, both having a value of 1. This is due to the binary representation of 257, which in 16-bit form is written 00000001 00000001.

The size specifier is in bits and need not be a multiplier of 8. The following snippet creates a binary by combining two 4-bit values:

```
iex(6)> <<1::4, 15::4>>
<<31>>
```

The resulting value has 1 byte and is represented in the output using the normalized form 31 (0001 1111).

If the total size of all the values isn't a multiple of 8, the binary is called a *bitstring* — a sequence of bits:

```
iex(7)> <<1::1, 0::1, 1::1>>
<<5::size(3)>>
```

You can also concatenate two binaries or bitstrings with the operator <>:

```
iex(8)> <<1, 2>> <> <<3, 4>>
<<1, 2, 3, 4>>
```

There's much more that can be done with binaries, but for the moment we'll put them aside. The most important thing you need to know about binaries is that they're consecutive sequences of bytes. Binaries play an important role in support for strings.

2.4.8 Strings

It may come as a surprise, but Elixir doesn't have a dedicated string type. Instead, strings are represented by using either a binary or a list type.

Binary strings

The most common way to use strings is to specify them with the familiar double-quotes syntax:

```
iex(1)> "This is a string"
"This is a string"
```

The result is printed as a string, but underneath it's a binary — nothing more than a consecutive sequence of bytes.

Elixir provides support for embedded string expressions. You can use #{} to place an Elixir expression in a string constant. The expression is immediately evaluated, and its string representation is placed at the corresponding location in the string:

```
iex(2)> "Embedded expression: #{3 + 0.14}"
"Embedded expression: 3.14"
```

Classical \ escaping works as you're used to:

```
iex(3)> "\r \n \" \\"
```

And strings don't have to finish on the same line:

```
iex(4)> "
       This is
       a multiline string
       "
```

Elixir provides another syntax for declaring strings, so-called *sigils*. In this approach, you enclose the string inside ~s():

```
iex(5)> ~s(This is also a string)
"This is also a string"
```

Sigils can be useful if you want to include quotes in a string:

```
iex(6)> ~s("Do... or do not. There is no try." -Master Yoda)
"\"Do... or do not. There is no try.\" -Master Yoda"
```

There's also an uppercase version ~S that doesn't handle interpolation or escape characters (\):

```
iex(7)> ~S(Not interpolated #{3 + 0.14})
"Not interpolated \#{3 + 0.14}"

iex(8)> ~S(Not escaped \n)
"Not escaped \\n"
```

Finally, there's a special *heredocs* syntax, which supports better formatting for multiline strings. Heredocs strings start with a triple double-quote. The ending triple double-quote must be on its own line:

```
iex(9)> """
       Heredoc must end on its own line """
       """
"Heredoc must end on its own line \"\"\"\n"
```

Because strings are binaries, you can concatenate them with the <> operator:

```
iex(10)> "String" <> " " <> "concatenation"
"String concatenation"
```

Many helper functions are available for working with binary strings. Most of them reside in the String module (https://hexdocs.pm/elixir/String.html).

Character lists

The alternative way of representing strings is to use single-quote syntax:

```
iex(1)> 'ABC'
'ABC'
```

This creates a *character list,* which is essentially a list of integers in which each element represents a single character.

The previous result is exactly the same as if you manually construct the list of integers:

```
iex(2)> [65, 66, 67]
'ABC'
```

As you can see, even the runtime doesn't distinguish between a list of integers and a character list. When a list consists of integers that represent printable characters, it's printed to the screen in the string form.

Just like with binary strings, there are syntax counterparts for various definitions of character lists:

```
iex(3)> 'Interpolation: #{3 + 0.14}'
'Interpolation: 3.14'

iex(4)> ~c(Character list sigil)
'Character list sigil'

iex(5)> ~C(Unescaped sigil #{3 + 0.14})
'Unescaped sigil \#{3 + 0.14}'

iex(6)> '''
        Heredoc
        '''
'Heredoc\n'
```

Character lists aren't compatible with binary strings. Most of the operations from the `String` module won't work with character lists. In general, you should prefer binary strings over character lists. Occasionally, some functions may work only with character lists. This mostly happens with pure Erlang libraries. In this case, you can convert a binary string to a character list

version using the `String.to_charlist/1` function:

```
iex(7)> String.to_charlist("ABC")
'ABC'
```

To convert a character list to a binary string, you can use `List.to_string/1`.

In general, you should prefer binary strings as much as possible, using character lists only when some third-party library (most often written in pure Erlang) requires it.

2.4.9 First-class functions

In Elixir, a function is a first-class citizen, which means it can be assigned to a variable. Here, assigning a function to a variable doesn't mean calling the function and storing its result to a variable. Instead, the function definition itself is assigned, and you can use the variable to call the function.

Let's look at some examples. To create a function variable, you can use the `fn` expression:

```
iex(1)> square = fn x ->
          x * x
        end
```

The variable `square` now contains a function that computes the square of a number. Because the function isn't bound to a global name, it's also called an *anonymous* function or a *lambda*.

Notice that the list of arguments isn't enclosed in parentheses. Technically, you can use parentheses here, but the prevalent convention, also enforced by the Elixir formatter, is to omit parentheses. In contrast, a list of arguments to a named function should be enclosed in parentheses. At first glance, this looks inconsistent, but there's a good reason for this convention, which will be explained in chapter 3.

You can call this function by specifying the variable name followed by a dot (`.`) and the arguments:

```
iex(2)> square.(5)
25
```

 Note

You may wonder why the dot operator is needed here. The motivation behind the dot operator is to make the code more explicit. When you encounter a `square.(5)` expression in the source code, you know an anonymous function is being invoked. In contrast, the expression `square(5)` is invoking a named function defined somewhere else in the module. Without the dot operator, you'd have to parse the surrounding code to understand whether you're calling a named or an anonymous function.

Because functions can be stored in a variable, they can be passed as arguments to other functions. This is often used to allow clients to parameterize generic logic. For example, the function `Enum.each/2` implements the generic iteration — it can iterate over anything enumerable, such as lists. The function `Enum.each/2` takes two arguments: an enumerable and a one-arity lambda (an anonymous function that accepts one argument). It iterates through the enumerable and calls the lambda for each element. The clients provide the lambda to specify what they want to do with each element.

The following snippet uses `Enum.each` to print each value of a list to the screen:

```
iex(3)> print_element = fn x -> IO.puts(x) end     #1
iex(4)> Enum.each(
          [1, 2, 3],
          print_element     #2
        )

1    #3
2    #3
3    #3

:ok     #4
```

Of course, you don't need a temp variable to pass the lambda to `Enum.each`:

```
iex(5)> Enum.each(
```

```
        [1, 2, 3],
        fn x -> IO.puts(x) end     #1
    )

1
2
3
```

Notice how the lambda just forwards all arguments to `IO.puts`, doing no other meaningful work. For such cases, Elixir makes it possible to directly reference the function and have a more compact lambda definition. Instead of writing `fn x -> IO.puts(x) end`, you can write `&IO.puts/1`.

The `&` operator, also known as the *capture* operator, takes the full function qualifier — a module name, a function name, and an arity — and turns that function into a lambda that can be assigned to a variable. You can use the capture operator to simplify the call to `Enum.each`:

```
iex(6)> Enum.each(
        [1, 2, 3],
        &IO.puts/1    #1
    )
```

The capture operator can also be used to shorten the lambda definition, making it possible to omit explicit argument naming. For example, you can turn this definition

```
iex(7)> lambda = fn x, y, z -> x * y + z end
```

into a more compact form:

```
iex(8)> lambda = &(&1 * &2 + &3)
```

This snippet creates a three-arity lambda. Each argument is referred to via the `&n` placeholder, which identifies the *n*th argument of the function. You can call this lambda like any other:

```
iex(9)> lambda.(2, 3, 4)
10
```

The return value 10 amounts to 2 * 3 + 4, as specified in the lambda definition.

Closures

A lambda can reference any variable from the outside scope:

```
iex(1)> outside_var = 5
5

iex(2)> my_lambda = fn ->
  IO.puts(outside_var)    #1
end

iex(3)> my_lambda.()
5
```

As long as you hold the reference to `my_lambda`, the variable `outside_var` is also accessible. This is also known as *closure*: by holding a reference to a lambda, you indirectly hold a reference to all variables it uses, even if those variables are from the external scope.

A closure always captures a specific memory location. Rebinding a variable doesn't affect the previously defined lambda that references the same symbolic name:

```
iex(1)> outside_var = 5
iex(2)> lambda = fn -> IO.puts(outside_var) end  #1
iex(3)> outside_var = 6     #2
iex(4)> lambda.()     #3
5
```

The preceding code illustrates another important point. Normally, after you have rebound `outside_var` to the value 6, the original memory location would be eligible for garbage collection. But because the `lambda` function captures the original location (the one that holds the number 5), and you're still referencing that lambda, the original location isn't available for garbage collection.

2.4.10 Other built-in types

There are a couple of types I still haven't presented. We won't deal with them in depth, but it's worth mentioning them for the sake of completeness:

- A *reference* is an almost unique piece of information in a BEAM instance. It's generated by calling `Kernel.make_ref/0` (or `make_ref`). According to the Elixir documentation, a reference will reoccur after approximately 2^82 calls. But if you restart a BEAM instance, reference generation starts from the beginning, so its uniqueness is guaranteed only during the lifetime of the BEAM instance.
- A *pid* (process identifier) is used to identify an Erlang process. Pids are important when cooperating between concurrent tasks, and you'll learn about them in chapter 5 when we deal with Erlang processes.
- The *port identifier* is important when using ports. It's a mechanism used in Erlang to talk to the outside world. File I/O and communication with external programs are done through ports. Ports are outside the scope of this book.

With that, we've covered all the basic data types. As you can see, Elixir has a simple type system consisting of only a handful of data types.

Of course, higher-level types are also available, which build on these basic types to provide additional functionality. Let's look at some of the most important ones that ship with Elixir.

2.4.11 Higher-level types

The built-in types just mentioned are inherited from the Erlang world. After all, Elixir code runs on BEAM, so its type system is heavily influenced by the Erlang foundations. But on top of these basic types, Elixir provides some higher-level abstractions. The ones most frequently used are `Range`, `Keyword`, `MapSet`, `Date`, `Time`, `NaiveDateTime`, and `DateTime`. Let's examine each of them.

Range

A *range* is an abstraction that allows you to represent a range of numbers. Elixir even provides a special syntax for defining ranges:

```
iex(1)> range = 1..2
```

You can ask whether a number falls in the range by using the `in` operator:

```
iex(2)> 2 in range
true

iex(3)> -1 in range
false
```

Ranges are enumerable, so functions from the `Enum` module know how to work with them. Earlier you met `Enum.each/2`, which iterates through an enumerable. The following example uses this function with a range to print the first three natural numbers:

```
iex(4)> Enum.each(
          1..3,
          &IO.puts/1
        )

1
2
3
```

Range isn't a special type. Internally it's represented as a map that contains range boundaries. Therefore, a range memory footprint is small and constant, regardless of the number of elements it represents. A million-number range is still just a small map.

For more information on ranges see the documentation for the `Range` module (https://hexdocs.pm/elixir/Range.html).

Keyword lists

A *keyword* list is a special case of a list, where each element is a two-element tuple, and the first element of each tuple is an atom. The second element can be of any type. Let's look at an example:

```
iex(1)> days = [{:monday, 1}, {:tuesday, 2}, {:wednesday, 3}]
```

Elixir supports a slightly shorter syntax for defining a keyword list:

```
iex(2)> days = [monday: 1, tuesday: 2, wednesday: 3]
```

Both expressions yield the same result: a list of pairs. Arguably, the second one is a bit more elegant.

Keyword lists are often used for small-size key/value structures, where keys are atoms. Many useful functions are available in the `Keyword` module (https://hexdocs.pm/elixir/Keyword.html). For example, you can use `Keyword.get/2` to fetch the value for a key:

```
iex(3)> Keyword.get(days, :monday)
1
```

```
iex(4)> Keyword.get(days, :noday)
nil
```

Just as with maps, you can use the operator `[]` to fetch a value:

```
iex(5)> days[:tuesday]
2
```

Don't let that fool you, though. Because you're dealing with a list, the complexity of a lookup operation is O(n).

Keyword lists are most often useful for allowing clients to pass an arbitrary number of optional arguments. For example, the result of the function `IO.inspect`, which prints a string representation of a term to the console, can be controlled by providing additional options through a keyword list:

```
iex(6)> IO.inspect([100, 200, 300])     #1
[100, 200, 300]    #1

iex(7)> IO.inspect([100, 200, 300], [width: 3])    #2
[100,       #2
 200,       #2
 300]       #2
```

In fact, this pattern is so frequent that Elixir allows you to omit the square brackets if the last argument is a keyword list:

```
iex(8)> IO.inspect([100, 200, 300], width: 3, limit: 1)
[100,
 ...]
```

Notice in this example that you're still sending two arguments to
`IO.inspect/2`: a list of numbers, and a two-element keyword list. But this
snippet demonstrates how to simulate optional arguments. You can accept a
keyword list as the last argument of your function, and make that argument
default to an empty list:

```
def my_fun(arg1, arg2, opts \\ []) do
  ...
end
```

Your clients can then pass options via the last argument. Of course, it's up to
you to check the contents in the `opts` argument and perform some conditional
logic depending on what the caller has sent you.

You may wonder if it's better to use maps instead of keywords for optional
arguments. A keyword list can contain multiple values for the same key. In
addition, you can control the ordering of keyword list elements — something
that isn't possible with maps. Finally, many functions in standard libraries of
Elixir and Erlang take their options as keyword lists. It's best to stick to the
existing convention and accept optional parameters via keyword lists.

MapSet

A `MapSet` is the implementation of a set — a store of unique values, where a
value can be of any type.

Let's look at some examples:

```
iex(1)> days = MapSet.new([:monday, :tuesday, :wednesday])   #1
#MapSet<[:monday, :tuesday, :wednesday]>

iex(2)> MapSet.member?(days, :monday)      #2
true      #2

iex(3)> MapSet.member?(days, :noday)       #3
false      #3

iex(4)> days = MapSet.put(days, :thursday)     #4
#MapSet<[:monday, :thursday, :tuesday, :wednesday]>
```

As you can see, you can manipulate the set using the function from the

MapSet module. For a detailed reference, refer to the official documentation at https://hexdocs.pm/elixir/MapSet.html.

A MapSet is also an enumerable, so you can pass it to functions from the Enum module:

```
iex(5)> Enum.each(days, &IO.puts/1)
monday
thursday
tuesday
wednesday
```

As you can tell from the output, MapSet doesn't preserve the ordering of the items.

Times and dates

Elixir has a couple of modules for working with date and time types: Date, Time, DateTime, and NaiveDateTime.

A date can be created with the ~D sigil. The following example creates a date that represents January 31, 2023:

```
iex(1)> date = ~D[2023-01-31]
~D[2023-01-31]
```

Once you've created the date, you can retrieve its individual fields:

```
iex(2)> date.year
2023

iex(3)> date.month
1
```

Similarly, you can represent a time with the ~T sigil, by providing hours, minutes, seconds, and microseconds:

```
iex(1)> time = ~T[11:59:12.00007]

iex(2)> time.hour
11
```

```
iex(3)> time.minute
59
```

There are also some useful functions available in the modules `Date` (https://hexdocs.pm/elixir/Date.html) and `Time` (https://hexdocs.pm/elixir/Time.html).

In addition to these two types, you can also work with datetimes using the `NaiveDateTime` and `DateTime` modules. The naive version can be created with the `~N` sigil:

```
iex(1)> naive_datetime = ~N[2023-01-31 11:59:12.000007]

iex(2)> naive_datetime.year
2023

iex(3)> naive_datetime.hour
11
```

The `DateTime` module can be used to work with datetimes in some timezone. Unlike with other types, no sigil is available. Instead, you can create a datetime by using `DateTime` functions:

```
iex(4)> datetime = DateTime.from_naive!(naive_datetime, "Etc/UTC"

iex(5)> datetime.year
2023

iex(6)> datetime.hour
11

iex(7)> datetime.time_zone
"Etc/UTC"
```

You can refer to the reference documentation, available at https://hexdocs.pm/elixir/NaiveDateTime.html and https://hexdocs.pm/elixir/DateTime.html, for more details on working with these types.

2.4.12 IO lists

An IO list is a special sort of list that's useful for incrementally building

output that will be forwarded to an I/O device, such as a network or a file. Each element of an IO list must be one of the following:

- An integer in the range of 0 to 255
- A binary
- An IO list

In other words, an IO list is a deeply nested structure in which leaf elements are plain bytes (or binaries, which are again a sequence of bytes). For example, here's "Hello world" represented as a convoluted IO list:

```
iex(1)> iolist = [[['H', 'e'], "llo,"], " worl", "d!"]
```

Notice how you can combine character lists and binary strings into a deeply nested list.

Many I/O functions can work directly and efficiently with such data. For example, you can print this structure to the screen:

```
iex(2)> IO.puts(iolist)
Hello, world!
```

Under the hood, the structure is flattened, and you can see the human-readable output. You'll get the same effect if you send an IO list to a file or a network socket.

IO lists are useful when you need to incrementally build a stream of bytes. Lists usually aren't good in this case, because appending to a list is an O(n) operation. In contrast, appending to an IO list is O(1), because you can use nesting. Here's a demonstration of this technique:

```
iex(3)> iolist = []      #1
iolist = [iolist, "This"]      #2
iolist = [iolist, " is"]       #2
iolist = [iolist, " an"]       #2
iolist = [iolist, " IO list."]  #2

[[[[[], "This"], " is"], " an"], " IO list."]    #3
```

Here, you append to an IO list by creating a new list with two elements: a previous version of the IO list and the suffix that's appended. Each such

operation is O(1), so this is performant. And, of course, you can send this
data to an IO function:

```
iex(4)> IO.puts(iolist)
This is an IO list.
```

This concludes our initial tour of the type system. We've covered most of the
basics, and we'll expand on this theory throughout the book as the need
arises. Next, it's time to learn a bit about Elixir operators.

2.5 Operators

You've been using various operators throughout this chapter, and in this
section we'll take a systematic look at the ones most commonly used in
Elixir. Most of the operators are defined in the `Kernel` module, and you can
refer to the module documentation for a detailed description.

Let's start with arithmetic operators. These include the standard `+`, `-`, `*`, and `/`,
and they work mostly as you'd expect, with the exception that the division
operator always returns a float, as explained earlier in this chapter when we
were dealing with numbers.

The comparison operators are more or less similar to what you're used to.
They're listed in table 2.1.

Table 2.1. Comparison operators

Operator	Description
`===`, `!==`	Strict equality/inequality
`==`, `!=`	Weak equality/inequality
`<`, `>`, `>=`, `<=`	Less than, greater than, less than or equal, greater than or equal

The only thing we need to discuss here is the difference between strict and weak equality. This is relevant only when comparing integers to floats:

```
iex(1)> 1 == 1.0    #1
true

iex(2)> 1 === 1.0   #2
false
```

Logical operators work on Boolean atoms. You saw them earlier in the discussion of atoms, but I'll repeat them once more: `and`, `or`, and `not`.

Unlike logical operators, short-circuit operators work with the concept of truthness: the atoms `false` and `nil` are treated as falsy, and everything else is treated as truthy. The `&&` operator returns the first expression if it's falsy; otherwise it returns the second expression. The `||` operator returns the first expression if it's truthy; otherwise it returns the second expression. The unary operator `!` returns `false` if the value is truthy; otherwise it returns `true`.

The operators presented here aren't the only ones available (for example, you've also seen the pipe operator `|>`). But these are the most common ones, so it was worth mentioning them in one place. You can find the detailed information on operators at https://hexdocs.pm/elixir/operators.html.

Many operators are functions

It's worth noting that many operators in Elixir are actually functions. For example, instead of calling `a+b`, you can call `Kernel.+(a,b)`. Of course, no one would ever want to write this kind of code, but operator functions have a benefit when turned into anonymous functions. For example, you can create a two-arity lambda that sums two numbers by calling `&Kernel.+/2` or the shorter `&+/2`. Such lambdas can then be used with various enumeration and stream functions, as I'll explain in chapter 3.

We've almost completed our initial tour of the language. One thing remains: Elixir macros.

2.6 Macros

Macros are arguably one of the most important features Elixir brings to the table, compared to plain Erlang. They make it possible to perform powerful code transformations at compile time, thus reducing boilerplate and providing elegant, mini-DSL expressions.

Macros are a fairly complex subject, and it would take a small book to treat them extensively. Because this book is more oriented toward runtime and BEAM, and macros are a somewhat advanced feature that should be used sparingly, I won't provide a detailed treatment. But you should have a general idea of how macros work because many Elixir features are powered by them.

A *macro* consists of Elixir code that can change the semantics of the input code. A macro is always called at compile time; it receives the parsed representation of the input Elixir code, and it has the opportunity to return an alternative version of that code.

Let's clear this up with an example. `unless` (an equivalent of `if not`) is a simple macro provided by Elixir:

```
unless some_expression do
  block_1
else
  block_2
end
```

`unless` isn't a special keyword. It's a macro (meaning an Elixir function) that transforms the input code into something like this:

```
if some_expression do
  block_2
else
  block_1
end
```

Such a transformation isn't possible with C-style macros, because the code of the expression can be arbitrarily complex and deeply nested. But in Elixir macros (which are heavily inspired by Lisp), you already work on a parsed

source representation, so you'll have access to the expression and both blocks in separate variables.

The end effect is that many parts of Elixir are written in Elixir with the help of macros. This includes the `unless` and `if` expressions, and also `defmodule` and `def`. Whereas other languages usually use keywords for such features, in Elixir they're built on top of a much smaller language core.

The main point to take away is that macros are compile-time code transformers. Whenever I say that something is a macro, the underlying implication is that it runs at compile time and produces alternative code.

Special forms

The Elixir compiler treats some language constructs in a special way. Such constructs are called *special forms* (https://hexdocs.pm/elixir/Kernel.SpecialForms.html). Some examples include the capture syntax `&(…)`, `for` comprehension (presented in chapter 3), `receive` expression (chapter 5), and `try` blocks (chapter 8).

For details, you may want to look at the official meta-programming guide (https://elixir-lang.org/getting-started/meta/quote-and-unquote.html). Meanwhile, we're done with our initial tour of the Elixir language. But before we finish this chapter, we should discuss some important aspects of the underlying runtime.

2.7 Understanding the runtime

As has been mentioned, the Elixir runtime is a BEAM instance. Once the compiling is done and the system is started, Erlang takes control. It's important to be familiar with some details of the virtual machine so you can understand how your systems work.

First, let's look at the significance of modules in the runtime.

2.7.1 Modules and functions in the runtime

Regardless of how you start the runtime, an OS process for the BEAM instance is started, and everything runs inside that process. This is true even when you're using the `iex` shell. If you need to find this OS process, you can look it up under the name `beam`.

Once the system is started, you run some code, typically by calling functions from modules. How does the runtime access the code? The VM keeps track of all modules loaded in memory. When you call a function from a module, BEAM first checks whether the module is loaded. If it is, the code of the corresponding function is executed. Otherwise the VM tries to find the compiled module file — the bytecode — on the disk and then load it and execute the function.

ⓘ Note

The previous description reveals that each compiled module resides in a separate file. A compiled module file has the extension .beam (for Bogdan/Björn's Erlang Abstract Machine). The name of the file corresponds to the module name.

Module names and atoms

Let's recall how modules are defined:

```
defmodule Geometry do
  ...
end
```

Also recall from the discussion about atoms that `Geometry` is an alias that corresponds to `:"Elixir.Geometry"`, as demonstrated in the following snippet:

```
iex(1)> Geometry == :"Elixir.Geometry"
true
```

This isn't an accident. When you compile the source containing the `Geometry` module, the file generated on the disk is named `Elixir.Geometry.beam`, regardless of the name of the input source file. In fact, if multiple modules are

defined in a single source file, the compiler will produce multiple .beam files that correspond to those modules. You can try this by calling the Elixir compiler (`elixirc`) from the command line:

```
$ elixirc source.ex
```

where the file source.ex defines a couple of modules. Assuming there are no syntax errors, you'll see multiple .beam files generated on the disk.

In the runtime, module names are aliases, and as I said, aliases are atoms. The first time you call the function of a module, BEAM tries to find the corresponding file on the disk. The VM looks for the file in the current folder and then in the code paths.

When you start BEAM with Elixir tools (such as `iex`), some code paths are predefined for you. You can add additional code paths by providing the `-pa` switch:

```
$ iex -pa my/code/path -pa another/code/path
```

You can check which code paths are used at runtime by calling the Erlang function `:code.get_path`.

If the module is loaded, the runtime doesn't search for it on the disk. This can be used when starting the shell, to auto-load modules:

```
$ iex my_source.ex
```

This command compiles the source file and then immediately loads all generated modules. Notice that in this case, .beam files aren't saved to disk. The `iex` tool performs an in-memory compilation.

Similarly, you can define modules in the shell:

```
iex(1)> defmodule MyModule do      #1
           def my_fun, do: :ok
         end

iex(2)> MyModule.my_fun
:ok
```

Again, the bytecode isn't saved to the disk in this case.

Pure Erlang modules

You've already seen how to call a function from a pure (non-Elixir) Erlang module. Let's talk a bit about this syntax:

```
:code.get_path   #1
```

In Erlang, modules also correspond to atoms. Somewhere on the disk is a file named code.beam that contains the compiled code of the `:code` module. Erlang uses simple filenames, which is the reason for this call syntax. But the rules are the same as with Elixir modules. In fact, Elixir modules are nothing more than Erlang modules with fancier names (such as `Elixir.MyModule`).

You can create modules with simple names in Elixir (although this isn't recommended):

```
defmodule :my_module do
  ...
end
```

Compiling the source file that contains such a definition will generate my_module.beam on the disk.

The important thing to remember from this discussion is that at runtime, module names are atoms. And somewhere on the disk is an *xyz*.beam file, where *xyz* is the expanded form of an alias (such as `Elixir.MyModule` when the module is named `MyModule`).

Dynamically calling functions

Somewhat related to this discussion is the ability to dynamically call functions at runtime. This can be done with the help of the `Kernel.apply/3` function:

```
iex(1)> apply(IO, :puts, ["Dynamic function call."])
Dynamic function call.
```

`Kernel.apply/3` receives three arguments: the module atom, the function atom, and the list of arguments passed to the function. Together, these three arguments, often called MFA (module, function, arguments), contain all the information needed to call an exported (public) function. `Kernel.apply/3` can be useful when you need to make a runtime decision about which function to call.

2.7.2 Starting the runtime

There are multiple ways of starting BEAM. So far, you've been using `iex`, and you'll continue to do so for some time. But let's quickly look at all the possible ways to start the runtime.

Interactive shell

When you start the shell, the BEAM instance is started underneath, and the Elixir shell takes control. The shell takes the input, *interprets* it, and prints the result.

It's important to be aware that input is interpreted, because that means it won't be as performant as the compiled code. This is generally fine, because you use the shell only to experiment with the language. But you shouldn't try to measure performance directly from `iex`.

On the other hand, modules are always compiled. Even if you define a module in the shell, it will be compiled and loaded in memory, so there will be no performance hit.

Running scripts

The `elixir` command can be used to run a single Elixir source file. Here's the basic syntax:

```
$ elixir my_source.ex
```

When you start this, the following actions take place:

1. The BEAM instance is started.
2. The file my_source.ex is compiled in memory, and the resulting modules are loaded to the VM. No .beam file is generated on the disk.
3. Whatever code resides outside of a module is interpreted.
4. Once everything is finished, BEAM is stopped.

This is generally useful for running scripts. In fact, it's recommended that such a script have an .exs extension, the trailing "s" indicating that it's a script.

The following listing shows a simple Elixir script.

Listing 2.5. Elixir script (script.exs)

```
defmodule MyModule do
  def run do
    IO.puts("Called MyModule.run")
  end
end

MyModule.run     #1
```

You can execute this script from the command line:

```
$ elixir script.exs
```

This call first does the in-memory compilation of the `MyModule` module and then calls `MyModule.run`. After the call to `MyModule.run` finishes, the BEAM instance is stopped.

If you don't want a BEAM instance to terminate, you can provide the `--no-halt` parameter:

```
$ elixir --no-halt script.exs
```

This is most often useful if your main code (outside a module) starts concurrent tasks that perform all the work. In this situation, your main call finishes as soon as the concurrent tasks are started, and BEAM is immediately terminated (and no work is done). Providing the `--no-halt` option keeps the entire system alive and running.

The mix tool

The `mix` tool is used to manage projects that are made up of multiple source files. Whenever you need to build a production-ready system, `mix` is your best option.

To create a new mix project, you can call `mix new project_name` from the command line:

```
$ mix new my_project
```

This creates a new folder named my_project containing a couple of subfolders and files. You can change to the my_project folder and compile the entire project:

```
$ cd my_project
$ mix compile

$ mix compile
Compiling 1 file (.ex)
Generated my_project app
```

The compilation goes through all the files from the lib folder and places the resulting .beam files in the ebin folder.

You can execute various `mix` commands on the project. For example, the generator created the module `MyProject` with the single function `hello/0`. You can invoke it with `mix run`:

```
$ mix run -e "IO.puts(MyProject.hello())"
world
```

The generator also create a couple of tests, which can be executed with `mix test`:

```
$ mix test
..

Finished in 0.03 seconds
2 tests, 0 failures
```

Regardless of how you start the mix project, it ensures that the ebin folder (where the .beam files are placed) is in the load path so the VM can find your modules.

You'll use `mix` a lot once you start creating more complex systems. For now, there's no need to go into any more detail.

2.8 Summary

- Elixir code is divided into modules and functions.
- Elixir is a dynamic language. The type of a variable is determined by the value it holds.
- Data is immutable — it can't be modified. A function can return the modified version of the input that resides in another memory location. The modified version shares as much memory as possible with the original data.
- The most important primitive data types are numbers, atoms, and binaries.
- There is no Boolean type. Instead, the atoms `true` and `false` are used.
- There is no nullability. The atom `nil` can be used for this purpose.
- There is no string type. Instead, you can use either binaries (recommended) or lists (when needed).
- The built-in complex types are tuples, lists, and maps. Tuples are used to group a small, fixed-size number of fields. Lists are used to manage variable-size collections. A map is a key/value data structure.
- `Range`, keyword lists, `MapSet`, `Date`, `Time`, `NaiveDateTime`, and `DateTime` are abstractions built on top of the existing built-in types.
- Functions are first-class citizens.
- Module names are atoms (or aliases) that correspond to .beam files on the disk.
- There are multiple ways of starting programs: `iex`, `elixir`, and the `mix` tool.

3 Control flow

This chapter covers

- Understanding pattern matching
- Working with multiclause functions
- Using conditional expressions
- Working with loops

Now that you're familiar with Elixir's basic building blocks, it's time to look at some typical low-level idioms of the language. In this chapter, we'll deal with conditionals and loops. As you'll see, these work differently than in many imperative languages.

Classical conditional constructs such as `if` and `case` are often replaced with *multiclause* functions, and there are no loop statements such as `while`. But you can still solve problems of arbitrary complexity in Elixir, and the resulting code is no more complicated than a typical OO solution.

All this may sound a bit radical, which is why conditionals and loops receive a detailed treatment in this chapter. But before we start discussing branching and looping, you need to learn about the important underlying mechanism: pattern matching.

3.1 Pattern matching

As mentioned in chapter 2, the = operator isn't an assignment. In the expression a = 1, we bind the variable a to the value 1. The operator = is called the *match operator*, and the assignment-like expression is an example of *pattern matching*.

Pattern matching is an important concept in Elixir. It's a feature that makes manipulations with complex variables (such as tuples and lists) a lot easier. Less obviously, it allows you to write elegant, declarative-like conditionals and loops. You'll see what this means by the end of the chapter; in this

section we'll look at the basic mechanical workings of pattern matching.

Let's begin by looking at the match operator.

3.1.1 The match operator

So far, you've seen the most basic use of the match operator:

```
iex(1)> person = {"Bob", 25}
```

We treated this as something akin to an assignment, but in reality something more complex is going on here. At runtime, the left side of the = operator is matched to the right side. The left side is called a *pattern*, whereas on the right side you have an expression that evaluates to an Elixir term.

In the example, you match the variable `person` to the right-side term `{"Bob", 25}`. A variable always matches the right-side term, and it becomes *bound* to the value of that term. This may seem a bit theoretical, so let's look at a slightly more complex use of the match operator that involves tuples.

3.1.2 Matching tuples

The following example demonstrates basic pattern matching of tuples:

```
iex(1)> {name, age} = {"Bob", 25}
```

This expression assumes that the right-side term is a tuple of two elements. When the expression is evaluated, the variables `name` and `age` are bound to the corresponding elements of the tuple. You can now verify that these variables are correctly bound:

```
iex(2)> name
"Bob"

iex(3)> age
25
```

This feature is useful when you call a function that returns a tuple and you want to bind individual elements of that tuple to separate variables. The

following example calls the Erlang function `:calendar.local_time/0` to get the current date and time:

```
iex(4)> {date, time} = :calendar.local_time()
```

The date and time are also tuples, which you can further decompose:

```
iex(5)> {year, month, day} = date
iex(6)> {hour, minute, second} = time
```

What happens if the right side doesn't correspond to the pattern? The match fails, and an error is raised:

```
iex(7)> {name, age} = "can't match"
** (MatchError) no match of right hand side value: "can't match"
```

ⓘ Note

We haven't yet covered the error-handling mechanisms — they'll be discussed in chapter 8. For now, suffice it to say that raising an error works somewhat like the classical exception mechanisms in mainstream languages. When an error is raised, control is immediately transferred to code somewhere up the call chain, which catches the error (assuming such code exists).

Finally, it's worth noting that, just like any other expression, the match expression also returns a value. The result of a match expression is always the right-side term you're matching against:

```
iex(8)> {name, age} = {"Bob", 25}     #1
{"Bob", 25}       #2
```

3.1.3 Matching constants

Matching isn't confined to destructuring tuple elements to individual variables. Surprisingly enough, even constants are allowed on the left side of the match expression:

```
iex(1)> 1 = 1
```

Recall that the match operator = tries to match the right-side term to the left-side pattern. In the example, you try to match the pattern `1` to the term `1`. Obviously this succeeds, and the result of the entire expression is the right-side term.

This example doesn't have much practical benefit, but it illustrates that you can place constants to the left of =, which proves that = is not an assignment operator.

Constants are much more useful in compound matches. For example, tuples are sometimes used to group various fields of a record. The following snippet creates a tuple that holds a person's name and age:

```
iex(2)> person = {:person, "Bob", 25}
```

The first element is a constant atom `:person`, which you use to denote that this tuple represents a person. Later you can rely on this knowledge and retrieve individual attributes of the person:

```
iex(3)> {:person, name, age} = person
{:person, "Bob", 25}
```

Here you expect the right-side term to be a three-element tuple, with its first element having a value of `:person`. After the match, the remaining elements of the tuple are bound to the variables `name` and `age`, which you can easily verify:

```
iex(4)> name
"Bob"
```

```
iex(5)> age
25
```

This is a common idiom in Elixir. Many functions from Elixir and Erlang return either `{:ok, result}` or `{:error, reason}`. For example, imagine that your system relies on a configuration file and expects it to always be available. You can read the file contents with the help of the `File.read/1` function:

```
{:ok, contents} = File.read("my_app.config")
```

In this single line of code, three distinct things happen:

1. An attempt to open and read the file my_app.config takes place.
2. If the attempt succeeds, the file contents are extracted to the variable `contents`.
3. If the attempt fails, an error is raised. This happens because the result of `File.read` is a tuple in the form `{:error, reason}`, so the match to `{:ok, contents}` fails.

By using constants in patterns, you tighten the match, making sure some part of the right side has a specific value.

3.1.4 Variables in patterns

Whenever a variable name exists in the left-side pattern, it always matches the corresponding right-side term. In addition, the variable is bound to the term it matches.

Occasionally we aren't interested in a value from the right-side term, but we still need to match on it. For example, let's say you want to get the current time of day. You can use the function `:calendar.local_time/0`, which returns a tuple `{date, time}`. But you aren't interested in a date, so you don't want to store it to a separate variable. In such cases, you can use the *anonymous variable* (_):

```
iex(1)> {_, time} = :calendar.local_time()
iex(2)> time

{20, 44, 18}
```

When it comes to matching, the anonymous variable works just like a named variable: it matches any right-side term. But the value of the term isn't bound to any variable.

You can also add a descriptive name after the underscore character:

```
iex(1)> {_date, time} = :calendar.local_time()
```

The _date is regarded as an anonymous variable, because its name starts with an underscore. Technically speaking, you could use that variable in the rest of the program, but the compiler will emit a warning.

Patterns can be arbitrarily nested. Taking the example further, let's say you only want to retrieve the current hour of the day:

```
iex(3)> {_, {hour, _, _}} = :calendar.local_time()

iex(4)> hour
20
```

A variable can be referenced multiple times in the same pattern. In the following expressions, you expect an RGB triplet with the same number for each component:

```
iex(5)> {amount, amount, amount} = {127, 127, 127}    #1
{127, 127, 127}

iex(6)> {amount, amount, amount} = {127, 127, 1}   #2
** (MatchError) no match of right hand side value: {127, 127, 1}
```

Occasionally, you'll need to match against the contents of the variable. For this purpose, the *pin operator* (^) is provided. This is best explained with an example:

```
iex(7)> expected_name = "Bob"    #1
"Bob"

iex(8)> {^expected_name, _} = {"Bob", 25}   #2
{"Bob", 25}

iex(9)> {^expected_name, _} = {"Alice", 30}    #2
** (MatchError) no match of right hand side value: {"Alice", 30}
```

Using ^expected_name in patterns says that you expect the *value* of the variable expected_name to be in the appropriate position in the right-side term. In this example, it would be the same as if you used the hard-coded pattern {"Bob", _}. Therefore, the first match succeeds, but the second one fails.

Notice that the pin operator doesn't bind the variable. You expect that the

variable is already bound to a value, and you try to match against that value.

3.1.5 Matching lists

List matching works similarly to tuples. The following example decomposes a three-element list:

```
iex(1)> [first, second, third] = [1, 2, 3]
[1, 2, 3]
```

And of course, the previously mentioned pattern techniques work as well:

```
[1, second, third] = [1, 2, 3]        #1
[first, first, first] = [1, 1, 1]     #2
[first, second, _ ] = [1, 2, 3]       #3
[^first, second, _ ] = [1, 2, 3]      #4
```

Matching lists is more often done by relying on their recursive nature. Recall from chapter 2 that each non-empty list is a recursive structure that can be expressed in the form [head | tail]. You can use pattern matching to put each of these two elements into separate variables:

```
iex(3)> [head | tail] = [1, 2, 3]
[1, 2, 3]

iex(4)> head
1

iex(5)> tail
[2, 3]
```

If you need only one element of the [head, tail] pair, you can use the anonymous variable. Here's an inefficient way of calculating the smallest element in the list:

```
iex(6)> [min | _] = Enum.sort([3,2,1])

iex(7)> min
1
```

First you sort the list, and then, with the pattern [min | _], you take only the head of the (sorted) list. Note that this could also be done with the hd function

mentioned in chapter 2. In fact, for this case, `hd` would be more elegant. The pattern `[head | _]` is more useful when pattern-matching function arguments, as you'll see in section 3.2.

3.1.6 Matching maps

To match a map, the following syntax can be used:

```
iex(1)> %{name: name, age: age} = %{name: "Bob", age: 25}
%{age: 25, name: "Bob"}

iex(2)> name
"Bob"

iex(3)> age
25
```

When matching a map, the left-side pattern doesn't need to contain all the keys from the right-side term:

```
iex(4)> %{age: age} = %{name: "Bob", age: 25}

iex(5)> age
25
```

You may wonder about the purpose of such a partial-matching rule. Maps are frequently used to represent structured data. In such cases, you're often interested in only some of the map's fields. For example, in the previous snippet, you just want to extract the age field, ignoring everything else. The partial-matching rule allows you to do exactly this.

Of course, a match will fail if the pattern contains a key that's not in the matched term:

```
iex(6)> %{age: age, works_at: works_at} = %{name: "Bob", age: 25}
** (MatchError) no match of right hand side value
```

3.1.7 Matching bitstrings and binaries

We won't deal with bitstrings and pure binaries much in this book, but it's worth mentioning some basic matching syntax. Recall that a *bitstring* is a

chunk of bits, and a *binary* is a special case of a bitstring that's always aligned to the byte size.

To match a binary, you use syntax similar to creating one:

```
iex(1)> binary = <<1, 2, 3>>
<<1, 2, 3>>

iex(2)> <<b1, b2, b3>> = binary     #1
<<1, 2, 3>>

iex(3)> b1
1

iex(4)> b2
2

iex(5)> b3
3
```

This example matches on a 3-byte binary and extracts individual bytes to separate variables.

The following example takes the binary apart by taking its first byte into one variable and the rest of the binary into another:

```
iex(6)> <<b1, rest :: binary>> = binary
<<1, 2, 3>>

iex(7)> b1
1

iex(8)> rest
<<2, 3>>
```

`rest::binary` states that you expect an arbitrary-sized binary. You can even extract separate bits or groups of bits. The following example splits a single byte into two 4-bit values:

```
iex(9)> <<a :: 4, b :: 4>> = << 155 >>
<< 155 >>

iex(10)> a
9
```

```
iex(11)> b
11
```

Pattern `a::4` states that you expect a four-bit value. In the example, you put the first four bits into variable `a` and the other four bits into variable `b`. Because the number 155 is in binary represented as 10011011, you get values of 9 (1001 binary) and 11 (1011 binary).

Matching bitstrings and binaries is immensely useful when you're trying to parse packed binary content that comes from a file, an external device, or a network. In such situations, you can use binary matching to extract separate bits and bytes elegantly.

As mentioned, the examples in this book won't need this feature. Still, you should make a mental note of binaries and pattern matching, in case the need arises at some point.

Matching binary strings

Recall that strings are binaries, so you can use binary matches to extract individual bits and bytes from a string:

```
iex(13)> <<b1, b2, b3>> = "ABC"
"ABC"

iex(13)> b1
65

iex(14)> b2
66

iex(15)> b3
67
```

Variables `b1`, `b2`, and `b3` hold corresponding bytes from the string you matched on. This isn't very useful, especially if you're dealing with Unicode strings. Extracting individual characters is better done using functions from the `String` module.

A more useful pattern is to match the beginning of the string:

```
iex(16)> command = "ping www.example.com"
"ping www.example.com"

iex(17)> "ping " <> url = command      #1
"ping www.example.com"

iex(18)> url
"www.example.com"
```

In this example, you construct a string that holds a `ping` command. When you write `"ping " <> url = command`, you state the expectation that a `command` variable is a binary string starting with `"ping "`. If this matches, the rest of the string is bound to the variable `url`.

3.1.8 Compound matches

You've already seen this, but let's make it explicit. Patterns can be arbitrarily nested, as in the following contrived example:

```
iex(1)> [_, {name, _}, _] = [{"Bob", 25}, {"Alice", 30}, {"John",
35}]
```

In this example, the term being matched is a list of three elements. Each element is a tuple representing a person, consisting of two fields: the person's name and age. The match extracts the name of the second person in the list.

Another interesting feature is match chaining. Before you see how that works, let's discuss match expressions in more detail.

A match expression has this general form:

```
pattern = expression
```

As you've seen in examples, you can place any expression on the right side:

```
iex(2)> a = 1 + 3
4
```

Let's break down what happens here:

1. The expression on the right side is evaluated.

2. The resulting value is matched against the left-side pattern.
3. Variables from the pattern are bound.
4. The result of the match expression is the result of the right-side term.

An important consequence of this is that match expressions can be chained:

```
iex(3)> a = (b = 1 + 3)
4
```

In this (not so useful) example, the following things happen:

1. The expression `1 + 3` is evaluated.
2. The result (4) is matched against the pattern `b`.
3. The result of the inner match (which is again 4) is matched against the pattern `a`.

Consequently, both a and b have the value `4`.

Parentheses are optional, and many developers omit them in this case:

```
iex(4)> a = b = 1 + 3
4
```

This yields the same result, due to the fact that the operator = is right-associative.

Now let's look at a more useful example. Recall the function `:calendar.local_time/0`:

```
iex(5)> :calendar.local_time()
{{2023, 11, 11}, {21, 28, 41}}
```

Let's say you want to retrieve the function's total result (datetime) as well as the current hour of the day. Here's the way to do it in a single compound match:

```
iex(6)> date_time = {_, {hour, _, _}} = :calendar.local_time()
```

You can even swap the ordering. It still gives the same result (assuming you call it in the same second):

```
iex(7)> {_, {hour, _, _}} = date_time = :calendar.local_time()
```

In any case, you get what you wanted:

```
iex(8)> date_time
{{2023, 11, 11}, {21, 32, 34}}

iex(9)> hour
21
```

This works because the result of a pattern match is always the result of the term being matched (whatever is on the right side of the match operator). You can successively match against the result of that term and extract different parts you're interested in.

3.1.9 General behavior

We're almost done with basic pattern-matching mechanics. We've worked through a lot of examples, so let's try to formalize the behavior a bit.

The pattern-matching expression consists of two parts: the *pattern* (left side) and the *term* (right side). In a match expression, the attempt to match the term to the pattern takes place.

If the match succeeds, all variables in the pattern are bound to the corresponding values from the term. The result of the entire expression is the entire term you matched. If the match fails, an error is raised.

Therefore, in a pattern-matching expression, you perform two different tasks:

- You assert your expectations about the right-side term. If these expectations aren't met, an error is raised.
- You bind some parts of the term to variables from the pattern.

Finally, it's worth mentioning that we didn't cover all possible patterns here. For the detailed reference you can refer to the official documentation (https://hexdocs.pm/elixir/patterns-and-guards.html).

The match operator = is just one example where pattern matching can be

used. Pattern matching powers many other kinds of expressions, and it's especially powerful when used in functions.

3.2 Matching with functions

The pattern-matching mechanism is used in the specification of function arguments. Recall the basic function definition:

```
def my_fun(arg1, arg2) do
  ...
end
```

The argument specifiers `arg1` and `arg2` are patterns, and you can use standard matching techniques.

Let's see this in action. As mentioned in chapter 2, tuples are often used to group related fields together. For example, if you do a geometry manipulation, you can represent a rectangle with a tuple, `{a, b}`, containing the rectangle's sides. The following listing shows a function that calculates a rectangle's area.

Listing 3.1. Pattern-matching function arguments (rect.ex)

```
defmodule Rectangle do
  def area({a, b}) do    #1
    a * b
  end
end
```

Notice how you pattern-match the argument. The function `Rectangle.area/1` expects that its argument is a two-element tuple. It then binds corresponding tuple elements into variables and returns the result.

You can see whether this works from the shell. Start the shell, and load the module:

```
$ iex rect.ex
```

Then try the function:

```
iex(1)> Rectangle.area({2, 3})
6
```

What happens here? When you call a function, the arguments you provide are matched against the patterns specified in the function definition. The function expects a two-element tuple and binds the tuple's elements to variables a and b.

When calling functions, the term being matched is the argument provided to the function call. The pattern you match against is the argument specifier, in this case {a, b}.

Of course, if you provide anything that isn't a two-element tuple, an error will be raised:

```
iex(2)> Rectangle.area(2)
** (FunctionClauseError) no function clause matching in Rectangle
```

Pattern-matching function arguments is an extremely useful tool. It underpins one of the most important features of Elixir: *multiclause functions*.

3.2.1 Multiclause functions

Elixir allows you to overload a function by specifying multiple clauses. A *clause* is a function definition specified by the def expression. If you provide multiple definitions of the same function with the same arity, it's said that the function has multiple clauses.

Let's see this in action. Extending the previous example, let's say you need to develop a Geometry module that can handle various shapes. You'll represent shapes with tuples and use the first element of each tuple to indicate which shape it represents:

```
rectangle = {:rectangle, 4, 5}
square = {:square, 5}
circle = {:circle, 4}
```

Given these shape representations, you can write the following function to calculate a shape's area.

Listing 3.2. Multiclause function (geometry.ex)

```
defmodule Geometry do
  def area({:rectangle, a, b}) do    #1
    a * b
  end

  def area({:square, a}) do       #2
    a * a
  end

  def area({:circle, r}) do     #3
    r * r * 3.14
  end
end
```

As you can see, you provide three clauses of the same function. Depending on which argument you pass, the appropriate clause is called. Let's try this from the shell:

```
iex(1)> Geometry.area({:rectangle, 4, 5})
20

iex(2)> Geometry.area({:square, 5})
25

iex(3)> Geometry.area({:circle, 4})
50.24
```

When you call the function, the runtime goes through each of its clauses, in the order they're specified in the source code, and tries to match the provided arguments. The first clause that successfully matches all arguments is executed.

Of course, if no clause matches, an error is raised:

```
iex(4)> Geometry.area({:triangle, 1, 2, 3})
** (FunctionClauseError) no function clause matching in Geometry.
```

From the caller's perspective, a multiclause function is a single function. You can't directly reference a specific clause. Instead, you always work on the entire function. Recall from chapter 2 that you can create a function value with the capture operator, &:

```
&Module.fun/arity
```

If you capture Geometry.area/1, you capture all of its clauses:

```
iex(4)> fun = &Geometry.area/1      #1

iex(5)> fun.({:circle, 4})
50.24

iex(6)> fun.({:square, 5})
25
```

This proves that the function is treated as a whole, even if it consists of multiple clauses.

Sometimes you'll want a function to return a term indicating a failure, rather than raising an error. You can introduce a *default* clause that always matches. Let's do this for the `area` function. The next listing adds a final clause that handles any invalid input.

Listing 3.3. Multiclause function (geometry_invalid_input.ex)

```
defmodule Geometry do
  def area({:rectangle, a, b}) do
    a * b
  end

  def area({:square, a}) do
    a * a
  end

  def area({:circle, r}) do
    r * r * 3.14
  end

  def area(unknown) do      #1
    {:error, {:unknown_shape, unknown}}
  end
end
```

If none of the first three clauses match, the final clause is called. This is because a variable pattern always matches the corresponding term. In this case, you return a two-element tuple `{:error, reason}`, to indicate that

something has gone wrong.

Try it from the shell:

```
iex(1)> Geometry.area({:square, 5})
25

iex(2)> Geometry.area({:triangle, 1, 2, 3})
{:error, {:unknown_shape, {:triangle, 1, 2, 3}}}
```

Tip

For this to work correctly, it's important to place the clauses in the appropriate order. The runtime tries to select the clauses using the order in the source code. If the `area(unknown)` clause was defined first, you'd always get the error result.

Notice that the `area(unknown)` clause works only for `area/1`. If you pass more than one argument, this clause won't be called. Recall from chapter 2 that functions differ in name and arity. Because functions with the same name but different arities are in reality two different functions, there's no way to specify an `area` clause that's executed regardless of how many arguments are passed.

One final note: you should always group clauses of the same function together, instead of scattering them in various places in the module. If a multiclause function is spread all over the file, it becomes increasingly hard to analyze the function's complete behavior. Even the compiler complains about this by emitting a compilation warning.

3.2.2 Guards

Let's say you want to write a function that accepts a number and returns an atom `:negative`, `:zero`, or `:positive`, depending on the number's value. This isn't possible with the simple pattern matching you've seen so far. Elixir gives you a solution for this in the form of *guards*.

Guards are an extension of the basic pattern-matching mechanism. They

allow you to state additional broader expectations that must be satisfied for the entire pattern to match.

A guard can be specified by providing the `when` clause after the arguments list. This is best illustrated by example. The following code tests whether a given number is positive, negative, or zero.

Listing 3.4. Using guards (test_num.ex)

```
defmodule TestNum do
  def test(x) when x < 0 do
    :negative
  end

  def test(x) when x == 0 do
    :zero
  end

  def test(x) when x > 0 do
    :positive
  end
end
```

The guard is a logical expression that adds further conditions to the pattern. In this example we have three clauses with the same pattern (`x`) that would normally always match. The additional guard refines the pattern, making sure that the clause is invoked only if the given condition is satisfied, as demonstrated in this shell session:

```
iex(1)> TestNum.test(-1)
:negative

iex(2)> TestNum.test(0)
:zero

iex(3)> TestNum.test(1)
:positive
```

Surprisingly enough, calling this function with a non-number yields strange results:

```
iex(4)> TestNum.test(:not_a_number)
:positive
```

What gives? The explanation lies in the fact that Elixir terms can be compared with the operators < and >, even if they're not of the same type. In this case, the type ordering determines the result:

```
number < atom < reference < fun < port < pid <
tuple < map < list < bitstring (binary)
```

A number is smaller than any other type, which is why `TestNum.test/1` always returns `:positive` if you provide a non-number. To fix this, you have to extend the guard by testing whether the argument is a number, as illustrated next.

Listing 3.5. Using guards (test_num2.ex)

```
defmodule TestNum do
  def test(x) when is_number(x) and x < 0 do
    :negative
  end

  def test(x) when x == 0 do
    :zero
  end

  def test(x) when is_number(x) and x > 0 do
    :positive
  end
end
```

This code uses the function `Kernel.is_number/1` to test whether the argument is a number. Now `TestNum.test/1` raises an error if you pass a non-number:

```
iex(1)> TestNum.test(-1)
:negative

iex(2)> TestNum.test(:not_a_number)
** (FunctionClauseError) no function clause matching in TestNum.t
```

The set of operators and functions that can be called from guards is very limited. In particular, you may not call your own functions, and most of the other functions won't work. These are some examples of operators and functions allowed in guards:

- Comparison operators (==, !=, ===, !==, >, <, <=, >=)
- Boolean operators (and, or) and negation operators (not, !)
     A. Arithmetic operators (+, -, *, /)
- Type-check functions from the Kernel module (for example, is_number/1, is_atom/1, and so on)

You can find the complete up-to-date list at https://hexdocs.pm/elixir/patterns-and-guards.html#guards.

In some cases, a function used in a guard may cause an error to be raised. For example, length/1 makes sense only on lists. Imagine you have the following function that calculates the smallest element of a non-empty list:

```
defmodule ListHelper do
  def smallest(list) when length(list) > 0 do
    Enum.min(list)
  end

  def smallest(_), do: {:error, :invalid_argument}
end
```

You may think that calling ListHelper.smallest/1 with anything other than a list will raise an error, but this won't happen. If an error is raised from inside the guard, it won't be propagated, and the guard expression will return false. The corresponding clause won't match, but some other might.

In the preceding example, if you call ListHelper.smallest(123), you'll get the result {:error, :invalid_argument}. This demonstrates that an error in the guard expression is internally handled.

3.2.3 Multiclause lambdas

Anonymous functions (lambdas) may also consist of multiple clauses. First, recall the basic way of defining and using lambdas:

```
iex(1)> double = fn x -> x * 2 end     #1

iex(2)> double.(3)       #2
6
```

The general lambda syntax has the following shape:

```
fn
  pattern_1, pattern_2 ->
    ...           #1

  pattern_3, pattern_4 ->
    ...           #2

  ...
end
```

Let's see this in action by reimplementing the `test/1` function that inspects whether a number is positive, negative, or zero:

```
iex(3)> test_num =
          fn
            x when is_number(x) and x < 0 -> :negative
            x when x == 0 -> :zero
            x when is_number(x) and x > 0 -> :positive
          end
```

Notice that there's no special ending terminator for a lambda clause. The clause ends when the new clause is started (in the form `pattern ->`) or when the lambda definition is finished with `end`.

Note

Because all clauses of a lambda are listed under the same `fn` expression, the parentheses for each clause are by convention omitted. In contrast, each clause of a named function is specified in a separate `def` (or `defp`) expression. As a result, parentheses around named function arguments are recommended.

You can now test this lambda:

```
iex(4)> test_num.(-1)
:negative

iex(5)> test_num.(0)
:zero

iex(6)> test_num.(1)
```

```
:positive
```

Multiclause lambdas come in handy when using higher-order functions, as you'll see later in this chapter. But for now, we're done with the basic theory behind multiclause functions. They play an important role in conditional runtime branching, which is our next topic.

3.3 Conditionals

Elixir provides some standard ways of doing conditional branching, with expressions such as `if` and `case`. Multiclause functions can be used for this purpose as well. In this section, we'll cover all the branching techniques, starting with multiclause functions.

3.3.1 Branching with multiclause functions

You've already seen how to do conditional logic with multiclauses, but let's repeat it once more:

```
defmodule TestNum do
  def test(x) when x < 0, do: :negative
  def test(0), do: :zero
  def test(x), do: :positive
end
```

The three clauses constitute three conditional branches. In a typical imperative language, such as JavaScript, you could write something like the following:

```
function test(x){
  if (x < 0) return "negative";
  if (x == 0) return "zero";
  return "positive";
}
```

Arguably, both versions are equally readable. But with multiclauses you can reap all the benefits of pattern matching, such as branching depending on the shape of the data. In the following example, a multiclause is used to test whether a given list is empty:

```
defmodule TestList do
  def empty?([]), do: true
  def empty?([_|_]), do: false
end
```

The first clause matches the empty list, whereas the second clause relies on the `head | tail` representation of a non-empty list.

By relying on pattern matching, you can implement polymorphic functions that do different things depending on the input type. The following example implements a function that doubles a variable. The function behaves differently depending on whether it's called with a number or with a binary (string):

```
iex(1)> defmodule Polymorphic do
          def double(x) when is_number(x), do: 2 * x
          def double(x) when is_binary(x), do: x <> x
        end

iex(2)> Polymorphic.double(3)
6

iex(3)> Polymorphic.double("Jar")
"JarJar"
```

The power of multiclauses starts to show in recursions. The resulting code seems declarative and is devoid of redundant `if`s and `return`s. Here's a recursive implementation of a factorial, based on multiclauses:

```
iex(4)> defmodule Fact do
          def fact(0), do: 1
          def fact(n), do: n * fact(n - 1)
        end

iex(5)> Fact.fact(1)
1

iex(6)> Fact.fact(3)
6
```

A multiclause-powered recursion is also used as a primary building block for looping. This will be thoroughly explained in the next section, but here's a simple example. The following function sums all the elements of a list:

```
iex(7)> defmodule ListHelper do
          def sum([]), do: 0
          def sum([head | tail]), do: head + sum(tail)
        end

iex(8)> ListHelper.sum([])
0

iex(9)> ListHelper.sum([1, 2, 3])
6
```

The solution implements the sum by relying on the recursive definition of a list. The sum of an empty list is always 0, and the sum of a non-empty list equals the value of its head plus the sum of its tail.

Everything that can be done with classical branching expressions can be accomplished with multiclauses. But the underlying pattern-matching mechanism can often be more expressive, allowing you to branch depending on values, types, and shapes of function arguments.

In some cases, though, the code looks better with the classical, imperative style of branching. Let's look at the other branching expressions we have in Elixir.

3.3.2 Classical branching expressions

Multiclause solutions may not always be appropriate. Using them requires creating a separate function and passing the necessary arguments. Sometimes it's simpler to use a classical branching expression in the function, and for such cases, the expressions `if`, `unless`, `cond`, and `case` are provided. These work roughly as you might expect, although there are a couple of twists. Let's look at each of them.

If and unless

The `if` expression has a familiar syntax:

```
if condition do
  ...
else
```

```
  ...
end
```

This causes one or the other branch to execute, depending on the truthiness of the condition. If the condition is anything other than `false` or `nil`, you end up in the main branch; otherwise the `else` part is called.

You can also condense this into a one-liner, much like a `def` expression:

```
if condition, do: something, else: another_thing
```

Recall that everything in Elixir is an expression that has a return value. The `if` expression returns the result of the executed block (that is, of the block's last expression). If the condition isn't met and the `else` clause isn't specified, the return value is the atom `nil`:

```
iex(1)> if 5 > 3, do: :one
:one

iex(2)> if 5 < 3, do: :one
nil

iex(3)> if 5 < 3, do: :one, else: :two
:two
```

Let's look at a more concrete example. The following code implements a `max` function that returns the larger of two elements (according to the semantics of the > operator):

```
def max(a, b) do
  if a >= b, do: a, else: b
end
```

The `unless` expression is also available, which is the equivalent of `if (not …)`. Consider the following `if` expression:

```
if result != :error, do: send_notification(...)
```

This can be also expressed as

```
unless result == :error, do: send_notification(...)
```

Cond

The `cond` expression can be thought of as equivalent to an `if-else-if` pattern. It takes a list of expressions and executes the block of the first expression that evaluates to a truthy value:

```
cond do
  expression_1 ->
    ...

  expression_2 ->
    ...

  ...
end
```

The result of `cond` is the result of the corresponding executed block. If none of the conditions is satisfied, `cond` raises an error.

The `cond` expression is a good fit if there are more than two branching choices:

```
def call_status(call) do
  cond do
    call.ended_at != nil -> :ended
    call.started_at != nil -> :started
    true -> :pending                            #1
  end
end
```

In this example, we're computing the status of a call. If the `ended_at` field is populated, the call has ended. Otherwise, if the `started_at` field is populated, the call has started. If neither of these two fields is populated, the call is pending. Notice the final clause (`true -> :pending`). Since the condition of this clause (`true`) is always satisfied, this effectively becomes the fallback clause that is invoked if none of the previously stated conditions in the `cond` expression are met.

Case

The general syntax of `case` is as follows:

```
case expression do
  pattern_1 ->
    ...

  pattern_2 ->
    ...

  ...
end
```

The term *pattern* here indicates that it deals with pattern matching. In the
case expression, the provided expression is evaluated, and then the result is
matched against the given clauses. The first one that matches is executed, and
the result of the corresponding block (its last expression) is the result of the
entire case expression. If no clause matches, an error is raised.

The case-powered version of the max function would then look like this:

```
def max(a,b) do
  case a >= b do
    true -> a
    false -> b
  end
end
```

The case expression is most suitable if you don't want to define a separate
multiclause function. Other than that, there are no differences between case
and multiclause functions. In fact, the general case syntax can be directly
translated into the multiclause approach:

```
defp fun(pattern_1), do: ...
defp fun(pattern_2), do: ...
...
```

This must be called using the fun(expression).

You can specify the default clause by using the anonymous variable to match
anything:

```
case expression do
  pattern_1 -> ...
  pattern_2 -> ...
  ...
```

```
  _ -> ...          #1
end
```

As you've seen, there are different ways of doing conditional logic in Elixir. Multiclauses offer a more declarative feel of branching, but they require you to define a separate function and pass all the necessary arguments to it. Classical expressions like `if` and `case` seem more imperative but can often prove simpler than the multiclause approach. Selecting an appropriate solution depends on the specific situation as well as your personal preferences.

3.3.3 The with expression

The final branching expression we'll discuss is the `with` expression, which can be very useful when you need to chain a couple of expressions and return the error of the first expression that fails. Let's look at a simple example.

Suppose you need to process registration data submitted by a user. The input is a map, with keys being strings ("login", "email", and "password"). Here's an example of one input map:

```
%{
  "login" => "alice",
  "email" => "some_email",
  "password" => "password",
  "other_field" => "some_value",
  "yet_another_field" => "...",
  ...
}
```

Your task is to normalize this map into a map that contains only the fields login, email, and password. Usually, if the set of fields is well-defined and known upfront, you can represent the keys as atoms. Therefore, for the given input, you can return the following structure:

```
%{login: "alice", email: "some_email", password: "password"}
```

But some required field might not be present in the input map. In this case, you want to report the error, so your function can have two different

outcomes. It can return the normalized user map, or it can return an error. An idiomatic approach in such cases is to make the function return `{:ok, some_result}` or `{:error, error_reason}`. In this exercise, the successful result is the normalized user map, whereas the error reason is descriptive text.

Start by writing the helper functions for extracting each field:

```
defp extract_login(%{"login" => login}), do: {:ok, login}
defp extract_login(_), do: {:error, "login missing"}

defp extract_email(%{"email" => email}), do: {:ok, email}
defp extract_email(_), do: {:error, "email missing"}

defp extract_password(%{"password" => password}), do: {:ok, passw
defp extract_password(_), do: {:error, "password missing"}
```

Here you're relying on pattern matching to detect the field's presence.

Now you need to write the top-level `extract_user/1` function, which combines these three functions. Here's one way to do it with `case`:

```
def extract_user(user) do
  case extract_login(user) do
    {:error, reason} ->
      {:error, reason}

    {:ok, login} ->
      case extract_email(user) do
        {:error, reason} ->
          {:error, reason}

        {:ok, email} ->
          case extract_password(user) do
            {:error, reason} ->
              {:error, reason}

            {:ok, password} ->
              %{login: login, email: email, password: password}
          end
      end
  end
end
```

This is quite noisy, given that the code composes three functions. Each time

you fetch something, you need to branch depending on the result, and you end up with three nested cases. In real life you usually have to perform many more validations, so the code can become quite nasty pretty quickly.

This is precisely where `with` can help you. The `with` special form allows you to use pattern matching to chain multiple expressions, verify that the result of each conforms to the desired pattern, and return the first unexpected result.

In its simplest form, `with` has the following shape:

```
with pattern_1 <- expression_1,
     pattern_2 <- expression_2,
     ...

do
  ...
end
```

You start from the top, evaluating the first expression and matching the result against the corresponding pattern. If the match succeeds, you move to the next expression. If all the expressions are successfully matched, you end up in the `do` block, and the result of the `with` expression is the result of the last expression in the `do` block.

If any match fails, however, `with` will not proceed to evaluate subsequent expressions. Instead, it will immediately return the result that couldn't be matched.

Let's look at an example:

```
iex(1)> with {:ok, login} <- {:ok, "alice"},
             {:ok, email} <- {:ok, "some_email"} do
          %{login: login, email: email}
        end

%{email: "some_email", login: "alice"}
```

Here you went through two pattern matches to extract the login and the email. Then the `do` block is evaluated. The result of the `with` expression is the last result of the expression in the `do` block. Superficially, this is no different than the following:

```
{:ok, login} = {:ok, "alice"}
{:ok, email} = {:ok, "email"}
%{login: login, email: email}
```

The benefit of `with` is that it returns the first term that fails to be matched against the corresponding pattern:

```
iex(2)> with {:ok, login} <- {:error, "login missing"},
             {:ok, email} <- {:ok, "email"} do
          %{login: login, email: email}
        end

{:error, "login missing"}
```

This is precisely what you need in your case. Armed with this new knowledge, refactor the top-level `extract_user` function, as shown in the next listing.

Listing 3.6. `with`-based user extraction (user_extraction.ex)

```
def extract_user(user) do
  with {:ok, login} <- extract_login(user),
       {:ok, email} <- extract_email(user),
       {:ok, password} <- extract_password(user) do
    {:ok, %{login: login, email: email, password: password}}
  end
end
```

As you can see, this code is much shorter and clearer. You extract desired pieces of data, moving forward only if you succeed. If something fails, you return the first error. Otherwise, you return the normalized structure. The complete implementation can be found in user_extraction.ex. Try it out:

```
$ iex user_extraction.ex

iex(1)> UserExtraction.extract_user(%{})
{:error, "login missing"}

iex(2)> UserExtraction.extract_user(%{"login" => "some_login"})
{:error, "email missing"}

iex(3)> UserExtraction.extract_user(%{
          "login" => "some_login",
          "email" => "some_email"
```

```
        })
{:error, "password missing"}

iex(4)> UserExtraction.extract_user(%{
          "login" => "some_login",
          "email" => "some_email",
          "password" => "some_password"
        })
{:ok, %{email: "some_email", login: "some_login",
  password: "some_password"}}
```

The `with` special form has a couple more features not presented here, and you're advised to study it in more detail at https://hexdocs.pm/elixir/Kernel.SpecialForms.html#with/1.

This concludes our tour of the branching expressions in Elixir. Now it's time to look at how you can perform loops and iterations.

3.4 Loops and iterations

Looping in Elixir works very differently than it does in mainstream languages. Constructs such as `while` and `do…while` aren't provided. Nevertheless, any serious program needs to do some kind of dynamic looping. So how do you go about it in Elixir?

The principal looping tool in Elixir is *recursion*, so we'll take a detailed look at how to use it.

ⓘ Note

Although recursion is the basic building block of any kind of looping, most production Elixir code uses it sparingly. That's because there are many higher-level abstractions that hide the recursion details. You'll learn about many of these abstractions throughout the book, but it's important to understand how recursion works in Elixir, because most of the complex code is based on this mechanism.

ⓘ Note

Most of the examples in this section deal with simple problems, such as calculating the sum of all the elements in a list — tasks Elixir allows you to do in an effective and elegant one-liner. The point of the examples, however, is to understand the different aspects of recursion-based processing on simple problems.

3.4.1 Iterating with recursion

Let's say you want to implement a function that prints the first *n* natural numbers (positive integers). Because there are no loops, you must rely on recursion. The basic approach is illustrated in the following listing.

Listing 3.7. Printing the first *n* natural numbers (natural_nums.ex)

```
defmodule NaturalNums do
  def print(1), do: IO.puts(1)

  def print(n) do
    print(n - 1)
    IO.puts(n)
  end
end
```

This code relies on recursion, pattern matching, and multiclause functions. If n is equal to 1, you print the number. Otherwise, you print the first n − 1 numbers and then the *n*th one.

Trying it in the shell gives satisfying results:

```
iex(1)> NaturalNums.print(3)
1
2
3
```

You may have noticed that the function won't work correctly if you provide a negative integer or a float. This could be resolved with additional guards and is left for you as an exercise.

The code in listing 3.7 demonstrates the basic way of doing a conditional loop. You specify a multiclause function, first providing the clauses that stop

the recursion. This is followed by more general clauses that produce part of the result and call the function recursively.

Next, let's look at computing something in a loop and returning the result. You've already seen this example when dealing with conditionals, but let's repeat it. The following code implements a function that sums all the elements in a given list.

Listing 3.8. Calculating the sum of the list (sum_list.ex)

```
defmodule ListHelper do
  def sum([]), do: 0

  def sum([head | tail]) do
    head + sum(tail)
  end
end
```

This code looks very declarative:

1. The sum of all the elements of an empty list is 0.
2. The sum of all the elements of a non-empty list equals the list's head plus the sum of the list's tail.

Let's see it in action:

```
iex(1)> ListHelper.sum([1, 2, 3])
6

iex(2)> ListHelper.sum([])
0
```

You probably know from other languages that a function call will lead to a stack push, and therefore will consume some memory. A very deep recursion might lead to a stack overflow and crash the entire program. This isn't necessarily a problem in Elixir, because of the tail-call optimization.

3.4.2 Tail function calls

If the last thing a function does is call another function (or itself), you're

dealing with a *tail call*:

```
def original_fun(...) do
  ...
  another_fun(...)          #1
end
```

Elixir (or, more precisely, Erlang) treats tail calls in a specific manner by performing a *tail-call optimization*. In this case, calling a function doesn't result in the usual stack push. Instead, something more like a goto or a jump statement happens. You don't allocate additional stack space before calling the function, which in turn means the tail function call consumes no additional memory.

How is this possible? In the previous snippet, the last thing done in `original_fun` is the call of `another_fun`. The final result of `original_fun` is the result of `another_fun`. This is why the compiler can safely perform the operation by jumping to the beginning of `another_fun` without doing additional memory allocation. When `another_fun` finishes, you return to whatever place `original_fun` was called from.

Tail calls are especially useful in recursive functions. A tail-recursive function — that is, a function that calls itself at the very end — can run virtually forever without consuming additional memory.

The following function is the Elixir equivalent of an endless loop:

```
def loop_forever(...) do
  ...
  loop_forever(...)
end
```

Because tail recursion doesn't consume additional memory, it's an appropriate solution for arbitrarily large iterations.

In the next listing, you'll convert the `ListHelper.sum/1` function to the tail-recursive version.

Listing 3.9. Tail-recursive sum of the first *n* natural numbers (sum_list_tc.ex)

```
defmodule ListHelper do
  def sum(list) do
    do_sum(0, list)
  end

  defp do_sum(current_sum, []) do
    current_sum
  end

  defp do_sum(current_sum, [head | tail]) do
    new_sum = head + current_sum
    do_sum(new_sum, tail)
  end
end
```

The first thing to notice is that you have two functions. The exported function
sum/1 is called by the module clients, and on the surface it works just like
before.

The recursion takes place in the private do_sum/2 function, which is
implemented as tail-recursive. It's a two-clause function, and we'll analyze it
clause by clause. The second clause is more interesting, so we'll start with it.
Here it is in isolation:

```
defp do_sum(current_sum, [head | tail]) do
  new_sum = head + current_sum      #1
  do_sum(new_sum, tail)             #2
end
```

This clause expects two arguments: the non-empty list to operate on, and the
sum you've calculated so far (current_sum). It then calculates the new sum
and calls itself recursively with the remainder of the list and the new sum.
Because the call happens at the very end, the function is tail-recursive, and
the call consumes no additional memory.

The variable new_sum is introduced here just to make things more obvious.
You could also inline the computation:

```
defp do_sum(current_sum, [head | tail]) do
  do_sum(head + current_sum, tail)
end
```

This function is still tail-recursive because it calls itself at the very end.

The remaining thing to see is the first clause of `do_sum/2`:

```
defp do_sum(current_sum, []) do
  current_sum
end
```

This clause is responsible for stopping the recursion. It matches on an empty list, which is the last step of the iteration. When you get here, there's nothing else to sum, so you return the accumulated result.

Finally, you have the function `sum/1`:

```
def sum(list) do
  do_sum(0, list)
end
```

This function is used by clients and is also responsible for initializing the value of the `current_sum` parameter that's passed recursively in `do_sum`.

You can think of tail recursion as a direct equivalent of a classical loop in imperative languages. The parameter `current_sum` is a classical accumulator: the value where you incrementally add the result in each iteration step. The `do_sum/2` function implements the iteration step and passes the accumulator from one step to the next. Elixir is an immutable language, so you need this trick to maintain the accumulated value throughout the loop. The first clause of `do_sum/2` defines the ending point of the iteration and returns the accumulator value.

In any case, the tail-recursive version of the list sum is now working, so you can try it from the shell:

```
iex(1)> ListHelper.sum([1, 2, 3])
6

iex(2)> ListHelper.sum([])
0
```

As you can see, from the caller's point of view, the function works exactly the same way. Internally, you rely on the tail recursion and can therefore process arbitrarily large lists without requiring extra memory for this task.

Tail vs. non-tail recursion

Given the properties of tail recursion, you might think it's always a preferred approach for doing loops. If you need to run an infinite loop, tail recursion is the only way that will work. Otherwise, aim for the version which seems more readable. And non-tail recursion can often produce the code which is more elegant and concise, and which sometimes even perform better than its tail recursive counterpart.

Recognizing tail calls

Tail calls can take different shapes. You've seen the most obvious case, but there are a couple of others. A tail call can also happen in a conditional expression:

```
def fun(...) do
  ...
  if something do
    ...
    another_fun(...)      #1
  end
end
```

The call to `another_fun` is a tail call because it's a last thing the function does. The same rule holds for `unless`, `cond`, `case`, and `with` expressions.

But the following code isn't a tail call:

```
def fun(...) do
  1 + another_fun(...)    #1
end
```

This is because the call to `another_fun` isn't the last thing done in the `fun` function. After `another_fun` finishes, you have to increment its result by 1 to compute the final result of `fun`.

Practicing

All this may seem complicated, but it's not that hard. If you're coming from

imperative languages, it's probably not what you're used to, and it will take some time to get accustomed to the recursive way of thinking, combined with the pattern-matching facility. You may want to take some time and experiment with recursion yourself. Here are a couple of functions you can write for practice:

- A `list_len/1` function that calculates the length of a list
- A `range/2` function that takes two integers, `from` and `to`, and returns a list of all integer numbers in the given range
- A `positive/1` function that takes a list and returns another list that contains only the positive numbers from the input list

Try to write these functions first in the non-tail-recursive form, and then convert them to the tail-recursive version. If you get stuck, the solutions are provided in the recursion_practice.ex and recursion_practice_tc.ex files (for the tail-recursive versions).

Recursion is the basic looping technique, and no loop can be done without it. Still, you won't need to write explicit recursion all that often. Many typical tasks can be performed by using higher-order functions.

3.4.3 Higher-order functions

A *higher-order function* is a fancy name for a function that takes one or more functions as its input or returns one or more functions (or both). The word *function* here means "function value."

You already made first contact with higher-order functions in chapter 2, when you used `Enum.each/2` to iterate through a list and print all of its elements. Let's recall how to do this:

```
iex(1)> Enum.each(
          [1, 2, 3],
          fn x -> IO.puts(x) end    #1
        )
1
2
3
```

The function `Enum.each/2` takes an enumerable (in this case, a list), and a lambda. It iterates through the enumerable, calling the lambda for each of its elements. Because `Enum.each/2` takes a lambda as its input, it's called a higher-order function.

You can use `Enum.each/2` to iterate over enumerable structures without writing the recursion. Under the hood, `Enum.each/2` is powered by recursion: there's no other way to do loops and iterations in Elixir. But the complexity of writing the recursion, the repetitive code, and the intricacies of tail recursion are hidden from you.

`Enum.each/2` is just one example of an iteration powered by a higher-order function. Elixir's standard library provides many other useful iteration helpers in the `Enum` module. You should spend some time researching the module documentation (https://hexdocs.pm/elixir/Enum.html). Here we'll look at some of the most frequently used `Enum` functions.

Enumerables

Most functions from the `Enum` module work on *enumerables*. You'll learn what this means in chapter 4. For now, it's sufficient to know that an enumerable is a data structure for which a certain contract is implemented, which makes that data structure suitable to be used by functions from the `Enum` module.

Some examples of enumerables are lists, ranges, maps, and `MapSet`. It's also possible to turn your own data structures into enumerables and thus harness all the features from the `Enum` module.

One manipulation you'll often need is a 1:1 transformation of a list to another list. For this purpose, `Enum.map/2` is provided. It takes an enumerable and a lambda that maps each element to another element. The following example doubles every element in the list:

```
iex(1)> Enum.map(
          [1, 2, 3],
          fn x -> 2 * x end
        )
[2, 4, 6]
```

Recall from chapter 2 that you can use the capture operator, `&`, to make the lambda definition a bit denser:

```
iex(2)> Enum.map(
          [1, 2, 3],
          &(2 * &1)
        )
```

The `&(…)` denotes a simplified lambda definition, where you use `&n` as a placeholder for the *n*th argument of the lambda.

Another useful function is `Enum.filter/2`, which can be used to extract only some elements of the list, based on certain criteria. The following snippet returns all odd numbers from a list:

```
iex(3)> Enum.filter(
          [1, 2, 3],
          fn x -> rem(x, 2) == 1 end
        )
[1, 3]
```

`Enum.filter/2` takes an enumerable and a lambda. It returns only those elements for which the lambda returns `true`.

Of course, you can use the capture syntax as well:

```
iex(3)> Enum.filter(
          [1, 2, 3],
          &(rem(&1, 2) == 1)
        )
[1, 3]
```

Let's play a bit more with `Enum`. Recall the example from section 3.3.3, where you used `with` to verify that the login, email, and password are submitted. In that example, you returned the first encountered error. Armed with this new knowledge, you can improve that code to report all missing fields immediately.

To briefly recap, your input is a map, and you need to fetch the keys "login", "email", and "password", and convert them into a map where keys are atoms. If a required field isn't provided, you need to report an error. In the previous

version, you simply reported the first missing field. A better user experience would be to return a list of all missing fields.

This is something you can do easily with the help of `Enum.filter/2`. The idea is to iterate through the list of required fields and take only those fields that aren't present in the map. You can easily check for the presence of a key with the help of `Map.has_key?/2`. The sketch of the solution then looks like the next listing.

Listing 3.10. Reporting all missing fields (user_extraction_2.ex)

```
case Enum.filter(            #1
        ["login", "email", "password"],    #1
        &(not Map.has_key?(user, &1))      #2
    ) do
  [] ->           #3
    ...

  missing_fields ->     #4
    ...
end
```

There are two possible outcomes of `Enum.filter/2`. If the result is an empty list, all the fields are provided, and you can extract the data. Otherwise, some fields are missing, and you need to report an error. The code for each branch is omitted here for the sake of brevity, but you can find the complete solution in user_extraction_2.ex.

Reduce

Probably the most versatile function from the `Enum` module is `Enum.reduce/3`, which can be used to transform an enumerable into anything. If you're coming from languages that support first-class functions, you may already know `reduce` under the name *inject* or *fold*.

Reducing is best explained with a specific example. You'll use `reduce` to sum all the elements in a list. Before doing it in Elixir, let's see how you could do this task in an imperative manner. Here's an imperative JavaScript example:

```
var sum = 0;        #1
[1, 2, 3].forEach(function(element) {
  sum += element;      #2
})
```

This is a standard imperative pattern. You initialize an accumulator (the variable sum) and then do some looping, adjusting the accumulator value in each step. After the loop is finished, the accumulator holds the final value.

In a functional language, you can't change the accumulator, but you can still calculate the result incrementally by using Enum.reduce/3. The function has the following shape:

```
Enum.reduce(
  enumerable,
  initial_acc,
  fn element, acc ->
    ...
  end
)
```

Enum.reduce/3 takes an enumerable as its first argument. The second argument is the initial value for the accumulator, the thing you compute incrementally. The final argument is a lambda that's called for each element. The lambda receives the element from the enumerable and the current accumulator value. The lambda's task is to compute and return the new accumulator value. When the iteration is done, Enum.reduce/3 returns the final accumulator value.

Let's use Enum.reduce/3 to sum up elements in the list:

```
iex(4)> Enum.reduce(
          [1, 2, 3],
          0,          #1
          fn element, sum -> sum + element end    #2
        )
6
```

That's all there is to it! Coming from an imperative background myself, it helps me to think of the lambda as the function that's called in each iteration step. Its task is to add a bit of the information to the result.

You may recall that I mentioned that many operators are functions, and you can turn an operator into a lambda by calling `&+/2`, `&*/2`, and so on. This combines nicely with higher-order functions. For example, the sum example can be written in a more compact form:

```
iex(5)> Enum.reduce([1,2,3], 0, &+/2)
6
```

It's worth mentioning that there's a function called `Enum.sum/1` that works exactly like this snippet. The point of the sum example was to illustrate how to iterate through a collection and accumulate the result.

Let's work a bit more with `reduce`. The previous example works only if you pass a list that consists exclusively of numbers. If the list contains anything else, an error is raised (because the + operator is defined only for numbers). The next example can work on any type of list and sums only its numeric elements:

```
iex(6)> Enum.reduce(
          [1, "not a number", 2, :x, 3],
          0,
          fn      #1
            element, sum when is_number(element) ->     #2
              sum + element

            _, sum ->
              sum       #3
          end
        )
```

This example relies on a multiclause lambda to obtain the desired result. If the element is a number, you add its value to the accumulated sum. Otherwise (if the element isn't a number), you return whatever sum you have at the moment, effectively passing it unchanged to the next iteration step.

Personally, I tend to avoid writing elaborate lambdas. If there's a bit more logic in the anonymous function, it's a sign that it will probably look better as a distinct function. In the following snippet, the lambda code is pushed to a separate private function:

```
defmodule NumHelper do
```

```
  def sum_nums(enumerable) do
    Enum.reduce(enumerable, 0, &add_num/2)     #1
  end

  defp add_num(num, sum) when is_number(num), do: sum + num   #2
  defp add_num(_, sum), do: sum
end
```

This is more or less similar to the approach you saw earlier. This example moves the iteration step to the separate, private function `add_num/2`. When calling `Enum.reduce`, you pass the lambda that delegates to that function, using the capture operator `&`.

Notice how when capturing the function, you don't specify the module name. That's because `add_num/2` resides in the same module, so you can omit the module prefix. In fact, because `add_num/2` is private, you can't capture it with the module prefix.

This concludes our basic showcase of the `Enum` module. Be sure to check the other functions that are available because you'll find a lot of useful helpers that can simplify loops, iterations, and manipulations of enumerables.

3.4.4 Comprehensions

The cryptic "comprehensions" name denotes another expression that can help you iterate and transform enumerables. The following example uses a comprehension to square each element of a list:

```
iex(1)> for x <- [1, 2, 3] do
          x*x
        end
```

The comprehension iterates through each element and runs the `do/end` block. The result is a list that contains all the results returned by the `do/end` block. In this basic form, `for` is no different than `Enum.map/2`.

Comprehensions have various other features that often make them elegant, compared to `Enum`-based iterations. For example, it's possible to perform nested iterations over multiple collections. The following example takes advantage of this feature to calculate a small multiplication table:

```
iex(2)> for x <- [1, 2, 3], y <- [1, 2, 3], do: {x, y, x*y}
[
  {1, 1, 1}, {1, 2, 2}, {1, 3, 3},
  {2, 1, 2}, {2, 2, 4}, {2, 3, 6},
  {3, 1, 3}, {3, 2, 6}, {3, 3, 9}
]
```

In this example, the comprehension performs a nested iteration, calling the provided block for each combination of input collections.

Just like functions from the Enum module, comprehensions can iterate through anything that's enumerable. For example, you can use ranges to compute a multiplication table for single-digit numbers:

```
iex(3)> for x <- 1..9, y <- 1..9, do: {x, y, x*y}
```

In the examples so far, the result of the comprehension has been a list. But comprehensions can return anything that's collectable. *Collectable* is an abstract term for a functional data type that can collect values. Some examples include lists, maps, MapSet, and file streams; you can even make your own custom type collectable (more on that in chapter 4).

In more general terms, a comprehension iterates through enumerables, calling the provided block for each value and storing the results in some collectable structure. Let's see this in action.

The following snippet makes a map that holds a multiplication table. Its keys are tuples of factors {x,y}, and the values contain products:

```
iex(4)> multiplication_table =
          for x <- 1..9,
              y <- 1..9,
              into: %{} do      #1
            {{x, y}, x*y}
          end

iex(5)> Map.get(multiplication_table, {7, 6})
42
```

Notice the into option, which specifies what to collect. In this case, it's an empty map %{} that will be populated with values returned from the do block. Notice how you return a {factors, product} tuple from the do block. You

use this format because map "knows" how to interpret it. The first element will be used as a key, and the second will be used as the corresponding value.

Another interesting comprehension feature is that you can specify filters. This gives you the possibility to skip some elements from the input. The following example computes a nonsymmetrical multiplication table for numbers `x` and `y`, where `x` is never greater than `y`:

```
iex(6)> multiplication_table =
          for x <- 1..9,
              y <- 1..9,
              x <= y,          #1
              into: %{} do
            {{x, y}, x*y}
          end

iex(7)> Map.get(multiplication_table, {6, 7})
42

iex(8)> Map.get(multiplication_table, {7, 6})
nil
```

The comprehension filter is evaluated for each element of the input enumerable, prior to block execution. If the filter returns `true`, the block is called and the result is collected. Otherwise the comprehension moves on to the next element.

As you can see, comprehensions are an interesting feature, allowing you to do some elegant transformations of the input enumerable. Although this can be done with `Enum` functions, most notably `Enum.reduce/3`, often the resulting code looks more elegant when comprehensions are used. This is particularly true when you have to perform a Cartesian product (cross-join) of multiple enumerables, or traverse a nested collection to produce a flat result.

ℹ️ **Note**

Comprehensions can also iterate through a binary. The syntax is somewhat different, and we won't treat it here. For details, it's best to look at the official `for` documentation at
https://hexdocs.pm/elixir/Kernel.SpecialForms.html#for/1.

3.4.5 Streams

Streams are a special kind of enumerable that can be useful for doing lazy composable operations over anything enumerable. To see what this means, let's look at one shortcoming of standard `Enum` functions.

Let's say you have a list of employees, and you need to print each one prefixed by their position in the list:

```
1. Alice
2. Bob
3. John
...
```

This is fairly simple to perform by combining various `Enum` functions. For example, there's a function `Enum.with_index/1` that takes an enumerable and returns a list of tuples, where the first element of the tuple is a member from the input enumerable and the second element is its zero-based index:

```
iex(1)> employees = ["Alice", "Bob", "John"]
["Alice", "Bob", "John"]

iex(2)> Enum.with_index(employees)
[{"Alice", 0}, {"Bob", 1}, {"John", 2}]
```

You can now feed the result of `Enum.with_index/1` to `Enum.each/2` to get the desired output:

```
iex(3)> employees
        |> Enum.with_index()
        |> Enum.each(fn {employee, index} ->
             IO.puts("#{index + 1}. #{employee}")
           end)

1. Alice
2. Bob
3. John
```

Here you rely on the pipe operator to chain together various function calls. This saves you from having to use intermediate variables and makes the code a bit cleaner.

So what's the problem with this code? The `Enum.with_index/1` function goes through the entire list to produce another list with tuples, and `Enum.each` then performs another iteration through the new list. It would be better if you could do both operations in a single pass without building another list. This is where streams can help you.

Streams are implemented in the `Stream` module (https://hexdocs.pm/elixir/Stream.html), which at first glance looks similar to the `Enum` module, containing functions like `map`, `filter`, and `take`. These functions take any enumerable as an input and give back a stream: an enumerable with some special powers.

A *stream* is a lazy enumerable, which means it produces the actual result on demand. Let's look at what this means.

The following snippet uses a stream to double each element in a list:

```
iex(4)> stream = Stream.map([1, 2, 3], fn x -> 2 * x end)    #1

#Stream<[enum: [1, 2, 3],      #2
  funs: [#Function<44.45151713/1 in Stream.map/2>]]>     #2
```

Because a stream is a lazy enumerable, the iteration over the input list (`[1, 2, 3]`) and the corresponding transformation (multiplication by 2) haven't yet happened. Instead, you get the structure that describes the computation.

To make the iteration happen, you have to pass the stream to an `Enum` function, such as `each`, `map`, or `filter`. You can also use the `Enum.to_list/1` function, which converts any kind of enumerable into a list:

```
iex(5)> Enum.to_list(stream)     #1
[2, 4, 6]
```

`Enum.to_list/1` (and any other `Enum` function, for that matter) is an eager operation. It immediately starts iterating through the input and creates the result. In doing so, `Enum.to_list/1` requests that the input enumerable start producing values. This is why the output of the stream is created when you send it to an `Enum` function.

The laziness of streams goes beyond iterating the list on demand. Values are

produced one by one when `Enum.to_list` requests another element. For example, you can use `Enum.take/2` to request only one element from the stream:

```
iex(6)> Enum.take(stream, 1)
[2]
```

Because `Enum.take/2` iterates only until it collects the desired number of elements, the input stream doubled only one element in the list. The others were never visited.

Going back to the example of printing employees, using a stream allows you to print employees in a single go. The change to the original code is simple enough. Instead of using `Enum.with_index/1`, you can rely on its lazy equivalent, `Stream.with_index/1`:

```
iex(7)> employees
        |> Stream.with_index()                          #1
        |> Enum.each(fn {employee, index} ->
            IO.puts("#{index + 1}. #{employee}")
          end)

1. Alice
2. Bob
3. John
```

The output is the same, but the list iteration is done only once. This becomes increasingly useful when you need to compose multiple transformations of the same list. The following example takes the input list and prints the square root of only those elements that represent a non-negative number, adding an indexed prefix at the beginning:

```
iex(1)> [9, -1, "foo", 25, 49]
        |> Stream.filter(&(is_number(&1) and &1 > 0))
        |> Stream.map(&{&1, :math.sqrt(&1)})
        |> Stream.with_index()
        |> Enum.each(fn {{input, result}, index} ->
            IO.puts("#{index + 1}. sqrt(#{input}) = #{result}")
          end)

1. sqrt(9) = 3.0
2. sqrt(25) = 5.0
3. sqrt(49) = 7.0
```

This code is dense, and it illustrates how concise you can be by relying only on functions as the abstraction tool. You start with the input list and filter only positive numbers. You transform each such number into an `{input_number, square_root}` tuple. Then you index the resulting tuples using `Stream.with_index/1`, and, finally, you print the result.

Even though you stack multiple transformations, everything is performed in a single pass when you call `Enum.each`. In contrast, if you used `Enum` functions everywhere, you'd have to run multiple iterations over each intermediate list, which would incur a memory-usage penalty.

This lazy property of streams can become useful for consuming slow and potentially large enumerable input. A typical case is when you need to parse each line of a file. Relying on eager `Enum` functions means you have to read the entire file into memory and then iterate through each line. In contrast, using streams makes it possible to read and immediately parse one line at a time. For example, the following function takes a filename and returns the list of all lines from that file that are longer than 80 characters:

```
def large_lines!(path) do
  File.stream!(path)
  |> Stream.map(&String.trim_trailing(&1, "\n"))
  |> Enum.filter(&(String.length(&1) > 80))
end
```

Here you rely on the `File.stream!/1` function, which takes the path of a file and returns a stream of its lines. Because the result is a stream, the iteration through the file happens only when you request it. After `File.stream!` returns, no byte from the file has been read yet. Then you remove the trailing newline character from each line, again in the lazy manner. Finally, you eagerly take only long lines, using `Enum.filter/2`. It's at this point that iteration happens. The consequence is that you never read the entire file in memory; instead, you work on each line individually.

Note

There are no special tricks in the Elixir compiler that allow these lazy enumerations. The real implementation is fairly involved, but the basic idea

behind streams is simple and relies on anonymous functions. In a nutshell, to make a lazy computation, you need to return a lambda that performs the computation. This makes the computation lazy, because you return its description rather than its result. When the computation needs to be materialized, the consumer code can call the lambda.

Infinite streams

So far you've produced streams by transforming an existing collection with the functions such as `Stream.map` or `Stream.filter`. Some functions from the `Stream` module allow you to create a stream from scratch.

One such function is `Stream.iterate/2`, which can be used to produce an infinite collection where each element is calculated based on the previous one. For example, the following snippet builds an infinite stream of natural numbers:

```
iex(1)> natural_numbers = Stream.iterate(
          1,
          fn previous -> previous + 1 end
        )
```

We can feed this infinite collection to other `Enum` and `Stream` functions to produce a finite sequence. For example to take the first 10 natural numbers, you can use `Enum.take/2`:

```
iex(2)> Enum.take(natural_numbers, 10)
[1, 2, 3, 4, 5, 6, 7, 8, 9, 10]
```

Another example is the function `Stream.repeatedly/1`, which repeatedly invokes the provided lambda to generate an element. In the following example we'll use it to repeatedly read the user's input from the console, stopping when the user submits the blank input:

```
iex(3)> Stream.repeatedly(fn -> IO.gets("> ") end)
        |> Stream.map(&String.trim_trailing(&1, "\n"))
        |> Enum.take_while(&(&1 != ""))

> Hello
> World
```

```
>
["Hello", "World"]
```

The `Stream` module contains a few more functions that produce an infinite stream, such as `Stream.unfold/2` or `Stream.resource/3`. Take a look at the official documentation for details.

Practicing

This style of coding takes some getting used to. You'll use the techniques presented here throughout the book, but you should try to write a couple such iterations yourself. Here are some exercise ideas that may help you get into the swing of things.

Using `large_lines!/1` as a model, write the following functions:

- A `lines_lengths!/1` that takes a file path and returns a list of numbers, with each number representing the length of the corresponding line from the file.
- A `longest_line_length!/1` that returns the length of the longest line in a file.
- A `longest_line!/1` that returns the contents of the longest line in a file.
- A `words_per_line!/1` that returns a list of numbers, with each number representing the word count in a file. Hint: to get the word count of a line, use `length(String.split(line))`.

Solutions are provided in the enum_streams_practice.ex file, but I strongly suggest that you spend some time trying to crack these problems yourself.

3.5 Summary

- Pattern matching is a mechanism that attempts to match a right-side term to the left-side pattern. In the process, variables from the pattern are bound to corresponding subterms from the term. If a term doesn't match the pattern, an error is raised.
- Function arguments are patterns. Calling a function tries to match the provided values to the patterns specified in the function definition.

- Functions can have multiple clauses. The first clause that matches all the arguments is executed.
- For conditional branching, you can use multiclause functions and expressions such as `if`, `unless`, `cond`, `case`, and `with`.
- Recursion is the main tool for implementing loops. Tail recursion is used when you need to run an arbitrarily long loop.
- Higher-order functions make writing loops much easier. There are many useful generic iteration functions in the `Enum` module. The `Stream` module additionally makes it possible to implement lazy and composable iterations.
- Comprehensions can also be used to iterate, transform, filter, and join various enumerables.

4 Data abstractions

This chapter covers

- Abstracting with modules
- Working with hierarchical data
- Polymorphism with protocols

This chapter deals with building higher-level data structures. In any complex system, there will be a need for abstractions such as `Money`, `Date`, `Employee`, and `OrderItem`, all textbook examples of higher-level abstractions that usually aren't directly supported by the language and are instead written on top of built-in types.

In Elixir, such abstractions are implemented with pure, stateless modules. In this chapter, you'll learn how to create and work with your own abstractions.

In a typical OO language, the basic abstraction building blocks are classes and objects. For example, there may be a `String` class that implements various string operations. Each string is then an instance of that class and can be manipulated by calling methods, as the following Ruby snippet illustrates:

```
"a string".upcase
```

This approach generally isn't used in Elixir. Being a functional language, Elixir promotes decoupling of data from the code. Instead of classes, you use modules, which are collections of functions. Instead of calling methods on objects, you explicitly call module functions and provide input data via arguments. The following snippet shows the Elixir way of uppercasing a string:

```
String.upcase("a string")
```

Another big difference from OO languages is that data is immutable. To modify data, you must call some function and take its result into a variable; the original data is left intact. The following examples demonstrate this

technique:

```
iex(1)> list = []
[]

iex(2)> list = List.insert_at(list, -1, :a)
[:a]

iex(3)> list = List.insert_at(list, -1, :b)
[:a, :b]

iex(4)> list = List.insert_at(list, -1, :c)
[:a, :b, :c]
```

In these examples, you're constantly keeping the result of the last operation and feeding it to the next one.

The important thing to notice in both Elixir snippets is that the module is used as the abstraction over the data type. When you need to work with strings, you reach for the `String` module. When you need to work with lists, you use the `List` module.

`String` and `List` are examples of modules that are dedicated to a specific data type. They're implemented in pure Elixir, and their functions rely on the predefined format of the input data. `String` functions expect a binary string as the first argument, whereas `List` functions expect a list.

Additionally, *modifier functions* (the ones that transform the data) return data of the same type. The function `String.upcase/1` returns a binary string, whereas `List.insert_at/3` returns a list.

Finally, a module also contains *query functions* that return some piece of information from the data, such as `String.length/1` and `List.first/1`. Such functions still expect an instance of the abstraction as the first argument, but they return another type of information.

The basic principles of abstraction in Elixir can be summarized as follows:

1. A module is in charge of abstracting some behavior.
2. The module's functions usually expect an instance of the abstraction as the first argument.

3. Modifier functions return a modified version of the abstraction.
4. Query functions return some other type of data.

Given these principles, it's fairly straightforward to create your own higher-level abstractions, as you'll see in the next section.

4.1 Abstracting with module

Lists and strings are admittedly lower-level types. But higher-level abstractions are based on the same principles just stated. In fact, you already saw examples of a higher-level abstraction in chapter 2. For example, a `MapSet` module implements a set. `MapSet` is implemented in pure Elixir and can serve as a good template for how to design an abstraction in Elixir.

Let's look at an example that uses `MapSet`:

```
iex(1)> days =
        MapSet.new() |>           #1
        MapSet.put(:monday) |>    #2
        MapSet.put(:tuesday)      #2

iex(2)> MapSet.member?(days, :monday)   #3
true
```

This approach more or less follows the principles stated earlier. The code is slightly simplified by using the pipe operator to chain operations together. This is possible because all the functions from the `MapSet` module take a set as the first argument. Such functions are pipe-friendly and can be chained with the `|>` operator.

Notice the `new/0` function that creates a new instance of the abstraction. There's nothing special about this function, and it could have been given any name. Its only purpose is to create a new data structure you can work on.

Because `MapSet` is an abstraction, you, as a client of this module, don't concern yourself with its internal workings or its data structure. You call `MapSet` functions, holding on to whatever result you get and passing that result back to functions from the same module.

You may think that abstractions like `MapSet` are something like user-defined types. Although there are many similarities, module-based abstractions aren't proper data types like the ones explained in chapter 2. Instead, they're implemented by composing built-in data types. For example, a `MapSet` instance is also a map, which you can verify by invoking `is_map(MapSet.new())`.

Given this template, let's try to build a simple abstraction.

4.1.1 Basic abstraction

The example in this section is a simple to-do list. The problem is admittedly not spectacular, but it's complex enough to give you something to work on while not being overly complicated. This will allow you to focus on techniques without spending too much time trying to grasp the problem itself.

The basic version of the to-do list will support the following features:

- Creating a new to-do list
- Adding new entries to the list
- Querying the list

Here's an example of the desired usage:

```
$ iex simple_todo.ex

iex(1)> todo_list =
          TodoList.new() |>
          TodoList.add_entry(~D[2023-12-19], "Dentist") |>
          TodoList.add_entry(~D[2023-12-20], "Shopping") |>
          TodoList.add_entry(~D[2023-12-19], "Movies")

iex(2)> TodoList.entries(todo_list, ~D[2023-12-19])
["Movies", "Dentist"]

iex(3)> TodoList.entries(todo_list, ~D[2023-12-18])
[]
```

This is fairly self-explanatory. You create an instance by calling
`TodoList.new/0`, then add some entries, and finally execute some queries.
The expression `~D[2023-12-19]`, as explained in section 2.4.11, creates a
date (December 19, 2023), powered by the `Date` module.

As the chapter progresses, you'll add additional features and modify the
interface slightly. You'll continue adding features throughout this book, and
by the end you'll have a fully working distributed web server that can
manage a large number of to-do lists.

For now, let's start with this simple interface. First you have to decide on the
internal data representation. In the preceding snippet, you can see that the
primary use case is finding all entries for a single date. Therefore, using a
map seems like a reasonable initial approach. You'll use dates as keys, with
values being lists of entries for given dates. With this in mind, the
implementation of the `new/0` function is straightforward.

Listing 4.1. Initializing a to-do list (simple_todo.ex)

```
defmodule TodoList do
  def new(), do: %{}
  ...
end
```

Next you have to implement the `add_entry/3` function. This function expects
a to-do list (which you know is a map) and has to add the entry to the list
under a given key (date). Of course, it's possible that no entries for that date
exist, so you have to cover that case as well. As it turns out, this can be done
with a single call to the function `Map.update/4`.

Listing 4.2. Adding an entry (simple_todo.ex)

```
defmodule TodoList do
  ...
  def add_entry(todo_list, date, title) do
    Map.update(
      todo_list,
      date,
      [title],        #1
      fn titles -> [title | titles] end    #2
```

```
    )
  end
...
end
```

The `Map.update/4` function receives a map, a key, an initial value, and an updater function. If no value exists for the given key, the initial value is used. Otherwise, the updater function is called. The function receives the existing value and returns the new value for that key. In this case, you push the new entry to the top of list. You may remember from chapter 2 that lists are most efficient when pushing new elements to the top. Therefore, you opt for a fast insertion operation but sacrifice ordering — more recently added entries are placed before the older ones in the list.

Finally, you need to implement the `entries/2` function that returns all entries for a given date, or an empty list if no task exists for that date. This is fairly straightforward, as you can see in the next listing.

Listing 4.3. Querying the to-do list (simple_todo.ex)

```
defmodule TodoList do
  ...
  def entries(todo_list, date) do
    Map.get(todo_list, date, [])
  end
end
```

You fetch a value for the given date from `todo_list`, which must be a map. The third argument to `Map.get/3` is a default value that's returned if a given key isn't present in the map.

4.1.2 Composing abstractions

Nothing stops you from creating one abstraction on top of another. In our initial take on the to-do list, there's an opportunity to move some of the code into a separate abstraction.

Look at the way you operate on a map, allowing multiple values to be stored under a single key, and retrieving all values for that key. This code could be moved to a separate abstraction. Let's call this `MultiDict`, which is

implemented in the next listing.

Listing 4.4. Implementing the `MultiDict` abstraction (todo_multi_dict.ex)

```
defmodule MultiDict do
  def new(), do: %{}

  def add(dict, key, value) do
    Map.update(dict, key, [value], &[value | &1])
  end

  def get(dict, key) do
    Map.get(dict, key, [])
  end
end
```

This is more or less a copy-and-paste from the initial to-do list implementation. The names are changed a bit, and you use the capture operator to shorten the updater lambda definition: `&[value | &1]`.

With this abstraction in place, the `TodoList` module becomes much simpler.

Listing 4.5. `TodoList` relying on a `MultiDict` (todo_multi_dict.ex)

```
defmodule TodoList do
  def new(), do: MultiDict.new()

  def add_entry(todo_list, date, title) do
    MultiDict.add(todo_list, date, title)
  end

  def entries(todo_list, date) do
    MultiDict.get(todo_list, date)
  end
end
```

This is a classical separation of concerns, where you extract a distinct responsibility into a separate abstraction, and then create another abstraction on top of it. A distinct `MultiDict` abstraction is now readily available to be used in other places in code if needed. Furthermore, you can extend `TodoList` with additional functions that are specific to to-do lists, and which therefore don't belong to `MultiDict`.

The point of this refactoring is to illustrate that the code organization isn't radically different from an OO approach. You use different tools to create abstractions (modules and functions instead of classes and methods), but the general idea is the same.

4.1.3 Structuring data with maps

`TodoList` now supports the basic features. You can insert entries into the structure and get all entries for a given date. But the interface is somewhat clumsy. When adding a new entry, you have to specify each field as a separate argument:

```
TodoList.add_entry(todo_list, ~D[2023-12-19], "Dentist")
```

If you want to extend an entry with another attribute — such as time — you must change the signature of the function, which will in turn break all the clients. Moreover, you have to change every place in the implementation where this data is being propagated.

An obvious solution to this problem is to somehow combine all entry fields as a single abstraction.

As explained in section 2.4.6, the most common way of doing this in Elixir is to use maps, with field names stored as keys of the atom type. The following snippet demonstrates how you can create and use an `entry` instance:

```
iex(1)> entry = %{date: ~D[2023-12-19], title: "Dentist"}

iex(2)> entry.date
~D[2023-12-19]

iex(3)> entry.title
"Dentist"
```

You can immediately adapt your code to represent entries with maps. As it turns out, this change is extremely simple. All you need to do is change the code of the `TodoList.add_entry` function to accept two arguments: a to-do list instance, and a map that describes an entry. The new version is presented in the following listing.

Listing 4.6. Representing entries with maps (todo_entry_map.ex)

```elixir
defmodule TodoList do
  ...
  def add_entry(todo_list, entry) do
    MultiDict.add(todo_list, entry.date, entry)
  end
  ...
end
```

That was simple! You assume an entry is a map and add it to `MultiDict`, using its date field as a key.

Let's see this in action. To add a new entry, clients now must provide a map:

```elixir
iex(1)> todo_list = TodoList.new() |>
TodoList.add_entry(%{date: ~D[2023-12-19], title: "Dentist"})
```

The client code is obviously more verbose, because it must provide field names. But because entries are now structured in a map, data retrieval is improved. The `TodoList.entries/2` function now returns complete entries, not just their titles:

```elixir
iex(2)> TodoList.entries(todo_list, ~D[2023-12-19])
[%{date: ~D[2023-12-19], title: "Dentist"}]
```

The current implementation of `TodoList` relies on a map. This means that at runtime, it's impossible to make a distinction between a map and a `TodoList` instance. In some situations, you may want to define and enforce a more precise structure definition. For such cases, Elixir provides a feature called *structs*.

4.1.4 Abstracting with structs

Let's say you need to deal with fractions in your program. A fraction is a part of a whole, represented in the form of `a/b`, where a and b are integers called a numerator and a denominator. Passing around these two values separately is noisy and error prone. Therefore it makes sense to introduce a small abstraction to help working with fractions. The following snippet demonstrates how such an abstraction could be used:

```
$ iex fraction.ex

iex(1)> Fraction.new(1, 2)
        |> Fraction.add(Fraction.new(1, 4))
        |> Fraction.value()
0.75
```

Here you sum one half (1/2) with one quarter (1/4) and return the numerical value of the resulting fraction. A fraction is created using `Fraction.new/2` and is then passed to various other functions that know how to work with it.

How can you implement this? There are many approaches, such as relying on plain tuples or using maps. In addition, Elixir provides a facility called *structs* that allows you to specify the abstraction structure up front and bind it to a module. Each module can define only one struct, which can then be used to create new instances and pattern-match on them.

A fraction has a well-defined structure, so you can use struct to specify and enforce data shape. Let's see this in action.

To define a struct, you use the `defstruct` macro (https://hexdocs.pm/elixir/Kernel.html#defstruct/1), as shown next.

Listing 4.7. Defining a structure (fraction.ex)

```
defmodule Fraction do
  defstruct a: nil, b: nil
    ...
end
```

A keyword list provided to `defstruct` defines the struct's fields together with their initial values. You can now instantiate a struct using this special syntax:

```
iex(1)> one_half = %Fraction{a: 1, b: 2}
%Fraction{a: 1, b: 2}
```

Notice how a struct bears the name of the module it's defined in. There's a tight relation between structs and modules. A struct may exist only in a module, and a single module can define only one struct.

Internally, a struct is a special kind of map. Therefore, individual fields are

accessed just like maps:

```
iex(2)> one_half.a
1

iex(3)> one_half.b
2
```

The nice thing about structs is that you can pattern-match on them:

```
iex(4)> %Fraction{a: a, b: b} = one_half
%Fraction{a: 1, b: 2}

iex(5)> a
1

iex(6)> b
2
```

This makes it possible to assert that some variable is really a struct:

```
iex(6)> %Fraction{} = one_half      #1
%Fraction{a: 1, b: 2}               #1

iex(7)> %Fraction{} = %{a: 1, b: 2}     #2
** (MatchError) no match of right hand side value: %{a: 1, b: 2}
```

Here you use a `%Fraction{}` pattern that matches any `Fraction` struct, regardless of its contents. Pattern matching with structs works much like it does with maps. This means in a pattern match, you need to specify only the fields you're interested in, ignoring all other fields.

Updating a struct works similarly to the way it works with maps:

```
iex(8)> one_quarter = %Fraction{one_half | b: 4}
%Fraction{a: 1, b: 4}
```

This code creates a new struct instance based on the original one (`one_half`), changing the value of the field `b` to 4.

The shape of the struct is defined at compile time. As a result, some errors can be caught by the Elixir compiler. For example, suppose we make a typo in the field name:

```
iex(9)> %Fraction{a: 1, d: 2}
** (KeyError) key :d not found
```

The struct doesn't specify the field :d, so the error is reported. In contrast, if you used a regular map, this code would succeed. But the program would fail in a distant place with a non-obvious reason, which would make the error more difficult to debug.

It's worth noting that this error is reported at compile time. If you make the same mistake in a source file, the code won't even compile.

Armed with this knowledge, let's add some functionality to the `Fraction` abstraction. First you need to provide the creation function.

Listing 4.8. Instantiating a fraction (fraction.ex)

```
defmodule Fraction do
  ...
  def new(a, b) do
    %Fraction{a: a, b: b}
  end
  ...
end
```

This is a simple wrapper around the `%Fraction{}` syntax. It makes the client code clearer and less coupled with the fact that structs are used.

Next, implement a `Fraction.value/1` function that returns a float representation of the fraction.

Listing 4.9. Calculating the fraction value (fraction.ex)

```
defmodule Fraction do
  ...

  def value(%Fraction{a: a, b: b}) do      #1
    a / b
  end

  ...
end
```

The `value/1` function matches a fraction, taking its fields into individual variables and using them to compute the final result. The benefit of pattern matching is that the input type is enforced. If you pass anything that isn't a fraction instance, you'll get a match error.

Instead of decomposing fields into variables, you could also use dot notation:

```
def value(fraction) do
  fraction.a / fraction.b
end
```

This version is arguably clearer, but on the flip side it accepts any map, not just `Fraction` structs, which might lead to subtle bugs. For example, suppose there's a `Rectangle` struct which has the same fields. You could accidentally pass such struct to this function, and instead of failing, the function would return some meaningless result.

One final thing left to do is to implement the `add` function.

Listing 4.10. Adding two fractions (fraction.ex)

```
defmodule Fraction
  ....

  def add(%Fraction{a: a1, b: b1}, %Fraction{a: a2, b: b2}) do
    new(
      a1 * b2 + a2 * b1,
      b2 * b1
    )
  end

  ....
end
```

You can now test your fraction:

```
iex(1)> Fraction.new(1, 2)
        |> Fraction.add(Fraction.new(1, 4))
        |> Fraction.value()
0.75
```

The code works as expected. By representing fractions with a struct, you can

provide the definition of your type, listing all fields and their default values. Furthermore, it's possible to distinguish struct instances from any other data type. This allows you to place `%Fraction{}` matches in function arguments, thus asserting that you only accept fraction instances.

Structs vs. maps

You should always be aware that structs are in reality just maps, so they have the same characteristics with respect to performance and memory usage. But a struct instance receives special treatment. Some things that can be done with maps don't work with structs. For example, you can't call `Enum` functions on a struct:

```
iex(1)> one_half = Fraction.new(1, 2)

iex(2)> Enum.to_list(one_half)
** (Protocol.UndefinedError) protocol Enumerable not implemented
   %Fraction{a: 1, b: 2}
```

Remember, a struct is a functional abstraction and should therefore behave according to the implementation of the module where it's defined. In the case of the `Fraction` abstraction, you must define whether `Fraction` is enumerable and, if so, in what way. If this isn't done, `Fraction` isn't an enumerable, so you can't call `Enum` functions on it.

In contrast, a plain map is an enumerable, so you can convert it to a list:

```
iex(3)> Enum.to_list(%{a: 1, b: 2})
[a: 1, b: 2]
```

On the other hand, because structs are maps, directly calling `Map` functions works:

```
iex(4)> Map.to_list(one_half)
[__struct__: Fraction, a: 1, b: 2]
```

Notice the `__struct__: Fraction` bit. This key/value pair is automatically included in each struct. It helps Elixir distinguish structs from plain maps and perform proper runtime dispatches from within polymorphic generic code. You'll learn more about this later when we describe protocols.

The `__struct__` field has an important consequence for pattern matching. A struct pattern can't match a plain map:

```
iex(5)> %Fraction{a: a, b: b} = %{a: 1, b: 2}
** (MatchError) no match of right hand side value: %{a: 1, b: 2}
```

But a plain map pattern can match a struct:

```
iex(5)> %{a: a, b: b} = %Fraction{a: 1, b: 2}
%Fraction{a: 1, b: 2}

iex(6)> a
1

iex(7)> b
2
```

This is due to the way pattern matching works with maps. Remember, all fields from the pattern must exist in the matched term. When matching a map to a struct pattern, this isn't the case, because `%Fraction{}` contains the field `struct`, which isn't present in the map being matched.

The opposite works, because you match a struct to the `%{a: a, b: b}` pattern. Because all these fields exist in the `Fraction` struct, the match is successful.

Records

In addition to maps and structs, there's another way to structure data: *records*. This is a feature that lets you use tuples and still be able to access individual elements by name. Records can be defined using the `defrecord` and `defrecordp` macros from the `Record` module (https://hexdocs.pm/elixir/Record.html).

Given that they're essentially tuples, records should be faster than maps (although the difference usually isn't significant in the grand scheme of things). On the flip side, the usage is more verbose, and it's not possible to access fields by name dynamically.

Records are present mostly for historical reasons. Before maps appeared,

records were one of the main tools for structuring data. In fact, many libraries from the Erlang ecosystem use records as their interface. If you need to interface an Erlang library using a record defined in that library, you must import that record into Elixir and define it as a record. This can be done with the `Record.extract/2` function in conjunction with the `defrecord` macro. This idiom isn't required often, so records won't be demonstrated here. Still, it may be useful to keep this information in the back of your head and research it if the need arises.

4.1.5 Data transparency

The modules you've devised so far are abstractions, because clients aren't aware of their implementation details. For example, as a client, you call `Fraction.new/2` to create an instance of the abstraction, and then send that instance back to some other function from that same module.

But the entire data structure is always visible. As a client, you can obtain individual fraction values, even if this was not intended by the library developer.

It's important to be aware that data in Elixir is always transparent. Clients can read any information from your structs (and any other data type), and there's no easy way of preventing that. In that sense, encapsulation works differently than in typical OO languages. In Elixir, modules are in charge of abstracting the data and providing operations to manipulate and query that data, but the data is never hidden.

Let's verify this in a shell session:

```
$ iex todo_entry_map.ex

iex(1)> todo_list = TodoList.new() |>
          TodoList.add_entry(%{date: ~D[2023-12-19], title: "Dent

%{~D[2023-12-19] => [%{date: ~D[2023-12-19], title: "Dentist"}]}
```

Looking at the return value, you can see the entire structure of the to-do list. From the output, you can immediately tell that the to-do list is powered by a map, and you can also find out details about how individual entries are kept.

Let's look at another example. A `MapSet` instance is also an abstraction, powered by the `MapSet` module and a corresponding struct. At first glance, this isn't visible:

```
iex(1)> mapset = MapSet.new([:monday, :tuesday])
MapSet.new([:monday, :tuesday])
```

Notice how the result of the expression is printed in a special way, using `MapSet.new(…)` output. This is due to the inspection mechanism in Elixir: whenever a result is printed in the shell, the function `Kernel.inspect/1` is called to transform the structure into an inspected string. For each abstraction you build, you can override the default behavior and provide your own inspected format. This is exactly what `MapSet` does, and you'll learn how to do this for your type later in this chapter when we discuss protocols.

Occasionally you may want to see the pure data structure, without this decorated output. This can be useful when you're debugging, analyzing, or reverse-engineering code. To do so, you can provide a special option to the `inspect` function:

```
iex(2)> IO.puts(inspect(mapset, structs: false))
%{__struct__: MapSet, map: %{monday: [], tuesday: []}, version: 2
```

The output now reveals the complete structure of a date, and you can "see through" the `MapSet` abstraction. This demonstrates that data privacy can't be fully enforced in Elixir. Remember from chapter 2 that the only complex types are tuples, lists, and maps. Any other abstraction, such as `MapSet` or your own `TodoList`, will ultimately be built on top of these types.

The benefit of data transparency is that the data can be easily inspected, which can be useful for debugging purposes. But as a client of an abstraction, you shouldn't rely on its internal representation, even though it's visible to you. You shouldn't pattern-match on the internal structure or try to extract or modify individual parts of it because a proper abstraction, such as `MapSet`, doesn't guarantee what the data will look like. The only guarantee is that the module's functions will work if you send them a properly structured instance that you already received from that same module.

Sometimes a module will publicly document some parts of its internal

structure. Good examples of this are the date and time modules, such as `Date`, `Time`, and `DateTime`. Looking at the documentation, you'll see explicit mention that the corresponding data is represented as a structure with fields such as year, month, hour, and so on. In this case, the structure of the data is publicly documented, and you can freely rely on it.

One final thing you should know, related to data inspection, is the `IO.inspect/1` function. This function prints the inspected representation of a structure to the screen and returns the structure itself. This is particularly useful when debugging a piece of code. Look at the following example:

```
iex(1)> Fraction.new(1, 4) |>
          Fraction.add(Fraction.new(1, 4)) |>
          Fraction.add(Fraction.new(1, 2)) |>
          Fraction.value()
1.0
```

This code relies on the pipe operator to perform a series of fraction operations. Let's say you want to inspect the entire structure after each step. You can easily insert the call to `IO.inspect/1` after every line:

```
iex(2)> Fraction.new(1, 4) |>
          IO.inspect() |>
          Fraction.add(Fraction.new(1, 4)) |>
          IO.inspect() |>
          Fraction.add(Fraction.new(1, 2)) |>
          IO.inspect() |>
          Fraction.value()

%Fraction{a: 1, b: 4}     #1
%Fraction{a: 8, b: 16}    #1
%Fraction{a: 32, b: 32}   #1
```

This works because `IO.inspect/1` prints the data structure and then returns that same data structure unchanged.

We're now done with the basic theory behind functional abstractions, but you'll practice some more by extending the to-do list.

4.2 Working with hierarchical data

In this section, you'll extend the `TodoList` abstraction to provide basic CRUD support. You already have the *C* and *R* parts resolved with the `add_entry/2` and `entries/2` functions, respectively. Now you need to add support for updating and deleting entries. To do this, you must be able to uniquely identify each entry in the to-do list, so you'll begin by adding unique ID values to each entry.

4.2.1 Generating IDs

When adding a new entry to the list, you'll autogenerate its ID value, using incremental integers for IDs. To implement this, you have to do a couple of things:

- Represent the to-do list as a struct — You need to do this because the to-do list now has to keep two pieces of information: the entries collection and the ID value for the next entry.
- Use the entry's ID as the key — So far, when storing entries in a collection, you used the entry's date as the key. You'll change this and use the entry's ID instead. This will make it possible to quickly insert, update, and delete individual entries. You'll now have exactly one value per each key, so you won't need the `MultiDict` abstraction anymore.

Let's start implementing this. The code in the following listing contains the module and struct definitions.

Listing 4.11. TodoList struct (todo_crud.ex)

```
defmodule TodoList do
  defstruct next_id: 1, entries: %{}      #1

  def new(), do: %TodoList{}      #2
  ...
end
```

The to-do list will now be represented as a struct with two fields. The field `next_id` contains the ID value that will be assigned to the new entry while it's being added to the structure. The field `entries` is the collection of entries. As has been mentioned, you're now using a map, and the keys are entry ID values.

During the struct definition, the default values for the `next_id` and `entries` fields are immediately specified. Therefore, you don't have to provide these when creating a new instance. The function `new/0` creates and returns an instance of the struct.

Next, it's time to reimplement the `add_entry/2` function. It has to do more work:

- Set the ID for the entry being added.
- Add the new entry to the collection.
- Increment the `next_id` field.

Here's the code.

Listing 4.12. Autogenerating ID values for new entries (todo_crud.ex)

```
defmodule TodoList do
  ...

  def add_entry(todo_list, entry) do
    entry = Map.put(entry, :id, todo_list.next_id)    #1

    new_entries = Map.put(        #2
      todo_list.entries,          #2
      todo_list.next_id,          #2
      entry                       #2
    )

    %TodoList{todo_list |         #3
      entries: new_entries,       #3
      next_id: todo_list.next_id + 1     #3
    }
  end

  ...
end
```

A lot of things happen here, so let's take them one at a time.

In the function body, you first update the entry's `id` value with the value stored in the `next_id` field. Notice how you use `Map.put/3` to update the entry map. The input map may not contain the `id` field, so you can't use the

standard `%{entry | id: next_id}` technique, which works only if the `id` field is already present in the map.

Once the entry is updated, you add it to the entries collection, keeping the result in the `new_entries` variable.

Finally, you must update the `TodoList` struct instance, setting its `entries` field to the `new_entries` collection and incrementing the `next_id` field. Essentially, you made a complex change in the struct, modifying multiple fields as well as the input entry (because you set its `id` field).

To the external caller, the entire operation will be atomic. Either everything will happen or, in case of an error, nothing at all. This is the consequence of immutability. The effect of adding an entry is visible to others only when the `add_entry/2` function finishes and its result is taken into a variable. If something goes wrong and you raise an error, the effect of any transformations won't be visible.

It's also worth repeating, as mentioned in chapter 2, that the new to-do list (the one returned by the `add_entry/2` function) will share as much memory as possible with the input to-do list.

With the `add_entry/2` function finished, you need to adjust the `entries/2` function. This will be more complicated because you changed the internal structure. Earlier, you kept a date-to-entries mapping. Now entries are stored using `id` as the key, so you have to iterate through all the entries and return the ones that fall on a given date. This code is shown next.

Listing 4.13. Filtering entries for a given date (todo_crud.ex)

```
defmodule TodoList do
  ...

  def entries(todo_list, date) do
    todo_list.entries
    |> Map.values()                                        #1
    |> Enum.filter(fn entry -> entry.date == date end)  #2
  end
  ...
end
```

This function first uses `Map.values/1` to take the entries from the `entries` map. Then, only the entries that fall on the given date are taken using `Enum.filter/2`.

You can check whether your new version of the to-do list works:

```
$ iex todo_crud.ex
iex(1)> todo_list = TodoList.new() |>
          TodoList.add_entry(%{date: ~D[2023-12-19], title: "Dent
          TodoList.add_entry(%{date: ~D[2023-12-20], title: "Shop
          TodoList.add_entry(%{date: ~D[2023-12-19], title: "Movi

iex(2)> TodoList.entries(todo_list, ~D[2023-12-19])
[
  %{date: ~D[2023-12-19], id: 1, title: "Dentist"},
  %{date: ~D[2023-12-19], id: 3, title: "Movies"}
]
```

This works as expected, and you can even see the ID value for each entry. Also note that the interface of the `TodoList` module is the same as the previous version. You've made a number of internal modifications, changed the data representation, and practically rewritten the entire module. And yet, the module's clients don't need to be altered, because you kept the same interface for your functions.

This is nothing revolutionary — it's a classical benefit of wrapping the behavior behind a properly chosen interface. But it demonstrates how you can construct and reason about higher-level types when working with stateless modules and immutable data.

4.2.2 Updating entries

Now that your entries have ID values, you can add additional modifier operations. Let's implement the `update_entry` operation, which can be used to modify a single entry in the to-do list.

This function will accept an entry id, and an updater lambda which will be invoked to update the entry. This will work similarly to `Map.update`. The lambda will receive the original entry and return its modified version. To keep things simple, the function won't raise an error if the entry with a given

ID doesn't exist.

The following snippet illustrates the usage. Here, you modify the date of an entry that has an ID value of 1:

```
iex(1)> TodoList.update_entry(
          todo_list,
          1,                          #1
          &Map.put(&1, :date, ~D[2023-12-20])      #2
        )
```

The implementation is presented in the following listing.

Listing 4.14. Updating an entry (todo_crud.ex)

```
defmodule TodoList do
  ...

  def update_entry(todo_list, entry_id, updater_fun) do
    case Map.fetch(todo_list.entries, entry_id) do
      :error ->          #1
        todo_list

      {:ok, old_entry} ->     #2
        new_entry = updater_fun.(old_entry)
        new_entries = Map.put(todo_list.entries, new_entry.id, ne
        %TodoList{todo_list | entries: new_entries}
    end
  end

  ...
end
```

Let's break down what happens here. First, you look up the entry with the given ID, using `Map.fetch/2`. The function will return `:error` if the entry doesn't exist, and `{:ok, value}` otherwise.

In the first case, if the entry doesn't exist, you return the original version of the list. Otherwise, you have to call the updater lambda to get the modified entry. Then you store the modified entry into the entries collection. Finally, you store the modified entries collection in the `TodoList` instance and return that instance.

4.2.3 Immutable hierarchical updates

You may not have noticed, but in the previous example you performed a deep update of an immutable hierarchy. Let's break down what happens when you call `TodoList.update_entry(todo_list, id, updater_lambda)`:

1. You take the target entry into a separate variable.
2. You call the updater that returns the modified version of the entry to you.
3. You call `Map.put` to put the modified entry into the entries collection.
4. You return the new version of the to-do list, which contains the new entries collection.

Notice that steps 2, 3, and 4 are the ones where you transform data. Each of these steps creates a new variable that contains the transformed data. In each subsequent step, you take this data and update its container, again by creating a transformed version of it.

This is how you work with immutable data structures. If you have hierarchical data, you can't directly modify part of it that resides deep in its tree. Instead, you have to walk down the tree to the particular part that needs to be modified, and then transform it and all of its ancestors. The result is a copy of the entire model (in this case, the to-do list). As mentioned, the two versions — new and previous — will share as much memory as possible.

Provided helpers

Although the technique presented works, it may become cumbersome for deeper hierarchies. Remember, to update an element deep in the hierarchy, you have to walk to that element and then update all of its parents. To simplify this, Elixir offers support for more elegant deep hierarchical updates.

Let's look at a basic example. Let's say that the to-do list is represented as a simple map, where keys are IDs and values are plain maps consisting of fields. Let's create one instance of such to-do list:

```
iex(1)> todo_list = %{
  1 => %{date: ~D[2023-12-19], title: "Dentist"},
```

```
    2 => %{date: ~D[2023-12-20], title: "Shopping"},
    3 => %{date: ~D[2023-12-19], title: "Movies"}
}
```

Now, let's say you change your mind and want to go to the theater instead of a movie. The original structure can be modified elegantly using the `Kernel.put_in/2` macro:

```
iex(2)> put_in(todo_list[3].title, "Theater")     #1

%{
  1 => %{date: ~D[2023-12-19], title: "Dentist"},
  2 => %{date: ~D[2023-12-20], title: "Shopping"},
  3 => %{date: ~D[2023-12-19], title: "Theater"}     #2
}
```

What happened here? Internally, `put_in/2` does something similar to what you did. It walks recursively to the desired element, transforms it, and then updates all the parents. Notice that this is still an immutable operation, meaning the original structure is left intact, and you have to take the result to a variable.

To be able to do a recursive walk, `put_in/2` needs to receive source data and a path to the target element. In the preceding example, the source is provided as `todo_list` and the path is specified as `[3].title`. The macro `put_in/2` then walks down that path, rebuilding the new hierarchy on the way up.

It's also worth noting that Elixir provides similar alternatives for data retrieval and updates in the form of the `get_in/2`, `update_in/2`, and `get_and_update_in/2` macros. The fact that these are macros means the path you provide is evaluated at compile time and can't be built dynamically.

If you need to construct paths at runtime, there are equivalent functions that accept data and path as separate arguments. For example, Elixir also includes the `put_in/3` macro, which can be used as follows:

```
iex(3)> path = [3, :title]

iex(4)> put_in(todo_list, path, "Theater")     #1
```

Functions and macros, such as `put_in`, rely on the `Access` module, which

allows you to work with key/value structures such as maps. You can also make your own abstraction work with `Access`. You need to implement a couple of functions required by the `Access` contract, and then `put_in` and related macros and functions will know how to work with your own abstraction. Refer to the official `Access` documentation (https://hexdocs.pm/elixir/Access.html) for more details.

Exercise: deleting an entry

Your `TodoList` module is almost complete. You've already implemented create (`add_entry/2`), retrieve (`entries/2`), and update (`update_entry/3`) operations. The last thing remaining is the `delete_entry/2` operation. This is straightforward, and it's left for you to do as an exercise. If you get stuck, the solution is provided in the source file todo_crud.ex.

4.2.4 Iterative updates

So far, you've been doing updates manually, one at a time. Now it's time to implement iterative updates. Imagine that you have a raw list describing the entries:

```
$ iex todo_builder.ex

iex(1)> entries = [
          %{date: ~D[2023-12-19], title: "Dentist"},
          %{date: ~D[2023-12-20], title: "Shopping"},
          %{date: ~D[2023-12-19], title: "Movies"}
        ]
```

Now you want to create an instance of the to-do list that contains all of these entries:

```
iex(2)> todo_list = TodoList.new(entries)
```

It's obvious that the function `new/1` performs an iterative build of the to-do list. How can you implement such a function? As it turns out, this is simple, as you can see in the following listing.

Listing 4.15. Iteratively building the to-do list (todo_builder.ex)

```
defmodule TodoList do
  ...

  def new(entries \\ []) do
    Enum.reduce(
      entries,
      %TodoList{},        #1
      fn entry, todo_list_acc ->     #2
        add_entry(todo_list_acc, entry)
      end
    )
  end
  ...
end
```

To build the to-do list iteratively, you're relying on `Enum.reduce/3`. Recall from chapter 3 that `reduce` is used to transform something enumerable to anything else. In this case, you're transforming a raw list of `Entry` instances into an instance of the `TodoList` struct. Therefore, you call `Enum.reduce/3`, passing the input list as the first argument, the new structure instance as the second argument (the initial accumulator value), and the lambda that's called in each step.

The lambda is called for each entry in the input list. Its task is to add the entry to the current accumulator (`TodoList` struct) and return the new accumulator value. To do this, the lambda delegates to the already-present `add_entry/2` function, reversing the argument order. The arguments need to be reversed because `Enum.reduce/3` calls the lambda, passing the iterated element (entry) and accumulator (`TodoList` struct). In contrast, `add_entry` accepts a struct and an entry.

Notice that you can make the lambda definition more compact with the help of the capture operator:

```
def new(entries \\ []) do
  Enum.reduce(
    entries,
    %TodoList{},
    &add_entry(&2, &1)     #1
  )
end
```

Whether you prefer this version or the previous one is entirely up to your personal taste.

4.2.5 Exercise: importing from a file

Now it's time for you to practice a bit. In this exercise, you'll create a `TodoList` instance from the comma-separated file.

Assume that you have a todos.csv file in the current folder. Each line in the file describes a single to-do entry:

```
2023-12-19,Dentist
2023-12-20,Shopping
2023-12-19,Movies
```

Your task is to create an additional module, `TodoList.CsvImporter`, that can be used to create a `TodoList` instance from the file contents:

```
iex(1)> todo_list = TodoList.CsvImporter.import("todos.csv")
```

To simplify the task, assume that the file is always available and in the correct format. Also assume that the comma character doesn't appear in the entry title.

This is generally not hard to do, but it might require some cracking and experimenting. Here are a couple of hints that will lead you in the right direction.

First, create a single file with the following layout:

```
defmodule TodoList do
  ...
end

defmodule TodoList.CsvImporter do
  ...
end
```

Always work in small steps. Implement part of the calculation, and then print the result to the screen using `IO.inspect/1`. I can't stress enough how important this is. This task requires some data pipelining. Working in small

steps will allow you to move gradually and verify that you're on the right track.

The general steps you should undertake are as follows:

1. Open a file and go through it, removing `\n` from each line. Hint: use `File.stream!/1`, `Stream.map/2`, and `String.trim_trailing/2`. You did this in chapter 3, when we talked about streams, in the example where you filtered lines longer than 80 characters.
2. Using `Stream.map`, transform each line obtained from the previous step into a to-do list entry.
   A. Convert the line into a `[date_string, title]` list using `String.split/2`.
   B. Convert the date string into a date using `Date.from_iso8601!` (https://hexdocs.pm/elixir/Date.html#from_iso8601!/2).
   C. Create the to-do list entry (a map in the shape of `%{date: date, title: title}`).

The output of step 2 is an enumerable that consists of to-do entries. Pass that enumerable to the `TodoList.new/1` function that you recently implemented.

In each of these steps, you'll receive an enumerable as an input, transform each element, and pass the resulting enumerable forward to the next step. In the final step, the resulting enumerable is passed to the already-implemented `TodoList.new/1`, and the to-do list is created.

If you work in small steps, it's harder to get lost. For example, you can start by opening a file and printing each line to the screen. Then try to remove the trailing newline from each line and print them to the screen, and so on.

While transforming the data in each step, you can work with `Enum` functions or functions from the `Stream` module. It will probably be simpler to start with eager functions from the `Enum` module and get the entire thing to work. Then try to replace as many of the `Enum` functions as possible with their `Stream` counterparts. Recall from chapter 3 that the `Stream` functions are lazy and composable, which can reduce the amount of intermediate memory required for the operation.

If you get lost, the solution is provided in the file todo_import.ex.

In the meantime, we're almost done with our exploration of higher-level abstractions. The final topic we'll briefly discuss is the Elixir way of doing polymorphism.

4.3 Polymorphism with protocols

Polymorphism is a runtime decision about which code to execute, based on the nature of the input data. In Elixir, the basic (but not the only) way of doing this is by using the language feature called *protocols*.

Before discussing protocols, let's see them in action. You've already seen polymorphic code. For example, the entire `Enum` module is generic code that works on anything enumerable, as the following snippet illustrates:

```
Enum.each([1, 2, 3], &IO.inspect/1)
Enum.each(1..3, &IO.inspect/1)
Enum.each(%{a: 1, b: 2}, &IO.inspect/1)
```

Notice how you use the same `Enum.each/2` function, sending it different data structures: a list, a range, and a map. How does `Enum.each/2` know how to walk each structure? It doesn't. The code in `Enum.each/2` is generic and relies on a contract. This contract, called a *protocol*, must be implemented for each data type you wish to use with `Enum` functions.

Let's see how to define and use protocols.

4.3.1 Protocol basics

A *protocol* is a module in which you declare functions without implementing them. Consider it a rough equivalent of an OO interface. The generic logic relies on the protocol and calls its functions. Then you can provide a concrete implementation of the protocol for different data types.

Let's look at an example. The protocol `String.Chars` is provided by the Elixir standard library and is used to convert data to a binary string. This is how the protocol is defined in the Elixir source:

```
defprotocol String.Chars do     #1
  def to_string(term)    #2
end
```

This resembles the module definition, with the notable difference that functions are declared but not implemented.

Notice the first argument of the function (the `term`). At runtime, the type of this argument determines the implementation that's called. Let's see this in action. Elixir already implements the protocol for atoms, numbers, and some other data types, so you can issue the following calls:

```
iex(1)> String.Chars.to_string(1)
"1"

iex(2)> String.Chars.to_string(:an_atom)
"an_atom"
```

If the protocol isn't implemented for the given data type, an error is raised:

```
iex(3)> String.Chars.to_string(TodoList.new())
** (Protocol.UndefinedError) protocol String.Chars not implemente
```

Usually you don't need to call the protocol function directly. More often, there's generic code that relies on the protocol. In the case of `String.Chars`, this is the auto-imported function `Kernel.to_string/1`:

```
iex(4)> to_string(1)
"1"

iex(5)> to_string(:an_atom)
"an_atom"

iex(6)> to_string(TodoList.new())
** (Protocol.UndefinedError) protocol String.Chars not implemente
```

As you can see, the behavior of `to_string/1` is exactly the same as that of `String.Chars.to_string/1`. This is because `Kernel.to_string/1` delegates to the `String.Chars` implementation.

In addition, you can send anything that implements `String.Chars` to `IO.puts/1`:

```
iex(7)> IO.puts(1)
1

iex(8)> IO.puts(:an_atom)
an_atom

iex(9)> IO.puts(TodoList.new())
** (Protocol.UndefinedError) protocol String.Chars not implemente
```

As you can see, an instance of the `TodoList` isn't printable because
`String.Chars` isn't implemented for the corresponding type.

4.3.2 Implementing a protocol

How do you implement a protocol for a specific type? Let's refer to the Elixir
source again. The following snippet implements `String.Chars` for integers:

```
defimpl String.Chars, for: Integer do
  def to_string(term) do
    Integer.to_string(term)
  end
end
```

You start the implementation by calling the `defimpl` macro. Then you specify
which protocol to implement and the corresponding data type. Finally, the
`do`/`end` block contains the implementation of each protocol function. In the
example, the implementation delegates to the existing standard library
function `Integer.to_string/1`.

The `for: Type` part deserves some explanation. The type is an atom and can
be any of following aliases: `Tuple`, `Atom`, `List`, `Map`, `BitString`, `Integer`,
`Float`, `Function`, `PID`, `Port`, or `Reference`. These values correspond to built-
in Elixir types.

In addition, the alias `Any` is allowed, which makes it possible to specify a
fallback implementation. If a protocol isn't defined for a given type, an error
will be raised, unless a fallback to `Any` is specified in the protocol definition
and an `Any` implementation exists. Refer to the protocol documentation
(https://hexdocs.pm/elixir/Protocol.html) for details.

Finally, and most importantly, the type can be any other arbitrary alias (but

not a regular, simple atom):

```
defimpl String.Chars, for: SomeAlias do
  ...
end
```

This implementation will be called if the first argument of the protocol function is a struct defined in the corresponding module.

For example, you can implement `String.Chars` for `TodoList`. Do this:

```
iex(1)> defimpl String.Chars, for: TodoList do
          def to_string(_) do
            "#TodoList"
          end
        end
```

Now you can pass a to-do list instance to `IO.puts/1`:

```
iex(2)> IO.puts(TodoList.new())
#TodoList
```

It's important to notice that the protocol implementation doesn't need to be part of any module. This has powerful consequences: you can implement a protocol for a type even if you can't modify the type's source code. You can place the protocol implementation anywhere in your own code, and the runtime will be able to take advantage of it.

4.3.3 Built-in protocols

Elixir comes with some predefined protocols. It's best to consult the online documentation for the complete reference (https://hexdocs.pm/elixir), but let's mention some of the more important ones.

You've already seen `String.Chars`, which specifies the contract for converting data into a binary string. There's also the `List.Chars` protocol, which converts input data to a character string (a list of characters).

If you want to control how your structure is printed in the debug output (via the `inspect` function), you can implement the `Inspect` protocol.

Arguably the most important protocol is `Enumerable`. By implementing it, you can make your data structure *enumerable*. This means you can use all the functions from the `Enum` and `Stream` modules for free! This is probably the best demonstration of protocol usefulness. Both `Enum` and `Stream` are generic modules that offer many useful functions, which can work on your custom data structures as soon as you implement the `Enumerable` protocol.

Closely related to enumeration is the `Collectable` protocol. Recall from chapter 3 that a collectable structure is one that you can repeatedly add elements to. A collectable can be used with comprehensions to collect results or with `Enum.into/2` to transfer elements of one structure (enumerable) to another (collectable).

And, of course, you can define your own protocols and implement them for any available data structure (your own or someone else's). See the `Kernel.defprotocol/2` documentation for more information.

Collectable to-do list

Let's look at a more involved example. You'll make your to-do list collectable so that you can use it as a comprehension target. This is a slightly more advanced example, so don't worry if you don't get every detail in the first go.

To make the abstraction collectable, you have to implement the corresponding protocol:

```
defimpl Collectable, for: TodoList do      #1
  def into(original) do                     #1
    {original, &into_callback/2}            #1
  end

  defp into_callback(todo_list, {:cont, entry}) do     #2
    TodoList.add_entry(todo_list, entry)               #2
  end                                                  #2
                                                       #2
  defp into_callback(todo_list, :done), do: todo_list #2
  defp into_callback(_todo_list, :halt), do: :ok       #2
end
```

The exported function `into/1` is called by the generic code (comprehensions, for example). Here you provide the implementation that returns the appender lambda. This appender lambda is then repeatedly invoked by the generic code to append each element to your data structure.

The appender function receives a to-do list and an instruction hint. If you receive `{:cont, entry}`, you must add a new entry. If you receive `:done`, you return the list, which at this point contains all appended elements. Finally, `:halt` indicates that the operation has been canceled, and the return value is ignored.

Let's see this in action. Copy and paste the previous code into the shell, and then try the following:

```
iex(1)> entries = [
          %{date: ~D[2023-12-19], title: "Dentist"},
          %{date: ~D[2023-12-20], title: "Shopping"},
          %{date: ~D[2023-12-19], title: "Movies"}
        ]

iex(2)> Enum.into(entries, TodoList.new())    #1
%TodoList{...}
```

By implementing the `Collectable` protocol, you essentially adapt the `TodoList` abstraction to any generic code that relies on that protocol, such as `Enum.into/2` or `for` comprehension.

4.4 Summary

- A module is used to create an abstraction. A module's functions create, manipulate, and query data. Clients can inspect the entire structure but shouldn't rely on it.
- Maps can be used to group different fields together in a single structure.
- Structs are special kinds of maps that allow you to define data abstractions related to a module.
- Polymorphism can be implemented with protocols. A protocol defines an interface that is used by the generic logic. You can then provide specific protocol implementations for a data type.

5 Concurrency primitives

This chapter covers

- Understanding BEAM concurrency principles
- Working with processes
- Working with stateful server processes
- Runtime considerations

Now that you have sufficient knowledge of Elixir and functional programming idioms, we'll turn our attention to BEAM concurrency, a feature that plays a central role in Elixir's and Erlang's support for scalability, fault-tolerance, and distribution.

In this chapter, we'll start our tour of BEAM concurrency by looking at basic techniques and tools. Before we explore the lower-level details, we'll take a look at higher-level principles.

5.1 Concurrency in BEAM

Erlang is all about writing highly available systems — systems that run forever and are always able to meaningfully respond to client requests. To make your system highly available, you have to tackle the following challenges:

- *Fault-tolerance* — Minimize, isolate, and recover from the effects of runtime errors.
- *Scalability* — Handle a load increase by adding more hardware resources without changing or redeploying the code.
- *Distribution* — Run your system on multiple machines so that others can take over if one machine crashes.

If you address these challenges, your systems can constantly provide service with minimal downtime and failures.

Concurrency plays an important role in achieving high availability. In BEAM, the unit of concurrency is a *process*: a basic building block that makes it possible to build scalable, fault-tolerant, distributed systems.

 Note

A BEAM process shouldn't be confused with an OS process. As you're about to learn, BEAM processes are much lighter and cheaper than OS processes. Because this book deals mostly with BEAM, the term *process* in the remaining text refers to a BEAM process.

In production, a typical server system must handle many simultaneous requests from many different clients, maintain a shared state (for example, caches, user session data, and server-wide data), and run some additional background processing jobs. For the server to work normally, all of these tasks should run reasonably quickly and be reliable.

Because many tasks are pending simultaneously, it's imperative to execute them in parallel as much as possible, thus taking advantage of all available CPU resources. For example, it's extremely bad if the lengthy processing of one request blocks all other pending requests and background jobs. Such behavior can lead to a constant increase in the request queue, and the system can become unresponsive.

Moreover, tasks should be as isolated from each other as possible. You don't want an unhandled exception in one request handler to crash another unrelated request handler, a background job, or, especially, the entire server. You also don't want a crashing task to leave behind an inconsistent memory state, which might later compromise another task.

That's exactly what the BEAM concurrency model does for us. Processes help us run things in parallel, allowing us to achieve *scalability* — the ability to address a load increase by adding more hardware power that the system automatically takes advantage of.

Processes also ensure isolation, which in turn gives us *fault-tolerance* — the ability to localize and limit the impact of unexpected runtime errors that

inevitably occur. If you can localize exceptions and recover from them, you can implement a system that truly never stops, even when unexpected errors occur.

In BEAM, a process is a concurrent thread of execution. Two processes run concurrently and may therefore run in parallel, assuming at least two CPU cores are available. Unlike OS processes or threads, BEAM processes are lightweight concurrent entities handled by the VM, which uses its own scheduler to manage their concurrent execution.

Figure 5.1. BEAM as a single OS process, using a few threads to schedule a large number of processes
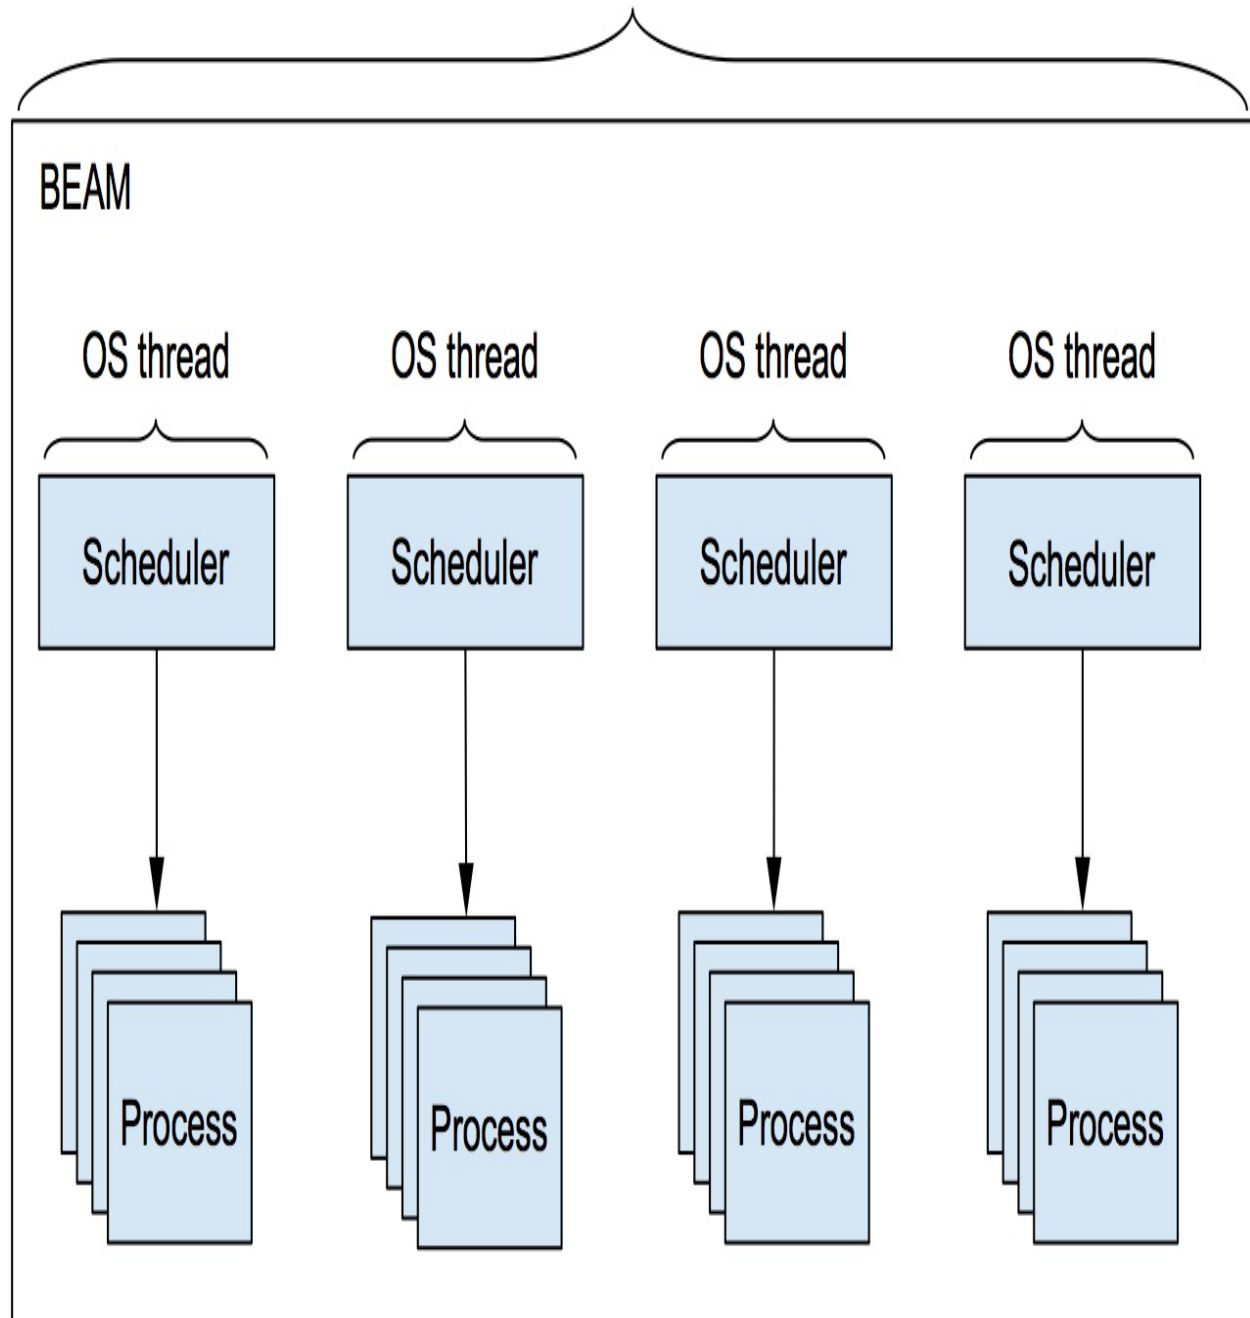
By default, BEAM uses as many schedulers as there are CPU cores available. For example, on a quad-core machine, four schedulers are used, as shown in figure 5.1.

Each scheduler runs in its own thread, and the entire VM runs in a single OS process. In figure 5.1, there's one OS process and four OS threads, and that's all you need to run a highly concurrent server system.

A scheduler is in charge of the interchangeable execution of processes. Each process gets an execution time slot; after the time is up, the running process is preempted, and the next one takes over.

Processes are light. It takes only a couple of microseconds to create a single process, and its initial memory footprint is a few kilobytes. By comparison, OS threads usually use a couple of megabytes just for the stack. Therefore, you can create a large number of processes: the theoretical limit imposed by the VM is roughly 134 million!

This feature can be exploited in server-side systems to manage various tasks that should run simultaneously. Using a dedicated process for each task, you can take advantage of all available CPU cores and parallelize the work as much as possible.

Moreover, running tasks in different processes improves the server's reliability and fault-tolerance. BEAM processes are completely isolated; they share no memory, and a crash of one process won't take down other processes. In addition, BEAM provides a means to detect a process crash and do something about it, such as restarting the crashed process. All this makes it easier to create systems that are more stable and can gracefully recover from unexpected errors, which inevitably occur in production.

Finally, each process can manage some state and can receive messages from other processes to manipulate or retrieve that state. As you saw in part 1 of this book, data in Elixir is immutable. To keep it alive, you have to hold on to it, constantly passing the result of one function to another. A process can be considered a container of this data — a place where an immutable structure is stored and kept alive for a longer time, possibly forever.

As you can see, there's more to concurrency than parallelization of the work. With this high-level view of BEAM processes in place, let's look at how you can create processes and work with them.

5.2 Working with processes

The benefits of processes are most obvious when you want to run something concurrently and parallelize the work as much as possible. For example, let's say you need to run a bunch of potentially long-running database queries. You could run those queries sequentially, one at a time, or you can try to run them concurrently, hoping that the total execution time will be reduced.

Concurrency vs. parallelism

It's important to realize that concurrency doesn't necessarily imply parallelism. Two concurrent programs have independent execution contexts, but this doesn't mean they will run in parallel. If you run two CPU-bound concurrent tasks and you only have one CPU core, parallel execution can't happen. You can achieve parallelism by adding more CPU cores and relying on an efficient concurrent framework. But you should be aware that concurrency itself doesn't necessarily speed things up.

To keep things simple, we'll use a simulation of a long-running database query, presented in the following snippet:

```
iex(1)> run_query =
          fn query_def ->
            Process.sleep(2000)     #1
            "#{query_def} result"
          end
```

Here, the code sleeps for two seconds to simulate a long-running operation. When you call the `run_query` lambda, the shell is blocked until the lambda is done:

```
iex(2)> run_query.("query 1")

"query 1 result"     #1
```

Consequently, if you run five queries, it will take 10 seconds to get all the results:

```
iex(3)> Enum.map(
          1..5,
          fn index ->
            query_def = "query #{index}"
            run_query.(query_def)
          end
        )

["query 1 result", "query 2 result", "query 3 result"     #1
 "query 4 result", "query 5 result"]
```

Obviously, this is neither performant nor scalable. Assuming that the queries are already optimized, the only thing you can do to try to make things faster is to run the queries concurrently. This won't speed up individual queries, but the total time required to run all the queries should be reduced. In the BEAM world, to run something concurrently, you have to create a separate process.

5.2.1 Creating processes

To create a process, you can use the auto-imported `spawn/1` function:

```
spawn(fn ->
  expression_1     #1
  ...              #1
  expression_n     #1
end)
```

The function `spawn/1` creates (spawns) a new process. The provided zero-arity lambda will run concurrently, in the spawned process. After the lambda finishes, the spawned process is stopped. As soon as the new process is spawned, `spawn/1` returns, and the caller process can continue its execution.

You can try this to run the query concurrently:

```
iex(4)> spawn(fn ->
          query_result = run_query.("query 1")
          IO.puts(query_result)
        end)
#PID<0.48.0>        #1
```

```
query 1 result    #2
```

As you can see, the call to `spawn/1` returns immediately, and you can do something else in the shell while the query runs concurrently. Then, after two seconds, the result is printed to the screen.

The funny-looking `#PID<0.48.0>` that's returned by `spawn/1` is the identifier of the created process, often called a *pid*. This can be used to communicate with the process, as you'll see later in this chapter.

In the meantime, let's play some more with concurrent execution. First, you'll create a helper lambda that concurrently runs the query and prints the result:

```
iex(5)> async_query =
          fn query_def ->
            spawn(fn ->
              query_result = run_query.(query_def)
              IO.puts(query_result)
            end)
          end

iex(6)> async_query.("query 1")
#PID<0.52.0>

query 1 result    #1
```

This code demonstrates an important technique: passing data to the created process. Notice that `async_query` takes one argument and binds it to the `query_def` variable. This data is then passed to the newly created process via the closure mechanism. The inner lambda — the one that runs in a separate process — references the variable `query_def` from the outer scope. This results in cross-process data passing — the contents of `query_def` are passed from the main process to the newly created one. When it's passed to another process, the data is deep-copied, because two processes can't share any memory.

Note

In BEAM, everything runs in a process. This also holds for the interactive shell. All expressions you enter in `iex` are executed in a single shell-specific process. In this example, the main process is the shell process.

Now that you have the `async_query` lambda in place, you can try to run five queries concurrently:

```
iex(7)> Enum.each(1..5, &async_query.("query #{&1}"))
:ok        #1

query 1 result      #2
query 2 result      #2
query 3 result      #2
query 4 result      #2
query 5 result      #2
```

As expected, the call to `Enum.each/2` now returns immediately (in the first sequential version you had to wait 10 seconds for it to finish). Moreover, all the results are printed at practically the same time, two seconds later, which is a five-fold improvement over the sequential version. This happens because you run each computation concurrently.

For the same reason, the order of execution isn't deterministic. The output results might be printed in any order.

In contrast to the sequential version, the caller process doesn't get the result of the spawned processes. The processes run concurrently, each one printing the result to the screen. At the same time, the caller process runs independently and has no access to any data from the spawned processes. Remember, processes are completely independent and isolated.

Often, a simple "fire and forget" concurrent execution, where the caller process doesn't receive any notification from the spawned ones, will suffice. Sometimes, though, you'll want to return the result of the concurrent operation to the caller process. For this purpose, you can use the message-passing mechanism.

5.2.2 Message passing

In complex systems, you often need concurrent tasks to cooperate in some way. For example, you may have a main process that spawns multiple concurrent calculations, and then you want to handle all the results in the main process.

Being completely isolated, processes can't use shared data structures to exchange knowledge. Instead, processes communicate via messages, as illustrated in figure 5.2.

Figure 5.2. Inter-process communication via messages

Asynchronous sending of messages.
A message is stored in the mailbox of the
receiver, and the sender continues its work.

Process A

Message

Message

Process B

Time

```
receive do
  ...
end
```

Message sending
can work in both
directions.

Takes one message from the mailbox
and processes it

When process A wants process B to do something, it sends an asynchronous message to B. The content of the message is an Elixir term — anything you can store in a variable. Sending a message amounts to storing it into the receiver's mailbox. The caller then continues with its own execution, and the receiver can pull the message in at any time and process it in some way. Because processes can't share memory, a message is deep-copied when it's sent.

The process mailbox is a FIFO queue limited only by the available memory. The receiver consumes messages in the order received, and a message can be removed from the queue only if it's consumed.

To send a message to a process, you need to have access to its process identifier (pid). Recall from the previous section that the pid of the newly created process is the result of the `spawn/1` function. In addition, you can obtain the pid of the current process by calling the auto-imported `self/0` function.

Once you have a receiver's pid, you can send it messages using the `Kernel.send/2` function:

```
send(pid, {:an, :arbitrary, :term})
```

The consequence of `send` is that a message is placed in the mailbox of the receiver. The caller process then continues running subsequent expressions.

On the receiver side, to pull a message from the mailbox, you have to use the `receive` expression:

```
receive do
  pattern_1 -> do_something
  pattern_2 -> do_something_else
end
```

The `receive` expression works similarly to the `case` expression you saw in chapter 3. It tries to pull one message from the process mailbox, match it against any of the provided patterns, and run the corresponding code. You can easily test this by making the shell process send messages to itself:

```
iex(1)> send(self(), "a message")            #1

iex(2)> receive do                           #2
          message -> IO.inspect(message)  #2
        end                                  #2

"a message"
```

If you want to handle a specific message, you can rely on pattern matching:

```
iex(3)> send(self(), {:message, 1})

iex(4)> receive do
          {:message, id} ->      #1
            IO.puts("received message #{id}")
        end

received message 1
```

If there are no messages in the mailbox, `receive` waits indefinitely for a new message to arrive. The following call blocks the shell, and you need to manually terminate it:

```
iex(5)> receive do
          message -> IO.inspect(message)
        end      #1
```

The same thing happens if no message in the mailbox matches the provided patterns:

```
iex(1)> send(self(), {:message, 1})

iex(2)> receive do
          {_, _, _} ->              #1
            IO.puts("received")
        end
#2
```

If you don't want `receive` to block, you can specify the `after` clause, which is executed if a message isn't received in a given time frame (in milliseconds):

```
iex(1)> receive do
          message -> IO.inspect(message)
```

```
        after
           5000 -> IO.puts("message not received")
        end

message not received      #1
```

Receive algorithm

Recall from chapter 3 that an error is raised when you can't pattern-match the given term. The `receive` expression is an exception to this rule. If a message doesn't match any of the provided clauses, it's put back into the process mailbox, and the next message is processed.

The `receive` expression works as follows:

1. Take the first message from the mailbox.
2. Try to match it against any of the provided patterns, going from top to bottom.
3. If a pattern matches the message, run the corresponding code.
4. If no pattern matches, take the next message and start from step 2.
5. If there are no more messages in the queue, wait for a new one to arrive. When a new message arrives, start from step 2.
6. If the `after` clause is specified and no message is matched in the given amount of time, run the code from the `after` block.

As you already know, each Elixir expression returns a value, and `receive` is no different. The result of `receive` is the result of the last expression in the appropriate clause:

```
iex(1)> send(self(), {:message, 1})

iex(2)> receive_result =
          receive do
            {:message, x} ->
              x + 2                    #1
          end

iex(3)> IO.inspect(receive_result)
3
```

To summarize, `receive` tries to find the first (oldest) message in the process

mailbox that can be matched against any of the provided patterns. If such a message is found, the corresponding code is executed. Otherwise, `receive` waits for such a message for a specified amount of time, or indefinitely if the `after` clause isn't provided.

Synchronous sending

The basic message-passing mechanism is the asynchronous "fire and forget" kind. A process sends a message and then continues to run, oblivious to what happens in the receiver. Sometimes a caller needs some kind of response from the receiver. There's no special support for doing this. Instead, you must program both parties to cooperate using the basic asynchronous messaging facility.

The caller must include its own pid in the message contents and then wait for a response from the receiver:

```
send(pid, {self(), some_message})   #1

receive do
  {:response, response} -> ...       #2
end
```

The receiver uses the included pid to send the response to the caller:

```
receive do
  {caller_pid, message} ->
    response = ...
    send(caller_pid, {:response, response})   #1
end
```

You'll see this in action a bit later, when we discuss server processes.

Collecting query results

Let's try message-passing with the concurrent queries developed in the previous section. In the previous attempt, you run queries in separate processes and print them to the screen from those processes. Let's recall how this works:

```
iex(1)> run_query =
          fn query_def ->
            Process.sleep(2000)
            "#{query_def} result"
          end

iex(2)> async_query =
          fn query_def ->
            spawn(fn ->
              query_result = run_query.(query_def)
              IO.puts(query_result)
            end)
          end
```

Now, instead of printing to the screen, let's collect all the results in the main process. First, you need to make the lambda send the query result to the caller process:

```
iex(3)> async_query =
          fn query_def ->
            caller = self()                                    #1

            spawn(fn ->
              query_result = run_query.(query_def)
              send(caller, {:query_result, query_result})    #2
            end)
          end
```

In this code, you first store the pid of the calling process to a distinct `caller` variable. This is necessary so the worker process (the one doing the calculation) can know the pid of the process that should receive the response.

Keep in mind that the result of `self/0` depends on the calling process. If you didn't store the result to the `caller` variable, and you tried to `send(self(), …)` from the inner lambda, it would have no effect. The spawned process would send the message to itself, because calling `self/0` returns the pid of the process that invoked the function.

The worker process can now use the `caller` variable to return the result of the calculation. The message is in the custom format `{:query_result, result}`. This makes it possible to distinguish between your messages and any others that might be sent to the caller process.

Now you can start your queries:

```
iex(4)> Enum.each(1..5, &async_query.("query #{&1}"))
```

This runs all the queries concurrently, and the result is stored in the mailbox of the caller process. In this case, this is the shell (`iex`) process.

Notice that the caller process is neither blocked nor interrupted while receiving messages. Sending a message doesn't disturb the receiving process in any way. If the process is performing computations, it continues to do so. The only thing affected is the content of the receiving process's mailbox. Messages remain in the mailbox until they're consumed or the process terminates.

Let's get the results. First you make a lambda that pulls one message from the mailbox and extracts the query result from it:

```
iex(5)> get_result =
          fn ->
            receive do
              {:query_result, result} -> result
            end
          end
```

Now you can pull all the messages from the mailbox into a single list:

```
iex(6)> results = Enum.map(1..5, fn _ -> get_result.() end)
["query 3 result", "query 2 result", "query 1 result",
 "query 5 result", "query 4 result"]
```

Notice the use of `Enum.map/2`, which maps anything enumerable to a list of the same length. In this case, you create a range of size 5 and then map each element to the result of the `get_result` lambda. This works because you know there are five messages waiting for you. Otherwise, the loop would get stuck waiting for new messages to arrive.

It's also worth repeating that results arrive in a nondeterministic order. Because all computations run concurrently, it's not certain in which order they'll finish.

This is a simple implementation of a parallel map technique that can be used

to process a larger amount of work in parallel and then collect the results into a list. This idea can be expressed with a pipeline:

```
iex(7)> 1..5
        |> Enum.map(&async_query.("query #{&1}"))  #1
        |> Enum.map(fn _ -> get_result.() end)     #2
```

5.3 Stateful server processes

Spawning processes to perform one-off tasks isn't the only use case for concurrency. In Elixir, it's common to create long-running processes that can serve various requests, sent in the form of messages. In addition, such processes may maintain some internal state — an arbitrary piece of data that may change over time.

We call such processes stateful server processes, and they are an important concept in Elixir/Erlang systems, so we'll spend some time exploring them.

5.3.1 Server processes

A *server process* is an informal name for a process that runs for a long time (or forever) and can handle various requests (messages). To make a process run forever, you have to use endless tail recursion. You may remember from chapter 3 that tail calls receive special treatment. If the last thing a function does is call another function (or itself), a simple jump takes place instead of a stack push. Consequently, a function that always calls itself will run forever, without causing a stack overflow or consuming additional memory.

This can be used to implement a server process. You need to run the endless loop and wait for a message in each step of the loop. When the message is received, you handle it and then continue the loop. Let's try this by creating a server process that can run a query on demand.

The basic sketch of a long-running server process is provided in the following listing.

Listing 5.1. Long-running server process (database_server.ex)

```
defmodule DatabaseServer do
  def start do
    spawn(&loop/0)     #1
  end

  defp loop do
    receive do       #2
      ...            #2
    end              #2

    loop()      #3
  end

  ...
end
```

start/0 is the so-called *interface function* that's used by clients to start the server process. When start/0 is called, it spawns the process that runs the loop/0 function. This function powers the infinite loop of the process. The function waits for a message, handles it, and then calls itself, thus ensuring that the process never stops.

Such implementation is what makes this process a server. Instead of actively running some computation, the process is mostly idle, awaiting for the message (request) to arrive. It's worth noting that this loop isn't CPU-intensive. Waiting for a message puts the process in a suspended state and doesn't waste CPU cycles.

Notice that functions in this module run in different processes. The function start/0 is called by clients and runs in a client process. The private function loop/0 runs in the server process. It's perfectly normal to have different functions from the same module running in different processes — there's no special relationship between modules and processes. A module is just a collection of functions, and these functions can be invoked in any process.

When implementing a server process, it usually makes sense to put all of its code in a single module. The functions of this module generally fall into two categories: interface and implementation. *Interface functions* are public and are executed in the caller process. They hide the details of process creation and the communication protocol. *Implementation functions* are usually private and run in the server process.

Note

As was the case with classical loops, you typically won't need to code the recursion loop yourself. A standard abstraction called `GenServer` (generic server process) is provided, which simplifies the development of stateful server processes. The abstraction still relies on recursion, but this recursion is implemented in `GenServer`. You'll learn about this abstraction in chapter 6.

Let's look at the full implementation of the `loop/0` function.

Listing 5.2. Database server loop (database_server.ex)

```
defmodule DatabaseServer do
  ...

  defp loop do
    receive do       #1
      {:run_query, caller, query_def} ->
        query_result = run_query(query_def)              #2
        send(caller, {:query_result, query_result})      #2
    end

    loop()
  end

  defp run_query(query_def) do       #3
    Process.sleep(2000)
    "#{query_def} result"
  end

  ...
end
```

This code reveals the communication protocol between the caller process and the database server. The caller sends a message in the format `{:run_query, caller, query_def}`. The server process handles such a message by executing the query and sending the query result back to the caller process.

Usually you want to hide these communication details from your clients. Clients shouldn't depend on knowing the exact structure of messages that must be sent or received. To hide this, it's best to provide a dedicated

interface function. Let's introduce a function called `run_async/2` that will be used by clients to request the operation — in this case, a query execution — from the server. This function makes the clients unaware of message-passing details — they just call `run_async/2` and get the result. The implementation is given in the following listing.

Listing 5.3. Implementation of run_async/2 (database_server.ex)

```
defmodule DatabaseServer do
  ...

  def run_async(server_pid, query_def) do
    send(server_pid, {:run_query, self(), query_def})
  end

  ...
end
```

The `run_async/2` function receives the pid of the database server and a query you want to execute. It sends the appropriate message to the server and then does nothing else. Calling `run_async/2` from the client requests that the server process run the query while the caller goes about its business.

Once the query is executed, the server sends a message to the caller process. To get this result, you need to add another interface function. It's called `get_result/0`.

Listing 5.4. Implementation of get_result/0 (database_server.ex)

```
defmodule DatabaseServer do
  ...

  def get_result do
    receive do
      {:query_result, result} -> result
    after
      5000 -> {:error, :timeout}
    end
  end

  ...
end
```

`get_result/0` is called when the client wants to get the query result. Here, you use `receive` to get the message. The `after` clause ensures that you give up after some time passes — for example, if something goes wrong during the query execution and a response never comes back.

The database server is now complete. Let's see how to use it:

```
iex(1)> server_pid = DatabaseServer.start()

iex(2)> DatabaseServer.run_async(server_pid, "query 1")
iex(3)> DatabaseServer.get_result()
"query 1 result"

iex(4)> DatabaseServer.run_async(server_pid, "query 2")
iex(5)> DatabaseServer.get_result()
"query 2 result"
```

Notice how you execute multiple queries in the same process. First you run query 1 and then query 2. This proves that the server process continues running after a message is received.

Because communication details are wrapped in functions, the client isn't aware of them. Instead, it communicates with the process with plain functions. Here, the server pid plays an important role. You receive the pid by calling `DatabaseServer.start/0`, and then you use it to issue requests to the server.

Of course, the request is handled asynchronously in the server process. After calling `DatabaseServer.run_async/2`, you can do whatever you want in the client (`iex`) process and collect the result when you need it.

Server processes are sequential

It's important to realize that a server process is internally sequential. It runs a loop that processes one message at a time. Thus, if you issue five asynchronous query requests to a single server process, they will be handled one by one, and the result of the last query will come after 10 seconds.

This is a good thing because it helps you reason about the system. A server

process can be considered a synchronization point. If multiple actions need to happen synchronously, in a serialized manner, you can introduce a single process and forward all requests to that process, which handles the requests sequentially.

Of course, in this case, a sequential property is a problem. You want to run multiple queries concurrently to get the result as quickly as possible. What can you do about it?

Assuming that the queries can be run independently, you can start a pool of server processes, and then for each query somehow choose one of the processes from the pool and have that process run the query. If the pool is large enough and you divide the work uniformly across each worker in the pool, you'll parallelize the total work as much as possible.

Here's a basic sketch of how can this be done. First, create a pool of database-server processes:

```
iex(1)> pool = Enum.map(1..100, fn _ -> DatabaseServer.start() en
```

Here you create 100 database-server processes and store their pids in a list. You may think that 100 processes is a lot, but recall that processes are lightweight. They take up a small amount of memory (~2 KB) and are created very quickly (in a few microseconds). Furthermore, because all of these processes wait for a message, they're effectively idle and don't waste CPU time.

Next, when you run a query, you need to decide which process will execute the query. The simplest way is to use the `:rand.uniform/1` function, which takes a positive integer `n` and returns a random number in the range `1..n`. Taking advantage of this, the following expression distributes five queries over a pool of processes:

```
iex(2)> Enum.each(
          1..5,
          fn query_def ->
            server_pid = Enum.at(pool, :rand.uniform(100) - 1)
            DatabaseServer.run_async(server_pid, query_def)
          end
        )
```

Note that this isn't efficient. You're using `Enum.at/2` to get the pid at a random position. Because you use a list to keep the processes, and a random lookup is an O(N) operation, selecting a random worker isn't very performant. You could do better if you used a map with process indexes as keys and pids as values; and there are other alternatives, such as using a round-robin approach. But for now, let's stick with this simple implementation.

Once you've queued the queries to the workers, you need to collect the responses. This is now straightforward, as illustrated in the following snippet:

```
iex(3)> Enum.map(1..5, fn _ -> DatabaseServer.get_result() end)
["5 result", "3 result", "1 result", "4 result", "2 result"]
```

Thanks to this, you get all the results much faster, because queries are again executed concurrently.

5.3.2 Keeping a process state

Server processes open the possibility of keeping some kind of process-specific state. For example, when you talk to a database, you need to maintain a connection handle.

To keep state in the process, you can extend the `loop` function with additional argument(s). Here's a basic sketch:

```
def start do
  spawn(fn ->
    initial_state = ...      #1
    loop(initial_state)      #2
  end)
end

defp loop(state) do
  ...
  loop(state)               #3
end
```

Let's use this technique to extend the database server with a connection. In this example, you'll use a random number as a simulation of the connection handle. First you need to initialize the connection while the process starts, as

demonstrated in the following listing.

Listing 5.5. Initializing the process state (stateful_database_server.ex)

```
defmodule DatabaseServer do
  ...

  def start do
    spawn(fn ->
      connection = :rand.uniform(1000)
      loop(connection)
    end)
  end

  ...
end
```

Here, you open the connection and then pass the corresponding handle to the `loop` function. In real life, instead of generating a random number, you'd use a database client library (such as ODBC) to open the connection.

Next you need to modify the `loop` function.

Listing 5.6. Using the connection while querying (stateful_database_server.ex)

```
defmodule DatabaseServer do
  ...

  defp loop(connection) do
    receive do
      {:run_query, from_pid, query_def} ->
        query_result = run_query(connection, query_def)     #1
        send(from_pid, {:query_result, query_result})
    end

    loop(connection)     #2
  end

  defp run_query(connection, query_def) do
    Process.sleep(2000)
    "Connection #{connection}: #{query_def} result"
  end

  ...
```

```
end
```

The `loop` function takes the state (connection) as the first argument. Every time the loop is resumed, the function passes on the state to itself, so it is available in the next step.

You have to additionally extend the `run_query` function to use the connection while querying the database. The connection handle (in this case, a number) is included in the query result.

With this, your stateful database server is complete. Notice that you didn't change the interface of its public functions, so the usage remains the same as it was. Let's see how it works:

```
iex(1)> server_pid = DatabaseServer.start()

iex(2)> DatabaseServer.run_async(server_pid, "query 1")
iex(3)> DatabaseServer.get_result()
"Connection 753: query 1 result"

iex(4)> DatabaseServer.run_async(server_pid, "query 2")
iex(5)> DatabaseServer.get_result()
"Connection 753: query 2 result"
```

The results for different queries are executed using the same connection handle, which is kept internally in the process loop and is completely invisible to other processes.

5.3.3 Mutable state

So far, you've seen how to keep constant process-specific state. It doesn't take much to make this state mutable. Here's the basic idea:

```
defp loop(state) do
  new_state =       #1
    receive do
      msg1 ->
        ...

      msg2 ->
        ...
    end
```

```
  loop(new_state)      #2
end
```

This is a standard stateful server technique in Elixir. The process determines the new state while handling the message. Then the loop function calls itself with the new state, which effectively changes the state. The next received message operates on the new state.

From the outside, stateful processes are mutable. By sending messages to a process, a caller can affect its state and the outcome of subsequent requests handled in that server. In that sense, sending a message is an operation with possible side effects. Still, the server relies on immutable data structures. A state can be any valid Elixir variable ranging from simple numbers to complex data abstractions, such as TodoList (which you built in chapter 4).

Let's see this in action. You'll start with a simple example: a stateful calculator process that keeps a number as its state. Initially the state of the process is 0, and you can manipulate it by issuing requests such as add, sub, mul, and div. You can also retrieve the process state with the value request.

Here's how you use the server:

```
iex(1)> calculator_pid = Calculator.start()        #1

iex(2)> Calculator.value(calculator_pid)          #2
0                                                  #2

iex(3)> Calculator.add(calculator_pid, 10)        #3
iex(4)> Calculator.sub(calculator_pid, 5)         #3
iex(5)> Calculator.mul(calculator_pid, 3)         #3
iex(6)> Calculator.div(calculator_pid, 5)         #3

iex(7)> Calculator.value(calculator_pid)        #4
3.0
```

In this code, you start the process and check its initial state. Then you issue some modifier requests and verify the result of the operations (((0 + 10) - 5) * 3) / 5, which is 3.0.

Now it's time to implement this. First, let's look at the server's inner loop.

Listing 5.7. Concurrent stateful calculator (calculator.ex)

```
defmodule Calculator do
  ...

  defp loop(current_value) do
    new_value =
      receive do
        {:value, caller} ->                      #1
          send(caller, {:response, current_value})   #1
          current_value                          #1

        {:add, value} -> current_value + value    #2
        {:sub, value} -> current_value - value    #2
        {:mul, value} -> current_value * value    #2
        {:div, value} -> current_value / value    #2

        invalid_request ->
          IO.puts("invalid request #{inspect invalid_request}")
          current_value
      end

    loop(new_value)
  end

  ...
end
```

The loop handles various messages. The `:value` message is used to retrieve the server's state. Because you need to send the response back, the caller must include its pid in the message. Notice that the last expression of this block returns `current_value`. This is needed because the result of `receive` is stored in `new_value`, which is then used as the server's new state. By returning `current_value`, you specify that the `:value` request doesn't change the process state.

The arithmetic operations compute the new state based on the current value and the argument received in the message. Unlike a `:value` message handler, arithmetic operation handlers don't send responses back to the caller. This makes it possible to run these operations asynchronously, as you'll see soon when you implement interface functions.

The final `receive` clause matches all the other messages. These are the ones

you're not supporting, so you log them to the screen and return `current_value`, leaving the state unchanged.

Next, you have to implement the interface functions that will be used by clients. This is done in the next listing.

Listing 5.8. Calculator interface functions (calculator.ex)

```
defmodule Calculator do
  def start do                     #1
    spawn(fn -> loop(0) end)       #1
  end                              #1

  def value(server_pid) do                       #2
    send(server_pid, {:value, self()})           #2
                                                 #2
    receive do                                   #2
      {:response, value} ->                      #2
        value                                    #2
      end                                        #2
    end                                          #2

  def add(server_pid, value), do: send(server_pid, {:add, value})
  def sub(server_pid, value), do: send(server_pid, {:sub, value})
  def mul(server_pid, value), do: send(server_pid, {:mul, value})
  def div(server_pid, value), do: send(server_pid, {:div, value})

  ...
end
```

The interface functions follow the protocol specified in the `loop/1` function. The `:value` request is an example of the synchronous message passing mentioned in section 5.2.2. The caller sends a message and then it awaits for the response. The caller is blocked until the response comes back, which makes the request-handling synchronous.

The arithmetic operations run asynchronously. There's no response message, so the caller doesn't have to wait for anything. Therefore, a caller can issue a number of these requests and continue doing its own work while the operations run concurrently in the server process. Keep in mind that the server handles messages in the order received, so requests are handled in the proper order.

Why make the arithmetic operations asynchronous? Because you don't care when they're executed. Until you request the server's state (via the `value/1` function), you don't want the client to block. This makes the client more efficient, because it doesn't block while the server is doing a computation.

Refactoring the loop

As you introduce multiple requests to your server, the `loop` function becomes more complex. If you have to handle many requests, it will become bloated, turning into a huge `switch/case`-like expression.

You can refactor this by relying on pattern matching and moving the message handling to a separate multiclause function. This keeps the code of the `loop` function very simple:

```
defp loop(current_value) do
  new_value =
    receive do
      message -> process_message(current_value, message)
    end

  loop(new_value)
end
```

Looking at this code, you can see the general workflow of the server. A message is first received and then processed. Message processing generally amounts to computing the new state based on the current state and the message received. Finally, you loop with this new state, effectively setting it in place of the old one.

`process_message/2` is a simple multiclause function that receives the current state and the message. Its task is to perform message-specific code and return the new state:

```
defp process_message(current_value, {:value, caller}) do
  send(caller, {:response, current_value})
  current_value
end

defp process_message(current_value, {:add, value}) do
  current_value + value
```

```
end

...
```

This code is a simple reorganization of the server process loop. It allows you to keep the loop code compact and to move the message-handling details to a separate multiclause function, with each clause handling a specific message.

5.3.4 Complex states

State is usually much more complex than a simple number. But the technique remains the same — you keep the mutable state using the private `loop` function. As the state becomes more complex, the code of the server process can become increasingly complicated. It's worth extracting the state manipulation to a separate module and keeping the server process focused only on passing messages and keeping the state.

Let's look at this technique using the `TodoList` abstraction developed in chapter 4. First, let's recall the basic usage of the structure:

```
iex(1)> todo_list =
          TodoList.new() |>
          TodoList.add_entry(%{date: ~D[2023-12-19], title: "Dent
          TodoList.add_entry(%{date: ~D[2023-12-20], title: "Shop
          TodoList.add_entry(%{date: ~D[2023-12-19], title: "Movi

iex(2)> TodoList.entries(todo_list, ~D[2023-12-19])
[
  %{date: ~D[2023-12-19], id: 1, title: "Dentist"},
  %{date: ~D[2023-12-19], id: 3, title: "Movies"}
]
```

As you may recall, a `TodoList` is a pure functional abstraction. To keep the structure alive, you must constantly hold on to the result of the last operation performed on the structure.

In this example, you'll build a `TodoServer` module that keeps this abstraction in the private state. Let's see how the server is used:

```
iex(1)> todo_server = TodoServer.start()
```

```
iex(2)> TodoServer.add_entry(
          todo_server,
          %{date: ~D[2023-12-19], title: "Dentist"}
        )

iex(3)> TodoServer.add_entry(
          todo_server,
          %{date: ~D[2023-12-20], title: "Shopping"}
        )

iex(4)> TodoServer.add_entry(
          todo_server,
          %{date: ~D[2023-12-19], title: "Movies"}
        )

iex(5)> TodoServer.entries(todo_server, ~D[2023-12-19])
[
  %{date: ~D[2023-12-19], id: 3, title: "Movies"},
  %{date: ~D[2023-12-19], id: 1, title: "Dentist"}
]
```

You start the server and then interact with it via the `TodoServer` API. In contrast to the pure functional approach, you don't need to take the result of a modification and feed it as an argument to the next operation. Instead, you constantly use the same `todo_server` variable to work with the to-do list.

Let's start implementing this server. First you need to place all the modules in a single file, as shown in the following listing.

Listing 5.9. `TodoServer` modules (todo_server.ex)

```
defmodule TodoServer do
  ...
end

defmodule TodoList do
  ...
end
```

Putting both modules in the same file ensures that you have everything available when you load the file while starting the `iex` shell. In more complicated systems, you'd use a proper `mix` project, as will be explained in chapter 7, but for now, this is sufficient.

The `TodoList` implementation is the same as in chapter 4. It has all the features you need to use it in a server process.

Now set up the basic structure of the to-do server process.

Listing 5.10. `TodoServer` basic structure (todo_server.ex)

```
defmodule TodoServer do
  def start do
    spawn(fn -> loop(TodoList.new()) end)     #1
  end

  defp loop(todo_list) do
    new_todo_list =
      receive do
        message -> process_message(todo_list, message)
      end

    loop(new_todo_list)
  end

  ...
end
```

There's nothing new here. You start the loop using a new instance of the `TodoList` abstraction as the initial state. In the loop, you receive messages and apply them to the state by calling the `process_message/2` function, which returns the new state. Finally, you loop with the new state.

For each request you want to support, you have to add a dedicated clause in the `process_message/2` function. Additionally, a corresponding interface function must be introduced. You'll begin by supporting the `add_entry` request.

Listing 5.11. The `add_entry` request (todo_server.ex)

```
defmodule TodoServer do
  ...

  def add_entry(todo_server, new_entry) do     #1
    send(todo_server, {:add_entry, new_entry})  #1
  end                                             #1
```

```
    ...

    defp process_message(todo_list, {:add_entry, new_entry}) do
      TodoList.add_entry(todo_list, new_entry)
    end

    ...
end
```

The interface function sends the new entry data to the server. This message will be handled in the corresponding clause of `process_message/2`. Here, you delegate to the `TodoList.add_entry/2` function and return the modified `TodoList` instance. This returned instance is used as the new server's state.

Using a similar approach, you can implement the `entries` request, keeping in mind that you need to wait for the response message. The code is shown in the next listing.

Listing 5.12. The `entries` request (todo_server.ex)

```
defmodule TodoServer do
  ...

  def entries(todo_server, date) do
    send(todo_server, {:entries, self(), date})

    receive do
      {:todo_entries, entries} -> entries
    after
      5000 -> {:error, :timeout}
    end
  end

  ...

  defp process_message(todo_list, {:entries, caller, date}) do
    send(caller, {:todo_entries, TodoList.entries(todo_list, date
    todo_list                    #2
  end

  ...
end
```

This is a synthesis of techniques you've seen previously. You send a message

and wait for the response. In the corresponding `process_message/2` clause, you delegate to `TodoList`, and then you send the response and return the unchanged to-do list. This is needed because `loop/2` takes the result of `process_message/2` as the new state.

In a similar way, you can add support for other to-do list requests such as `update_entry` and `delete_entry`. The implementation of these requests is left for you as an exercise.

Concurrent vs. functional approach

A process that maintains mutable state can be regarded as a kind of mutable data structure. But you shouldn't abuse processes to avoid using the functional approach of transforming immutable data.

The data should be modeled using pure functional abstractions, just as you did with `TodoList`. A pure functional data structure provides many benefits, such as integrity, atomicity, reusability, and testability.

A stateful process serves as a container of such data structure. The process keeps the state alive, and allows other processes in the system to interact with this data via the exposed API.

With such separation of responsibilities, building a highly concurrent system becomes straightforward. For example, if you're implementing a web server that manages multiple to-do lists, you could run one server process for each to-do list. While handling an HTTP request, you can find the corresponding to-do server and have it perform the requested operation. Each to-do list manipulation runs concurrently, thus making your server scalable and more performant. Moreover, there are no synchronization problems, because each to-do list is managed in a dedicated process. Recall that a single process is always sequential, so multiple competing requests that manipulate the same to-do list are serialized and handled sequentially in the corresponding process. Don't worry if this seems vague — you'll see it in action in chapter 7.

5.3.5 Registered processes

In order for a process to cooperate with other processes, it must know their whereabouts. In BEAM, a process is identified by the corresponding pid. To make process A send messages to process B, you have to bring the pid of process B to process A.

Sometimes it can be cumbersome to keep and pass pids. If you know there will always be only one instance of some type of server, you can give the process a *local name* and use that name to send messages to the process. The name is called *local* because it has meaning only in the currently running BEAM instance. This distinction becomes important when you start connecting multiple BEAM instances to a distributed system, as you'll see in chapter 12.

Registering a process can be done with `Process.register(pid, name)`, where a name must be an atom. Here's a quick illustration:

```
iex(1)> Process.register(self(), :some_name)      #1

iex(2)> send(:some_name, :msg)         #2

iex(3)> receive do             #3
           msg -> IO.puts("received #{msg}")
         end

received msg
```

The following constraints apply to registered names:

- The name can only be an atom.
- A single process can have only one name.
- Two processes can't have the same name.

If these constraints aren't met, an error is raised.

For practice, try to change the to-do server to run as a registered process. The interface of the server will then be simplified, because you don't need to keep and pass the server's pid.

Here's an example of how such a server can be used:

```
iex(1)> TodoServer.start()

iex(2)> TodoServer.add_entry(%{date: ~D[2023-12-19], title: "Dent
iex(3)> TodoServer.add_entry(%{date: ~D[2023-12-20], title: "Shop
iex(4)> TodoServer.add_entry(%{date: ~D[2023-12-19], title: "Movi

iex(5)> TodoServer.entries(~D[2023-12-19])
[%{date: ~D[2023-12-19], id: 3, title: "Movies"},
 %{date: ~D[2023-12-19], id: 1, title: "Dentist"}]
```

To make this work, you have to register a server process under a name (such as `:todo_server`). Then you change all the interface functions to use the registered name when sending a message to the process. If you get stuck, the solution is provided in the registered_todo_server.ex file.

Using the registered server is much simpler because you don't have to store the server's pid and pass it to the interface functions. Instead, the interface functions internally use the registered name to send messages to the process.

Local registration plays an important role in process discovery. Registered names provide a way of communicating with a process without knowing its pid. This becomes increasingly important when you start dealing with restarting processes (as you'll see in chapters 8 and 9) and distributed systems (discussed in chapter 12).

This concludes our initial exploration of stateful processes. They play an important role in Elixir-based systems, and you'll continue using them throughout the book. Now, though, we'll look at some important runtime properties of BEAM processes.

5.4 Runtime considerations

You've learned a great deal about how to work with processes. Now its time to discuss some important runtime properties of BEAM concurrency.

5.4.1 A process is sequential

It has already been mentioned, but it's very important, so I'll stress it again: a single process is a sequential program — it runs expressions in a sequence,

one by one.

Multiple processes run concurrently, and so they may run in parallel to each other. But if many processes send messages to a single process, that single process may become a bottleneck which significantly affects overall throughput of the system.

Let's look at an example. The code in the following listing implements a slow echo server.

Listing 5.13. Demonstration of a process bottleneck (process_bottleneck.ex)

```
defmodule Server do
  def start do
    spawn(fn -> loop end)
  end

  def send_msg(server, message) do
    send(server, {self(), message})

    receive do
      {:response, response} -> response
    end
  end

  defp loop do
    receive do
      {caller, msg} ->
        Process.sleep(1000)     #1
        send(caller, {:response, msg})   #2
    end

    loop()
  end
end
```

Upon receiving a message, the server sends the message back to the caller. Before that, it sleeps for a second to simulate a long-running request.

To test its behaviour in a concurrent setting, start the server and fire up five concurrent clients:

```
iex(1)> server = Server.start()
```

```
iex(2)> Enum.each(
          1..5,
          fn i ->
            spawn(fn ->        #1
              IO.puts("Sending msg ##{i}")
              response = Server.send_msg(server, i)    #2
              IO.puts("Response: #{response}")
            end)
          end
        )
```

As soon as you start this, you'll see the following lines printed:

```
Sending msg #1
Sending msg #2
Sending msg #3
Sending msg #4
Sending msg #5
```

So far, so good. Five processes have been started and are running concurrently. But now the problems begin — the responses come back slowly, one by one, a second apart:

```
Response: 1     #1
Response: 2     #2
Response: 3     #3
Response: 4     #4
Response: 5     #5
```

What happened? The echo server can handle only one message per second. Because all other processes depend on the echo server, they're constrained by its throughput.

What can you do about this? Once you identify the bottleneck, you should try to optimize the process internally. Generally, a server process has a simple flow. It receives and handles messages one by one. So the goal is to make the server handle messages at least as fast as they arrive. In this example, server optimization amounts to removing the `Process.sleep/1` call.

If you can't make message handling fast enough, you can try to split the server into multiple processes, effectively parallelizing the original work and hoping that doing so will boost performance on a multicore system. This

should be your last resort, though. Parallelization isn't a remedy for a poorly structured algorithm.

5.4.2 Unlimited process mailboxes

Theoretically, a process mailbox has an unlimited size. In practice, the mailbox size is limited by available memory. Thus, if a process constantly falls behind, meaning messages arrive faster than the process can handle them, the mailbox will constantly grow and increasingly consume memory. Ultimately, a single slow process may cause an entire system to crash by consuming all the available memory.

A more subtle version of the same problem occurs if a process doesn't handle some messages at all. Consider the following server loop:

```
defp loop
  receive do
    {:message, msg} -> do_something(msg)
  end

  loop()
end
```

A server powered by this loop handles only messages that are in the form `{:message, something}`. All other messages remain in the process mailbox forever, taking up memory space for no reason.

Large mailbox contents can significantly affect performance. It puts the extra pressure on the garbage collector, and it can lead to slow pattern matches in `receive`.

To avoid this, you should introduce a *match-all* receive clause that deals with unexpected kinds of messages. Typically, you'll emit a warning that a process has received the unknown message, and do nothing else about it:

```
defp loop
  receive
    {:message, msg} -> do_something(msg)
    other -> warn_about_unknown_message(other)    #1
  end
```

```
  loop()
end
```

Since the process handles every kind of message, the uncontrolled growth of its mailbox is less likely to happen.

It's also worth noting that BEAM gives you tools for analyzing processes at runtime. In particular, you can query each process for its mailbox size and thus detect the ones for which the mailbox-queue buildup occurs. We'll discuss this feature in chapter 13.

5.4.3 Shared-nothing concurrency

As already mentioned, processes share no memory. Thus, sending a message to another process results in a deep copy of the message contents:

```
send(target_pid, data)   #1
```

Less obviously, closing on a variable from a `spawn` also results in deep-copying the closed variable:

```
data = ...

spawn(fn ->
  ...
  some_fun(data)     #1
  ...
end)
```

This is something you should be aware of when moving code into a separate process. Deep-copying is an in-memory operation, so it should be reasonably fast, and occasionally sending a big message shouldn't present a problem. But having many processes frequently send big messages may affect system performance. The notion of small versus big is subjective. Simple data, such as a number, an atom, or a tuple with few elements, is obviously small. A list of a million complex structures is big. The border lies somewhere in between and depends on your specific case.

There are a couple of special cases where the data is copied by reference. This happens with binaries (including strings) that are larger than 64 bytes,

hardcoded constants (also known as literals), and terms created via the `:persistent_term` API (www.erlang.org/doc/man/persistent_term.html).

Shared-nothing concurrency ensures complete isolation between processes: one process can't compromise the internal state of another. This promotes the integrity and fault-tolerance of the system.

In addition, because processes share no memory, garbage collection can take place on a process level. Each process gets an initial small chunk of heap memory (~2 KB on 64-bit BEAM). When more memory is needed, garbage collection for that process takes place. As a result, garbage collection is concurrent and distributed. Instead of one large "stop the entire system" collection, you end up running many smaller, typically faster collections. This prevents unwanted long complete blockages and keeps the entire system more responsive.

5.4.4 Scheduler inner workings

In general, you can assume that there are *n* schedulers that run *m* processes, with *m* most often being significantly larger than *n*. This is called *m:n* threading, and it reflects the fact that you run a large number of logical microthreads using a smaller number of OS threads, as illustrated in figure 5.3.

Figure 5.3. *m:n* threading: a small number of schedulers running a large number of BEAM processes
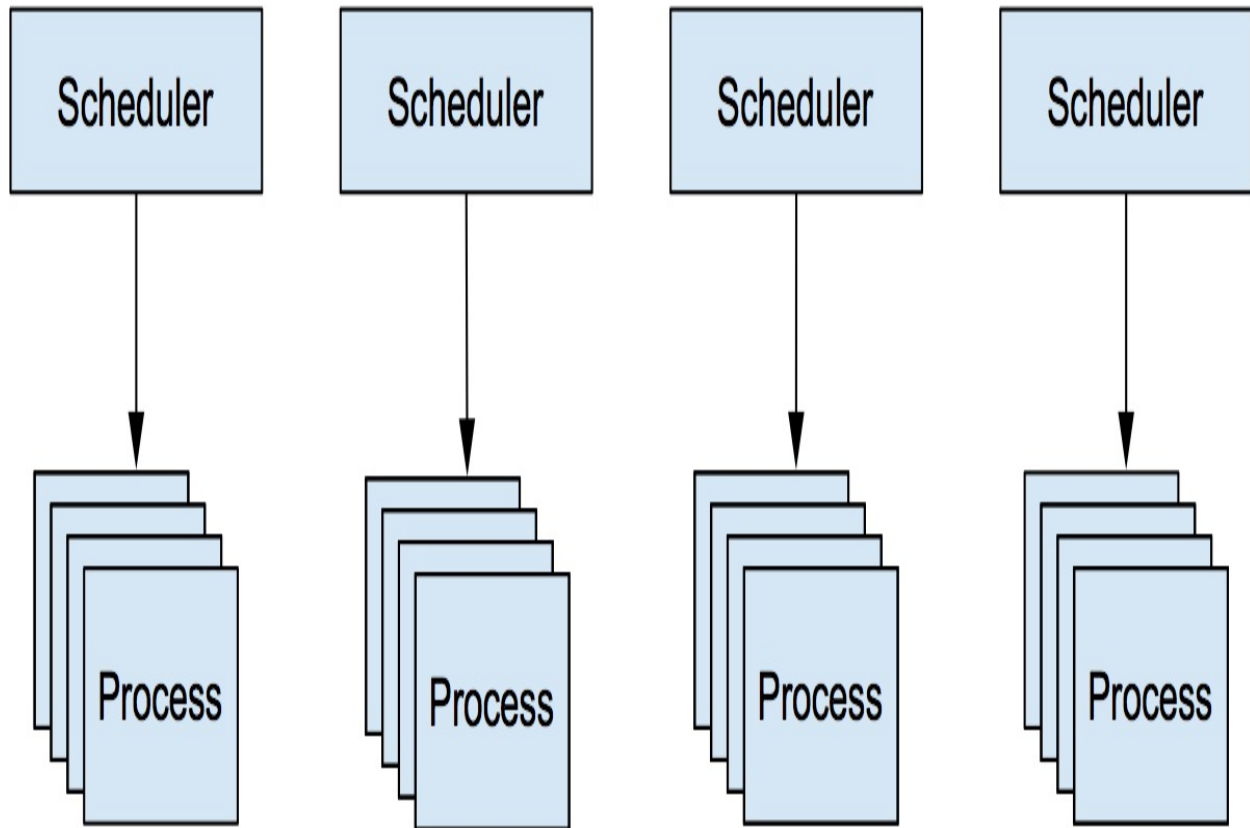
Each BEAM scheduler is in reality an OS thread that manages the execution of BEAM processes. By default, BEAM uses only as many schedulers as there are logical processors available. You can change these settings via various Erlang emulator flags.

To provide Erlang flags, you can use the following syntax:

```
$ iex --erl "put Erlang emulator flags here"
```

A list of all Erlang flags can be found at erlang.org/doc/man/erl.html.

For example, to use only one scheduler thread, you can provide the `+S 1` Erlang flag:

```
$ iex --erl "+S 1"
Erlang/OTP 25 [erts-13.0] [source] [64-bit] [smp:1:1] [ds:1:1:10]
```

Notice the `smp:1:1` part in the output. This informs us that only one scheduler thread is used. You can also check the number of schedulers programmatically:

```
iex(1)> System.schedulers()
1
```

If you're running other external services on the system, you could consider reducing the number of BEAM scheduler threads. Doing this will leave more computational resources for non-BEAM services.

Internally, each scheduler picks one process, runs it for a while, and then pick another process. While in the scheduler, the process gets a small execution window of approximately 2,000 function calls, after which it's preempted. It's also worth mentioning that in some cases a longer CPU bound work or a larger garbage collection might be performed on another thread (called dirty scheduler).

If a process is doing a network I/O or waiting for a message, it yields the execution to the scheduler. The same thing happens when `Process.sleep` is invoked. As a result, you don't have to care about the nature of the work performed in a process. If you want to separate the execution of one function from the rest of the code, you just need to run that function in a separate process, regardless of whether the work is CPU- or IO-bound.

As a consequence of all of this, the context switching is performed frequently. Typically, a single process is in the scheduler for less than one millisecond. This promotes the responsiveness of BEAM-powered systems. If one process performs a long CPU-bound operation, such as computing the value of pi to a billion decimals, it won't block the entire scheduler, and other processes shouldn't be affected.

This can easily be proven. Start an iex session with just one scheduler thread:

```
$ iex --erl "+S 1"
```

Spawn a process which runs an infinite CPU-bound loop:

```
iex(1)> spawn(fn ->
          Stream.repeatedly(fn -> :rand.uniform() end)
          |> Stream.run()
        end)
```

This code uses `Stream.repeatedly/1` to create a lazy infinite stream of

random numbers. The stream is executed using `Stream.run/1` function, which will effectively run an infinite CPU-bound loop. To avoid blocking the iex shell session, the work is done in a separate process.

As soon as you start this computation, you should notice the CPU usage going to 100%, which proves that we're now running an intensive long-running CPU work.

Still, even though BEAM is using only one scheduler thread, the iex session is responsive, and you can evaluate other expressions. For example, let's sum the first 1,000,000,000 integers:

```
iex(2)> Enum.sum(1..1_000_000_000)
500000000500000000
```

We were able to run another job on an already busy scheduler thread, and that job finished almost immediately. This is the consequence of frequent context switching, which ensures that an occasional long-running job won't significantly affect the responsiveness of the entire system.

5.5 Summary

- A BEAM process is a lightweight concurrent unit of execution. Processes are completely isolated and share no memory.
- Processes can communicate with asynchronous messages. Synchronous sends and responses are manually built on top of this basic mechanism.
- A server process is a process that runs for a long time (possibly forever) and handles various messages. Server processes are powered by endless recursion.
- Server processes can maintain their own private state using the arguments of the endless recursion.

6 Generic server processes

This chapter covers

- Building a generic server process
- Using `GenServer`

In chapter 5, you became familiar with basic concurrency techniques: you learned how to create processes and communicate with them. I also explained the idea behind stateful server processes — long-running processes that react to messages and maintain state.

Server processes play an important role and are used frequently when building highly concurrent systems in Elixir and Erlang, so we'll spend some time exploring them in detail. In this chapter, you'll learn how to reduce some of the boilerplate associated with server processes, such as infinite recursion, state management, and message passing.

Erlang provides a helper module for implementing server processes — it's part of the framework called *Open Telecom Platform (OTP)*. Despite its misleading name, the framework has nothing to do with telecoms; rather, it provides patterns and abstractions for tasks such as creating components, building releases, developing server processes, handling and recovering from runtime errors, logging, event handling, and upgrading code.

You'll learn about various parts of OTP throughout this book, but in this chapter we'll focus on one of its most important parts: `GenServer` — the module that simplifies the implementation of server processes. Before we look at `GenServer`, though, you'll implement a simplified version of it, based on the message-passing primitives you saw in chapter 5.

6.1 Building a generic server process

You saw a few examples of server processes in chapter 5. Although those processes serve different purposes, there are some commonalities in their

implementations. In particular, all code that implements a server process needs to do the following:

- Spawn a separate process
- Run an infinite loop in the process
- Maintain the process state
- React to messages
- Send a response back to the caller

No matter what kind of server process you run, you'll always need to do these tasks, so it's worth moving this code to a single place. Concrete implementations can then reuse this code and focus on their specific needs. Let's look at how you can implement such generic code.

6.1.1 Plugging in with modules

The generic code will perform various tasks common to server processes, leaving the specific decisions to concrete implementations. For example, the generic code will spawn a process, but the concrete implementation must determine the initial state. Similarly, the generic code will run the loop, receive messages, and optionally send the responses, but the concrete implementation must decide how each message is handled and what the response is.

In other words, the generic code drives the entire process, and the specific implementation must fill in the missing pieces. Therefore, you need a plug-in mechanism that lets the generic code call into the concrete implementation when a specific decision needs to be made.

The simplest way to do this is to use modules. Remember that a module name is an atom. You can store that atom in a variable and use the variable later to invoke functions on the module:

```
iex(1)> some_module = IO      #1

iex(2)> some_module.puts("Hello")    #2
Hello
```

You can use this feature to provide callback hooks from the generic code. In

particular, you can take the following approach:

1. Make the generic code accept a plug-in module as the argument. That module is called a *callback module*.
2. Maintain the module atom as part of the process state.
3. Invoke callback-module functions when needed.

Obviously, for this to work, a callback module must implement and export a well-defined set of functions, which I'll gradually introduce as we implement the generic code.

6.1.2 Implementing the generic code

Let's start building a generic server process. First you need to start the process and initialize its state, as shown in the following listing.

Listing 6.1. Starting the server process (server_process.ex)

```
defmodule ServerProcess do
  def start(callback_module) do
    spawn(fn ->
      initial_state = callback_module.init()     #1
      loop(callback_module, initial_state)
    end)
  end

  ...
end
```

`ServerProcess.start/1` takes a module atom as the argument and then spawns the process. In the spawned process, the callback function `init/0` is invoked to create the initial state. Obviously, for this to work, the callback module must export the `init/0` function.

Finally, you enter the loop that will power the server process and maintain this state. The return value of `ServerProcess.start/1` is a pid, which can be used to send messages to the request process.

Next, you need to implement the loop code that powers the process, waits for messages, and handles them. In this example, you'll implement a

synchronous send-and-response communication pattern. The server process must receive a message, handle it, send the response message back to the caller, and change the process state.

The generic code is responsible for receiving and sending messages, whereas the specific implementation must handle the message and return the response and the new state. The idea is illustrated in the following listing.

Listing 6.2. Handling messages in the server process (server_process.ex)

```
defmodule ServerProcess do
  ...

  defp loop(callback_module, current_state) do
    receive do
      {request, caller} ->
        {response, new_state} =
          callback_module.handle_call(        #1
            request,                          #1
            current_state                     #1
          )                                   #1

        send(caller, {:response, response})   #2

        loop(callback_module, new_state)      #3
    end
  end

  ...
end
```

Here, you expect a message in the form of a `{request, caller}` tuple. The `request` is data that identifies the request and is meaningful to the specific implementation. The callback function `handle_call/2` takes the request payload and the current state, and it must return a `{response new_state}` tuple. The generic code can then send the response back to the caller and continue looping with the new state.

There's only one thing left to do: you need to provide a function to issue requests to the server process, as shown in the following listing.

Listing 6.3. Helper for issuing requests (server_process.ex)

```
defmodule ServerProcess do
  ...

  def call(server_pid, request) do
    send(server_pid, {request, self()})      #1

    receive do
      {:response, response} ->      #2
        response                    #3
    end
  end
end
```

At this point you have the abstraction for the generic server process in place. Let's see how it can be used.

6.1.3 Using the generic abstraction

To test the server process, you'll implement a simple key/value store. It will be a process that can be used to store mappings between arbitrary terms.

Remember that the callback module must implement two functions: `init/0`, which creates the initial state, and `handle_call/2`, which handles specific requests. The code is shown next.

Listing 6.4. Key/value store implementation (server_process.ex)

```
defmodule KeyValueStore do
  def init do
    %{}          #1
  end

  def handle_call({:put, key, value}, state) do      #2
    {:ok, Map.put(state, key, value)}                #2
  end                                                #2

  def handle_call({:get, key}, state) do      #3
    {Map.get(state, key), state}              #3
  end                                         #3
end
```

That's all it takes to create a specific server process. Because the infinite loop and message-passing boilerplate are pushed to the generic code, the specific

implementation is more concise and focused on its main task.

Take particular notice of how you use a multiclause in `handle_call/2` to handle different types of requests. This is the place where the specific implementation decides how to handle each request. The `ServerProcess` module is generic code that blindly forwards requests from client processes to the callback module.

Let's test the process:

```
iex(1)> pid = ServerProcess.start(KeyValueStore)

iex(2)> ServerProcess.call(pid, {:put, :some_key, :some_value})
:ok

iex(3)> ServerProcess.call(pid, {:get, :some_key})
:some_value
```

Notice how you start the process with `ServerProcess.start(KeyValueStore)`. This is where you plug the specific `KeyValueStore` into the generic code of `ServerProcess`. All subsequent invocations of `ServerProcess.call/2` will send messages to that process, which will in turn call `KeyValueStore.handle_call/2` to perform the handling.

It's beneficial to make clients completely oblivious to the fact that the `ServerProcess` abstraction is used. This can be achieved by introducing helper functions, as shown here.

Listing 6.5. Wrapping `ServerProcess` function calls (server_process.ex)

```
defmodule KeyValueStore do
  def start do
    ServerProcess.start(KeyValueStore)
  end

  def put(pid, key, value) do
    ServerProcess.call(pid, {:put, key, value})
  end

  def get(pid, key) do
    ServerProcess.call(pid, {:get, key})
```

```
    end

  ...
end
```

Clients can now use `start/0`, `put/3`, and `get/2` to manipulate the key/value store. These functions are informally called *interface functions*. Clients use the interface functions of `KeyValueStore` to start and interact with the process.

In contrast, `init/0` and `handle_call/2` are callback functions used internally by the generic code. Note that interface functions run in client processes, whereas callback functions are always invoked in the server process.

6.1.4 Supporting asynchronous requests

The current implementation of `ServerProcess` supports only synchronous requests. Let's expand on this and introduce support for asynchronous fire-and-forget requests, where a client sends a message and doesn't wait for a response.

In the current code, we use the term *call* for synchronous requests. For asynchronous requests, we'll use the term *cast*. This is the naming convention used in OTP, so it's good to adopt it.

Because you're introducing the second request type, you need to change the format of messages that are passed between client processes and the server. This will allow you to determine the request type in the server process and handle different types of requests in different ways.

This can be as simple as including the request-type information in the tuple being passed from the client process to the server, as shown next.

Listing 6.6. Including the request type in the message (server_process_cast.ex)

```
defmodule ServerProcess do
  ...

  def call(server_pid, request) do
    send(server_pid, {:call, request, self()})     #1
```

```
      ...
  end

  defp loop(callback_module, current_state) do
    receive do
      {:call, request, caller} ->      #2
        ...
    end
  end

  ...
end
```

Now you can introduce support for cast requests. In this scenario, when the message arrives, the specific implementation handles it and returns the new state. No response is sent back to the caller, so the callback function must return only the new state. The code is provided in the following listing.

Listing 6.7. Supporting casts in the server process (server_process_cast.ex)

```
defmodule ServerProcess do
  ...

  def cast(server_pid, request) do          #1
    send(server_pid, {:cast, request})      #1
  end                                        #1

  defp loop(callback_module, current_state) do
    receive do
      {:call, request, caller} ->
        ...

      {:cast, request} ->        #2
        new_state =
          callback_module.handle_cast(
            request,
            current_state
          )

        loop(callback_module, new_state)
    end
  end

  ...
end
```

To handle a cast request, you need the callback function `handle_cast/2`. This function must handle the message and return the new state. In the server loop, you then invoke this function and loop with the new state. That's all it takes to support cast requests.

Finally, you'll change the implementation of the key/value store to use casts. Keep in mind that a cast is a fire-and-forget type of request, so it's not suitable for all requests. In this example, the `get` request must be a call, because the server process needs to respond with the value associated with a given key. In contrast, the `put` request can be implemented as a cast because the client doesn't need to wait for the response.

Listing 6.8. Implementing `put` as a cast (server_process_cast.ex)

```
defmodule KeyValueStore do
  ...

  def put(pid, key, value) do
    ServerProcess.cast(pid, {:put, key, value})     #1
  end

  ...

  def handle_cast({:put, key, value}, state) do     #2
    Map.put(state, key, value)
  end

  ...
end
```

Now you can try the server process:

```
iex(1)> pid = KeyValueStore.start()

iex(2)> KeyValueStore.put(pid, :some_key, :some_value)

iex(3)> KeyValueStore.get(pid, :some_key)
:some_value
```

With a simple change in the generic implementation, you added another feature to the service processes. Specific implementations can now decide whether each concrete request should be implemented as a call or as a cast.

6.1.5 Exercise: refactoring the to-do server

An important benefit of the generic `ServerProcess` abstraction is that it lets you easily create various kinds of processes that rely on this common code. For example, in chapter 5, you developed a simple to-do server that maintains a to-do list abstraction in its internal state. This server can also be powered by the generic `ServerProcess`.

This is the perfect opportunity for you to practice a bit. Take the complete code from todo_server.ex from the chapter 5 source, and save it to a different file. Then add the last version of the `ServerProcess` module to the same file. Finally, adapt the code of the `TodoServer` module to work with `ServerProcess`.

Once you have everything working, compare the code between the two versions. The new version of `TodoServer` should be smaller and simpler, even for such a simple server process that supports only two different requests. If you get stuck, you can find the solution in the server_process_todo.ex file.

Note

It's clumsy to place multiple modules in a single file and maintain multiple copies of the `ServerProcess` code in different files. In chapter 7, you'll start using a better approach to code organization powered by the `mix` tool. But for the moment, let's stick with the current simplistic approach.

You're now finished implementing a basic abstraction for generic server processes. The current implementation is simple and leaves a lot of room for improvement, but it demonstrates the basic technique of generic server processes. Now it's time to use the full-blown OTP abstraction for generic server processes: `GenServer`.

6.2 Using GenServer

When it comes to production-ready code, it doesn't make much sense to build

and use the manually baked `ServerProcess` abstraction. That's because Elixir ships with a much better support for generic server processes, called `GenServer`. In addition to being much more feature-rich than `ServerProcess`, `GenServer` also handles various kinds of edge cases and is battle-tested in production in complex concurrent systems.

Some of the compelling features provided by `GenServer` include the following:

- Support for calls and casts
- Customizable timeouts for call requests
- Propagation of server-process crashes to client processes waiting for a response
- Support for distributed systems

Note that there's no special magic behind `GenServer`. Its code relies on concurrency primitives explained in chapter 5 and fault-tolerance features explained in chapter 9. After all, `GenServer` is implemented in plain Erlang and Elixir. The heavy lifting is done in the `:gen_server` module, which is included in the Erlang standard library. Some additional wrapping is performed in the Elixir standard library, in the `GenServer` module.

In this section, you'll learn how to build your server processes with `GenServer`. But first, let's examine the concept of OTP *behaviours*.

Note

Notice the British spelling of the word *behaviour*: this is the preferred spelling both in OTP code and official documentation. This book uses the British spelling to specifically denote an OTP behaviour but retains the American spelling (*behavior*) for all other purposes.

6.2.1 OTP behaviours

In Erlang terminology, a behaviour is generic code that implements a common pattern. The generic logic is exposed through the behaviour module, and you can plug into it by implementing a corresponding callback module.

The callback module must satisfy a contract defined by the behaviour, meaning it must implement and export a set of functions. The behaviour module then calls into these functions, allowing you to provide your own specialization of the generic code.

This is exactly what `ServerProcess` does. It powers a generic server process, requiring specific implementations to provide the callback module that implements the `init/0`, `handle_call/2`, and `handle_cast/2` functions. `ServerProcess` is a simple example of a behaviour.

It's even possible to specify the behaviour contract and verify that the callback module implements required functions during compilation. For details, see the official documentation (hexdocs.pm/elixir/Module.html#module-behaviour).

The Erlang standard library includes the following OTP behaviours:

- `gen_server` — Generic implementation of a stateful server process
- `supervisor` — Provides error handling and recovery in concurrent systems
- `application` — Generic implementation of components and libraries
- `gen_event` — Provides event-handling support
- `gen_statem` — Runs a finite state machine in a stateful server process

Elixir provides its own wrappers for the most frequently used behaviours via the modules `GenServer`, `Supervisor`, and `Application`. This book focuses on these behaviours. The `GenServer` behaviour receives detailed treatment in this chapter and chapter 7, `Supervisor` is discussed in chapters 8 and 9, and `Application` is presented in chapter 11.

The remaining behaviours, although useful, are used less often and won't be discussed in this book. Once you get a grip on `GenServer` and `Supervisor`, you should be able to research other behaviours on your own and use them when the need arises. You can find more about `gen_event` and `gen_statem` in the Erlang documentation (erlang.org/doc/design_principles/des_princ.html).

6.2.2 Plugging into GenServer

Using `GenServer` is roughly similar to using `ServerProcess`. There are some differences in the format of the returned values, but the basic idea is the same.

The `GenServer` behaviour requires seven callback functions, but frequently you'll need only a subset of those. You can get some sensible default implementations of all required callback functions if you use the `GenServer` module:

```
iex(1)> defmodule KeyValueStore do
          use GenServer
        end
```

The `use` macro is a language feature you haven't seen previously. During compilation, when this instruction is encountered, the specific macro from the `GenServer` module is invoked. That macro then injects a bunch of functions into the calling module (`KeyValueStore`, in this case). You can verify this in the shell:

```
iex(2)> KeyValueStore.__info__(:functions)
[child_spec: 1, code_change: 3, handle_call: 3, handle_cast: 2,
 handle_info: 2, init: 1, terminate: 2]
```

Here you use the `__info__`/1 function that's automatically injected into each Elixir module during compilation. It lists all exported functions of a module (except `__info__`/1).

As you can see in the output, many functions are automatically included in the module due to `use GenServer`. These are all callback functions that need to be implemented for you to plug into the `GenServer` behaviour.

Of course, you can then override the default implementation of each function as required. If you define a function of the same name and arity in your module, it will overwrite the default implementation you get through `use`.

At this point, you can plug your callback module into the behaviour. To start the process, use the `GenServer.start/2` function:

```
iex(3)> GenServer.start(KeyValueStore, nil)
{:ok, #PID<0.51.0>}
```

This works roughly like `ServerProcess`. The server process is started, and the behaviour uses `KeyValueStore` as the callback module. The second argument of `GenServer.start/2` is a custom parameter that's passed to the process during its initialization. For the moment, you don't need this, so you send the `nil` value. Finally, notice that the result of `GenServer.start/2` is a tuple of the form `{:ok, pid}`.

6.2.3 Handling requests

Now you can convert the `KeyValueStore` to work with `GenServer`. To do this, you need to implement three callbacks: `init/1`, `handle_cast/2`, and `handle_call/3`, shown in the following listing:

Listing 6.9. Implementing `GenServer` callbacks (key_value_gen_server.ex)

```
defmodule KeyValueStore do
  use GenServer

  def init(_) do
    {:ok, %{}}
  end

  def handle_cast({:put, key, value}, state) do
    {:noreply, Map.put(state, key, value)}
  end

  def handle_call({:get, key}, _, state) do
    {:reply, Map.get(state, key), state}
  end
end
```

These callbacks work similarly to the ones in `ServerProcess`, with a couple of differences:

- `init/1` accepts one argument. This is the second argument provided to `GenServer.start/2`, and you can use it to pass data to the server process while starting it.
- The result of `init/1` must be in the format `{:ok, initial_state}`.
- `handle_cast/2` accepts the request and the state and should return the result in the format `{:noreply, new_state}`.
- `handle_call/3` takes the request, the caller information, and the state. It

should return the result in the format `{:reply, response, new_state}`.

The second argument to `handle_call/3` is a tuple that contains the request ID (used internally by the `GenServer` behaviour) and the pid of the caller. This information is in most cases not needed, so in this example you ignore it.

With these callbacks in place, the only things missing are interface functions. To interact with a `GenServer` process, you can use functions from the `GenServer` module. In particular, you can use `GenServer.start/2` to start the process and `GenServer.cast/2` and `GenServer.call/2` to issue requests. The code is shown in the next listing:

Listing 6.10. Adding interface functions (key_value_gen_server.ex)

```
defmodule KeyValueStore do
  use GenServer

  def start do
    GenServer.start(KeyValueStore, nil)
  end

  def put(pid, key, value) do
    GenServer.cast(pid, {:put, key, value})
  end

  def get(pid, key) do
    GenServer.call(pid, {:get, key})
  end

  ...
end
```

That's it! With only a few changes, you've moved from a basic `ServerProcess` to a full-blown `GenServer`. Let's test the server:

```
iex(1)> {:ok, pid} = KeyValueStore.start()

iex(2)> KeyValueStore.put(pid, :some_key, :some_value)

iex(3)> KeyValueStore.get(pid, :some_key)
:some_value
```

It works as expected.

There are many differences between `ServerProcess` and `GenServer`, but a couple points deserve special mention.

First, `GenServer.start/2` returns only after the `init/1` callback has finished in the server process. Consequently, the client process that starts the server is blocked until the server process is initialized.

Second, `GenServer.call/2` doesn't wait indefinitely for a response. By default, if the response message doesn't arrive in five seconds, an error is raised in the client process. You can alter this by using `GenServer.call(pid, request, timeout)`, where the timeout is given in milliseconds. In addition, if the receiver process happens to terminate while you're waiting for the response, `GenServer` detects it and raises a corresponding error in the caller process.

6.2.4 Handling plain messages

Messages sent to the server process via `GenServer.call` and `GenServer.cast` contain more than just a request payload. Those functions include additional data in the message sent to the server process. This is something you did in the `ServerProcess` example in section 6.1:

```
defmodule ServerProcess do
  ...

  def call(server_pid, request) do
    send(server_pid, {:call, request, self()})     #1
    ...
  end

  def cast(server_pid, request) do
    send(server_pid, {:cast, request})     #2
  end

  ...

  defp loop(callback_module, current_state) do
    receive do
      {:call, request, caller} ->        #3
```

```
        ...

      {:cast, request} ->      #4
        ...
      end
  end
  ...
end
```

Notice that you don't send the plain `request` payload to the server process; you include additional data, such as the request type and the caller for call requests.

`GenServer` uses a similar approach, using `:"$gen_cast"` and `:"$gen_call"` atoms to decorate cast and call messages. You don't need to worry about the exact format of those messages, but it's important to understand that `GenServer` internally uses particular message formats and handles those messages in a specific way.

Occasionally you may need to handle messages that aren't specific to `GenServer`. For example, imagine that you need to do a periodic cleanup of the server process state. You can use the Erlang function `:timer.send_interval/2`, which periodically sends a message to the caller process. Because this message isn't a `GenServer`-specific message, it's not treated as a cast or a call. Instead, for such plain messages, `GenServer` calls the `handle_info/2` callback, giving you a chance to do something with the message.

Here's a sketch of this technique:

```
iex(1)> defmodule KeyValueStore do
          use GenServer

          def init(_) do
            :timer.send_interval(5000, :cleanup)     #1
            {:ok, %{}}
          end

          def handle_info(:cleanup, state) do     #2
            IO.puts "performing cleanup..."
            {:noreply, state}
          end
```

```
        end

iex(2)> GenServer.start(KeyValueStore, nil)

performing cleanup...       #3
performing cleanup...
performing cleanup...
```

During process initialization, you make sure a `:cleanup message` is sent to the process every five seconds. This message is handled in the `handle_info/2` callback, which essentially works like `handle_cast/2`, returning the result as `{:noreply, new_state}`.

6.2.5 Other GenServer features

There are various other features and subtleties I haven't mentioned in this basic introduction to `GenServer`. You'll learn about some of them elsewhere in this book, but you should definitely take the time to look over the documentation for the `GenServer` module (hexdocs.pm/elixir/GenServer.html) and its Erlang foundation (erlang.org/doc/man/gen_server.html).

A couple of points still deserve special mention.

Compile-time checking

One problem with the callbacks mechanism is that it's easy to make a subtle mistake when defining a callback function. Consider the following example:

```
iex(1)> defmodule EchoServer do
          use GenServer

          def handle_call(some_request, server_state) do
            {:reply, some_request, server_state}
          end
        end
```

Here you have a simple echo server, which handles every call request by sending the request back to the client. Try it out:

```
iex(2)> {:ok, pid} = GenServer.start(EchoServer, nil)
```

```
{:ok, #PID<0.96.0>}

iex(3)> GenServer.call(pid, :some_call)
** (exit) exited in: GenServer.call(#PID<0.96.0>, :some_call, 500
   ** (EXIT) an exception was raised:
      ** (RuntimeError) attempted to call GenServer #PID<0.96.0>
                        no handle_call/3 clause was provided
```

Issuing a call caused the server to crash with an error that no `handle_call/3` clause is provided, although the clause is listed in the module. What happened? If you look closely at the definition of `EchoServer`, you'll see that you defined `handle_call/2`, while `GenServer` requires `handle_call/3`.

You can get a compile-time warning here if you tell the compiler that the function being defined is supposed to satisfy a contract by some behaviour. To do this, you need to provide the `@impl` module attribute immediately before the first clause of the callback function:

```
iex(1)> defmodule EchoServer do
          use GenServer

          @impl GenServer     #1
          def handle_call(some_request, server_state) do
            {:reply, some_request, server_state}
          end
        end
```

The `@impl GenServer` tells the compiler that the function about to be defined is a callback function for the `GenServer` behaviour. As soon as you execute this expression in the shell, you'll get a warning:

```
warning: got "@impl GenServer" for function handle_call/2 but thi
behaviour does not specify such callback.
```

The compiler tells you that `GenServer` doesn't deal with `handle_call/2`, so already during compilation you get a hint that something is wrong. It's a good practice to always specify the `@impl` attribute for every callback function you define in your modules.

Name registration

Recall from chapter 5 that a process can be registered under a local name (an atom), where *local* means the name is registered only in the currently running BEAM instance. This allows you to create a singleton process that you can access by name without needing to know its pid.

Local registration is an important feature because it supports patterns of fault-tolerance and distributed systems. You'll see exactly how this works in later chapters, but it's worth mentioning that you can provide the process name as an option to `GenServer.start`:

```
GenServer.start(
  CallbackModule,
  init_param,
  name: :some_name      #1
)
```

You can then issue calls and casts using the name:

```
GenServer.call(:some_name, ...)
GenServer.cast(:some_name, ...)
```

The most frequent approach is to use the same name as the module name. As explained in section 2.4.2, module names are atoms, so you can safely pass them as the `:name` option. Here's a sketch of this approach:

```
defmodule KeyValueStore do
  def start() do
    GenServer.start(KeyValueStore, nil, name: KeyValueStore)    #
  end

  def put(key, value) do
    GenServer.cast(KeyValueStore, {:put, key, value})      #2
  end

  ...
end
```

Notice how `KeyValueStore.put` now doesn't need to take the `pid`. It will simply issue a request to the registered process.

You can also replace `KeyValueStore` with the special form `__MODULE__`. During compilation, `__MODULE__` is replaced with the name of the module

where the code resides:

```
defmodule KeyValueStore do
  def start() do
    GenServer.start(__MODULE__, nil, name: __MODULE__)      #1
  end

  def put(key, value) do
    GenServer.cast(__MODULE__, {:put, key, value})     #2
  end

  ...
end
```

After compilation, this code is completely equivalent to the previous version, but some future refactoring is made easier. If, for example, you rename `KeyValueStore` to `KeyValue.Store`, you need to do it only in one place in the module.

Stopping the server

Different callbacks can return various types of responses. So far, you've seen the most common cases:

- `{:ok, initial_state}` from `init/1`
- `{:reply, response, new_state}` from `handle_call/3`
- `{:noreply, new_state}` from `handle_cast/2` and `handle_info/2`

There are additional possibilities, the most important one being the option to stop the server process.

In `init/1`, you can decide against starting the server. In this case, you can either return `{:stop, reason}` or `:ignore`. In both cases, the server won't proceed with the loop, and will instead terminate.

If `init/1` returns `{:stop, reason}`, the result of `start/2` will be `{:error, reason}`. In contrast, if `init/1` returns `:ignore`, the result of `start/2` will also be `:ignore`. The difference between these two return values is in their intention. You should opt for `{:stop, reason}` when you can't proceed further due to some error. In contrast, `:ignore` should be used when stopping

the server is the normal course of action.

Returning `{:stop, reason, new_state}` from `handle_*` callbacks causes `GenServer` to stop the server process. If the termination is part of the standard workflow, you should use the atom `:normal` as the stoppage reason. If you're in `handle_call/3` and also need to respond to the caller before terminating, you can return `{:stop, reason, response, new_state}`.

You may wonder why you need to return a new state if you're terminating the process. The reason is that just before the termination, `GenServer` calls the callback function `terminate/2`, sending it the termination reason and the final state of the process. This can be useful if you need to perform cleanup.

Finally, you can also stop the server process by invoking `GenServer.stop/3` from the client process. This invocation will issue a synchronous request to the server. The behaviour will handle the stop request itself by stopping the server process.

6.2.6 Process lifecycle

It's important to always be aware of how `GenServer`-powered processes tick and where (in which process) various functions are executed. Let's do a quick recap by looking at figure 6.1, which shows the lifecycle of a typical server process.

Figure 6.1. Lifecycle of a `GenServer`-powered process

A client process starts the server by calling `GenServer.start` and providing the callback module (**1**). This creates the new server process, which is powered by the `GenServer` behaviour.

Requests can be issued by client processes using various `GenServer` functions or plain `send`. When a message is received, `GenServer` invokes callback functions to handle it. Therefore, callback functions are always executed in the server process.

The process state is maintained in the `GenServer` loop but is defined and manipulated by the callback functions. It starts with `init/1`, which defines the initial state that's then passed to subsequent `handle_*` callbacks (**2**). Each of these callbacks receives the current state and must return its new version, which is used by the `GenServer` loop in place of the old one.

The Actor model

Erlang is an accidental implementation of the Actor model originally described by Carl Hewitt. An actor is a concurrent computational entity that encapsulates state and can communicate with other actors. When processing a single message, an actor can designate the new state that will be used when processing the next message. This is roughly similar to how `GenServer`-based processes work in Erlang. Note, though, that as Robert Virding (one of Erlang's co-inventors) has repeatedly stated, Erlang developers arrived at this idea on their own and learned about the existence of the Actor model much later.

There are some disagreements about whether Erlang is a proper implementation of the Actor model, and the term actor isn't used much in the Erlang community. This book doesn't use this terminology either. Still, it's worth keeping in mind that in the context of Erlang, an actor corresponds to a server process, most frequently a `GenServer`.

6.2.7 OTP-compliant processes

For various reasons, once you start building production systems, you should avoid using plain processes started with `spawn`. Instead, all of your processes should be so-called *OTP-compliant processes*. Such processes adhere to OTP conventions, they can be used in supervision trees (described in chapter 9), and errors in those processes are logged with more details.

All processes powered by OTP behaviours, such as `GenServer` and `Supervisor`, are OTP-compliant. Elixir also includes other modules that can be used to run OTP-compliant processes. For example, the `Task` module (hexdocs.pm/elixir/Task.html) is perfect to run one-off jobs that process some input and then stop. The `Agent` module (hexdocs.pm/elixir/Agent.html) is a simpler (but less powerful) alternative to `GenServer`-based processes and is appropriate if the single purpose of the process is to manage and expose state. Both `Task` and `Agent` are discussed in chapter 10.

In addition, there are various other OTP-compliant abstractions available via third-party libraries. For example, `GenStage` (hexdocs.pm/gen_stage) can be used for back-pressure and load control. The `Phoenix.Channel` module (hexdocs.pm/phoenix/Phoenix.Channel.html), which is part of the Phoenix web framework (phoenixframework.org), is used to facilitate bidirectional communication between a client and a web server over protocols such as WebSocket or HTTP.

There isn't enough space in this book to treat every possible OTP-compliant abstraction, so you'll need to do some research of your own. But it's worth pointing out that most such abstractions follow the ideas of `GenServer`. Except for the `Task` module, all of the OTP abstractions mentioned in this section are internally implemented on top of `GenServer`. Therefore, in my personal opinion, `GenServer` is likely the most important part of OTP. If you properly understand the principles of `GenServer`, most other abstractions should be much easier to grasp.

6.2.8 Exercise: GenServer-powered to-do server

Let's wrap up this chapter with a simple but important exercise. For practice, try to change the to-do server, implemented earlier in this chapter, to work with the `GenServer` behaviour. This should be a straightforward task, but if you get stuck, the solution is in the todo_server.ex file.

Be sure to either finish this exercise or analyze and understand the solution, because in future chapters you'll gradually expand on this simple server process and build a highly concurrent distributed system.

6.3 Summary

- A generic server process is an abstraction that implements tasks common to any kind of server process, such as recursion-powered looping and message passing.
- A generic server process can be implemented as a behaviour. A behaviour drives the process, whereas specific implementations can plug into the behaviour via callback modules.
- The behaviour invokes callback functions when the specific implementation needs to make a decision.
- `GenServer` is a behaviour that implements a generic server process.
- A callback module for `GenServer` must implement various functions. The most frequently used ones are `init/1`, `handle_cast/2`, `handle_call/3`, and `handle_info/2`.
- You can interact with a `GenServer` process with the `GenServer` module.
- Two types of requests can be issued to a server process: calls and casts.
- A cast is a fire-and-forget type of request — a caller sends a message and immediately moves on to do something else.
- A call is a synchronous send-and-respond request — a caller sends a message and waits until the response arrives, the timeout occurs, or the server crashes.

7 Building a concurrent system

This chapter covers

- Working with the `mix` project
- Managing multiple to-do lists
- Persisting data
- Reasoning with processes

The concurrent examples you've seen so far have relied on a single server-process instance. But typical Elixir/Erlang systems are powered by a multitude of processes, many of which are stateful server processes. It's not uncommon for a moderately complex system to run a few thousand processes, whereas larger systems may be powered by hundreds of thousands or even millions of processes. Remember that processes are cheap, so you can create them in abundance. And owing to message-passing concurrency, it's still fairly easy to reason about highly concurrent systems. Therefore, it's useful to run different tasks in separate processes. Such a highly concurrent approach can often improve the scalability and reliability of your systems.

In this chapter, you'll see an example of a more involved system powered by many processes that cooperate to provide the full service. Your ultimate goal is to build a distributed HTTP server that can handle many end users who are simultaneously manipulating many to-do lists. You'll do this throughout the remaining chapters and reach the final goal in chapter 12. In this chapter, you'll develop an infrastructure for handling multiple to-do lists and persisting them to disk.

But first, let's look at how you can manage more complex projects with the `mix` tool.

7.1 Working with the mix project

As code gets more involved, placing all the modules in a single file becomes

increasingly clumsy. This is the right time to start working with multifile projects.

Chapter 2 briefly mentioned that Elixir ships with the `mix` tool, which you can use to create, build, and run projects as well as manage their dependencies, run tests, and create custom project-based tasks. Here you'll learn just enough about `mix` to create and run a project. Additional `mix` features will be introduced as the need arises.

You'll use `mix` to create a project for the to-do list. Type the following at the command line:

```
$ mix new todo
```

This creates the todo folder and a project structure under it. The result is a folder that contains only a handful of files, including a readme, unit-test support files, and the .gitignore file. `mix` projects are extremely simple and don't introduce a plethora of autogenerated files.

Tip

This book doesn't provide a detailed treatment of the `mix` tool. Instead, essential features are introduced when needed. To find out more about `mix`, check the online *Introduction to Mix* guide (https://elixir-lang.org/getting-started/mix-otp/introduction-to-mix.html). In addition, from the command line, you can run `mix help` to get a list of available commands and `mix help command` to get detailed help for a particular command. Finally, the online reference for `mix` is available at https://hexdocs.pm/mix.

Once the project is in place, you can go to its folder and run `mix` tasks from there. For example, you can compile the project with the `mix compile` command, or you can run tests with `mix test`.

You can also use a special way of starting `iex`, which is useful when you want to play with `mix` projects in the Elixir shell. When you run `iex -S mix`, two things happen. First, the project is compiled (just as with `mix compile`). If this is successful, the shell is started, and all modules from the project can

be referenced and used.

Using `mix`, it's possible to organize your code into multiple files and folders. You can place .ex files under the lib folder, and they'll automatically be included in the next build. You can also use arbitrarily nested subfolders under the lib folder.

There are no hard rules regarding how files should be named and organized, but there are some preferred conventions:

- You should place your modules under a common top-level alias. For example, modules might be called `Todo.List`, `Todo.Server`, and so on. This reduces the chance of module names conflicting when you combine multiple projects into a single system.
- In general, one file should contain one module. Occasionally, if a helper module is small and used only internally, it can be placed in the same file as the module using it. If you want to implement protocols for the module, you can do this in the same file as well.
- A filename should be an underscore case (aka *snake* case) version of the main module name it implements. For example, a `TodoServer` module would reside in a todo_server.ex file in the lib folder.
- The folder structure should correspond to multipart module names. A module called `Todo.Server` should reside in the file lib/todo/server.ex.

These aren't strict rules, but they're the ones used by the Elixir project as well as many third-party libraries.

With this out of the way, you can start adding code to the project.

In chapter 4 you've developed the module `TodoList`. In chapter 6, as a part of the exercises, you've developed the module `TodoServer`, which implements a server process that maintains the state of a single to-do list. The final version of both modules resides in the file todo_server.ex from chapter 6.

Now you'll add the code of those modules to the newly generated `todo` project. Here's what you need to do:

1. Remove the file todo/lib/todo.ex.

2. Remove the file todo/test/todo_test.exs.
3. Place the `TodoList` code in the todo/lib/todo/list.ex file. Rename the module to `Todo.List`.
4. Place the `TodoServer` code in the todo/lib/todo/server.ex file. Rename the module to `Todo.Server`.
5. Replace all references to `TodoServer` with `Todo.Server` and all references to `TodoList` with `Todo.List`.

The final version is available in the todo folder. Now you can start the system with `iex -S mix` and verify that it works:

```
$ iex -S mix

iex(1)> {:ok, todo_server} = Todo.Server.start()

iex(2)> Todo.Server.add_entry(
          todo_server,
          %{date: ~D[2023-12-19], title: "Dentist"}
        )

iex(3)> Todo.Server.entries(todo_server, ~D[2023-12-19])
[%{date: ~D[2023-12-19], id: 1, title: "Dentist"}]
```

At this point, the to-do code is in the `mix` project, and you can continue to extend it with additional features.

7.2 Managing multiple to-do lists

This section introduces support for managing multiple to-do lists. Before starting, let's recap what you've built so far:

- A pure functional `Todo.List` abstraction
- A to-do server process that can be used to manage one to-do list for a long time

There are two approaches to extending this code to work with multiple lists:

- Implement a pure functional abstraction to work with multiple to-do lists. Modify `Todo.Server` to use the new abstraction as its internal state.
- Run one instance of the existing to-do server for each to-do list.

The problem with the first approach is that you'll end up having only one process to serve all users. This approach isn't very scalable. If the system is used by many different users, they'll frequently block each other, competing for the same resource — a single server process that performs all tasks.

The alternative is to use as many processes as there are to-do lists. With this approach, each list is managed concurrently, and the system should be more responsive and scalable.

To run multiple to-do server processes, you need another entity — something you'll use to create `Todo.Server` instances or fetch the existing ones. That "something" must manage a state — essentially a key/value structure that maps to-do list names to to-do server pids. This state will of course be mutable (the number of lists changes over time) and must be available during the system's lifetime.

Therefore, you'll introduce another server process: a to-do cache. You'll run only one instance of this process, and it will be used to create and return a pid of a to-do server process that corresponds to the given name. The module will export only two functions: `start/0`, which starts the process, and `server_process/2`, which retrieves a to-do server process (its pid) for a given name, optionally starting the process if it isn't already running.

7.2.1 Implementing a cache

Let's begin implementing the cache process. First, copy the entire mix project (the todo folder) to the todo_cache folder. Then add the new file todo_cache/lib/todo/cache.ex, which is where the code for `Todo.Cache` will reside.

Now you need to decide what the process state will be. Remember, the process will provide to-do server processes. You give it a name, and it returns the pid of the corresponding process. In this case, it seems reasonable to use a map that associates to-do list names with to-do server pids. This is implemented in the following listing.

Listing 7.1. Cache initialization (todo_cache/lib/todo/cache.ex)

```
defmodule Todo.Cache do
  use GenServer

  def init(_) do
    {:ok, %{}}
  end

  ...
end
```

With this in place, you can begin introducing the `server_process` request.
You need to decide whether this request will be a call or a cast. Because this
request must return a result to the caller (a to-do server pid), there are no
options — it needs to be a call. The implementation is shown next.

Listing 7.2. Handling the `server_process` request (todo_cache/lib/todo/cache.ex)

```
defmodule Todo.Cache do
  ...

  def handle_call({:server_process, todo_list_name}, _, todo_serv
    case Map.fetch(todo_servers, todo_list_name) do
      {:ok, todo_server} ->              #1
        {:reply, todo_server, todo_servers}

      :error ->          #2
        {:ok, new_server} = Todo.Server.start()      #3

        {
          :reply,
          new_server,
          Map.put(todo_servers, todo_list_name, new_server)
        }
    end
  end

  ...
end
```

In this example, you use `Map.fetch/2` to query the map. If there's something
for the given key, you return the value to the caller, leaving the state
unchanged. Otherwise, you must start a server, return its pid, and insert an
appropriate name/value pair in the process state.

Finally, you shouldn't forget to include interface functions.

Listing 7.3. Interface functions (todo_cache/lib/todo/cache.ex)

```
defmodule Todo.Cache do
  ...

  def start do
    GenServer.start(__MODULE__, nil)
  end

  def server_process(cache_pid, todo_list_name) do
    GenServer.call(cache_pid, {:server_process, todo_list_name})
  end

  ...
end
```

Notice how `__MODULE__` is passed as the first argument to `GenServer.start/2`. During compilation, this expression is replaced with the name of the current module. This is a simple convenience; you could write `Todo.Cache` instead, but this approach removes this minor duplication and guards the code against a possible change of the module name.

At this point, the to-do cache is complete, and you can try it. Start the shell with `iex -S mix`, and do the following:

```
iex(1)> {:ok, cache} = Todo.Cache.start()

iex(2)> Todo.Cache.server_process(cache, "Bob's list")   #1
#PID<0.69.0>                                              #1

iex(3)> Todo.Cache.server_process(cache, "Bob's list")   #2
#PID<0.69.0>                                              #2

iex(4)> Todo.Cache.server_process(cache, "Alice's list") #3
#PID<0.72.0>                                              #3
```

The returned pid represents a to-do server process that manages a single to-do list. You can use it in the familiar way to manipulate the list:

```
iex(5)> bobs_list = Todo.Cache.server_process(cache, "Bob's list"
```

```
iex(6)> Todo.Server.add_entry(
          bobs_list,
          %{date: ~D[2023-12-19], title: "Dentist"}
        )

iex(7)> Todo.Server.entries(bobs_list, ~D[2023-12-19])
[%{date: ~D[2023-12-19], id: 1, title: "Dentist"}]
```

Of course, Alice's list isn't affected by these manipulations:

```
iex(8)> Todo.Cache.server_process(cache, "Alice's list") |>
          Todo.Server.entries(~D[2023-12-19])
[]
```

Having the cache in place makes it possible for you to manage many to-do lists independently. The following session creates 100,000 to-do list servers and verifies that you have that many processes running:

```
iex(1)> {:ok, cache} = Todo.Cache.start()

iex(2)> length(Process.list())
65

iex(3)> Enum.each(
          1..100_000,
          fn index ->
            Todo.Cache.server_process(cache, "to-do list #{index}
          end
        )

iex(4)> length(Process.list())
100065
```

Here you use the `Process.list/0` function to get the list of currently running processes.

You might be puzzled as to why you initially have 65 processes running, even though you started just 1. The remaining processes are those started and used internally by Elixir and Erlang.

7.2.2 Writing tests

Now that the code is organized in the mix project, you can write automated

tests. The testing framework for Elixir is called `ex_unit`, and it's included in the Elixir distribution. Running tests is as easy as invoking `mix test`. All you need to do is write the test code.

Let's look at a quick example by testing the behavior of `Todo.Cache.server_process/2`. First you need to create the test file. The sketch is provided in the following listing.

Listing 7.4. Test file skeleton (todo_cache/test/todo/cache_test.exs)

```
defmodule Todo.CacheTest do
  use ExUnit.Case      #1

  ...
end
```

Take note of the file location and the name. A test file must reside in the test folder, and its name must end with `_test.exs` to be included in the test execution. As explained in chapter 2, the .exs extension stands for Elixir script, and it's used to indicate that a file isn't compiled to disk. Instead, `mix` will interpret this file every time the tests are executed.

The script file must define the test module that contains the tests. The expression `use ExUnit.Case` prepares the test module for testing. This expression injects some boilerplate that makes the module compliant with `ex_unit` and imports some helper test macros to the module.

One such macro is `test`, which can be used to define tests. You'll use it to test the behavior of `Todo.Cache.server_process/2`. The code is provided in the following listing.

Listing 7.5. `Testing server_process` (todo_cache/test/todo/cache_test.exs)

```
defmodule Todo.CacheTest do
  use ExUnit.Case

  test "server_process" do      #1
    {:ok, cache} = Todo.Cache.start()
    bob_pid = Todo.Cache.server_process(cache, "bob")

    assert bob_pid != Todo.Cache.server_process(cache, "alice")
```

```
    assert bob_pid == Todo.Cache.server_process(cache, "bob")
  end

  ...
end
```

To define a test, you need to write `test test_description do … end`. The
test description is a string that's included in the output if the test fails. The
code of the test itself is included in the `do` block.

The `test` macro is an example of metaprogramming capabilities in Elixir.
This macro will generate a function that contains some boilerplate and the
code provided in the `do` block. This function will then be invoked by `ex_unit`
when you execute tests.

In this particular test, you first start the cache process and then fetch one
server process. Then you verify the expected behavior. This is done with the
help of the `assert` macro, which takes an expression and verifies its outcome.
If the expression fails, `assert` will raise an error with a descriptive output.
This error will be caught by `ex_unit` and displayed.

For example, take a look at the first assertion:

```
assert bob_pid != Todo.Cache.server_process(cache, "alice")
```

In this assertion, you're verifying that Alice's and Bob's to-do lists are
powered by different processes.

Just like `test`, `assert` is a macro and therefore is invoked during
compilation. The macro introspects the expression and transforms it into
different code. An approximation of the generated code could be something
like this:

```
left_value = bob_pid
right_value = Todo.Cache.server_process(cache, "alice")

if left_value == right_value do
  # raise an error
end
```

In other words, the `assert` macro generates the code that will fail if the

expression `bob_pid != Todo.Cache.server_process(cache, "alice")` returns `false`.

A great benefit of the way `assert` works is that you don't need to learn a completely new set of functions, such as `assert_equal`, `assert_not_equal`, or `assert_gt`, to write your assertions. Instead, you use the same expressions as in the regular code to verify the desired behavior. You can assert on standard comparisons such as ==, !=, >, <, and so on.

You can even assert that a pattern matching expression succeeded. Let's look at a quick example. You'll add another test that verifies the behavior of to-do server operations. To keep things simple, you'll include the test in the same file. The code is provided in the following listing.

Listing 7.6. Testing to-do server operations (todo_cache/test/todo/cache_test.exs)

```
defmodule Todo.CacheTest do
  use ExUnit.Case

  ...

  test "to-do operations" do
    {:ok, cache} = Todo.Cache.start()

    alice = Todo.Cache.server_process(cache, "alice")
    Todo.Server.add_entry(alice, %{date: ~D[2023-12-19], title: "

    entries = Todo.Server.entries(alice, ~D[2023-12-19])
    assert [%{date: ~D[2023-12-19], title: "Dentist"}] = entries
  end
end
```

Here you create one to-do server, add a single entry, and then fetch the entries for the given date. Finally, using pattern matching, you assert that the list of entries has exactly one element, with date and title fields having proper values. Relying on pattern matching allowed you to check only the relevant fields, and to verify the size of the result in a single expression.

At this point, you've created a single test file with a couple of tests. The test project in the todo_cache folder also includes another test file called test/todo/list_test.exs, which verifies the behavior of the `Todo.List` module.

For the sake of brevity, that code isn't presented here.

ⓘ Note

It's worth noting that the example projects in this book aren't test-driven or particularly well tested. In this book, the focus is on extremely simple code that illustrates a point. Such code is often not very testable, and some improvisations have been used to ensure basic correctness.

Now you can run all the tests with `mix test`:

```
$ mix test
.......

Finished in 0.05 seconds
7 tests, 0 failures
```

There are many other features available in `ex_unit`, but we'll stop here. To learn more about unit testing in Elixir, check out the official `ExUnit` reference at https://hexdocs.pm/ex_unit, and the `mix test` documentation at https://hexdocs.pm/mix/Mix.Tasks.Test.html.

7.2.3 Analyzing process dependencies

Let's reflect a bit on the current system. You've developed support for managing many to-do list instances, and the end goal is to use this infrastructure in an HTTP server. In the Elixir/Erlang world, HTTP servers typically use a separate process for each request. Thus, if you have many simultaneous end users, you can expect many BEAM processes accessing your to-do cache and to-do servers. The dependency between processes is illustrated in figure 7.1.

Figure 7.1. Cooperation of processes
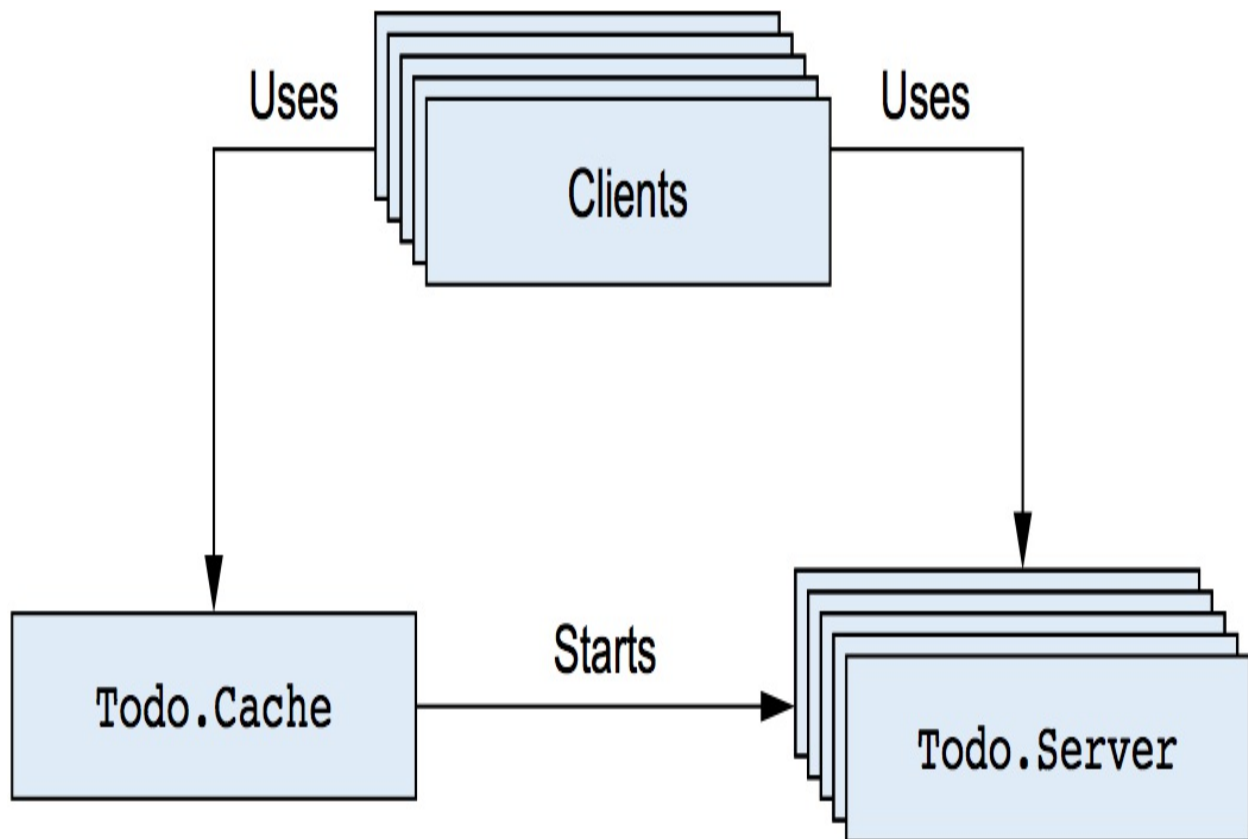
Here, each box represents a single process. The Clients boxes are arbitrary clients, such as HTTP request-handler processes. Looking at this diagram, you can immediately spot some characteristics of your system's concurrent behavior:

- Multiple clients (possibly a large number of them) issue requests to the single to-do cache process.
- Multiple clients communicate with multiple to-do server processes.

The first point identifies a possible source of a bottleneck. Because you have only one to-do cache process, you can handle only one `server_process` request simultaneously, regardless of how many CPU resources you have.

Notice that this problem may not be significant in practice. If your `server_process` takes, for example, one microsecond, the to-do cache could handle a load of up to 1,000,000 requests per second, which should be sufficient for most needs. But if request handling takes 100 milliseconds, you could process only 10 requests per second, and your system wouldn't be able

to handle higher loads.

It's easy to reason about an individual process. Regardless of how many concurrent requests are coming in, a single process can handle only one request at a time. Thus, a process is good enough if its request-handling rate is at least equal to the incoming rate. Otherwise, you have to either optimize the process or do other interventions.

For this specific case, the to-do cache performs a very simple operation — a map lookup followed by an optional process creation and map update. According to a quick test on my machine, for 1 million to-do lists, it takes about 5 microseconds to start a new to-do server and put it in the map, or 1.5 microseconds to fetch the existing one. This should be sufficient for a load of at least 200,000 requests/sec, which seems like reasonable performance for this initial attempt. If you want to repeat the test on your machine, take a look at the instructions in `todo_cache/lib/load_test.ex`.

I need to say a few words about client interactions with to-do servers. Once a client gets a to-do server pid, the list manipulation runs concurrently to all other activities in the system. Because you can expect list manipulations to be fairly involved, it's beneficial to run those operations concurrently. This is where your system is concurrent and scalable — it can manipulate multiple lists, using as many resources as possible.

Also recall from chapter 5 that a process waiting for a message is suspended and doesn't waste CPU resources. Thus, regardless of the number of processes, only those processes that are actually doing computations consume CPU. In this case, that means a client process doesn't use CPU while it waits for a to-do server to finish.

Finally, you can be sure that a single list can't be modified by two simultaneous clients. Recall that the list is managed by a single process. Even if a million clients try to modify the same list, their requests will be serialized in the corresponding to-do server and handled one by one. Because a process runs only one request at a time, its internal state is consistent. You know there can't be multiple simultaneous updates of the process state, which makes race conditions in a single process impossible.

If you need to make sure part of the code is synchronized — that is, that there are no multiple simultaneous executions of critical code — it's best to run that code in a dedicated process. When multiple clients want this code to run, they issue a request to that process. The process then serves as a synchronization point, making sure the critical code is run in a single process.

Now you have a basic system that you can use to manipulate many to-do lists. It's time to include basic persistence so your data can outlive server restarts.

7.3 Persisting data

In this section, you'll extend the to-do cache and introduce basic data persistence. The focus here isn't so much on the persistence itself, but rather on exploring the process model — how you can organize your system into various server processes, analyze dependencies, and identify and address bottlenecks. You'll start with the code from the todo_cache folder and extend it gradually. For data persistence, you'll use simple disk-based persistence, encoding the data into the Erlang external term format. The complete solution is in the persistable_todo_cache folder.

7.3.1 Encoding and persisting

To encode an arbitrary Elixir/Erlang term, you use the `:erlang.term_to_binary/1` function, which accepts an Erlang term and returns an encoded byte sequence as a binary value. The input term can be of arbitrary complexity, including deep hierarchies of nested lists and tuples. The result can be stored to disk, retrieved at a later point, and decoded to an Erlang term with the inverse function `:erlang.binary_to_term/1`.

Equipped with this knowledge, you'll introduce another process: a database powered by the `Todo.Database` module. This will be a server process that supports two requests: `store` and `get`. While storing data, clients will provide a key and the corresponding data. The data will be stored in the file that bears

the same name as the key. This approach is far from perfect and is error-prone, but it's simple enough to let us focus on the concurrency aspect of the problem.

The full implementation of the database process is given in the following listing.

Listing 7.7. Database process (persistable_todo_cache/lib/todo/database.ex)

```elixir
defmodule Todo.Database do
  use GenServer

  @db_folder "./persist"

  def start do
    GenServer.start(__MODULE__, nil,
      name: __MODULE__                          #1
    )
  end

  def store(key, data) do
    GenServer.cast(__MODULE__, {:store, key, data})
  end

  def get(key) do
    GenServer.call(__MODULE__, {:get, key})
  end

  def init(_) do
    File.mkdir_p!(@db_folder)          #2
    {:ok, nil}
  end

  def handle_cast({:store, key, data}, state) do      #3
    key                                                 #3
    |> file_name()                                      #3
    |> File.write!(:erlang.term_to_binary(data))        #3
                                                        #3
    {:noreply, state}                                   #3
  end                                                   #3

  def handle_call({:get, key}, _, state) do                #4
    data = case File.read(file_name(key)) do               #4
      {:ok, contents} -> :erlang.binary_to_term(contents)  #4
      _ -> nil                                             #4
```

```
      end                                                     #4
                                                              #4
      {:reply, data, state}                                   #4
    end

  defp file_name(key) do
    Path.join(@db_folder, to_string(key))
  end
end
```

This is mostly a synthesis of techniques mentioned earlier. First, you set the module attribute `@db_folder` to the hardcoded value of the database folder. As explained in section 2.3.6, this works as a compile-time constant, allowing you to encode the knowledge about the database folder in a single place in code.

The database server is locally registered under a name; this keeps things simple and relieves you from passing around the `Todo.Database` pid. Of course, a downside is that you can run only one instance of the database process.

It's worth noting that the `store` request is a cast, whereas `get` is a call. In this implementation, I decided to turn `store` into a cast because the client isn't interested in a response. Using casts promotes scalability of the system because the caller issues a request and goes about its business.

A huge downside of a cast is that the caller can't know whether the request was successfully handled. In fact, the caller can't even be sure that the request reached the target process. This is a property of casts. Casts promote overall availability by allowing client processes to move on immediately after a request is issued. But this comes at the cost of consistency, because you can't be confident about whether a request has succeeded.

In this example, you'll start with the `store` request being a cast. This makes the entire system more scalable and responsive, with the downside being that you can't guarantee that all changes have been persisted.

During initialization, you use `File.mkdir_p!/1` to create the specified folder if it doesn't exist. The exclamation mark at the end of the name indicates a function that raises an error if the folder can't be created for some reason.

The data is stored by encoding the given term to the binary and then persisting it to the disk. Data fetching is an inverse of storing. If the given file doesn't exist on the disk, you return nil.

7.3.2 Using the database

With the database process in place, it's time to use it from your existing system. You have to do three things:

1. Ensure that a database process is started
2. Persist the list on every modification
3. Try to fetch the list from disk during the first retrieval

To start the server, you'll plug into the `Todo.Cache.init/1` function. This is a quick hack, but it's sufficient for the moment. The modification is provided next.

Listing 7.8. Starting the database (persistable_todo_cache/lib/todo/cache.ex)

```
defmodule Todo.Cache do
  ...

  def init(_) do
    Todo.Database.start()
    {:ok, %{}}
  end

  ...
end
```

Here you use the persist subfolder of the current folder as the place to store data.

Storing the data

Next you have to persist the list after it's modified. Obviously, this must be done from the to-do server. But remember that the database's `store` request requires a key. For this purpose, you'll use the to-do list name. As you may recall, this name is currently maintained only in the to-do cache, so you must

propagate it to the to-do server as well. This means extending the to-do server state to be in the format `{list_name, todo_list}`. The code isn't shown here, but these are the corresponding changes:

- `Todo.Server.start` now accepts the to-do list name and passes it to `GenServer.start/2`.
- `Todo.Server.init/1` uses this parameter and keeps the list name in the process state.
- `Todo.Server.handle` callbacks are updated to work with the new state format.

While starting the new to-do server, the cache process passes the list name.

After these modifications, the to-do server knows its own name. Now it's trivial to persist the data, as shown in the following listing.

Listing 7.9. Persisting the data(persistable_todo_cache/lib/todo/server.ex)

```
defmodule Todo.Server do
  ...

  def handle_cast({:add_entry, new_entry}, {name, todo_list}) do
    new_list = Todo.List.add_entry(todo_list, new_entry)
    Todo.Database.store(name, new_list)              #1
    {:noreply, {name, new_list}}
  end

  ...
end
```

You can immediately test whether this works. Run `iex -S mix`, and try the following:

```
iex(1)> {:ok, cache} = Todo.Cache.start()

iex(2)> bobs_list = Todo.Cache.server_process(cache, "bobs_list")

iex(3)> Todo.Server.add_entry(
          bobs_list,
          %{date: ~D[2023-12-19], title: "Dentist"}
        )
```

If all goes well, there should be a file named persist/bobs_list on the disk.

Reading the data

All that's left to do is to read the data from the disk when the server is started. The first idea that comes to mind is to perform this in the `init/1` callback:

```
def init(name) do
  todo_list = Todo.Database.get(name) || Todo.List.new()
  {:ok, {name, todo_list}}
end
```

Here you try to fetch the data from the database, and you resort to the empty list if there's nothing on disk.

While this approach would work, in general you should be careful about possibly long-running `init/1` callbacks. Recall that `GenServer.start` returns only after the process has been initialized. Consequently, a long-running `init/1` function will cause the client (starter) process to block. In this case, a long initialization of a to-do server will block the cache process. And since cache process is used by many clients, this can in turn block a larger part of the system.

Thankfully, `GenServer` comes with a solution to this problem, by allowing you to split the initialization in two phases: one which blocks the client process, and another one which can be performed after the `GenServer.start` invocation in the client has finished.

To do this, `init/1` must return the result in the shape of `{:ok, initial_state, {:continue, some_arg}}`. In this case, the following things happen:

- The initial state of the server process is set.
- The `GenServer.start` invocation in the caller process is unblocked.
- The `handle_continue` callback is invoked in the server process. The callback receives the provided argument (from the `{:continue, some_arg}` tuple) and the server state.

The `handle_continue/2` function is the first callback invoked immediately

after `init/1`. Therefore it can be used as the second phase of the process initialization.

At the time `handle_continue/2` is invoked, the `GenServer.start` invocation in the client process has already finished. Therefore, this phase of the initialization doesn't block the client anymore, so its a fitting place for performing a potentially longer initialization work, such as reading from the database. The idea can be seen in the following listing:

Listing 7.10. Two-phase initialization (persistable_todo_cache/lib/todo/server.ex)

```
defmodule Todo.Server do
  ...

  def init(name) do
    {:ok, {name, nil}, {:continue, :init}}                      #1
  end

  def handle_continue(:init, {name, nil}) do
    todo_list = Todo.Database.get(name) || Todo.List.new()  #2
    {:noreply, {name, todo_list}}
  end

  ...
end
```

The execution of `init/1` is kept as short as possible, and the to-do list is set to `nil`. There's no point in setting it to anything else, because it's going to be overwritten in `handle_continue/2`, which is the first callback invoked immediately after `init/2`.

In any case, the to-do server now reads data from the database on creation. You can immediately test this. If you have the previous shell session open, close it, and start the new one. Then try the following:

```
iex(1)> {:ok, cache} = Todo.Cache.start()

iex(2)> bobs_list = Todo.Cache.server_process(cache, "bobs_list")

iex(3)> Todo.Server.entries(bobs_list, ~D[2023-12-19])
[%{date: ~D[2023-12-19], id: 1, title: "Dentist"}]
```
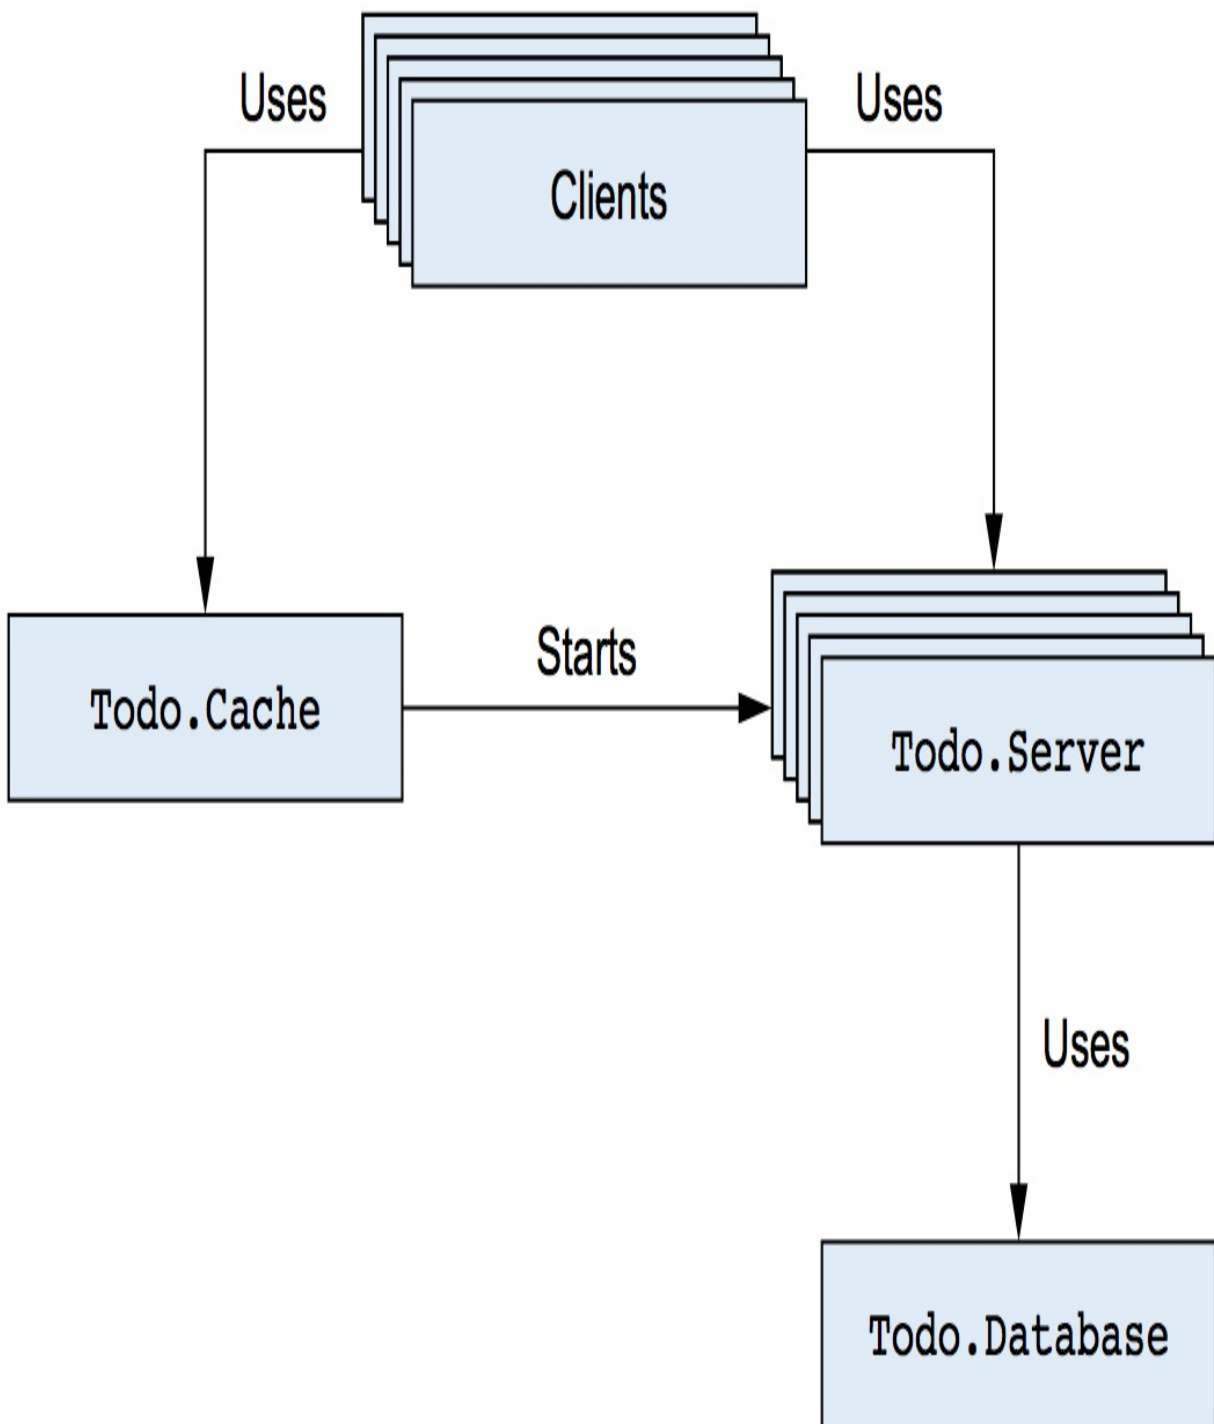
As you can see, your to-do list contains data, which proves that deserialization works.

7.3.3 Analyzing the system

Let's analyze how the new version of the system works. The process interaction is presented in figure 7.2.

Figure 7.2. Process dependencies

You introduced just one process, but it can have a negative impact on the entire system. Recall that the database performs term encoding/decoding and, even worse, disk I/O operations. Depending on the load and list sizes, this can affect performance badly. Let's recall all the places where database

requests are issued:

```
defmodule Todo.Server do
  ...

  def handle_continue(:init, {name, nil}) do
    todo_list = Todo.Database.get(name) || Todo.List.new()  #1
    ...
  end

  ...

  def handle_cast({:add_entry, new_entry}, {name, todo_list}) do
    ...
    Todo.Database.store(name, todo_list)       #2
    ...
  end

  ...
end
```

The `store` request may not seem problematic from the client-side perspective, because it's an asynchronous cast. A client issues a `store` request and then goes about its business. But if requests to the database come in faster than they can be handled, the process mailbox will grow and increasingly consume memory. Ultimately, the entire system may experience significant problems, resulting in the possible termination of the BEAM OS process.

The `get` request can cause additional problems. It's a synchronous call, so the to-do server waits while the database returns the response. While it's waiting for the response, this to-do server can't handle new messages.

It's worth repeating that the synchronous call won't block indefinitely. Recall that `GenServer.call` has a default timeout of five seconds, and you can configure it to be less for better responsiveness. Still, when a request times out, it isn't removed from the receiver's mailbox. A request is a message that's placed in the receiver's mailbox. A timeout means you give up waiting on the response, but the message remains in the receiver's mailbox and will be processed at some point.

7.3.4 Addressing the process bottleneck

There are many approaches to addressing the bottleneck introduced by the singleton database process. Here we'll discuss a few of them.

Bypassing the process

The simplest possible way to eliminate the process bottleneck is to bypass the process. You should ask yourself — does this need to be a process, or can it be a plain module?

There are various reasons for running a piece of code in a dedicated server process:

- The code must manage a long-living state.
- The code handles a kind of a resource that can and should be reused between multiple invocations, such as a TCP connection, database connection, file handle, and so on.
- A critical section of the code must be synchronized. Only one instance of this code may be running in any moment.

If none of these conditions are met, you probably don't need a process and can run the code in client processes, which will completely eliminate the bottleneck and promote parallelism and scalability.

In the current code, you could indeed store to the file directly from the to-do server process. All operations on the same list are serialized in the same process, so there are no race conditions. But the problem with this approach is that concurrency is unbound. If you have 100,000 simultaneous clients, then you'll issue that many concurrent I/O operations, which may negatively affect the entire system.

Handling requests concurrently

Another option is to keep the database process and make it handle database operations concurrently. This is useful when requests depend on a common state but can be handled independently. The idea is illustrated in figure 7.3.

Figure 7.3. Handling requests concurrently
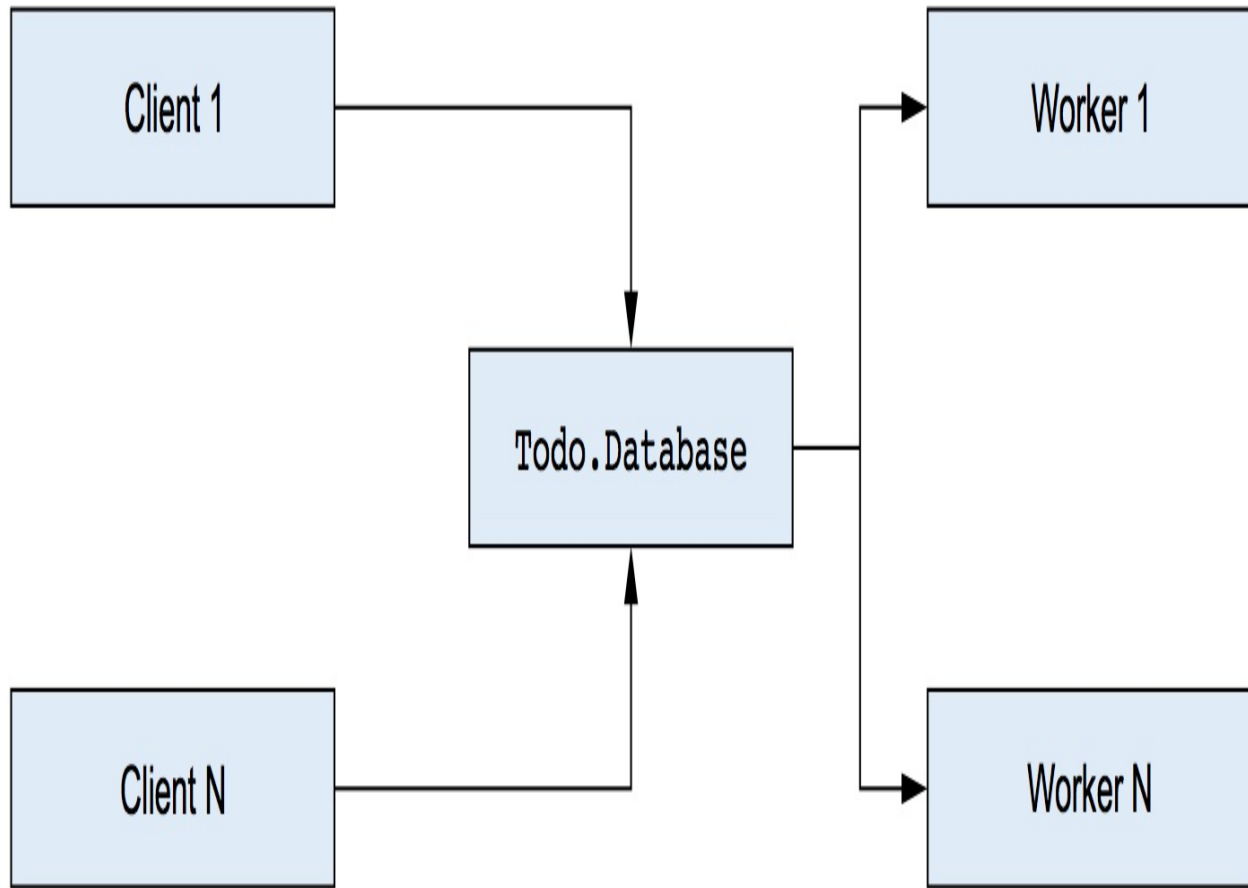
As you can see, each request is still serialized through the central server process, but this server process spawns one-off worker processes that perform the actual request handling. If you keep the code in the database process short and fast, you'll get to keep a high degree of scalability with many workers running concurrently.

To implement this, you must run each database operation in a spawned one-off process. For casts, this means transforming the body of the handler:

```
def handle_cast({:store, key, data}, state) do
  spawn(fn ->
    key                                         #1
    |> file_name()                              #1
    |> File.write!(:erlang.term_to_binary(data))  #1
  end)

  {:noreply, state}
```

```
end
```

The handler function spawns the new worker process and immediately returns. While the worker is running, the database process can accept new requests.

For synchronous calls, this approach is slightly more complicated because you have to return the response from the spawned worker process:

```
def handle_call({:get, key}, caller, state) do
  spawn(fn ->        #1
    data = case File.read(file_name(key)) do
      {:ok, contents} -> :erlang.binary_to_term(contents)
      _ -> nil
    end

    GenServer.reply(caller, data)     #2
  end)

  {:noreply, state}           #3
end
```

The server process spawns another worker process and then returns `{:noreply, state}`, indicating to `GenServer` that you won't reply at this point. In the meantime, the spawned process handles the request and reports back to the caller with `GenServer.reply/2`. This is one situation where you need to use the second argument of `handle_call/3`: the caller pid and the unique ID of the request. This information is used in the spawned process to send the response message to the caller.

This technique keeps the processing in the database process short while still allowing concurrent execution of database operations. This approach has the same drawbacks as the previous idea. The concurrency is still unbound, so too many simultaneous clients might overload the disk I/O to the point where the entire system becomes unresponsive.

Limiting concurrency with pooling

A typical remedy for this problem is to introduce pooling. For example, your database process might create three worker processes and keep their pids in

its internal state. When a request arrives, it's delegated to one of the worker processes, perhaps in a round-robin fashion or with some other load-distribution strategy. The idea is presented in figure 7.4.

Figure 7.4. Pooling database operations

All requests still arrive at the database process first, but they're quickly forwarded to one of the workers. Essentially, this technique keeps the

concurrency level under control, and it works best when dealing with resources that can't handle unbound concurrency.

This approach will work correctly in this example, so it's the one you'll use. In a different situation, some other approach might work better. The point of this analysis is to illustrate how you can think in terms of processes. Always keep in mind that multiple processes run concurrently, whereas a single process handles requests sequentially. If computations can safely run in parallel, you should consider running them in separate processes. In contrast, if an operation must be synchronized, you'll want to run it in a single process.

Database connection pool

In this example, increasing the number of concurrent disk-based operations doesn't yield significant improvements. In this sense, the optimizations serve more as a didactic example than an efficient solution. But in real life, you'd probably talk to a database that's able to handle multiple concurrent requests efficiently. In such a case, you'd typically need to constrain the number of simultaneous database operations. And this is the purpose of a pool of processes.

There's no need to implement such a pool yourself. A couple of generic pool libraries are available for the Elixir/Erlang ecosystem, one of the more popular being Poolboy (https://github.com/devinus/poolboy). Depending on which database library you're using, you'll either need to combine it with Poolboy (or another pooling solution), or this will be done by the library (as is the case, for example, with the Ecto library https://github.com/elixir-lang/ecto). In chapter 11, when you learn how to manage application dependencies, you'll replace the custom implementation with Poolboy.

7.3.5 Exercise: pooling and synchronizing

Now it's time for you to practice a bit. This exercise introduces pooling and makes the database internally delegate to three workers that perform the actual database operations. Moreover, there should be per-key (to-do list name) synchronization on the database level. Data with the same key should always be treated by the same worker.

Here are some pointers for doing this:

- Start with the existing solution, and migrate it gradually. Of the existing code, the only thing that needs to change is the `Todo.Database` implementation. You don't have to touch any of the other modules.
- Introduce a `Todo.DatabaseWorker` module. It will be almost a copy of the current `Todo.Database`, but the process must not be registered under a name, because you need to run multiple instances.
- `Todo.DatabaseWorker.start` should receive the database folder as its argument, and pass it as the second argument to `GenServer.start/2`. This argument is received in the `init/1` callback, and it should be stored in the worker state.
- `Todo.Database` will receive a significant rewrite, but its interface must remain the same. This means it still implements a locally registered process that's used via the functions `start/0`, `get/1`, and `store/2`.
- During the `Todo.Database` initialization, start three workers and store their pids in a map, using zero-based indexes as keys.
- In `Todo.Database`, implement a single request, `choose_worker`, that will return a worker's pid for a given key.
- `choose_worker` should always return the same worker for the same key. The easiest way to do this is to compute the key's numerical hash and normalize it to fall in the range [0, 2]. This can be done by calling `:erlang.phash2(key, 3)`.
- The interface functions `get` and `store` of `Todo.Database` internally call `choose_worker` to obtain the worker's pid and then forward to interface functions of `DatabaseWorker` using the obtained pid as the first argument.

Always try to work in small steps, and test as often as possible. For example, once you implement `Todo.DatabaseWorker`, you can immediately start `iex -S mix` and try it in isolation.

The same goes for `Todo.Database`. First you can initialize the state without implementing a request handler. Call `IO.inspect` from `init/1` to verify that the state is correct. Then implement `choose_worker`, and test that it works in the shell. Finally, add interface functions and test the entire system.

How can you be sure that requests for different keys are running in different

processes? You can use `IO.inspect` and, from within the worker, print the pid and the key using something like `IO.inspect "#{inspect(self())}: storing #{inspect(key)}"`. Use `IO.inspect` extensively. It's your friend and can help you significantly during development.

If you get stuck, the complete solution is in the todo_cache_pooling folder. Make sure you understand the solution, because you'll continue extending this version in subsequent chapters.

7.4 Reasoning with processes

You've now seen various examples of server processes in action. The point of these examples has been to demonstrate how to reason about an involved concurrent system.

From within, a server process is a sequential program that accepts and handles requests, optionally managing internal state. From the outside, it's a concurrent agent that exposes a well-defined communication interface.

Another way to look at server processes is to think of them as services. Each process is like a small service that's responsible for a single task. In the to-do example, there's a to-do server that handles a distinct to-do list. Different lists are handled by different to-do servers, which makes the system more efficient. But a single list is always handled by the same process, which eliminates race conditions and keeps consistency. The to-do cache is a service that maps to-do names to corresponding to-do servers. Finally, the database process is a service that handles database requests. Internally, it distributes the work over a limited pool of workers, making sure the same item is always handled by the same worker.

Those services (processes) are mostly independent, but in some cases they need to cooperate. For this purpose, you can use calls and casts. Obviously, when a client needs a response, you should use calls. But even when a response isn't needed, calls can sometimes be a better fit. The main problem with a cast is that it's a fire-and-forget kind of request, so the caller doesn't get any guarantees. You can't be sure that the request has reached the target, and you most certainly don't know about its outcome.

Essentially, both types have benefits and downsides. Casts promote system responsiveness (because a caller isn't blocked) at the cost of reduced consistency (because a caller doesn't know about the outcome of the request). On the other hand, calls promote consistency (a caller gets a response) but reduce system responsiveness (a caller is blocked while waiting for a response).

Finally, calls can also be used to apply back-pressure to client processes. Because a call blocks a client, it prevents the client from generating too much work. The client becomes synchronized with the server and can never produce more work than the server can handle. In contrast, if you use casts, clients may overload the server, and requests may pile up in the message box and consume memory. Ultimately, you may run out of memory, and the entire system may be terminated.

Which approach is a better fit depends on the specific situation and circumstances. If you're unsure, it's probably better to start with a call, because it's more consistent. You can then consider switching to casts in places where you establish that calls hurt performance and system responsiveness.

7.5 Summary

- When a system needs to perform various tasks, it's often beneficial to run different tasks in separate processes. Doing so promotes the scalability and fault-tolerance of the system.
- A process is internally sequential and handles requests one by one. A single process can thus keep its state consistent, but it can also cause a performance bottleneck if it serves many clients.
- Carefully consider calls versus casts. Calls are synchronous and therefore block the caller. If the response isn't needed, casts may improve performance at the expense of reduced guarantees, because a client process doesn't know the outcome.
- You can use `mix` projects to manage projects that consist of multiple modules.

8 Fault-tolerance basics

This chapter covers

- Runtime errors
- Errors in concurrent systems
- Supervisors

Fault-tolerance is a first-class concept in BEAM. The ability to develop reliable systems that can operate even when faced with runtime errors is what brought us Erlang in the first place.

The aim of fault-tolerance is to acknowledge the existence of failures, minimize their impact, and ultimately recover without human intervention. In a sufficiently complex system, many things can go wrong. Occasional bugs will happen, components you're depending on may fail, and you may experience hardware failures. A system may also become overloaded and fail to cope with an increased incoming request rate. Finally, if a system is distributed, you can experience additional issues such as a remote machine becoming unavailable, perhaps due to a crash or a broken network link.

It's hard to predict everything that can go wrong, so it's better to face the harsh reality that anything can fail. Regardless of which part of the system happens to fail, it shouldn't take down the entire system; you want to be able to provide at least some service. For example, if the database server becomes unreachable, you can still serve data from the cache. You might even queue incoming store requests and try to resolve them later, when the connection to the database is reestablished.

Another thing you need to do is detect failures and try to recover from them. In the previous example, the system may try to reconnect to the database until it succeeds, and then resume providing full service.

These are the properties of a resilient, self-healing system. Whatever goes wrong (and remember, anything can go wrong), the system should keep

providing as much service as possible and fully recover as soon as possible.

Such thinking significantly changes the approach to error handling. Instead of obsessively trying to reduce the number of errors, your priority should be to minimize their effects and recover from them automatically. In a system that has to run continuously, it's better to experience many isolated errors than to encounter a single error that takes down the entire system.

It's somewhat surprising that the core tool for error handling is concurrency. In the BEAM world, two concurrent processes are completely separated; they share no memory, and a crash in one process can't by default compromise the execution flow of another. Process isolation allows you to confine the negative effects of an error to a single process or a small group of related processes, which keeps most of the system functioning normally.

Of course, when a process crashes, you'll usually want to detect this state and do something about it. In this chapter, you'll learn the basic techniques of detecting and handling errors in a concurrent system. Then, in chapter 9, you'll expand on this knowledge and implement fine-grained error isolation. Let's start with a bit of theory about runtime errors.

8.1 Runtime errors

In previous chapters, I loosely mentioned that in various situations an error is raised. One of the most common examples is a failed pattern match. If a match fails, an error is raised. Another example is a synchronous `GenServer.call`. If the response message doesn't arrive in a given time interval (five seconds by default), a runtime error happens. There are many other examples, such as invalid arithmetic operations (such as division by zero), invocation of a nonexistent function, and explicit error signaling.

When a runtime error happens, execution control is transferred up the call stack to the error-handling code. If you didn't specify such code, the process where the error happened is terminated. All other processes by default run unaffected.

8.1.1 Error types

BEAM distinguishes three types of runtime errors: *errors, exits,* and *throws.* Here are some typical examples of errors:

```
iex(1)> 1/0          #1
** (ArithmeticError) bad argument in arithmetic expression

iex(1)> Module.nonexistent_function()      #2
** (UndefinedFunctionError) function Module.nonexistent_function/
   undefined or private

iex(1)> List.first({1,2,3})        #3
** (FunctionClauseError) no function clause matching in List.firs
```

You can also *raise* your own error by using the `raise/1` macro, passing an error string:

```
iex(1)> raise("Something went wrong")
** (RuntimeError) Something went wrong
```

If your function explicitly raises an error, you should append the `!` character to its name. This is a convention used in Elixir standard libraries. For example, `File.open!` raises an error if a file can't be opened:

```
iex(1)> File.open!("nonexistent_file")
** (File.Error) could not open non_existing_file: no such file or
```

In contrast, `File.open` (notice the lack of `!`) just returns the information that the file couldn't be opened:

```
iex(1)> File.open("nonexistent_file")
{:error, :enoent}
```

Notice that in this snippet there's no runtime error. `File.open` returns a result, which the caller can handle in some way.

Another type of a runtime error is the *exit*, which is used to deliberately terminate a process. To exit the current process, you can call `exit/1`, providing an exit reason:

```
iex(2)> spawn(fn ->
          exit("I'm done")        #1
          IO.puts("This doesn't happen")
        end)
```

The exit reason is an arbitrary term that describes why you're terminating the process. As you'll see later, it's possible for some other process to detect a process crash and obtain this exit reason.

The final runtime error type is a *throw*. To issue a throw, you can call `throw/1`:

```
iex(3)> throw(:thrown_value)
** (throw) :thrown_value
```

The purpose of throws is to allow nonlocal returns. As you saw in chapters 3 and 4, Elixir programs are organized in many nested function calls. In particular, loops are implemented as recursions. The consequence is that there are no constructs such as `break`, `continue`, and `return`, which you've probably seen in other languages. When you're deep in a loop, it's not trivial to stop the loop and return a value, and throws can help with this. You can throw a value and catch it up the call stack. But using throws for control flow is hacky, somewhat reminiscent of `goto`, and you should avoid this technique as much as possible.

8.1.2 Handling errors

It is, of course, possible to intercept any kind of runtime error (error, exit, or throw) and do something about it. The main tool for this is the `try` expression. Here's how to run some code and catch errors:

```
try do
  ...
catch error_type, error_value ->
  ...
end
```

This works much like what you've probably seen in other languages. The code in the `do` block is executed, and, if an error happens, execution is transferred to the `catch` block.

Notice that two things are specified in the `catch`. The `error_type` will contain an atom `:error`, `:exit`, or `:throw`, indicating the type of error that

has occurred. The `error_value` will contain error-specific information such as a value that was thrown or an error that was raised.

Let's play with this a bit by writing a helper lambda to make it easier to experiment with errors:

```
iex(1)> try_helper = fn fun ->
          try do
            fun.()
            IO.puts("No error.")

          catch type, value ->
            IO.puts("""
              Error
                #{inspect(type)}
                #{inspect(value)}
              """)
          end
        end
```

This helper lambda takes a function as its argument, calls this function in a `try`, and reports the type of error and the corresponding value. Since the output spans multiple lines, the heredoc syntax (`"""`) is used, which has briefly mentioned in chapter 2.

Let's try it out:

```
iex(2)> try_helper.(fn -> raise("Something went wrong") end)
Error
  :error       #1
  %RuntimeError{message: "Something went wrong"}  #2
```

Notice how the string message is wrapped in a `RuntimeError` struct. This is an Elixir-specific decoration done from within the `raise/1` macro. If you want to raise a plain, undecorated error, you can use Erlang's `:erlang.error/1` and provide an arbitrary term. The resulting error value will be the term you've raised.

If you attempt to throw a value, you'll get a different error type:

```
iex(3)> try_helper.(fn -> throw("Thrown value") end)
Error
```

```
  :throw
  "Thrown value"
```

Calling `exit/1` has its own type:

```
iex(4)> try_helper.(fn -> exit("I'm done") end)
Error
  :exit
  "I'm done"
```

Remember that in Elixir, everything is an expression that has a return value. With `try`, the return value is the result of the last executed expression — either from the `do` block or, if an error was raised, from the `catch` block:

```
iex(5)> result =
          try do
            throw("Thrown value")
          catch type, value -> {type, value}
          end

iex(6)> result
{:throw, "Thrown value"}
```

It's also worth noting that the `type` and `value` specified in the `catch` block are patterns. If you want to handle a specific type of error, you can do this by providing corresponding patterns.

For example, let's say you want to immediately return a value from inside a deep nested loop. You could invoke the following:

```
throw({:result, some_result})
```

Then, somewhere up the call stack, you would handle this particular throw:

```
try do
  ...
catch
  :throw, {:result, x} -> x
end
```

In this example, you only match for a specific runtime error: a throw in the form `{:result, x}`. If anything else is raised, you won't catch it, and an error will be propagated further up the call stack. If the error isn't handled,

the process terminates.

Because `catch` is a pattern match, multiple clauses can be specified, just as you've seen with `case` and `receive` expressions:

```
try do
  ...

catch
  type_pattern_1, error_value_1 ->
    ...

  type_pattern_2, error_value_2 ->
    ...

  ...
end
```

The block under the first pattern that matches a raised error is invoked, and the result of the last expression is returned.

If you want to catch anything, you can use the `type, value` pattern, or `_, _` if you're not interested in values. These patterns will handle any error that can occur.

It's also possible to specify code that should always be executed after the `try` block, regardless of whether an error was raised:

```
iex(7)> try do
          raise("Something went wrong")
        catch
          _,_ -> IO.puts("Error caught")
        after
          IO.puts("Cleanup code")    #1
        end

Error caught
Cleanup code
```

Because it's always executed, the `after` block is useful for cleaning up resources — for example, to close an open file.

It's worth noting that the `after` clause doesn't affect the result of the entire

`try` expression. The result of `try` is the result of the last expression either from the `do` block or from the corresponding `catch` block if something was caught.

Try and tail calls

You may recall the tail-call optimization from chapter 3. If the last thing a function does is call another function (or itself), then a simple jump will occur without a stack push. This optimization isn't possible if the function call resides in a `try` expression. This is fairly obvious, because the last thing a function does is a `try` block, and it won't finish until its `do` or `catch` block is done. Consequently, whatever is called in `try` isn't the last thing a function does and is therefore not available for tail-call optimization.

There's much more to signaling and handling runtime errors. Elixir provides some abstractions on top of this basic mechanism. You can define custom errors via a `defexception` macro (see https://hexdocs.pm/elixir/Kernel.html#defexception/1) and handle them in a slightly more elegant fashion. The `try` special form also has a couple other features we haven't discussed. You should definitely research the official `try` documentation (https://hexdocs.pm/elixir/Kernel.SpecialForms.html#try/1) as well as the corresponding "Getting Started" section (https://elixir-lang.org/getting-started/try-catch-and-rescue.html).

What I've presented here are the core concepts of runtime errors. All other extensions supported by Elixir eventually boil down to these concepts and have the same properties:

- A runtime error has a type, which can be `:error`, `:exit`, or `:throw`.
- A runtime error also has a value, which can be any arbitrary term.
- If a runtime error isn't handled, the corresponding process will terminate.

Compared to languages such as C++, C#, Java, and JavaScript, there's much less need to catch runtime errors. A more common idiom is to let the process crash and then do something about it (usually, restart the process). This approach may seem hacky, but there's reasoning behind it. In a complex system, most bugs are flushed out in the testing phase. The remaining bugs

mostly fall into a so-called *Heisenbug category* — unpredictable errors that occur irregularly in special circumstances and are hard to reproduce. The cause of such errors usually lies in corruptness of the state. Therefore, a reasonable remedy for such errors is to let the process crash and start another one.

This may help, because you're getting rid of the process state (which may be corrupt) and starting with a clean state. In many cases, doing so resolves the immediate problem. Of course, the error should be logged so you can analyze it later and detect the root cause. But in the meantime, you can recover from an unexpected failure and continue providing service.

Don't worry if this discussion seems vague. This approach to error handling, also known as *letting it crash*, will be explained in detail throughout this chapter and the next. In the following section, we'll look at the basics of error handling in concurrent systems.

8.2 Errors in concurrent systems

Concurrency plays a central role in building fault-tolerant, BEAM-based systems. This is due to the total isolation and independence of individual processes. A crash in one process won't affect the others (unless you explicitly want it to).

Here's a quick demonstration:

```
iex(1)> spawn(fn ->    #1
          spawn(fn ->     #2
            Process.sleep(1000)
            IO.puts("Process 2 finished")
          end)

          raise("Something went wrong")    #3
        end)
```

Running this yields the following output:

```
17:36:20.546 [error] Process #PID<0.116.0> raised an exception
...
Process 2 finished      #2
```

As you can see, the execution of process 2 goes on despite the fact that process 1 crashes. Information about the crash of process 1 is printed to the screen, but the rest of the system — including process 2 and the `iex` shell prompt — runs normally.

Furthermore, because processes share no memory, a crash in one process won't leave memory garbage that might corrupt another process. Therefore, by running independent actions in separate processes, you automatically ensure isolation and protection.

You already benefit from process isolation in this book's example to-do system. Recall the current architecture, shown in figure 8.1.

Figure 8.1. Isolating errors in the to-do system

All the boxes in the figure are BEAM processes. A crash in a single to-do server doesn't affect operations on other to-do lists. A crash in `Todo.Database` doesn't block cached reads that take place in to-do server processes.

Of course, this isolation isn't enough by itself. As you can see in figure 8.1, processes often communicate with each other. If a process isn't running, its clients can't use its services. For example, if the database process goes down, the to-do servers can't query it. What's worse, modifications to the to-do list won't be persisted. Obviously this isn't desirable behavior, and you must have a way of detecting a process crash and somehow recovering from it.

8.2.1 Linking processes

A basic primitive for detecting a process crash is the concept of *links*. If two processes are linked, and one of them terminates, the other process receives an *exit signal* — a notification that a process has crashed.

An exit signal contains the pid of the crashed process and the *exit reason* — an arbitrary Elixir term that provides a description of why the process has terminated. In the case of a normal termination (when the spawned function has finished), the exit reason is the atom `:normal`. By default, when a process receives an exit signal from another process, and that signal is anything other than `:normal`, the linked process terminates as well. In other words, when a process terminates abnormally, the linked process is also taken down.

One link connects exactly two processes and is always bidirectional. To create a link, you can use `Process.link/1`, which connects the current process with another process. More often, a link is created when you start a process. You can do this by using `spawn_link/1`, which spawns a process and links it to the current one.

Let's verify this. In the following example, you again spawn two processes, this time linking them together. Then you take down one process:

```
iex(1)> spawn(fn ->
          spawn_link(fn ->        #1
            Process.sleep(1000)
```

```
        IO.puts("Process 2 finished")
      end)

      raise("Something went wrong")
    end)
```

Not surprisingly, this example gives the following output:

```
17:36:20.546 [error] Process #PID<0.116.0> raised an exception
```

Notice in particular that you don't see the output from process 2. This is because process 1 terminated abnormally, which caused an exit signal to be emitted to process 2.

One process can be linked to an arbitrary number of other processes, and you can create as many links in the system as you want, as shown in figure 8.2.

Figure 8.2. Example of links with multiple processes

This illustrates the transitive nature of process links. In this structure, the crash of a single process will emit exit signals to all of its linked processes. If the default behavior isn't overridden, those processes will crash as well. Ultimately, the entire tree of linked processes will be taken down.

Trapping exits

You may be puzzled by the consequences of links. Earlier I explained how process isolation makes it possible to isolate the effect of a runtime error. Links break this isolation and propagate errors over process boundaries. You can think of a link as a communication channel for providing notifications about process terminations.

Usually you don't want a linked process to crash. Instead, you want to detect the process crash and do something about it. This can be done by *trapping exits*. When a process is trapping exits, it isn't taken down when a linked process crashes. Instead, an exit signal is placed in the surviving process's message queue, in the form of a standard message. A trapping process can receive this message and do something about the crash.

To set up an exit trap, you call `Process.flag(:trap_exit, true)`, which makes the current process trap exit signals. Let's look at how this works:

```
iex(1)> spawn(fn ->
          Process.flag(:trap_exit, true)      #1

          spawn_link(fn -> raise("Something went wrong") end)

          receive do                    #3
            msg -> IO.inspect(msg)      #3
          end                           #3
        end)
```

Here you make the parent process trap exits and then spawn a linked process that will crash. Then you receive a message and print it to the screen. The shell session produces the following output:

```
{:EXIT, #PID<0.118.0>,
 {%RuntimeError{message: "Something went wrong"},
  [{:elixir_eval, :__FILE__, `, [file: 'iex', line: 3]}]}}
```

The general format of the exit signal message is `{:EXIT, from_pid, exit_reason}`, where `from_pid` is the pid of the crashed process and `exit_reason` is an arbitrary term that describes the reason for process termination. If a process is terminated due to a throw or an error, the exit reason is a tuple in the form `{reason, where}`, with `where` containing the stack trace. Otherwise, if a process is terminated due to an exit, the reason is a term provided to `exit/1`.

8.2.2 Monitors

As mentioned earlier, links are always bidirectional. Most of the time, this is exactly what you need, but in some cases unidirectional propagation of a process crash works better. Sometimes you need to connect two processes, A and B, in such a way that process A is notified when B terminates, but not the other way around. In such cases, you can use a *monitor*, which is something like a unidirectional link.

To monitor a process, you use `Process.monitor`:

```
monitor_ref = Process.monitor(target_pid)
```

This makes the current process monitor the target process. The result is a unique reference that identifies the monitor. A single process can create multiple monitors.

If the monitored process dies, your process receives a message in the format `{:DOWN, monitor_ref, :process, from_pid, exit_reason}`. If you want to, you can also stop the monitor by calling `Process.demonitor(monitor_ref)`.

Here's a quick example:

```
iex(1)> target_pid = spawn(fn ->      #1
          Process.sleep(1000)         #1
        end)                          #1

iex(2)> Process.monitor(target_pid)   #2

iex(3)> receive do                    #3
          msg -> IO.inspect(msg)      #3
```

```
      end                               #3
```

```
{:DOWN, #Reference<0.1398266903.3291480065.256365>, :process,   #
  #PID<0.111.0>, :noproc}        #4
```

There are two main differences between monitors and links. First, monitors are unidirectional — only the process that created a monitor receives notifications. In addition, unlike a link, the observer process won't crash when the monitored process terminates. Instead, a message is sent, which you can handle or ignore.

Exits are propagated through GenServer calls

When you issue a synchronous request via `GenServer.call`, if a server process crashes, an exit signal will occur in your client process. This is a simple but very important example of cross-process error propagation. Internally, `GenServer` sets up a temporary monitor that targets the server process. While waiting for a response from the server, if a `:DOWN` message is received, `GenServer` can detect that a process has crashed and raise a corresponding exit signal in the client process.

Links, exit traps, and monitors make it possible to detect errors in a concurrent system. You can introduce a process whose responsibility is to receive links and monitor notifications, and do something when other processes in the system crash. Such processes, called *supervisors*, are the primary tools of error recovery in concurrent systems.

8.3 Supervisors

A supervisor is a generic process that manages the lifecycle of other processes in a system. A supervisor process can start other processes, which are then considered to be its children. Using links, monitors, and exit traps, a supervisor detects possible terminations of any child, and can restart it if needed.

Processes that aren't supervisors are called *workers*. These are the processes that provide the actual services of the system. Your current version of the to-do system consists only of worker processes, such as the to-do cache and to-

do server processes.

If any of the worker processes crashes, perhaps due to a bug, some part of your system will be gone forever. This is where supervisors can help. By running workers under a supervisor, you can ensure that a failing process is restarted, and the service of your system is restored.

In order to do that, you need at least one supervisor process in the system. In Elixir, this can be done using the `Supervisor` module (https://hexdocs.pm/elixir/Supervisor.html). By invoking `Supervisor.start_link/2`, you can start the supervisor process, which then works as follows:

1. The supervisor process traps exits, and then starts the child processes.
2. If at any point in time a child terminates, the supervisor process receives a corresponding exit message and performs corrective actions, such as restarting the crashed process.
3. If the supervisor process terminates, its children are also taken down.

There are two different ways of starting a supervisor. In a basic approach, you invoke the function `Supervisor.start_link`, passing it a list that describes each child to be started under the supervisor, together with some additional supervisor options. Alternatively, you can pass a module defining a callback function that returns this information. We'll start with the basic approach, and explain the second version a bit later.

Let's introduce one supervisor to the to-do system. Figure 8.3 recaps the processes in the system:

Figure 8.3. Processes in the to-do system

- `Todo.Server` — Allows multiple clients to work on a single to-do list
- `Todo.Cache` — Maintains a collection of to-do servers and is responsible for their creation and discovery
- `Todo.DatabaseWorker` — Performs read/write operations on the database
- `Todo.Database` — Manages a pool of database workers, and forwards database requests to them

The to-do cache process is the system entry point. When you start the cache, all the needed processes are started, so the cache can be considered the root of the system. Now we'll introduce a new supervisor process that will supervise the to-do cache process.

8.3.1 Preparing the existing code

Before you start working with the supervisor, you'll need to make a couple of changes to the cache.

First, you'll register the cache process. This will allow you to interact with the process without needing to know its pid.

You'll also need to create a link while starting the to-do cache process. This is required if you want to run the process under a supervisor. Why is the supervisor using links rather than monitors? Because links work in both directions, so the termination of a supervisor means all of its children will be automatically taken down. This, in turn, allows you to properly terminate any part of the system without leaving behind dangling processes. You'll see how this works in this chapter and the next, when you work with finer-grained supervision.

Creating a link to the caller process is as simple as using `GenServer.start_link` in place of `GenServer.start`. While you're at it, you can also rename the corresponding `Todo.Cache` interface function to `start_link`.

Finally, you'll make the `start_link` function take one argument and ignore it. This seems confusing, but it makes starting a supervised process a bit easier. The reasons will be explained later, when we discuss child

specifications.

The changes are shown in the following listing.

Listing 8.1. Changes in the to-do cache (supervised_todo_cache/lib/todo/cache.ex)

```
defmodule Todo.Cache do
  use GenServer

  def start_link(_) do              #1
    GenServer.start_link(__MODULE__, nil, name: __MODULE__)      #2
  end

  def server_process(todo_list_name) do
    GenServer.call(__MODULE__, {:server_process, todo_list_name})
  end

  def init(_) do
    IO.puts("Starting to-do cache.") #4
    ...
  end

  ...
end
```

Notice that you also call `IO.puts/1` from the `init/1` callback for debugging purposes. This debug expression is included in all other `GenServer` callback modules (`Todo.Database`, `Todo.DatabaseWorker`, and `Todo.Server`).

8.3.2 Starting the supervisor process

With these changes in place, you can immediately try to start the supervisor process with to-do cache as its only child. Change the current folder to supervised_todo_cache, and start the shell (`iex -S mix`). Now you can start the supervisor:

```
iex(1)> Supervisor.start_link([Todo.Cache], strategy: :one_for_on

Starting to-do cache.
Starting database server.
Starting database worker.
Starting database worker.
Starting database worker.
```

As you can see from the console output, invoking `Supervisor.start_link/2` caused the to-do cache to start. The cache process then started the database processes.

Let's take a closer look at the invocation of `Supervisor.start_link/2`:

```
Supervisor.start_link(
  [Todo.Cache],              #1
  strategy: :one_for_one    #2
)
```

As the function name hints, `Supervisor.start_link/2` starts a supervisor process and links it to the caller.

The first argument is the list of desired children. More precisely, each element of this list is a child specification that describes how the child should be started and managed. We'll discuss child specifications in detail a bit later. In this simple form, the provided child specification is a module name. In this case, the child is described by some callback function in the `Todo.Cache` module.

When the supervisor process is started, it will go through this list and start each child according to the specification. In this example, the supervisor will invoke `Todo.Cache.start_link/1`. Once all the children are started, `Supervisor.start_link/2` returns `{:ok, supervisor_pid}`.

The second argument to `Supervisor.start_link/2` is the list of supervisor-specific options. The option `:strategy`, also known as *restart strategy*, is mandatory. This option specifies how a supervisor should handle termination of its children. The `one_for_one` strategy states that if a child terminates, another child should be started in its place. There are a couple of other strategies (for example, "Restart all children if a single child crashes"), and we'll discuss them in chapter 9.

![info icon] **Note**

The term *restart* is used casually here. Technically, a process can't be restarted. It can only be terminated; then another process, powered by the

same module, can be started in its place. The new process has a different pid and doesn't share any state with the old one.

In any case, after `Supervisor.start_link/2` returns, all the required processes in the system are running, and you can interact with the system. For example, you can start one to-do server:

```
iex(2)> bobs_list = Todo.Cache.server_process("Bob's list")
Starting to-do server for Bob's list.
#PID<0.161.0>
```

The cache process is started as the child of the supervisor process, so we say that it's supervised. This means that if the cache process crashes, its supervisor will restart it.

You can quickly verify this by provoking a crash of the cache process. First, you need to get the pid of the cache. As mentioned, the cache is now registered under a name (its own module name), so getting its pid is easily done with the help of `Process.whereis/1`:

```
iex(3)> cache_pid = Process.whereis(Todo.Cache)
#PID<0.155.0>
```

Now, you can kill the process using the `Process.exit/2` function, which accepts a pid and the exit reason, and sends the corresponding exit signal to the given process. The exit reason can be an arbitrary term. Here, you'll use the atom `:kill`, which is treated in a special way. The exit reason `:kill` ensures that the target process is unconditionally taken down, even if the process is trapping exits. Let's see it in action:

```
iex(4)> Process.exit(cache_pid, :kill)
Starting to-do cache.
```

As you can see from the output, the process is immediately restarted. You can also prove that the to-do cache is now a process with a different pid:

```
iex(5)> Process.whereis(Todo.Cache)
#PID<0.164.0>
```

And you can use the new process, just as you did the old one:

```
iex(6)> bobs_list = Todo.Cache.server_process("Bob's list")
Starting to-do server for Bob's list.
#PID<0.167.0>
```

This brief experiment proves some basic fault-tolerance capabilities. After the crash, the system healed itself and resumed the full service.

Names allow process discovery

It's important to explain why you register the to-do cache under a local name. You should always keep in mind that in order to talk to a process, you need to have its pid. In chapter 7, you used a naive approach where you created a process and then passed around its pid. This works fine until you enter the supervisor realm.

The problem is that supervised processes can be restarted. Remember that restarting boils down to starting another process in place of the old one — the new process has a different pid. This means any reference to the pid of the crashed process becomes invalid, identifying a nonexistent process.

That's why registered names are important. They provide a reliable way of finding a process and talking to it, regardless of possible process restarts.

8.3.3 Child specification

In order to manage a child process, a supervisor needs some information, such as

- How should the child be started?
- What should be done if the child terminates?
- What term should be used to uniquely identify each child?

These pieces of information are collectively called the *child specification*. Recall that when invoking `Supervisor.start_link/2`, you sent a list of child specifications. In its basic shape, a specification is a map with a couple of fields configuring the properties of the child.

For example, here's what the specification for the to-do cache could look

like:

```
%{
  id: Todo.Cache,      #1
  start: {Todo.Cache, :start_link, [nil]},     #2
}
```

The `:id` field is an arbitrary term that's used to distinguish this child from any other child of the same supervisor.

The `:start` field is a triplet in the shape of `{module, start_function, list_of_arguments}`. When starting the child, the generic supervisor code will use `apply(module, start_function, list_of_arguments)` to invoke the function described by this tuple. The invoked function must start and link the process.

There are some other fields that you can omit from the specification, in which case some sensible defaults are chosen. We'll discuss some of them later, in chapter 9. You can also refer to the official documentation at https://hexdocs.pm/elixir/Supervisor.html#module-child-specification for more details.

In any case, you can pass the specification map directly to `Supervisor.start_link`. Here's an example:

```
Supervisor.start_link(
  [
    %{
      id: Todo.Cache,
      start: {Todo.Cache, :start_link, [nil]}
    }
  ],
  strategy: :one_for_one
)
```

This will instruct the supervisor to invoke `Todo.Cache.start_link(nil)` to start the child. Recall that you changed `Todo.Cache.start_link` to take one argument (which is ignored), so you need to pass some value, in this example `nil`.

One problem with this approach is that it's error-prone. If something changes

in the implementation of the cache, such as the signature of the start function, you need to remember to adapt the specification in the code starting the supervisor.

To address this issue, `Supervisor` allows you to pass a tuple `{module_name, arg}` in the child specification list. In this case, `Supervisor` will first invoke `module_name.child_spec(arg)` to get the actual specification. This function must return the specification map. The supervisor then proceeds to start the child according to the returned specification.

The `Todo.Cache` module already has `child_spec/1` defined, even though you didn't write it yourself. The default implementation is injected by `use GenServer`. Therefore, you can also start the supervisor in the following way:

```
Supervisor.start_link(
  [{Todo.Cache, nil}],
  strategy: :one_for_one
)
```

As a consequence, `Supervisor` will invoke `Todo.Cache.child_spec(nil)` and start the child according to the returned specification. It's easy to verify what the injected implementation of `child_spec/1` returns:

```
iex(1)> Todo.Cache.child_spec(nil)
%{id: Todo.Cache, start: {Todo.Cache, :start_link, [nil]}}
```

In other words, the generated `child_spec/1` returns a specification that invokes the module's `start_link/1` function with the argument passed to `child_spec/1`. This is precisely why you made `Todo.Cache.start_link` take one argument, even though the argument is ignored:

```
defmodule Todo.Cache do
  use GenServer      #1

  def start_link(_) do    #2
    ...
  end

  ...
end
```

By doing this, you made `Todo.Cache` compatible with the generated `child_spec/1`, which means you can include `Todo.Cache` in the list of children without needing to do any extra work.

If you don't like that approach, you can provide some options to `use` `GenServer` to tweak the output of the generated `child_spec/1`. Refer to the official documentation (https://hexdocs.pm/elixir/GenServer.html#module-how-to-supervise) for more details. If you need even more control, you can simply define `child_spec/1` yourself, which will override the default implementation.

Finally, if you don't care about the argument passed to `child_spec/1`, you can include just the module name in the child specification list. In this case, `Supervisor` will pass the empty list `[]` to `child_spec/1`. Therefore, you can also start `Todo.Cache` like this:

```
Supervisor.start_link(
  [Todo.Cache],
  strategy: :one_for_one
)
```

Before going further, let's recap how supervisor starting works. When you invoke `Supervisor.start_link(child_specs, options)`, the following happens:

1. The new process is started, powered by the `Supervisor` module.
2. The supervisor process goes through the list of child specifications and starts each child, one by one.
3. Each specification is resolved, if needed, by invoking `child_spec/1` from the corresponding module.
4. The supervisor starts the child process, according to the `:start` field of the child specification.

8.3.4 Wrapping the supervisor

So far, you've played with the supervisor in the shell. But in real life, you'll want to work with supervisor in the code. Just like with `GenServer`, it's advised to wrap the `Supervisor` in a module.

The following listing implements the module for your first supervisor:

Listing 8.2. To-do system supervisor (supervised_todo_cache/lib/todo/system.ex)

```
defmodule Todo.System do
  def start_link do
    Supervisor.start_link(
      [Todo.Cache],
      strategy: :one_for_one
    )
  end
end
```

With this simple addition, starting the whole system becomes easy:

```
$ iex -S mix

iex(1)> Todo.System.start_link()

Starting to-do cache.
Starting database server.
Starting database worker.
Starting database worker.
Starting database worker.
```

The name `Todo.System` is chosen to describe the purpose of the module. By invoking `Todo.System.start_link()` you start the entire to-do system, with all the required services, such as the cache and database.

8.3.5 Using a callback module

Another way of starting a supervisor is by providing a callback module. This works similarly to `GenServer`. You develop the module that must implement the `init/1` function. This function must return the list of child specifications and additional supervisor options, such as its strategy.

Here's how you could rewrite `Todo.System` to use this approach:

```
defmodule Todo.System do
  use Supervisor            #1

  def start_link do
```

```
    Supervisor.start_link(__MODULE__, nil)      #2
  end

  def init(_) do              #3
    Supervisor.init([Todo.Cache], strategy: :one_for_one)    #3
  end                         #3

end
```

As with `GenServer`, you start with `use Supervisor` to get some common boilerplate in your module.

The crucial bit happens when you invoke `Supervisor.start_link/2`. Instead of the list of child specifications, you're now passing the callback module. In this case, the supervisor process will invoke the `init/1` function to provide the supervisor specification. The argument passed to `init/1` is the second argument you pass to `Supervisor.start_link/2`.

Finally, in `init/1`, you describe the supervisor with the help of the `Supervisor.init/2` function, passing it the list of children and the supervisor options.

The preceding code is a more elaborate equivalent of `Supervisor.start_link([Todo.Cache], strategy: :one_for_one)`. Clearly, you need more lines of code to get the same effect. On the upside, this approach gives you more control. For example, if you need to perform some extra initialization before starting the children, you can do it in `init/1`. Moreover, the callback module is more flexible with respect to hot-code reloading, allowing you to modify the list of children without needing to restart the entire supervisor.

In most cases, the simple approach of passing the list of child specifications directly will be sufficient. Moreover, as you've seen in the preceding examples, if you wrap the use of `Supervisor` in a dedicated module, it's easy to switch from one approach to the other. Therefore, in this book, you'll exclusively use the simple approach without a callback module.

8.3.6 Linking all processes

At this point, you're supervising the to-do cache process, so you get some basic fault-tolerance. If the cache process crashes, a new process is started, and the system can resume providing the service.

However, there's a problem in your current implementation. When the supervisor restarts the to-do cache, you'll get a completely separate process hierarchy, and there will be a new set of to-do server processes that are in no way related to the previous ones. The previous to-do servers will be unused garbage that's still running and consuming both memory and CPU resources.

Let's demonstrate this issue. First, start the system and request one to-do server:

```
iex(1)> Todo.System.start_link()

iex(2)> Todo.Cache.server_process("Bob's list")
Starting to-do server for Bob's list.
#PID<0.159.0>
```

A cached to-do server isn't started on subsequent requests:

```
iex(3)> Todo.Cache.server_process("Bob's list")
#PID<0.159.0>
```

Check the number of running processes:

```
iex(4)> length(Process.list())
71
```

Now terminate the to-do cache:

```
iex(5)> Process.exit(Process.whereis(Todo.Cache), :kill)
Starting to-do cache.
```

Finally, request a to-do server for Bob's list:

```
iex(6)> Todo.Cache.server_process("Bob's list")
Starting to-do server for Bob's list.
#PID<0.165.0>
```

As you can see, after you restart the to-do cache, retrieving a previously fetched server creates a new process. This isn't surprising because you killed

the previous cache process, which also destroyed the process state.

When a process terminates, its state is released, and the new process starts with the fresh state. If you want to preserve the state, you must handle it yourself; we'll discuss this in chapter 9.

After the cache process is restarted, you have a completely new process that has no notion of what was previously cached. At the same time, your old cache structure (to-do servers) isn't cleaned up. You can see this by rechecking the number of running processes:

```
iex(7)> length(Process.list())
72
```

You have one additional process, which is the previously started to-do server for Bob's list. This obviously isn't good. Terminating a to-do cache destroys its state, so you should also take down all existing to-do servers. This way, you ensure proper process termination.

To do this, you must establish links between processes. Each to-do server must be linked to the cache. Going further, you'll also link the database server to the to-do cache and the database workers to the database server. This will effectively ensure that the entire structure is linked, as illustrated in figure 8.4.

Figure 8.4. Linking all processes in the to-do system
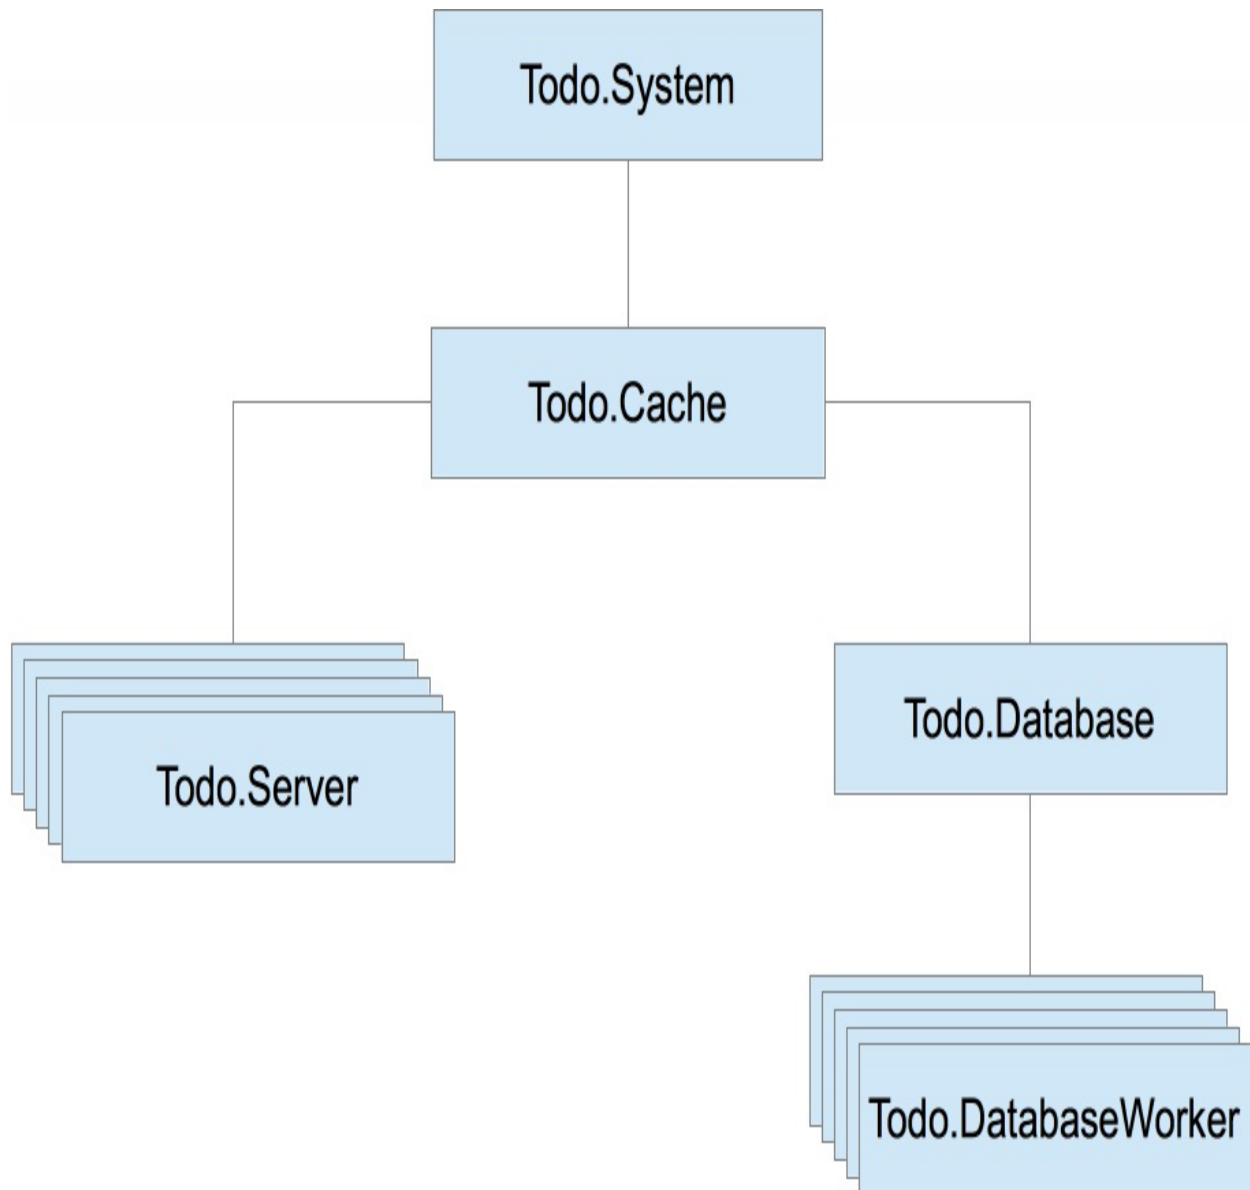
By linking a group of interdependent processes, you can ensure that the crash of one process takes down its dependencies as well. Regardless of which process crashes, links make sure the entire structure is terminated. Because this will lead to the termination of the cache process, it will be noticed by the supervisor, which will start a new system.

With this approach you can detect an error in any part of the system and recover from it without leaving behind dangling processes. On the downside, you're allowing errors to have a wide impact. An error in a single database worker or a single to-do server will take down the entire structure. This is far from perfect, and you'll make improvements in chapter 9.

For now, let's stick with this simple approach and implement the required code. In your present system, you have a to-do supervisor that starts and supervises the cache. You must ensure that the cache is directly or indirectly linked to all other worker processes.

The change is simple. All you need to do is switch from `start` to `start_link` for all the processes in the project. In the corresponding modules, you currently have something like this:

```
def start(...) do
  GenServer.start(...)
end
```

This snippet must be transformed into the following:

```
def start_link(...) do
  GenServer.start_link(...)
end
```

And of course, every `module.start` invocation must be replaced with `module.start_link`. These changes are mechanical, and the code isn't presented here. The complete solution resides in the todo_links folder.

Let's see how the new system works:

```
iex(1)> Todo.System.start_link()

iex(2)> Todo.Cache.server_process("Bob's list")
Starting to-do server for Bob's list.

iex(3)> length(Process.list())
71

iex(4)> Process.exit(Process.whereis(Todo.Cache), :kill)    #1

iex(5)> bobs_list = Todo.Cache.server_process("Bob's list")
Starting to-do server for Bob's list.

iex(6)> length(Process.list())
71      #2
```

When you crash a process, the entire structure is terminated, and the new process starts in its place. Links ensure that dependent processes are

terminated as well, which keeps the system consistent.

8.3.7 Restart frequency

It's important to keep in mind that a supervisor won't restart a child process forever. The supervisor relies on the *maximum* restart frequency, which defines how many restarts are allowed in a given time period. By default, the maximum restart frequency is three restarts in five seconds. You can change these parameters by passing `:max_restarts` and `:max_seconds` options to `Supervisor.start_link/2`. If this frequency is exceeded, the supervisor gives up and terminates itself together with all of its children.

Let's verify this in the shell. First, start the supervisor:

```
iex(1)> Todo.System.start_link()
Starting the to-do cache.
```

Now you need to perform frequent restarts of the to-do cache process:

```
iex(1)> for _ <- 1..4 do
          Process.exit(Process.whereis(Todo.Cache), :kill)
          Process.sleep(200)
        end
```

Here you terminate the cache process and sleep for a short while, allowing the supervisor to restart the process. This is done four times, meaning that in the last iteration, you'll exceed the default maximum restart frequency (three restarts in five seconds).

Here's the output:

```
Starting the to-do cache.      #1
Starting database server.      #1
...                            #1

** (EXIT from #PID<0.149.0>) :shutdown    #2
```

After the maximum restart frequency was exceeded, the supervisor gave up and terminated, taking down the child processes as well.

You may wonder about the reason for this mechanism. When a critical

process in the system crashes, its supervisor tries to bring it back online by starting the new process. If this doesn't help, there's no point in infinite restarting. If too many restarts occur in a given time interval, it's clear that the problem can't be fixed. In this case, the only sensible thing a supervisor can do is give up and terminate itself, which also terminates all of its children.

This mechanism plays an important role in so-called *supervision trees*, where supervisors and workers are organized in a deeper hierarchy that allows you to control how the system recovers from errors. This will be thoroughly explained in the next chapter, where you'll build a fine-grained supervision tree.

8.4 Summary

- There are three types of runtime errors: throws, errors, and exits.
- When a runtime error occurs, execution moves up the stack to the corresponding `try` block. If an error isn't handled, a process will crash.
- Process termination can be detected in another process. To do this, you can use links or monitors.
- Links are bidirectional — a crash of either process is propagated to the other process.
- By default, when a process terminates abnormally, all processes linked to it terminate as well. By trapping exits, you can react to the crash of a linked process and do something about it.
- A supervisor is a process that manages the lifecycle of other processes. It can start, supervise, and restart crashed processes.
- The `Supervisor` module is used to start supervisors and work with them.
- A supervisor is defined by the list of child specifications and the supervision strategy. You can provide these as the arguments to `Supervisor.start_link/2`, or you can implement a callback module.

9 Isolating error effects

This chapter covers

- Understanding supervision trees
- Starting workers dynamically
- "Let it crash"

In chapter 8 you learned about the basic theory behind error handling in concurrent systems based on the concept of supervisors. The idea is to have a process whose only job is to supervise other processes and to restart them if they crash. This gives you a way to deal with all sorts of unexpected errors in your system. Regardless of what goes wrong in a worker process, you can be sure that the supervisor will detect an error and restart the worker.

In addition to providing basic error detection and recovery, supervisors play an important role in isolating error effects. By placing individual workers directly under a supervisor, you can confine an error's impact to a single worker. This has an important benefit: it makes your system more available to its clients. Unexpected errors will occur no matter how hard you try to avoid them. Isolating the effects of such errors allows other parts of the system to run and provide service while you're recovering from the error.

For example, a database error in this book's example to-do system shouldn't stop the cache from working. While you're trying to recover from whatever went wrong in the database part, you should continue to serve existing cached data, thus providing at least partial service. Going even further, an error in an individual database worker shouldn't affect other database operations. Ultimately, if you can confine an error's impact to a small part of the system, your system can provide most of its service all of the time.

Isolating errors and minimizing their negative effects is the topic of this chapter. The main idea is to run each worker under a supervisor, which makes it possible to restart each worker individually. You'll see how this works in the next section, where you start to build a fine-grained supervision

tree.

9.1 Supervision trees

In this section, we'll discuss how to reduce the effect of an error on the entire system. The basic tools are processes, links, and supervisors, and the general approach is fairly simple. You always have to consider what will happen to the rest of the system if a process crashes due to an error, and you should take corrective measures when an error's impact is too wide (when the error affects too many processes).

9.1.1 Separating loosely dependent parts

Let's look at how errors are propagated in the to-do system. Links between processes are depicted in figure 9.1.

Figure 9.1. Process links in the to-do system

As you can see in the diagram, the entire structure is connected. Regardless of which process crashes, the exit signal will be propagated to its linked processes. Ultimately, the to-do cache process will crash as well, and this will be noticed by the `Todo.System`, which will in turn restart the cache process.

This is a correct error-handling approach because you restart the system and don't leave behind any dangling processes. But such a recovery approach is too coarse. Wherever an error happens, the entire system is restarted. In the case of a database error, the entire to-do cache will terminate. Similarly, an error in one to-do server process will take down all the database workers.

This coarse-grained error recovery is due to the fact that you're starting worker processes from within other workers. For example, a database server is started from the to-do cache. To reduce error effects, you need to start individual workers from the supervisor. Such a scheme makes it possible for the supervisor to supervise and restart each worker separately.

Let's see how to do this. First, you'll move the database server so it's started directly from the supervisor. This will allow you to isolate database errors from those that happen in the cache.

Placing the database server under supervision is simple enough. You must remove the call to `Todo.Database.start_link` from `Todo.Cache.init/1`. Then you have to add another child specification when invoking `Supervisor.start_link/2`, as illustrated in the following listing.

Listing 9.1. Supervising database server (supervise_database/lib/todo/system.ex)

```
defmodule Todo.System do
  def start_link do
    Supervisor.start_link(
      [
        Todo.Database,      #1
        Todo.Cache
      ],
      strategy: :one_for_one
    )
  end
end
```

There's one more small change that needs to be done. Just like you did with `Todo.Cache`, you need to adapt `Todo.Database.start_link` to take exactly one argument and ignore it. This will make it possible to rely on the autogenerated `Todo.Database.child_spec/1`, obtained by `use GenServer`.

Listing 9.2. Adapting `start_link` (supervise_database/lib/todo/system.ex)

```
defmodule Todo.Database do
  ...

  def start_link(_) do
    ...
```

```
    end

    ...
end
```

These changes ensure that the cache and the database are separated, as shown in figure 9.2. Running both the database and cache processes under the supervisor makes it possible to restart each worker individually. An error in the database worker will crash the entire database structure, but the cache will remain undisturbed. This means all clients reading from the cache will be able to get their results while the database part is restarting.

Figure 9.2. Separated supervision of database and cache
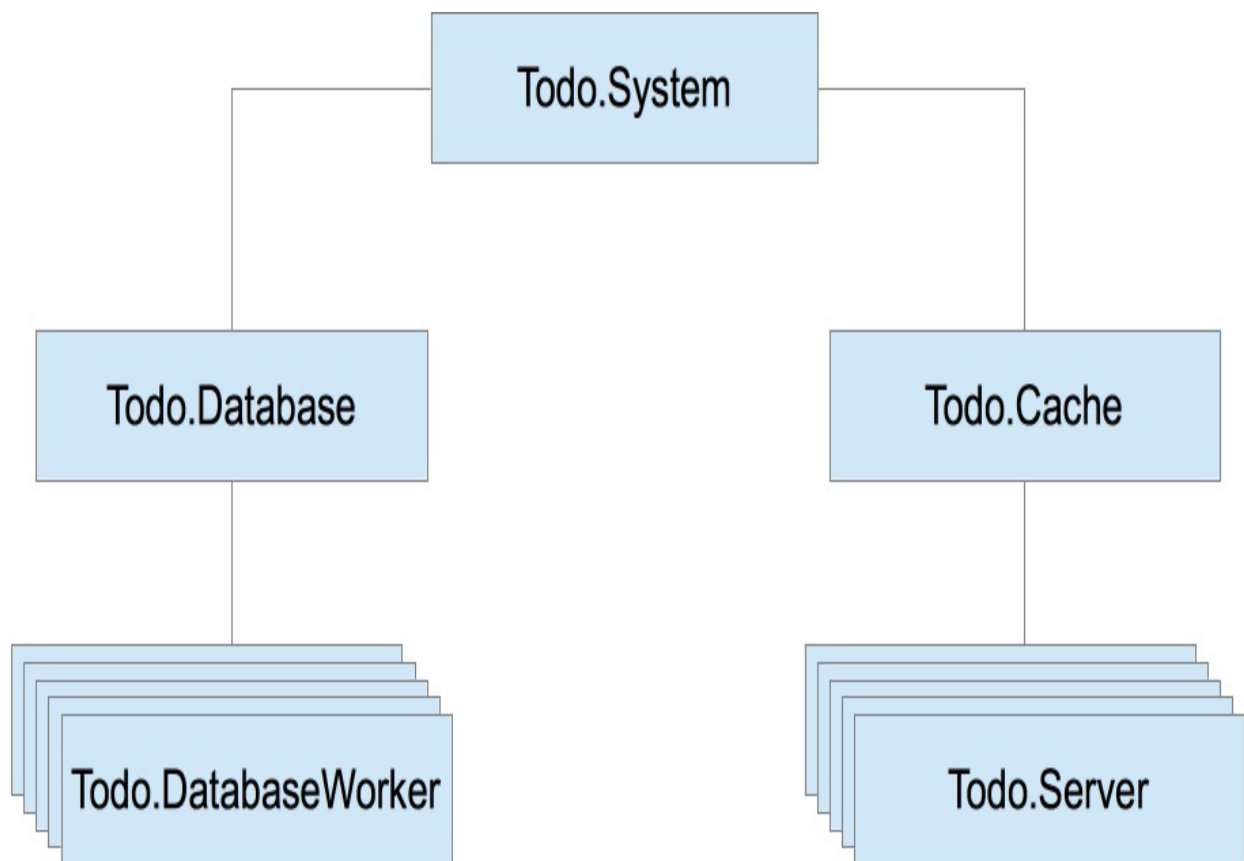
Let's verify this. Go to the supervise_database folder, and start the shell (`iex -S mix`). Then start the system:

```
iex(1)> Todo.System.start_link()
```

```
Starting database server.
Starting database worker.
Starting database worker.
Starting database worker.
Starting to-do cache.
```

Now, kill the database server:

```
iex(2)> Process.exit(Process.whereis(Todo.Database), :kill)

Starting database server.
Starting database worker.
Starting database worker.
Starting database worker.
```

As you can see from the output, only database-related processes are restarted. The same is true if you terminate the to-do cache. By placing both processes under a supervisor, you localize the negative impact of an error. A cache error will have no effect on the database part, and vice versa.

Recall chapter 8's discussion of process isolation. Because each part is implemented in a separate process, the database server and the to-do cache are isolated and don't affect each other. Of course, these processes are indirectly linked via the supervisor, but the supervisor is trapping exit signals, thus preventing further propagation. This is a property of one_for_one supervisors in particular—they confine an error's impact to a single worker and take the corrective measure (restart) only on that process.

Child processes are started synchronously

In this example, the supervisor starts two child processes. It's important to be aware that children are started synchronously, in the order specified. The supervisor starts a child, waits for it to finish, and then moves on to start the next child. When the worker is a GenServer, the next child is started only after the init/1 callback function for the current child is finished.

You may recall from chapter 7 that init/1 shouldn't run for a long time. This is precisely why. If Todo.Database was taking, say, five minutes to start, you wouldn't have the to-do cache available all that time. Always make sure your init/1 functions run quickly, and use the technique mentioned in

chapter 7 (post-init continuation via the `handle_continue/2` callback) when you need more complex initialization.

9.1.2 Rich process discovery

Although you now have some basic error isolation, there's still a lot to be desired. An error in one database worker will crash the entire database structure and terminate all running database operations. Ideally you want to confine a database error to a single worker. This means each database worker has to be directly supervised.

There's one problem with this approach. Recall that in the current version, the database server starts the workers and keeps their pids in its internal list. But if a process is started from a supervisor, you don't have access to the pid of the worker-process. This is a property of supervisors. You can't keep a worker's pid for a long time because that process might be restarted, and its successor will have a different pid.

Therefore, you need a way to give symbolic names to supervised processes and access each process via this name. When a process is restarted, the successor will register itself under the same name, which will allow you to reach the right process even after multiple restarts.

You could use registered names for this purpose. The problem is that names can only be atoms, and in this case you need something more elaborate that will allow you to use arbitrary terms, such as `{:database_worker, 1}`, `{:database_worker, 2}`, and so on. What you need is a process registry that maintains a key-value map, where the keys are names and the values are pids. A process registry differs from standard local registration in that names can be arbitrarily complex.

Every time a process is created, it can register itself to the registry under a name. If a process is terminated and restarted, the new process will re-register itself. Having a registry will give you a fixed point where you can discover processes (their pids). The idea is illustrated in figure 9.3.

Figure 9.3. Discovering processes through a registry

In step 1, the worker process registers itself, usually during initialization. Some time later, the client process will query the registry for the pid of the desired worker. The client can then issue a request to the server process.

Elixir's standard library includes the implementation of a process registry in the `Registry` module. This module allows you to associate a process with one or more arbitrary complex keys, and then find the process (its pid) by doing a key-based lookup.

Let's look at a couple of examples. The process registry is itself a process. You can start it by invoking `Registry.start_link/1`:

```
iex(1)> Registry.start_link(name: :my_registry, keys: :unique)
```

The single argument is a keyword list of registry options. The mandatory options are `:name` and `:keys`.

The `:name` option is an atom, and it specifies the name of the registry process. You'll use this name to interact with the registry.

The `:keys` option can either be `:unique` or `:duplicate`. In a unique registry, names are unique—only one process can be registered under any key. This is useful when you want to assign a unique role to processes. For example, in your system, only one process could be associated with `{:database_worker, 1}`. In contrast, in a duplicate registry, multiple processes can have the same name. Duplicate registry is useful in scenarios where a single publisher process needs to send notifications to a dynamic number of subscriber

processes, which tend to come and go over time.

Once you have the registry started, you can register a process under some key. Let's try it out. You'll spawn a mock `{:database_worker, 1}` process that waits for a message and then prints it to the console:

```
iex(2)> spawn(fn ->
    Registry.register(:my_registry, {:database_worker, 1}, nil)

    receive do
      msg -> IO.puts("got message #{inspect(msg)}")
    end
  end)
```

The crucial bit happens when invoking `Registry.register/3`. Here, you're passing the name of the registry (`:my_registry`), the desired name of the spawned process (`{:database_worker, 1}`), and an arbitrary value. The `Registry` will then store a mapping of the name to the provided value and the pid of the caller process.

At this point, the registered process can be discovered by other processes. Notice how in the preceding snippet, you didn't take the pid of the database worker. That's because you don't need it. You can look it up in the registry by invoking `Registry.lookup/2`:

```
iex(3)> [{db_worker_pid, _value}] =
          Registry.lookup(
            :my_registry,
            {:database_worker, 1}
          )
```

`Registry.lookup/2` takes the name of the registry and the key (process name), and returns a list of `{pid, value}` tuples. When the registry is unique, this list can be either empty (no process is registered under the given key), or it can have one element. For a duplicate registry, this list can have any number of entries. The `pid` element in each tuple is the pid of the registered process, whereas the `value` is the value provided to `Registry.register/3`.

Now that you've discovered the mock database worker, you can send it a message:

```
iex(4)> send(db_worker_pid, :some_message)
got message :some_message
```

A very useful property of `Registry` is that it links to all the registered processes. This allows the registry to notice the termination of these processes and remove the corresponding entry from its internal structure.

You can immediately verify this. The database worker mock was a one-off process. It received a message, printed it, and then stopped. Try to discover it again:

```
iex(5)> Registry.lookup(:my_registry, {:database_worker, 1})
[]
```

As you can see, no entry is found under the given key because the database worker terminated.

 Note

It's worth mentioning that `Registry` is implemented in plain Elixir. You can think of `Registry` as something like a `GenServer` that holds the map of names to pids in its state. In reality, the implementation is more sophisticated and relies on the ETS table feature, which you'll learn about in chapter 10. ETS tables allow `Registry` to be very efficient and scalable. Lookups and writes are very fast, and in many cases they won't block each other, meaning that multiple operations on the same registry may run in parallel.

`Registry` has more features and properties, which we won't discuss here. You can take a look at the official documentation at https://hexdocs.pm/elixir/Registry.html for more details. But there's one very important feature of OTP processes that you need to learn about, called *via tuple*.

9.1.3 Via tuples

A *via tuple* is a mechanism that allows you to use an arbitrary third-party registry to register OTP-compliant processes, such as `GenServer` and supervisor. Recall that you can provide a `:name` option when starting a

```
GenServer:

GenServer.start_link(callback_module, some_arg, name: some_name)
```

So far, you've only passed atoms as the `:name` option, which caused the started process to be registered locally. But the `:name` option can also be provided in the shape of `{:via, some_module, some_arg}`. Such a tuple is also called a *via tuple*.

If you provide a via tuple as the name option, `GenServer` will invoke a well-defined function from `some_module` to register the process. Likewise, you can pass a via tuple as the first argument to `GenServer.cast` and `GenServer.call`, and `GenServer` will discover the pid using `some_module`. In this sense, `some_module` acts like a custom third-party process registry, and the via tuple is the way of connecting such a registry with `GenServer` and similar OTP abstractions.

The third element of the via tuple, `some_arg`, is a piece of data that's passed to functions of `some_module`. The exact shape of this data is defined by the registry module. At the very least, this piece of data must contain the name under which the process should be registered and looked up.

In the case of `Registry`, the third argument should be a pair, `{registry_name, process_key}`, so the entire via tuple then has the shape of `{:via, Registry, {registry_name, process_key}}`.

Let's look at an example. We'll revisit our old friend from chapter 6, the `EchoServer`. This is a simple `GenServer` that handles a call request by returning the request payload. Now you'll add registration to the echo server. When you start the server, you'll provide the server ID—an arbitrary term that uniquely identifies the server. When you want to send a request to the server, you'll pass this ID, instead of the pid.

Here's the full implementation:

```
defmodule EchoServer do
  use GenServer

  def start_link(id) do
```

```
    GenServer.start_link(__MODULE__, nil, name: via_tuple(id))
  end

  def init(_), do: {:ok, nil}

  def call(id, some_request) do
    GenServer.call(via_tuple(id), some_request)      #2
  end

  defp via_tuple(id) do
    {:via, Registry, {:my_registry, {__MODULE__, id}}}    #3
  end

  def handle_call(some_request, _, state) do
    {:reply, some_request, state}
  end
end
```

Here you consolidate the shaping of the via tuple in the `via_tuple/1` helper function. The registered name of the process will be `{__MODULE__, id}`, or in this case, `{EchoServer, id}`.

Try it out. Start the `iex` session, copy and paste the module definition, and then start `:my_registry`:

```
iex(1)> defmodule EchoServer do ... end

iex(2)> Registry.start_link(name: :my_registry, keys: :unique)
```

Now you can start and interact with multiple echo servers without needing to keep track of their pids:

```
iex(3)> EchoServer.start_link("server one")
iex(4)> EchoServer.start_link("server two")

iex(5)> EchoServer.call("server one", :some_request)
:some_request

iex(6)> EchoServer.call("server two", :another_request)
:another_request
```

Notice that the IDs here are strings, and also recall that the whole registered key is in fact `{EchoServer, some_id}`, which proves that you're using arbitrary complex terms to register processes and discover them.

9.1.4 Registering database workers

Now that you've learned the basics of `Registry`, you can implement registration and discovery of your database workers. First you need to create the `Todo.ProcessRegistry` module. The code is presented in the following listing.

Listing 9.3. Todo process registry (pool_supervision/lib/todo/process_registry.ex)

```
defmodule Todo.ProcessRegistry do
  def start_link do
    Registry.start_link(keys: :unique, name: __MODULE__)
  end

  def via_tuple(key) do
    {:via, Registry, {__MODULE__, key}}
  end

  def child_spec(_) do                          #1
    Supervisor.child_spec(                      #1
      Registry,                                 #1
      id: __MODULE__,                           #1
      start: {__MODULE__, :start_link, []}      #1
    )                                           #1
  end
end
```

The interface functions are straightforward. The `start_link` function simply forwards to the `Registry` module to start a unique registry. The `via_tuple/1` function can be used by other modules, such as `Todo.DatabaseWorker`, to create the appropriate via tuple that registers a process with this registry.

Because the registry is a process, it should be supervised. Therefore you include `child_spec/1` in the module. Here you're using `Supervisor.child_spec/2` to adjust the default specification from the `Registry` module. This invocation essentially states that you'll use whatever child specification is provided by `Registry`, with `:id` and `:start` fields changed. By doing this, you don't need to know about the internals of the `Registry` implementation, such as whether the registry process is a worker or a supervisor.

With this in place, you can immediately put the registry under the `Todo.System` supervisor. The code is presented in the following listing.

Listing 9.4. Supervising registry (pool_supervision/lib/todo/system.ex)

```
defmodule Todo.System do
  def start_link do
    Supervisor.start_link(
      [
        Todo.ProcessRegistry,      #1
        Todo.Database,
        Todo.Cache
      ],
      strategy: :one_for_one
    )
  end
end
```

Keep in mind that processes are started synchronously, in the order you specify. Thus, the order in the child specification list matters and isn't chosen arbitrarily. A child must always be specified after its dependencies. In this case, you must start the registry first because database workers will depend on it.

With `Todo.ProcessRegistry` in place, you can start adapting the database workers. The relevant changes are presented in the following listing.

Listing 9.5. Registering workers (pool_supervision/lib/todo/database_worker.ex)

```
defmodule Todo.DatabaseWorker do
  use GenServer

  def start_link({db_folder, worker_id}) do
    GenServer.start_link(
      __MODULE__,
      db_folder,
      name: via_tuple(worker_id)     #1
    )
  end

  def store(worker_id, key, data) do
    GenServer.cast(via_tuple(worker_id), {:store, key, data})   #
  end
```

```
  def get(worker_id, key) do
    GenServer.call(via_tuple(worker_id), {:get, key})       #2
  end

  defp via_tuple(worker_id) do
    Todo.ProcessRegistry.via_tuple({__MODULE__, worker_id})
  end

  ...
end
```

This code introduces the notion of a `worker_id`, which is an integer in the range `1..pool_size`. The `start_link` function now takes this parameter together with `db_folder`. But notice that the function takes both parameters as a single `{db_folder, worker_id}` tuple. The reason is again in conformance with the autogenerated `child_spec/1`, which forwards to `start_link/1`. To manage a worker under a supervisor, you can now use the `{Todo.DatabaseWorker, {db_folder, worker_id}}` child specification.

When invoking `GenServer.start_link`, you provide the via tuple as the name option. The exact shape of the tuple is wrapped in the internal `via_tuple/1` function, which takes the worker ID and returns the corresponding via tuple. This function just delegates to `Todo.ProcessRegistry`, passing it the desired name in the form `{__MODULE__, worker_id}`. Therefore, a worker is registered with the key `{Todo.DatabaseWorker, worker_id}`. Such a name eliminates possible clashes with other types of processes that might be registered with the same registry.

Similarly, you use the `via_tuple/1` helper to discover the processes when invoking `GenServer.call` and `GenServer.cast`. Notice that `store/3` and `get/2` functions now receive a worker ID as the first argument. This means that their clients don't need to keep track of the pids anymore.

9.1.5 Supervising database workers

Now you can create a new supervisor that will manage the pool of workers. Why introduce a separate supervisor? Theoretically, you could place workers under `Todo.System`, and this would work fine. But remember from the

previous chapter that if restarts happen too often, the supervisor gives up at some point and terminates all of its children. If you keep too many children under the same supervisor, you might reach the maximum restart intensity sooner, in which case all processes are restarted. In other words, problems in a single process could easily trip over to the majority of the system.

In this case, I made an arbitrary decision to place a distinct part of the system (the database) under a separate supervisor. This approach may limit the impact of a failed restart to database operations. If restarting one database worker fails, the supervisor will terminate, which means the parent supervisor will try to restart the entire database service without touching other processes in the system.

Either way, the consequence of these changes is that you don't need the database `GenServer` anymore. The purpose of this server was to start a pool of worker processes and to manage the mapping of a worker ID to pid. With these new changes, the workers are started by the supervisor; the mapping is already handled by the registry. Therefore, the database `GenServer` is redundant.

You can still keep the `Todo.Database` module. It will now implement a supervisor of database worker processes and retain the same interface functions as before. As a result, you don't need to change the code of the client `Todo.Server` module at all, and you can still keep `Todo.Database` in the list of `Todo.System` children.

Next you'll convert the database into a supervisor. The code is presented in the next listing.

Listing 9.6. Supervising workers (pool_supervision/lib/todo/database.ex)

```
defmodule Todo.Database do
  @pool_size 3
  @db_folder "./persist"

  def start_link do
    File.mkdir_p!(@db_folder)

    children = Enum.map(1..@pool_size, &worker_spec/1)
    Supervisor.start_link(children, strategy: :one_for_one)
```

```
  end

  defp worker_spec(worker_id) do
    default_worker_spec = {Todo.DatabaseWorker, {@db_folder, work
    Supervisor.child_spec(default_worker_spec, id: worker_id)
  end

  ...
end
```

You start off by creating a list of three child specifications, each of them describing one database worker. Then you pass this list to `Supervisor.start_link/2`.

The specification for each worker is created in `worker_spec/1`. You start off with the default specification for the database worker, `{Todo.DatabaseWorker, {@db_folder, worker_id}}`. Then you use `Supervisor.child_spec/2` to set the unique ID for the worker.

Without that, you'd end up with multiple children having the same ID. Recall from chapter 8 that a default `child_spec/1`, generated via `use GenServer`, provides the name of the module in the `:id` field. Consequently, if you use that default specification and try to start two database workers, they'll both get the same ID of `Todo.DatabaseWorker`. The `Supervisor` module will complain about it and raise an error.

You also need to implement `Todo.Database.child_spec/1`. You just converted the database into a supervisor, so the module doesn't contain `use GenServer` anymore, meaning `child_spec/1` isn't auto-generated. The code is given in the following listing.

Listing 9.7. Database operations (pool_supervision/lib/todo/database.ex)

```
defmodule Todo.Database do
  ...

  def child_spec(_) do
    %{
      id: __MODULE__,
      start: {__MODULE__, :start_link, []},
      type: :supervisor
    }
```

```
    end

  ...
end
```

The specification contains the field `:type`, which hasn't been mentioned before. This field can be used to indicate the type of the started process. The valid values are `:supervisor` (if the child is a supervisor process) or `:worker` (for any other kind of process). If you omit this field, the default value of `:worker` is used.

The `child_spec/1` in listing 9.7 therefore specifies that `Todo.Database` is a supervisor, and that it can be started by invoking `Todo.Database.start_link/0`.

This is a nice example of how `child_spec/1` helps you keep implementation details in the module that powers a process. You just turned the database into a supervisor, and you changed the arity of its `start_link` function (it now takes zero arguments), but nothing needs to be changed in the `Todo.System` module.

Next, you need to adapt the `store/2` and `get/1` functions.

Listing 9.8. Database operations (pool_supervision/lib/todo/database.ex)

```
defmodule Todo.Database do
  ...

  def store(key, data) do
    key
    |> choose_worker()
    |> Todo.DatabaseWorker.store(key, data)
  end

  def get(key) do
    key
    |> choose_worker()
    |> Todo.DatabaseWorker.get(key)
  end

  defp choose_worker(key) do
    :erlang.phash2(key, @pool_size) + 1
```

```
    end

    ...
end
```

The only difference from the previous version is in the `choose_worker/1`
function. Previously this function issued a call to the database server. Now it
just selects the worker ID in the range `1..@pool_size`. This ID is then passed
to `Todo.DatabaseWorker` functions, which will perform a registry lookup and
forward the request to the corresponding database worker.

At this point, you can test how the system works. Start everything:

```
iex(1)> Todo.System.start_link()

Starting database server.
Starting database worker.
Starting database worker.
Starting database worker.
Starting to-do cache.
```

Now verify that you can restart individual workers correctly. To do that, you
need to get the pid of a worker. Because you know the internals of the
system, this can easily be done by looking it up in the registry. Once you
have the pid, you can terminate the worker:

```
iex(2)> [{worker_pid, _}] =
          Registry.lookup(
            Todo.ProcessRegistry,
            {Todo.DatabaseWorker, 2}
          )

iex(3)> Process.exit(worker_pid, :kill)
Starting database worker.
```

The worker is restarted, as expected, and the rest of the system is
undisturbed.

It's worth repeating how the registry supports proper behavior in the system
regarding restarted processes. When a worker is restarted, the new process
has a different pid. But owing to the registry, the client code doesn't care
about that. You resolve the pid at the latest possible moment, doing a registry

lookup prior to issuing a request to the database worker. Therefore, in most cases the lookup will succeed, and you'll talk to the proper process.

In some cases, the discovery of the database worker might return an invalid value, such as if the database worker crashes after the client process found its pid, but before the request is sent. In this case, the client process has a stale pid, so the request will fail. A similar problem can occur if a client wants to find a database worker that has just crashed. Restarting and registration run concurrently with the client, so the client might not find the worker pid in the registry.

Both scenarios lead to the same result: the client process, in this case a to-do server, will crash, and the error will be propagated to the end user. This is a consequence of the highly concurrent nature of the system. A failure recovery is performed concurrently in the supervisor process, so some part of the system might not be in a consistent state for a brief period of time.

9.1.6 Organizing the supervision tree

Let's stop for a moment and reflect on what you've done so far. The relationship between processes is presented in figure 9.4.

This is an example of a simple *supervision tree* — a nested structure of supervisors and workers. The tree describes how the system is organized into a hierarchy of services. In this example, the system consists of three services: the process registry, the database, and the cache.

Figure 9.4. Supervision tree
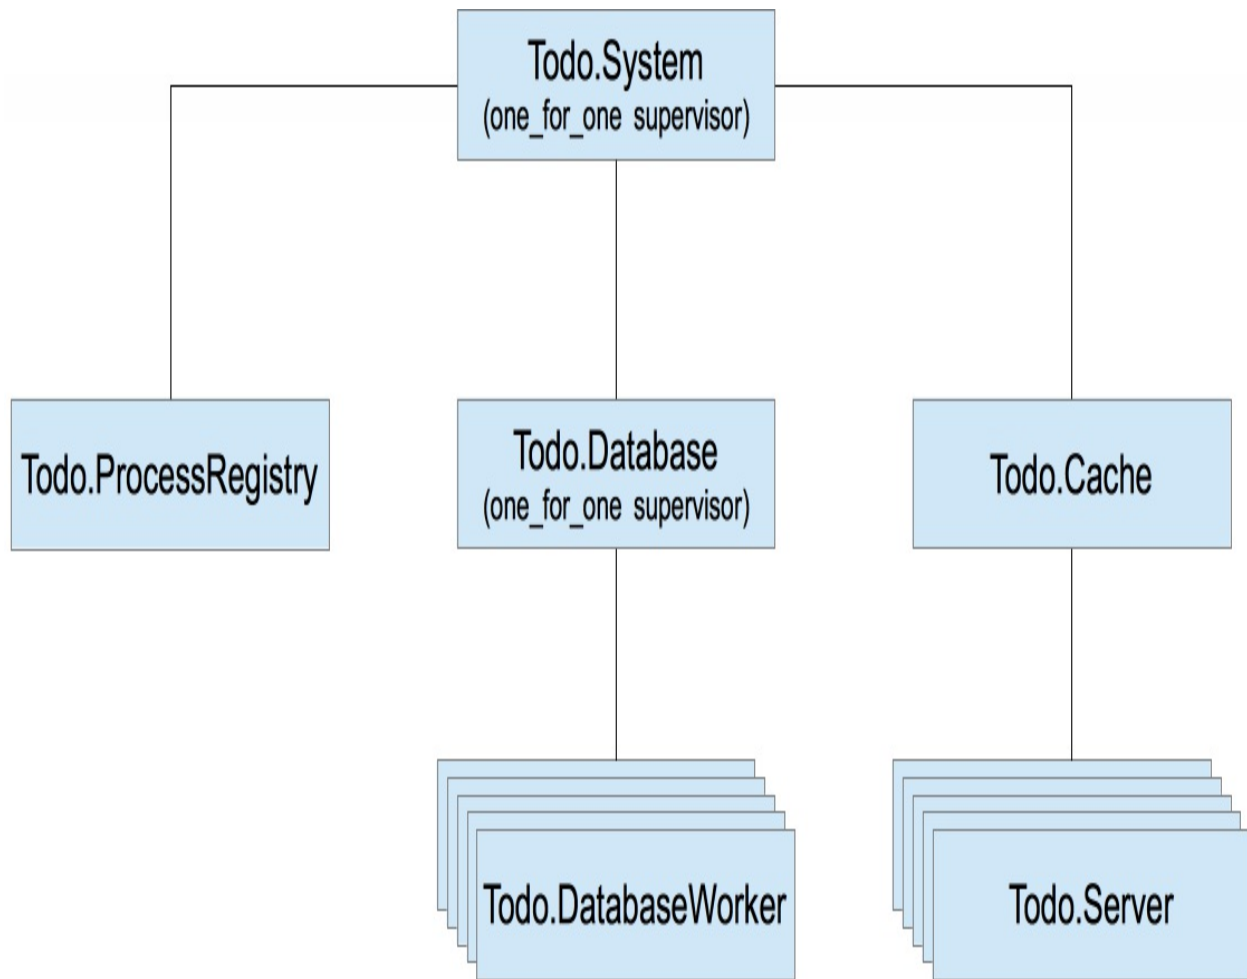
Each service can be further subdivided into subservices. For example, the database is composed of multiple workers and the cache is composed of multiple to-do servers. Even the registry is further subdivided into multiple processes, but that's an implementation detail of the `Registry` module, so it's not shown on the diagram.

Although supervisors are frequently mentioned in the context of fault-tolerance and error recovery, defining the proper starting order is their most essential role. The supervision tree describes how the system is started and how it's taken down.

A more granular tree allows you to take down an arbitrary part of the system, without touching anything else. In the current version, stopping the database service is as easy as asking its parent (`Todo.System`) to stop the `Todo.Database` child, using the `Supervisor.terminate_child/2` function.

This will take down the database process together with its descendants.

If worker processes are small services in a system, you can think of supervisors as being service managers—a built-in equivalent of systemd, Windows Service Manager, and the like. They're responsible for the lifecycle of services they directly manage. If any critical service stops, its parent will try to restart it.

Looking at the supervision tree, you can reason about how errors are handled and propagated throughout the system. If a database worker crashes, the database supervisor will restart it, leaving the rest of the system alone. If that doesn't help, you'll exceed the maximum restart frequency, and the database supervisor will terminate all database workers and then itself.

This will be noticed by the system supervisor, which will then start a fresh database pool in hopes of solving the problem. What does all this restarting get you? By restarting an entire group of workers, you effectively terminate all pending database operations and begin clean. If that doesn't help, there's nothing more you can do, so you propagate the error up the tree (in this case, killing everything). This is how error recovery works in supervision trees— you try to recover from an error locally, affecting as few processes as possible. If that doesn't work, you move up and try to restart the wider part of the system.

OTP-compliant processes

All processes that are started directly from a supervisor should be *OTP-compliant*. To implement an OTP-compliant process, it's not enough to spawn or link a process; you also must handle some OTP-specific messages in a particular way. The details of what exactly must be done are provided in the Erlang documentation at https://www.erlang.org/doc/design_principles/spec_proc.html#special-processes.

Luckily, you usually won't need to implement an OTP-compliant process from scratch. Instead, you can use various higher-level abstractions, such as `GenServer`, `Supervisor`, and `Registry`. The processes started with these

modules will be OTP-compliant. Elixir also ships with `Task` and `Agent` modules that can be used to run OTP-compliant processes. You'll learn about tasks and agents in the next chapter.

Plain processes started by `spawn_link` aren't OTP-compliant, so such processes shouldn't be started directly from a supervisor. You can freely start plain processes from workers such as `GenServer`, but it's generally better to use OTP-compliant processes wherever possible.

Shutting down processes

An important benefit of supervision trees is the ability to stop the entire system without leaving dangling processes. When you terminate a supervisor, all of its immediate children are also terminated. If all other processes are directly or indirectly linked to those children, they will eventually be terminated as well. Consequently, you can stop the entire system by terminating the top-level supervisor process.

Most often, a supervisor subtree is terminated in a controlled manner. A supervisor process will instruct its children to terminate gracefully, thus giving them the chance to do final cleanup. If some of those children are themselves supervisors, they will take down their own trees in the same way. Graceful termination of a `GenServer` worker involves invoking the `terminate/2` callback, but only if the worker process is trapping exits. Therefore, if you want to do some cleanup from a `GenServer` process, make sure you set up an exit trap from an `init/1` callback.

Because graceful termination involves the possible execution of cleanup code, it may take longer than desired. The `:shutdown` option in a child specification lets you control how long the supervisor will wait for the child to terminate gracefully. If the child doesn't terminate in this time, it will be forcefully terminated. You can choose the shutdown time by specifying `shutdown: shutdown_strategy` in `child_spec/1` and passing an integer representing a time in milliseconds. Alternatively, you can pass the atom `:infinity`, which instructs the supervisor to wait indefinitely for the child to terminate. Finally, you can pass the atom `:brutal_kill`, telling the supervisor to immediately terminate the child in a forceful way. The forceful

termination is done by sending a `:kill` exit signal to the process, like you did with `Process.exit(pid, :kill)`.

The default value of the `:shutdown` option is 5000 for a worker process or `:infinity` for a supervisor process.

Avoiding process restarting

By default, a supervisor restarts a terminated process regardless of the exit reason. Even if the process terminates with the reason `:normal`, it will be restarted. Sometimes you may want to alter this behavior.

For example, consider a process that handles an HTTP request or a TCP connection. If such a process fails, the socket will be closed, and there's no point in restarting the process (the remote party will be disconnected anyway). You still want to have such processes under a supervision tree because this makes it possible to terminate the entire supervisor subtree without leaving dangling processes. In this situation, you can set up a *temporary* worker by providing `restart: :temporary` in `child_spec/1`. A temporary worker isn't restarted on termination.

Another option is a *transient* worker, which is restarted only if it terminates abnormally. Transient workers can be used for processes that may terminate normally, as part of the standard system workflow. A typical example for this is a one-off job that you want to execute when the system is started. You could start the corresponding process (usually powered by the `Task` module) in the supervision tree, and configure it as transient. A transient worker can be specified by providing `restart: :transient` in `child_spec/1`.

Restart strategies

So far, you've been using only the `:one_for_one` restart strategy. In this mode, a supervisor handles a process termination by starting a new process in its place, leaving other children alone. There are two additional restart strategies:

- `:one_for_all` — When a child crashes, the supervisor terminates all

other children and then starts all children.
- `:rest_for_one` — When a child crashes, the supervisor terminates all younger siblings of the crashed child. Then the supervisor starts new child processes in place of the terminated ones.

These strategies are useful if there's tight coupling between siblings, where the service of some child doesn't make any sense without its siblings. One example is when a process keeps the pid of some sibling in its own state. In this case, the process is tightly coupled to a particular instance of the sibling. If the sibling terminates, so should the dependent process.

By opting for `:one_for_all` or `:rest_for_one`, you can make that happen. The former is useful when there's tight dependency in all directions (every sibling depends on other siblings). The latter is appropriate if younger siblings depend on the older ones.

For example, in the to-do system, you could use `:rest_for_one` to take down database workers if the registry process terminates. Without the registry, these processes can't serve any purpose, so taking them down would be a proper thing to do. In this particular case, however, you don't need to do that because `Registry` links each registered process to the registry process. As a result, a termination of the registry process is properly propagated to the registered processes. Any such process that doesn't trap exits will be taken down automatically; processes that trap exits will receive a notification message.

This concludes our initial look at fine-grained supervision. You've made a number of changes that minimize the effects of errors, but there's still a lot of room for improvement. You'll continue extending the system in the next section, where you'll learn how to start workers dynamically.

9.2 Starting processes dynamically

With the changes you made in the previous section, the impact of a database-worker error is now confined to a single worker. It's time to do the same thing for to-do servers. You'll use roughly the same approach as you did with database workers: you'll run each to-do server under a supervisor and register

the servers in the process registry.

9.2.1 Registering to-do servers

You'll start off by adding registration to to-do servers. The change is simple, as shown in the following listing.

Listing 9.9. Registering to-do servers (dynamic_workers/lib/todo/server.ex)

```
defmodule Todo.Server do
  use GenServer, restart: :temporary

  def start_link(name) do
    GenServer.start_link(__MODULE__, name, name: via_tuple(name))
  end

  defp via_tuple(name) do
    Todo.ProcessRegistry.via_tuple({__MODULE__, name})
  end

  ...
end
```

This is the same technique you used with database workers. You pass the via tuple as the name option. The via tuple will state that the server should be registered with the `{__MODULE__, name}` key to the process registry. Using this form of the key avoids possible collisions between to-do server keys and database worker keys.

The functions `add_entry/2` and `entries/2` are unchanged, and they still take the pid as the first argument, so the usage remains the same. A client process first obtains the pid of the to-do server by invoking `Todo.Cache.server_process/1`, and then it invokes `Todo.Server` functions.

9.2.2 Dynamic supervision

Next, you need to supervise to-do servers. There's a twist, though. Unlike database workers, to-do servers are created dynamically when needed. Initially, no to-do server is running; each is created on demand when you call `Todo.Cache.server_process/1`. This effectively means you can't specify

supervisor children up front because you don't know how many children you'll need.

For such cases, you need a dynamic supervisor that can start children on demand. In Elixir, this feature is available via the `DynamicSupervisor` module.

`DynamicSupervisor` is similar to `Supervisor`, but where `Supervisor` is used to start a predefined list of children, `DynamicSupervisor` is used to start children on demand. When you start a dynamic supervisor, you don't provide a list of child specifications, so only the supervisor process is started. Then, whenever you want to, you can start a supervised child using `DynamicSupervisor.start_child/2`.

Let's see this in action. You'll convert `Todo.Cache` into a dynamic supervisor, much like what you did with the database. The relevant code is presented in the following listing.

Listing 9.10. To-do cache as a supervisor (dynamic_workers/lib/todo/cache.ex)

```
defmodule Todo.Cache do
  def start_link() do
    IO.puts("Starting to-do cache")

    DynamicSupervisor.start_link(      #1
      name: __MODULE__,
      strategy: :one_for_one
    )
  end

  ...
end
```

You start the supervisor using `DynamicSupervisor.start_link/1`. This will start the supervisor process, but no children are specified at this point. Notice that when starting the supervisor, you're also passing the `:name` option. This will cause the supervisor to be registered under a local name.

By making the supervisor locally registered, it's easier for you to interact with the supervisor and ask it to start a child. You can immediately use this

by adding the `start_child/1` function, which starts the to-do server for the given to-do list:

```
defmodule Todo.Cache do
  ...

  defp start_child(todo_list_name) do
    DynamicSupervisor.start_child(
      __MODULE__,
      {Todo.Server, todo_list_name}
    )
  end

  ...
```

Here, you're invoking `DynamicSupervisor.start_child/2`, passing it the name of your supervisor and the child specification of the child you want to start. The specification `{Todo.Server, todo_list_name}` will lead to the invocation of `Todo.Server.start_link(todo_list_name)`. The to-do server will be started as the child of the `Todo.Cache` supervisor.

It's worth noting that `DynamicSupervisor.start_child/2` is a cross-process synchronous call. A request is sent to the supervisor process, which then starts the child. If multiple client processes simultaneously try to start a child under the same supervisor, the requests will be serialized.

For more details on dynamic supervisors, refer to the official documentation at https://hexdocs.pm/elixir/DynamicSupervisor.html.

One small thing left to do is the implementation of `child_spec/1`:

```
defmodule Todo.Cache do
  ...

  def child_spec(_arg) do
    %{
      id: __MODULE__,
      start: {__MODULE__, :start_link, []},
      type: :supervisor
    }
  end

  ...
```

```
end
```

At this point, the to-do cache is converted into a dynamic supervisor.

9.2.3 Finding to-do servers

The final thing left to do is to change the discovery
`Todo.Cache.server_process/1` function. This function takes a name and
returns the pid of the to-do server, starting it if it's not running. The
implementation is provided in the following listing.

Listing 9.11. Finding a to-do server (dynamic_workers/lib/todo/cache.ex)

```
defmodule Todo.Cache do
  ...

  def server_process(todo_list_name) do
    case start_child(todo_list_name) do
      {:ok, pid} -> pid                       #1
      {:error, {:already_started, pid}} -> pid    #2
    end
  end

  defp start_child(todo_list_name) do
    DynamicSupervisor.start_child(
      __MODULE__,
      {Todo.Server, todo_list_name}
    )
  end
end
```

The function first invokes the local `start_child/1` function, which you
prepared in the previous section, and which is a simple wrapper around
`DynamicSupervisor.start_child/2`.

This invocation can have two successful outcomes. In the most obvious case,
the function returns `{:ok, pid}` with the pid of the newly started to-do
server.

The second outcome is more interesting. If the result is `{:error,
{:already_started, pid}}`, the to-do process failed to register because

another process is already registered with the same name—a to-do server for the list with the given name is already running. For the to-do example, this outcome is also a success. You tried to start the server, but it was already running. That's fine. You have the pid of the server, and you can interact with it.

The result `{:error, {:already_started, pid}}` is returned due to the inner workings of `GenServer` registration. When a `:name` option is provided to `GenServer.start_link`, the registration is performed in the started process before `init/1` is invoked. This registration can fail if some other process is already registered under the same key. In this case, `GenServer.start_link` doesn't resume to run the server loop. Instead, it returns `{:error, {:already_started, pid}}` where the pid points to the process that's registered under the same key. This result is then returned by `DynamicSupervisor.start_child`.

It's worth briefly discussing how `server_process/1` behaves in a concurrent scenario. Consider the case of two processes invoking this function at the same time. The execution moves to `DynamicSupervisor.start_child/2`, so you might end up with two simultaneous executions of `start_child` on the same supervisor. Recall that a child is started in the supervisor process. Therefore, the invocations of `start_child` are serialized, and `server_process/1` doesn't suffer from race conditions.

On the flip side, the way `start_child` is used here is not very efficient. Every time you want to work with a to-do list, you issue a request to the supervisor, so the supervisor process can become a bottleneck. Even if the to-do server is already running, the supervisor will briefly start a new child, which will immediately stop. This can easily be improved, but we'll leave it for now because the current implementation is behaving properly. We'll revisit this issue in chapter 12 when we move to a distributed registration.

9.2.4 Using temporary restart strategy

There's one thing left to do. You'll configure the to-do server to be a `:temporary` child. As a result, if a to-do server stops, say due to a crash, it won't be restarted.

Why choose this approach? Servers are started on demand, so when a user tries to interact with a to-do list, if the server process isn't running, it will be started. If a to-do list server crashes, it will be started on the next use, so there's no need to restart it automatically.

Opting for the `:temporary` strategy also means that the parent supervisor won't be restarted due to too many failures in its children. Even if there are frequent crashes in one to-do server, say due to corrupt state, you'll never take down the entire cache, which should improve the availability of the entire system.

Changing the restart strategy is easily done by providing the `:restart` option to use `GenServer`:

Listing 9.12. Changing to-do server restart strategy (dynamic_workers/lib/todo/server.ex)

```
defmodule Todo.Server do
  use GenServer, restart: :temporary

  ...
end
```

The `:temporary` value will be included under the `:restart` key in the result of `child_spec/1`, so the parent supervisor will treat the child as temporary. If the child terminates, it won't be restarted.

You might wonder why to-do servers are supervised if they're not restarted. There are two important benefits. First, this structure ensures that the failure of a single to-do server doesn't affect any other process in the system. In addition, as explained in section 9.1.6, this allows you to properly take down the system, or some service in the system, without leaving any dangling processes behind. To stop all to-do servers, you need to stop the `Todo.Cache` supervisor. In other words, supervision isn't just about restarting crashed processes, but also about isolating individual crashes and enabling proper termination.

9.2.5 Testing the system

At this point, the to-do servers are supervised and you can test the code.

Notice that you didn't have to make any change in the `Todo.System` supervisor. The `Todo.Cache` was already listed as a child, and you only changed its internals. Let's see if this works.

Start the shell and the entire system:

```
iex(1)> Todo.System.start_link()

Starting database server.
Starting database worker.
Starting database worker.
Starting database worker.
Starting to-do cache.
```

Now, you can get one to-do server:

```
iex(2)> bobs_list = Todo.Cache.server_process("Bob's list")
Starting to-do server for Bob's list
#PID<0.118.0>
```

Repeating the request doesn't start another server:

```
iex(3)> bobs_list = Todo.Cache.server_process("Bob's list")
#PID<0.118.0>
```

In contrast, using a different to-do list name creates another process:

```
iex(4)> alices_list = Todo.Cache.server_process("Alice's list")
Starting to-do server for Alice's list
#PID<0.121.0>
```

Crash one to-do server:

```
iex(5)> Process.exit(bobs_list, :kill)
```

The subsequent call to `Todo.Cache.server_process/1` will return a different pid:

```
iex(6)> Todo.Cache.server_process("Bob's list")
Starting to-do server for Bob's list
#PID<0.124.0>
```

Of course, Alice's server remains undisturbed:

```
iex(7)> Todo.Cache.server_process("Alice's list")
#PID<0.121.0>
```

The supervision tree of the new code is presented in figure 9.5. The diagram depicts how you supervise each process, limiting the effect of unexpected errors.

Figure 9.5. Supervising to-do servers

With this, you're finished making your to-do system fault-tolerant. You've introduced additional supervisor processes to the system, and you've also managed to simplify some other parts (removing the to-do cache and database server processes). You'll make many more changes to this system, but for now let's leave it and look at some important practical considerations.

9.3 "Let it crash"

In general, when you develop complex systems, you should employ supervisors to do your error handling and recovery. With properly designed supervision trees, you can limit the impact of unexpected errors, and the system will hopefully recover. I can personally testify that supervisors have helped me in occasional weird situations in production, keeping the running system stable and saving me from unwanted phone calls in the middle of the night. It's also worth noting that OTP provides logging facilities, so process crashes are logged and you can see that something went wrong. It's even possible to set up an event handler that will be triggered on every process crash, thus allowing you to perform custom actions, such as sending an email or reporting to an external system.

An important consequence of this style of error handling is that the worker code is liberated from paranoid, defensive `try/catch` constructs. Usually these aren't needed because you use supervisors to handle error recovery. Joe Armstrong, one of the inventors of Erlang, described such a style in his PhD thesis ("Making reliable distributed systems in the presence of software errors," https://erlang.org/download/armstrong_thesis_2003.pdf) as *intentional programming*. Using this approach, the code states the programmer's intention, rather than being cluttered with all sorts of defensive constructs.

This style is also known as *let it crash*. In addition to making the code shorter and more focused, "let it crash" promotes clean-slate recovery. Remember, when a new process starts, it starts with new state, which should be consistent. Furthermore, the message queue (mailbox) of the old process is thrown away. This will cause some requests in the system to fail. But the new process starts fresh, which gives it a better chance to resume normal operation.

"Let it crash" can initially seem confusing, and people may mistake it for the "let everything crash" approach. There are two important situations in which you should explicitly handle an error:

- In critical processes that shouldn't crash

- When you expect an error that can be dealt with in a meaningful way

Let's look at each of these.

9.3.1 Processes that shouldn't crash

Processes that shouldn't crash are informally called a system's *error kernel* — processes that are critical for the entire system to work and whose state can't be restored in a simple and consistent way. Such processes are the heart of your system, and you generally don't want them to crash, because without them the system can't provide any service.

You should keep the code of such important processes as simple as possible. The less logic that happens in the process, the smaller the chance of a process crash. If the code of your error-kernel process is complex, consider splitting it into two processes: one that holds state, and another that does the actual work. The former process then becomes extremely simple and is unlikely to crash, whereas the worker process can be removed from the error kernel (because it no longer maintains critical state).

In addition, you could consider including defensive `try/catch` statements in each `handle_*` callback of a critical process, to prevent a process from crashing. Here's a simple sketch of the idea:

```
def handle_call(message, _, state) do
  try
    new_state =
      state
      |> transformation_1()
      |> transformation_2()
      ...

    {:reply, response, new_state}

  catch _, _ ->
    {:reply, {:error, reason}, state}     #1
  end
end
```

This snippet illustrates how immutable data structures allow you to implement a fault-tolerant server. While processing a request, you make a

series of transformations on the state. If anything bad happens, you use the initial state, effectively performing a rollback of all changes. This preserves state consistency while keeping the process constantly alive.

Keep in mind that this technique doesn't completely guard against a process crash. For example, you can always kill a process by invoking `Process.exit(pid, :kill)`, because a `:kill` exit reason can't be intercepted even if you're trapping exits. Therefore, you should always have a recovery plan for the crash of a critical process. Set up a proper supervision hierarchy to ensure the termination of all dependent processes in the case of an error-kernel process crash.

9.3.2 Handling expected errors

The whole point of the let-it-crash approach is to leave recovery of unexpected errors to supervisors. But if you can predict an error and you have a way to deal with it, there's no reason to let the process crash.

Here's a simple example. Look at the `:get` request in the database worker:

```
def handle_call({:get, key}, _, db_folder) do
  data =
    case File.read(file_name(db_folder, key)) do
      {:ok, contents} -> :erlang.binary_to_term(contents)
      _ -> nil          #1
    end

  {:reply, data, db_folder}
end
```

When handling a `get` request, you try to read from a file, covering the case when this read fails. If it doesn't succeed, you return `nil`, treating this case as if an entry for the given key isn't in the database.

But you can do better. Consider using an error only when the file isn't available. This error is identified with `{:error, :enoent}`, so the corresponding code would look like this:

```
case File.read(...) do
  {:ok, contents} -> do_something_with(contents)
```

```
  {:error, :enoent} -> nil
end
```

Notice how you rely on pattern matching here. If neither of these two expected situations happens, a pattern match will fail, and so will your process. This is the idea of "let it crash." You deal with expected situations (the file is either available or doesn't exist), crashing if anything else goes wrong (for example, you don't have permissions).

In contrast, when storing data, you use `File.write!/2` (notice the exclamation mark), which may throw an exception and crash the process. If you don't succeed in saving the data, your database worker has failed, and there's no point in hiding this fact. Better to fail fast, which will cause an error that will be logged and (hopefully) noticed and fixed.

Of course, restarting may not help. In this case, the supervisor will give up and crash itself, and the system will quickly come to a halt, which is probably a good thing. No point in working if you can't persist the data.

As a general rule, if you know what to do with an error, you should definitely handle it. Otherwise, for anything unexpected, let the process crash and ensure proper error isolation and recovery via supervisors.

9.3.3 Preserving the state

Keep in mind that state isn't preserved when a process is restarted. Remember from chapter 5 that a process's state is its own private affair. When a process crashes, the memory it occupied is reclaimed, and the new process starts with new state. This has the important advantage of starting clean. Perhaps a process crashed due to inconsistent state, and starting fresh may fix the error.

That said, in some cases you'll want the process's state to survive the crash. This isn't provided out of the box; you need to implement it yourself. The general approach is to save the state outside of the process (for example, in another process or to a database) and then restore the state when the successor process is started.

You already have this functionality in the to-do server. Recall that you have a simple database system that persists to-do lists to disk. When the to-do server is started, the first thing it tries to do is to restore the data from the database. This makes it possible for the new process to inherit the state of the old one.

In general, be careful when preserving state. As you learned in chapter 4, a typical change in a functional data abstraction goes through chained transformations:

```
new_state =
  state
  |> transformation_1(...)
  ...
  |> transformation_n(...)
```

As a rule, the state should be persisted after all transformations are completed. Only then can you be certain that your state is consistent, so this is a good opportunity to save it. For example, you do this in the to-do server after you modify the internal data abstraction:

```
def handle_cast({:add_entry, new_entry}, {name, todo_list}) do
  new_list = Todo.List.add_entry(todo_list, new_entry)
  Todo.Database.store(name, new_list)      #1
  {:noreply, {name, new_list}}
end
```

💡 **Tip**

Persistent state can have a negative effect on restarts. Let's say an error is caused by state that's somehow invalid (perhaps due to a bug). If this state is persisted, your process can never restart successfully, because the process will restore the invalid state and then crash again (either on starting or when handling a request). You should be careful when persisting state. If you can afford to, it's better to start clean and terminate all other dependent processes.

9.4 Summary

- Supervisors allow you to localize the impact of an error, keeping unrelated parts of the system undisturbed.

- The registry helps you find processes without needing to track their pids. This is very helpful if a process is restarted.
- Each process should reside somewhere in a supervision tree. This makes it possible to terminate the entire system (or an arbitrary sub-part of it) by terminating the supervisor.
- `DynamicSupervisor` is used for on-demand starting.
- When a process crashes, its state is lost. You can deal with this by storing state outside the process, but more often than not, it's best to start with clean state.
- In general, you should handle unexpected errors through a proper supervision hierarchy. Explicit handling through a `try` construct should be used only when you have a meaningful way to deal with an error.

10 Beyond GenServer

This chapter covers

- Tasks
- Agents
- ETS tables

Chapters 8 and 9 introduced the distinction between worker and supervisor processes. Workers are the processes that provide some part of your service, whereas supervisors organize the worker processes into a tree. This allows you to start and stop processes in the desired order, and also to restart critical processes if they fail.

As was mentioned in section 9.1.6, all processes that are started directly from a supervisor should be OTP-compliant processes. Processes started with plain `spawn` and `spawn_link` are not OTP-compliant, so you should refrain from running such processes in production. Modules such as `Supervisor`, `GenServer`, and `Registry` allow you to start OTP-compliant processes that can be placed into a supervision tree.

In this chapter you'll learn about two additional modules that also allow you to run OTP-compliant workers: `Task` and `Agent`. Tasks can be very useful when you need to run one-off jobs, whereas agents can be used to manage state and provide concurrent access to it. Finally, we'll discuss a related feature called ETS tables, which under some conditions can serve as more efficient alternatives to `GenServer` and `Agent`.

There's a lot of new ground to cover, so let's start by discussing tasks.

10.1 Tasks

The `Task` module can be used to concurrently run a job — a process that takes some input, performs some computation, and then stops. In this sense, task-powered processes have a different flow than server processes. Whereas

a `GenServer` process acts as a long-running server, a `Task`-powered process starts its work immediately, doesn't serve requests, and stops when the work is done.

The `Task` module can be used in two different ways, depending on whether the task process needs to send a result back to the process that started it or not. The former case is also called an *awaited task*, because the starter process waits for the task to send the result back. Let's discuss this option first.

10.1.1 Awaited tasks

An *awaited task* is a process that executes some function, sends the function result back to the starter process, and then terminates. Let's look at a basic example.

Suppose you want to start a concurrent, possibly long-running, job, and get its result back. You can simulate a long-running job with the following function:

```
iex(1)> long_job =
          fn ->
            Process.sleep(2000)
            :some_result
          end
```

This lambda, when invoked, sleeps for two seconds and then returns `:some_result`.

To run this lambda concurrently, you can use `Task.async/1`:

```
iex(2)> task = Task.async(long_job)
```

The function `Task.async/1` takes a zero-arity lambda, spawns a separate process, and invokes the lambda in the spawned process. The return value of the lambda will be sent as a message back to the starter process.

Because the computation runs in a separate process, `Task.async/1` returns immediately, even if the lambda itself takes a long time to finish. This means

that the starter process isn't blocked and can perform some additional work concurrently with the task process.

The return value of `Task.async/1` is a struct that describes the running tasks. This struct can be passed to `Task.await/1` to await the result of the task:

```
iex(3)> Task.await(task)
:some_result
```

The function `Task.await/1` waits for the response message from the task process. This message will contain the result of the lambda. When the message arrives, `Task.await/1` returns the lambda's result. If the message doesn't arrive within five seconds, `Task.await/1` will raise an exception. You can provide a different timeout as the second parameter to `Task.await/2`.

Awaited tasks can be very useful when you need to run a couple of mutually independent one-off computations and wait for all the results. To illustrate this, we'll reuse the example from section 5.2.2. In that example, you needed to execute multiple independent queries and collect all the results. Because queries are mutually independent, you can improve the total execution time by running each query in a separate process and sending the result as a message to the starter process. The starter process then needs to await all the results.

Back in chapter 5, you implemented this from scratch using `spawn`, `send`, and `receive`. Here you'll rely on `Task.async/1` and `Task.await/1`.

First, define a helper lambda that simulates a long-running query execution:

```
iex(1)> run_query =
          fn query_def ->
            Process.sleep(2000)
            "#{query_def} result"
          end
```

Now you can start five queries, each in a separate task:

```
iex(2)> queries = 1..5
```

```
iex(3)> tasks =
          Enum.map(
            queries,
            &Task.async(fn -> run_query.("query #{&1}") end)
          )
```

Here you create five queries, and then start each query execution in a separate task. The result in the `tasks` variable is a list of five `%Task{}` structs, each describing one task executing a query.

To wait for all the results, you pass each task from the `tasks` variable to the `Task.await/1` function:

```
iex(4)> Enum.map(tasks, &Task.await/1)
["query 1 result", "query 2 result", "query 3 result", "query 4 r
 "query 5 result"]
```

Using the pipe operator, you can write this code in a slightly shorter way:

```
iex(5)> 1..5
        |> Enum.map(&Task.async(fn -> run_query.("query #{&1}") e
        |> Enum.map(&Task.await/1)

["query 1 result", "query 2 result", "query 3 result", "query 4 r
 "query 5 result"]      #1
```

The fact that all the results are collected in two seconds proves that each task is running in a separate process.

It should be noted that `Task.async/1` links the new task to the starter process. Therefore, if any task process crashes, the starter process will crash too (unless it's trapping exits). The crash of the starter process will, in turn, cause all the other tasks started by the same process to crash. In other words, starting multiple tasks with `Task.async/1` has all-or-nothing semantics. The crash of a single task takes down all other tasks as well as the starter process.

If you want to explicitly handle failures of individual tasks, you'll need to trap exits and handle corresponding exit messages in the starter process. There are some functions available in the `Task` module that can help you here, most notably `Task.async_stream/3`. You can refer to the official documentation at hexdocs.pm/elixir/Task.html for more details.

In the meantime, let's take a look at how you can work with tasks when the starter process doesn't need to wait for their result.

10.1.2 Non-awaited tasks

If you don't want to send the result message back to the starter process, you can use `Task.start_link/1`. This function can be thought of as an OTP-compliant wrapper around plain `spawn_link`. The function starts a separate process and links it to the caller. Then the provided lambda is executed in the started process. Once the lambda finishes, the process terminates with the reason `:normal`. Unlike `Task.async/1`, `Task.start_link/1` won't send any message to the starter process. Here's a basic example:

```
iex(1)> Task.start_link(fn ->
          Process.sleep(1000)
          IO.puts("Hello from task")
        end)

{:ok, #PID<0.89.0>}     #1

Hello from task!    #2
```

Let's look at a more concrete example. Suppose you want to gather some metrics about your system and report them at regular intervals. This is an example of a non-responsive job. You don't really need a `GenServer` here because you don't need to serve requests from other client processes. Instead, you want a process that sleeps for a while and then gathers relevant metrics and reports them.

Let's start implementing this in your to-do system. First you'll implement a sequential loop that periodically gathers metrics and prints them to the screen. The code is provided in the following listing.

Listing 10.1. Reporting system metrics (todo_metrics/lib/todo/metrics.ex)

```
defmodule Todo.Metrics do
  ...

  defp loop() do
    Process.sleep(:timer.seconds(10))
```

```
      IO.inspect(collect_metrics())
      loop()
    end

    defp collect_metrics() do
      [
        memory_usage: :erlang.memory(:total),
        process_count: :erlang.system_info(:process_count)
      ]
    end
end
```

In real life you'd likely want to collect much more data and send it to an external service, but in this example we'll keep things simple.

You want to run this loop as a part of your system. To do this, you need to start a task.

Listing 10.2. Metrics reporter as task (todo_metrics/lib/todo/metrics.ex)

```
defmodule Todo.Metrics do
  use Task

  def start_link(_arg), do: Task.start_link(&loop/0)

  ...
end
```

First you specify `use Task`, which will inject the `child_spec/1` function into the `Todo.Metrics` module. Just like with `GenServer`, the injected specification will invoke `start_link/1`, so you need to define `start_link/1` even if you don't use the argument. The implementation of `start_link/1` simply invokes `Task.start_link/1` to start a task process where the loop is running.

With these two simple lines of code, the `Todo.Metrics` module is ready to be injected into the supervision tree, as shown in the following listing.

Listing 10.3. Starting a supervised metrics task (todo_metrics/lib/todo/system.ex)

```
defmodule Todo.System do
  def start_link do
```

```
    Supervisor.start_link(
      [
        Todo.Metrics,
        ...
      ],
      strategy: :one_for_one
    )
  end
end
```

This is the main purpose of `Task.start_link/1` — it allows you to start an OTP-compliant process that you can safely run as a child of some supervisor.

Try it out:

```
$ iex -S mix

iex(1)> Todo.System.start_link()

[memory_usage: 48110864, process_count: 74]     #1
[memory_usage: 48505592, process_count: 74]     #2
```

This was a simple way of implementing a periodic job in your system, without needing to run multiple OS processes and use external schedulers such as `cron`.

In more complex scenarios, it's worth separating scheduling from the job logic. The idea is to use one process for periodic scheduling, and then start each job instance in a separate one-off process. Such an approach improves fault-tolerance, because the crash of a job process won't disturb the scheduling process. You can try to implement this approach as an exercise, but when it comes to production, it's better to rely on battle-tested third-party libraries, such as Quantum (github.com/quantum-elixir/quantum-core).

10.1.3 Supervising dynamic tasks

In many situations you'll want to start non-awaited tasks dynamically. A frequent example is when you need to communicate with a remote service, such as a payment gateway, while handling a web request.

A naive approach would be to perform this communication synchronously,

while the request is being handled. However, this approach can lead to poor user experience. If there are some intermittent networking issues, communicating with a remote service might be slow, or it might completely fail.

A better approach is to perform this communication asynchronously, from a separate task process. You accept the incoming request, start a task that communicates with the remote service, and immediately respond that the request has been accepted. Once the task is done, you issue a notification about the outcome, perhaps via WebSocket or an email. This improves system responsiveness, and enhances system resilience to various networking issues. You could retry the failed communication over a longer period of time, without needing to block the request handler, or even keep the user connected.

This is an example of a dynamically started independent task. The task is started on-demand, and its lifecycle must be decoupled from the lifecycle of the process which started it (the process which handles the incoming request).

In such situations it's best to run the task under a dedicated supervisor. You could use `DynamicSupervisor` for this purpose, but Elixir includes a task-specific wrapper around it called `Task.Supervisor` (hexdocs.pm/elixir/Task.Supervisor.html).

To run dynamically supervised tasks, start the task supervisor:

```
iex(1)> Task.Supervisor.start_link(name: MyTaskSupervisor)
```

Now you can use `Task.Supervisor.start_child/2` to start a task under that supervisor:

```
iex(2)> Task.Supervisor.start_child(
          MyTaskSupervisor,
          fn ->
            IO.puts("Task started")
            Process.sleep(2000)
            IO.puts("Task stopping")
          end
        )
```

```
{:ok, #PID<0.118.0>}      #1

Task started              #2
Task stopping             #3
```

It's important to understand the distinction between the logical and the actual
starter process. The shell process is the process which initiated the task
creation. However, the task is actually started as a child of the task
supervisor. As a result of this process structure, the lifecycles of the logical
starter (the iex shell process) and the task are separated. A crash of one
process won't affect the other.

This concludes our brief tour of tasks. We haven't covered all the nuances, so
I advise you to study the official module documentation in more detail at
hexdocs.pm/elixir/Task.html.

Next, we'll take a look at agents.

10.2 Agents

The `Agent` module provides an abstraction that's similar to `GenServer`.
Agents require a bit less ceremony and can therefore eliminate some
boilerplate associated with `GenServer`. On the flip side, `Agent` doesn't
support all the scenarios that `GenServer` does. As a mechanical rule, if a
`GenServer` implements only `init/1`, `handle_cast/2`, and `handle_call/3`, it
can be replaced with an `Agent`. But if you need to use `handle_info/2` or
`terminate/1`, `Agent` won't suffice, and you'll need to use `GenServer`.

Let's explore this further, starting with the basic use of agents.

10.2.1 Basic use

To start an agent, you can use `Agent.start_link/1`:

```
iex(1)> {:ok, pid} = Agent.start_link(fn -> %{name: "Bob", age: 3
{:ok, #PID<0.86.0>}
```

`Agent.start_link/1` will start a new process and execute the provided
lambda in that process. Unlike a task, an agent process doesn't terminate

when the lambda is finished. Instead, an agent uses the return value of the lambda as its state. Other processes can access and manipulate an agent's state using various functions from the `Agent` module.

To fetch the agent's state, or some part of it, you can use `Agent.get/2`:

```
iex(2)> Agent.get(pid, fn state -> state.name end)
"Bob"
```

`Agent.get/2` takes the pid of the agent and a lambda. The lambda is invoked in the agent's process, and it receives the agent's state as the argument. The return value of the lambda is sent back to the caller process as a message. This message is received in `Agent.get/2`, which then returns the result to its caller.

To modify the agent's state you can use `Agent.update/2`:

```
iex(3)> Agent.update(pid, fn state -> %{state | age: state.age +
:ok
```

This will cause the internal state of the agent process to change. You can verify the change with `Agent.get/2`:

```
iex(2)> Agent.get(pid, fn state -> state end)
%{age: 31, name: "Bob"}
```

It's worth mentioning that `Agent.update/2` is synchronous. The function only returns after the update has succeeded. An asynchronous update can be performed with `Agent.cast/2`.

There are some other functions available in the `Agent` module, so you're advised to study the official documentation at hexdocs.pm/elixir/Agent.html. In the meantime, let's discuss how agents work in a concurrent setting.

10.2.2 Agents and concurrency

A single agent, being a process, can be used by multiple client processes. A change made by one process can be observed by other processes in subsequent agent operations. Let's demonstrate this.

You'll start an agent that's used as a counter:

```
iex(1)> {:ok, counter} = Agent.start_link(fn -> 0 end)
```

The initial state of the agent is 0. Now manipulate the agent's state from another process:

```
iex(2)> spawn(fn -> Agent.update(counter, fn count -> count + 1 e
```

Finally, let's check the agent's state from the shell process:

```
iex(3)> Agent.get(counter, fn count -> count end)
1
```

This example demonstrates that the state is associated with the agent process. When one client process changes the state of the agent, subsequent operations issued by other processes will see the new state.

An agent process works exactly like a GenServer. If multiple clients try to work with the same agent at the same time, the operations will be serialized and executed one by one. In fact, the Agent module is implemented in plain Elixir on top of GenServer. To demystify this, let's sketch a naive implementation of an Agent-like module.

Here's how you can implement the agent-like state initialization:

```
defmodule MyAgent do
  use GenServer          #1

  def start_link(init_fun) do
    GenServer.start_link(__MODULE__, init_fun)     #2
  end

  def init(init_fun) do
    {:ok, init_fun.()}           #3
  end

  ...
end
```

Recall from chapter 5 that any term can be sent as a message. This includes anonymous functions, and the agent implementation takes advantage of that

fact. Agent interface functions take an anonymous function as an argument and pass the function to the server process, which in turn invokes the function and does something with its result.

The same approach is used to provide `get` and `update` operations:

```elixir
defmodule MyAgent do
  ...

  def get(pid, fun) do
    GenServer.call(pid, {:get, fun})
  end

  def update(pid, fun) do
    GenServer.call(pid, {:update, fun})
  end

  def handle_call({:get, fun}, _from, state) do
    response = fun.(state)
   {:reply, response, state}
  end

  def handle_call({:update, fun}, _from, state) do
    new_state = fun.(state)
    {:reply, :ok, new_state}
  end

  ...
end
```

The real implementation of the `Agent` module is more sophisticated and feature-rich, but the basic idea is the same as in the preceding example. The `Agent` module is a plain `GenServer` that can be controlled by sending lambdas to the process. Therefore, concurrent reasoning about agents is exactly the same as with `GenServer`.

10.2.3 Agent-powered to-do server

Because `Agent` can be used to manage concurrent state, it's a perfect candidate to power your to-do list server. Converting a `GenServer` into an agent is a fairly straightforward job. You need to replace a pair of interface functions and the corresponding `GenServer` callback clause with a single

function that uses the `Agent` API.

The full code of the `Todo.Server` as an agent is provided in the following listing.

Listing 10.4. Agent-powered to-do server (todo_agent/lib/todo/server.ex)

```
defmodule Todo.Server do
  use Agent, restart: :temporary

  def start_link(name) do
    Agent.start_link(
      fn ->
        IO.puts("Starting to-do server for #{name}")
        {name, Todo.Database.get(name) || Todo.List.new()}
      end,
      name: via_tuple(name)
    )
  end

  def add_entry(todo_server, new_entry) do
    Agent.cast(todo_server, fn {name, todo_list} ->
      new_list = Todo.List.add_entry(todo_list, new_entry)
      Todo.Database.store(name, new_list)
      {name, new_list}
    end)
  end

  def entries(todo_server, date) do
    Agent.get(
      todo_server,
      fn {_name, todo_list} -> Todo.List.entries(todo_list, date)
    )
  end

  defp via_tuple(name) do
    Todo.ProcessRegistry.via_tuple({__MODULE__, name})
  end
 end
```

It's worth noting that the interface of the module remains unchanged, so there's no need to modify the code of any other module.

There are two things worth discussing in this code. The first is the expression

use `Agent` at the start of the module. Just like with `GenServer` and `Task`, this expression will inject the default implementation of `child_spec/1`, allowing you to list the module in a child specification list.

In addition, the implementation of `add_entry/2` uses `Agent.cast/2`. This function is the asynchronous version of `Agent.update/2`, which means that the function returns immediately and the update is performed concurrently. `Agent.cast/2` is used here to keep the same behavior as in the previous version, where `GenServer.cast/2` was used.

Always wrap agent code in a module

One problem with agents is that they completely open the process's state. Recall that with `GenServer`, the state is private to the server and can only be manipulated via well-defined messages. With an `Agent`, though, the state can be manipulated in an arbitrary way through lambdas passed to `Agent` functions, which means the state is prone to accidental corruption. To guard against this problem, you're advised to always wrap an agent in a dedicated module, and to only manipulate the agent process through functions of that module. This is precisely what you did when you converted `Todo.Server` into an agent.

The new version of `Todo.Server` requires only 29 lines of code, which is somewhat shorter than the previous 41 lines of code in the `GenServer`. An agent definitely seems like an appealing alternative to `GenServer`.

But agents can't handle all the scenarios that `GenServer` can, so they're not always appropriate. In the next section we'll take a look at those limitations.

10.2.4 Limitations of agents

The `Agent` module can't be used if you need to handle plain messages, or if you want to run some logic on termination. In such cases, you need to use `GenServer`. Let's look at an example.

In the current version of your system, you never expire items from the to-do cache. This means that when a user manipulates a single to-do list, the list will remain in memory until the system is terminated. This is clearly not

good, because as users work with different to-do lists, you'll consume more and more memory until the whole system runs out of memory and blows up.

Let's introduce a simple expiry of to-do servers. You'll stop to-do servers that have been idle for a while.

One way to implement this is to create a single cleaning process that would terminate an idle to-do server. In this approach, each to-do server would need to notify the cleaning process every time it's been used, and that would cause the cleaning process to become a possible bottleneck. You'd end up with one process that needs to handle a possibly high load of messages from many other processes, and it might not be able to keep up.

A better approach is to make each to-do server decide on its own when it wants to terminate. This will simplify the logic and also avoid any performance bottlenecks. This is an example of something that can be done with `GenServer`, but can't be implemented with `Agent`.

Detecting an idle period in a `GenServer` can be done in a couple of ways, and here you'll use a simple approach. In values returned from `GenServer` callbacks, you can include one extra element at the end of the return tuple. This element, if it's an integer, represents an idle time after which the timeout message is sent to the `GenServer` process.

For example, in `init/1`, instead of returning `{:ok, initial_state}`, you can return `{:ok, initial_state, 1000}`. The value of 1000 states that if no call, cast, or plain message arrives to the server process in 1,000 milliseconds, the `handle_info/2` callback will be invoked, and the first argument will have the value of `:timeout`.

The same thing holds true for other callbacks, such as `handle_cast/2` and `handle_call/3`, where you can return `{:noreply, new_state, timeout}` and `{:reply, response, new_state, timeout}` respectively.

Therefore, to make the to-do server stop itself after a period of inactivity, you need to do the following:

1. Convert the implementation of the to-do server back to `GenServer`.

2. Include the idle timeout integer in all result tuples of all callback functions.
3. Add `handle_info/2` and stop the server if the `:timeout` message arrives.

Starting with the last `GenServer`-powered version of the `Todo.Server`, you're going to include the idle timeout integer in the callback functions' results. The code is provided in the following listing.

Listing 10.5. Specifying idle timeout (todo_cache_expiry/lib/todo/server.ex)

```
defmodule Todo.Server do
  ...

  @expiry_idle_timeout :timer.seconds(10)      #1

  def init(name) do
    IO.puts("Starting to-do server for #{name}.")
    {:ok, {name, nil}, {:continue, :init}}
  end

  def handle_continue(:init, {name, nil}) do
    todo_list = Todo.Database.get(name) || Todo.List.new()

    {
      :noreply,
      {name, todo_list},
      @expiry_idle_timeout #2
    }
  end

  def handle_cast({:add_entry, new_entry}, {name, todo_list}) do
    new_list = Todo.List.add_entry(todo_list, new_entry)
    Todo.Database.store(name, new_list)
    {:noreply, {name, new_list}, @expiry_idle_timeout}      #2
  end

  def handle_call({:entries, date}, _, {name, todo_list}) do
    {
      :reply,
      Todo.List.entries(todo_list, date),
      {name, todo_list},
      @expiry_idle_timeout          #2
    }
  end
```

```
    ...
end
```

First you declare a module attribute, `@expiry_idle_timeout`, which will contain the value of 10000 (obtained by invoking `:timer.seconds(10)`). This attribute serves as a module-level constant, which you include as the last element of each return tuple of almost every callback function.

The only exception is `init/1` where the continuation tuple (`{:continue, …}`) is provided. As explained in chapter 7, this is used to perform a potentially long-running server initialization without blocking the caller. `GenServer` doesn't support returning both, the continue info and the timeout. Therefore, in `init/1` you only include the continuation, while the timeout is provided in `handle_continue`. This is perfectly fine, because `handle_continue` is unconditionally invoked immediately after `init/1`.

These changes ensure that `handle_info(:timeout, state)` will be invoked when there's no activity in the server process for 10 seconds.

Finally, you need to handle the `:timeout` message and stop the server, as shown next.

Listing 10.6. Stopping an idle to-do server (todo_cache_expiry/lib/todo/server.ex)

```
defmodule Todo.Server do
  ...

  def handle_info(:timeout, {name, todo_list}) do
    IO.puts("Stopping to-do server for #{name}")
    {:stop, :normal, {name, todo_list}}      #1
  end
end
```

Now quickly verify if expiration works properly. Go to the todo_cache_expiry folder, start the system, and start one to-do server:

```
$ iex -S mix

iex(1)> Todo.System.start_link()
iex(2)> pid = Todo.Cache.server_process("bobs_list")
```

Now wait for a while, and you should see the debug message:

```
Stopping to-do server for bobs_list
```

Finally, verify whether the process is still alive:

```
iex(3)> Process.alive?(pid)
false
```

This is an example of a scenario where agents just won't suffice, and you need to use `GenServer`. But until you wanted to implement expiry, agents were just as appropriate a solution as `GenServer`. As long as you don't need to handle plain messages or you don't need to run some termination code in `terminate/1`, you can use `Agent`.

Personally, I mostly don't use `Agent` and start immediately with `GenServer`. Because converting an `Agent` into a `GenServer` requires some work, I'd much rather start with `GenServer` immediately. As an added bonus, this keeps the code more uniform, because all the server processes are implemented using the same abstraction. If you feel confused and aren't sure whether to use `Agent` or `GenServer`, my advice is to always go for `GenServer` because it covers more scenarios and it's not much more complicated than `Agent`.

This concludes the story about agents. Next, we'll take a look at a feature called ETS tables.

10.3 ETS tables

ETS (Erlang Term Storage) tables are a mechanism that allows you to share some state between multiple processes in a more efficient way. ETS tables can be thought of as an optimization tool. Whatever you can do with an ETS table can also be done with `GenServer` or `Agent`, but the ETS version can often perform much better. But ETS tables can handle only limited scenarios, so often they can't replace server processes.

Typical situations where ETS tables can be useful are shared key/value structures and counters. Although these scenarios can also be implemented with `GenServer` (or `Agent`), such solutions might lead to performance and

scalability issues.

Let's look at a simple demonstration of those issues by implementing a concurrent key/value store with `GenServer`. First, let's look at the example use of such a store:

```
iex(1)> KeyValue.start_link()
{:ok, #PID<0.118.0>}

iex(2)> KeyValue.put(:some_key, :some_value)
:ok

iex(3)> KeyValue.get(:some_key)
:some_value
```

The full implementation of the `KeyValue` module is provided in the following listing.

Listing 10.7. GenServer-powered key/value store (key_value/lib/key_value.ex)

```
defmodule KeyValue do
  use GenServer

  def start_link do
    GenServer.start_link(__MODULE__, nil, name: __MODULE__)
  end

  def put(key, value) do
    GenServer.cast(__MODULE__, {:put, key, value})
  end

  def get(key) do
    GenServer.call(__MODULE__, {:get, key})
  end

  def init(_) do
    {:ok, %{}}
  end

  def handle_cast({:put, key, value}, store) do
    {:noreply, Map.put(store, key, value)}
  end

  def handle_call({:get, key}, _, store) do
    {:reply, Map.get(store, key), store}
```

```
    end
end
```

Nothing new happens here. The `KeyValue` module is a simple `GenServer` that holds a map in its state. The `put` and `get` requests boil down to invoking `Map.put/3` and `Map.get/2` in the server process.

Next you'll do some quick and inconclusive performance measurements of this key/value store. Go to the key_value folder and run the following command:

```
mix run -e "Bench.run(KeyValue)"
```

You haven't seen the `mix run` command before. This command compiles the project, starts a BEAM instance, and then executes the expression provided via the `-e` argument, which means that `Bench.run/1` is invoked. Once the function is done, the BEAM instance stops.

The `Bench` module, available in key_value/lib/bench.ex conducts a simple load test. It starts the `KeyValue` server and then performs operations on one million keys. For each key, the bench program executes 10 `put` operations. Each `put` is followed by a `get`, so in total the program performs 20,000,000 operations.

Once the test is done, the function prints the observed throughput:

```
mix run -e "Bench.run(KeyValue)"
953182 operations/sec
```

The throughput of about 950,000 operations/sec seems decent enough. But let's see how the key/value server performs when it must serve multiple client processes. You can verify this by providing the `:concurrency` option to `Bench.run`:

```
mix run -e "Bench.run(KeyValue, concurrency: 10000)"
735369 operations/sec
```

Somewhat unexpectedly, with 10,000 client processes you got worse throughput. What happened? The main problem is that despite having so many processes, there's just one key/value server process, so all of the

key/value operations are synchronized, as shown in figure 10.1.

Figure 10.1. Single process bottleneck
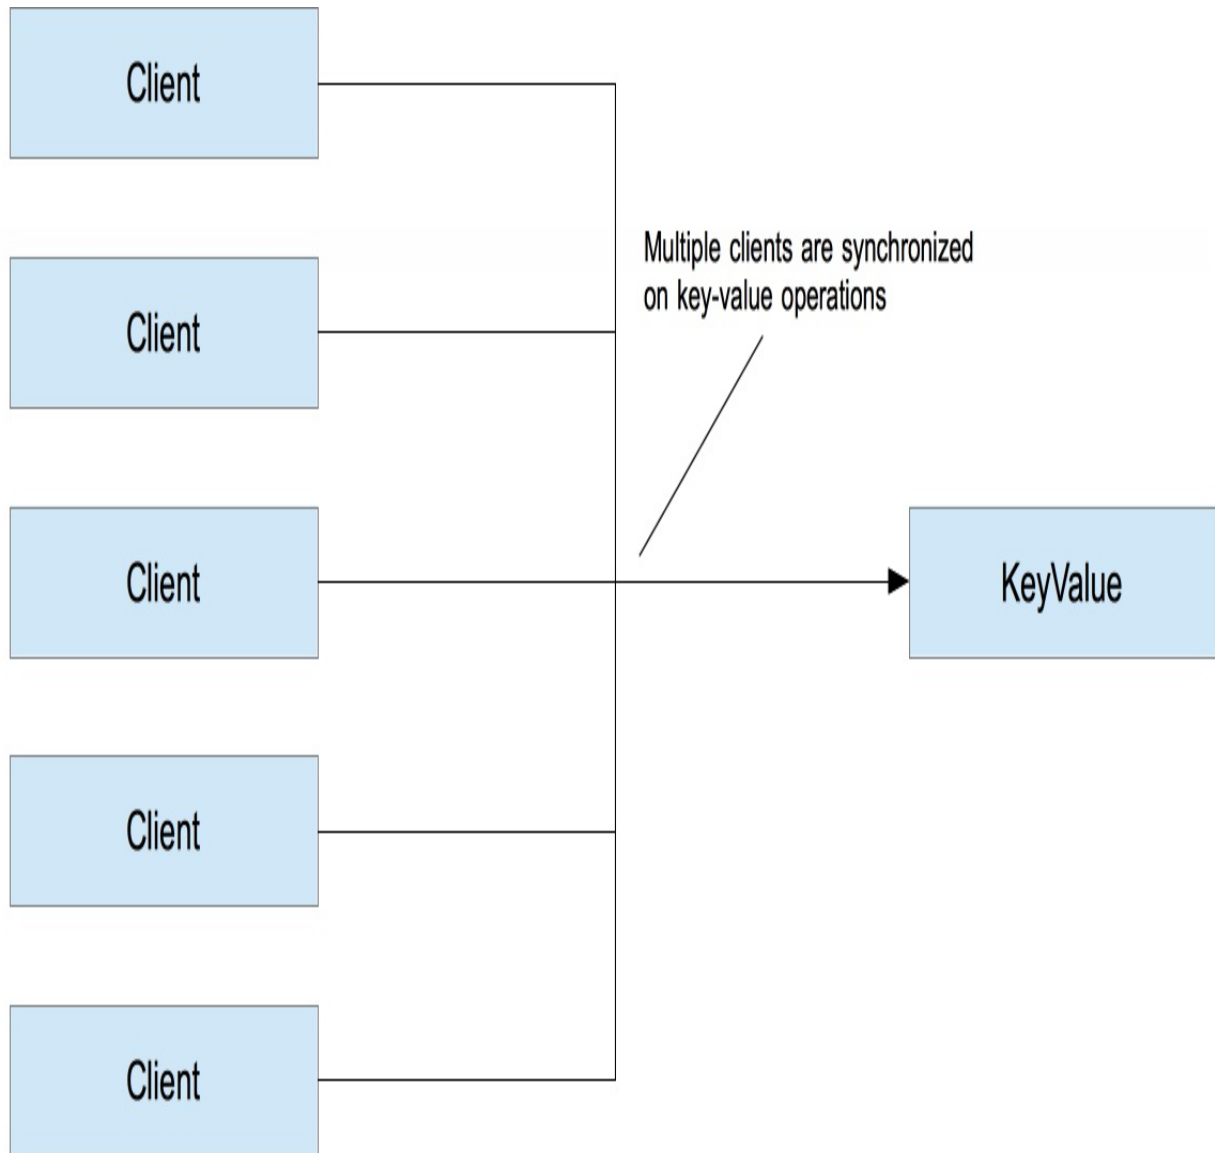
The key/value server therefore becomes a performance bottleneck and a scalability killer. The system isn't able to efficiently utilize all the hardware resources. Even though there are 10,000 client processes running, all the key/value operations are serialized.

Moreover, keep in mind that even in moderately concurrent systems, you usually run many more processes than there are CPU cores. In this case you

have 10,000 clients, which is many more than the number of CPU cores available. Consequently, not all processes can run at the same time — some have to wait their turn.

As explained in chapter 5, the VM goes to great lengths to use CPUs as well as possible, but the fact remains that you have many processes competing for limited resources. As a result, the key/value server doesn't get a single CPU core all to itself. The process must sometimes wait for its turn if BEAM schedulers run other processes in the system. Because the key/value server has fewer CPU resources for doing its job, it will take more time to compute the results.

This isn't to say that processes are bad. In general, you should strive to run independent tasks concurrently to improve scalability and fault tolerance. Processes should also be your first choice for maintaining state that changes over time. The problem isn't the many processes running in the system, but the single process on which many other processes depend.

In this particular scenario, you can do much better with ETS tables, so let's see what they are and how you can work with them.

10.3.1 Basic operations

ETS table is a special in-memory structure where you can store Erlang terms. This makes it possible to share some state between multiple processes without introducing a dedicated server process.

Compared to other data structures, ETS tables have some unusual characteristics:

- There's no specific ETS data type. A table is identified by its ID (a reference type) or a global name (an atom).
- ETS tables are mutable. A write to a table will affect subsequent read operations.
- Multiple processes can write to or read from a single ETS table. Writes and reads might be performed simultaneously.
- Minimum concurrency safety is ensured. Multiple processes can safely write to the same row of the same table. The last write wins.

- An ETS table resides in a separate memory space. Any data coming in or out is deep-copied.
- ETS doesn't put pressure on the garbage collector. Overwritten or deleted data is immediately released.
- An ETS table is deeply connected to its owner process (by default, the process that created the table). If the owner process terminates, the ETS table is reclaimed.
- Other than on owner-process termination, there's no automatic garbage collection of an ETS table. Even if you don't hold a reference to the table, it still occupies memory.

These characteristics mean ETS tables somewhat resemble processes. In fact, it's often said that ETS tables have process semantics. You could implement ETS table with processes, but such an implementation would be much less efficient. In BEAM, ETS tables are powered by C code, which ensures better speed and efficiency.

The fifth point in the previous list is especially interesting. Because data is deep-copied to and from an ETS table, there's no classical mutability problem. Once you read data from an ETS table, you have a snapshot that no one can change. Regardless of other processes possibly modifying the contents of those rows in the ETS table, the data you read remains unaffected.

Let's look at some examples. All functions related to ETS tables are contained in the Erlang `:ets` module (erlang.org/doc/man/ets.html). To create a table, you can call `:ets.new/2`:

```
iex(1)> table = :ets.new(:my_table, [])
#Reference<0.970221231.4117102596.53103>
```

The first argument is a table name, which is important only if you want to register the table (we'll discuss this in a minute). In addition, you can pass various options, some of which are discussed shortly. You should definitely spend some time researching the official documentation about `:ets.new/2` to see which options are possible.

The result of `:ets.new/2` is a reference, a unique opaque term that represents the ETS table in the running system.

Because the structure is a table, you can store multiple rows into it. Each row is an arbitrarily sized tuple (with at least one element), and each tuple element can contain any Erlang term, including a deep hierarchy of nested lists, tuples, maps, or anything else you can store in a variable.

To store data, you can use `:ets.insert/2`:

```
iex(2)> :ets.insert(table, {:key_1, 1})     #1
true

iex(3)> :ets.insert(table, {:key_2, 2})     #1
true

iex(4)> :ets.insert(table, {:key_1, 3})     #2
true
```

The first element of the tuple represents the *key* — something you can use for a fast lookup into the table. By default, ETS tables are of the *set* type, which means you can't store multiple tuples with the same key. Consequently, your last write overwrites the row from the first write.

To verify this, you can use `:ets.lookup/2`, which returns a list of rows for a given key:

```
iex(5)> :ets.lookup(table, :key_1)
[key_1: 3]

iex(6)> :ets.lookup(table, :key_2)
[key_2: 2]
```

You may wonder why the list is returned if you can have only one row per distinct key. The reason is that ETS tables support other table types, some of which allow duplicate rows. In particular, the following table types are possible:

- `:set` — The default. One row per distinct key is allowed.
- `:ordered_set` — Just like `:set`, but rows are in term order (comparison via the < and > operators).
- `:bag` — Multiple rows with the same key are allowed, but two rows can't be completely identical.
- `:duplicate_bag` — Just like `:bag`, but allows duplicate rows.

Another important option is the table's access permissions. The following values are possible:

- `:protected` — The default. The owner process can read from and write to the table. All other processes can read from the table.
- `:public` — All processes can read from and write to the table.
- `:private` — Only the owner process can access the table.

To create a table of a different type, or to use a different access level, you can simply include the desired option in the list passed to `:ets.new/2`. For example, to create a public duplicate bag list, you can invoke this:

```
:ets.new(:some_table, [:public, :duplicate_bag])
```

Finally, it's worth discussing the table name. This argument must be an atom, and by default it serves no purpose (although, strangely enough, you must still provide it). You can create multiple tables with the same name, and they're still different tables.

But if you provide a `:named_table` option, the table becomes accessible via its name:

```
iex(1)> :ets.new(:my_table, [:named_table])    #1
:my_table

iex(2)> :ets.insert(:my_table, {:key_1, 3})    #2

iex(3)> :ets.lookup(:my_table, :key_2)         #2
[]
```

In this sense, a table name resembles a locally registered process name. It's a symbolic name of a table, and it relieves you of having to pass around the ETS reference.

Trying to create a duplicate named table will result in an error:

```
iex(4)> :ets.new(:my_table, [:named_table])
** (ArgumentError) errors were found at the given arguments:
  * 2nd argument: invalid options
```

10.3.2 ETS powered key/value store

Equipped with this new knowledge, you're going to implement key/value store with ETS tables. Since a table has to be owned by some process, we'll still keep a `GenServer` around. Its sole purpose is to create the table in `init/1`, and keep it alive.

All interactions with the table (reads and writes) will be done directly from the caller process, without needing to issue a request to the server process.

The relevant code is contained in the same project as the initial attempt. First you need to start and initialize the table owner process, as shown in the next listing.

Listing 10.8. Creating the ETS table (key_value/lib/ets_key_value.ex)

```
defmodule EtsKeyValue do
  use GenServer

  def start_link do
    GenServer.start_link(__MODULE__, nil, name: __MODULE__)   #1
  end

  def init(_) do
    :ets.new(          #2
      __MODULE__,
      [:named_table, :public, write_concurrency: true]
    )

    {:ok, nil}
  end

  ...
end
```

In `start_link`, you start a `GenServer`. Then, in the `init/1` callback, the new ETS table is created. The table is configured as named, so the client processes can access it by its name (the name of the module). The access is set to public, which allows client processes to write to the table. The table type isn't provided, so it will default to `set`.

Notice the `:write_concurrency` option provided to `:ets.new`. This option allows you to issue concurrent writes to the table, which is exactly what you want in this case. There's also a `:read_concurrency` option, which can improve read performance in some cases. This option isn't set here, because the `Bench` module performs a lot of interleaved reads and writes, and in such cases `:read_concurrency` can lead to worse performance. Instead of randomly setting these options, it's always good to measure and observe their effects.

At this point, you can implement the operations:

```
defmodule EtsKeyValue do
  ...

  def put(key, value) do
    :ets.insert(__MODULE__, {key, value})     #1
  end

  def get(key) do
    case :ets.lookup(__MODULE__, key) do     #2
      [{^key, value}] -> value               #3
      [] -> nil     #4
    end
  end

  ...
end
```

The preceding code is a simple application of the presented `:ets` functions. To store an entry, you invoke `:ets.insert/2`. To perform a lookup, you invoke `:ets.lookup/2`. Because the ETS table is a `set`, the result list can contain at most one element — the key/value pair for the given key. If there's no row for the given key, the result is an empty list.

The crucial thing to notice here is that `get` and `put` operations now don't go through the server process. This means that multiple clients can work with the key/value store simultaneously, without blocking each other, as shown in figure 10.2.

Figure 10.2. Concurrency in an ETS-powered key/value store

As you can see, operations working on different keys can be executed in parallel. Multiple operations working on the same key will be properly synchronized to prevent possible race conditions. When you have a lot of keys, the chances for collisions are small, so you can expect better scheduler use, and therefore better scalability.

Now verify that the new key/value store works correctly:

```
iex(1)> EtsKeyValue.start_link()
{:ok, #PID<0.109.0>}

iex(2)> EtsKeyValue.put(:some_key, :some_value)
true
```

```
iex(3)> EtsKeyValue.get(:some_key)
:some_value
```

The key/value store seems to be working. Now let's see how it performs. You'll start with a sequential bench:

```
mix run -e "Bench.run(EtsKeyValue)"
5700062 operations/sec
```

On my machine, I obtained a projected throughput of about 5.7 million requests per second. Recalling that the pure `GenServer` version managed about 950,000 requests per second, this is an almost 6x increase in throughput!

There are a couple of reasons for this improvement. First, ETS operations are handled immediately in the client process. In contrast, a cross-process request involves putting a message in the mailbox of the receiver and then waiting for the receiver to be scheduled in and to handle the request. If the request is a synchronous call, the client process also has to wait for the response message to arrive.

In addition, changes to ETS tables are destructive. If a value under some key is changed, the old value is immediately released. Therefore, data managed in ETS tables doesn't put any pressure on a garbage collector. In contrast, transforming standard immutable data generates garbage. In a `GenServer`-based key/value store, frequent writes will generate a lot of garbage, and this means that the server process is occasionally blocked while it's being garbage-collected.

In this case, even in a plain sequential scenario, you get a significant improvement. But how does ETS hold up against multiple clients? Let's see:

```
mix run -e "Bench.run(EtsKeyValue, concurrency: 10000, num_update
52009220 operations/sec
```

Notice the `num_updates: 100` option. Because the ETS-based implementation is much faster, you're passing this option to run a longer test. This test will perform 100 `put` (and therefore also 100 `get`) operations on each key.

Using 10,000 client processes yields a 9x greater throughput. Compared to the plain `GenServer` solution, the improvement is 70x (52,000,000 vs. 735,000 requests per second). The single-process key/value server starts to slow down with an increase in the total number of running processes in the system. In contrast, the ETS-based cache scales better.

The main reason for this scaling lies in the fact that cache operations are executed in the client process, so you don't need to perform `GenServer`-based serialization. The atomic operations provided by the `:ets` module are properly synchronized and can safely run simultaneously in multiple processes. Operations working on different keys can run in parallel. Even the reads of the same key can work in parallel. Only writes will block other operations on the same key.

On the flip side, the vocabulary of write operations is small. You can perform key/value writes with `:ets.insert/2`, delete a row with `:ets.delete/2`, modify a row with `:ets.update_element/3`, and atomically update an integer in a row with `:ets.update_counter/4`. For more complex scenarios, you'll likely need to channel writes through a `GenServer`. Therefore, you can think of ETS tables as being an optimization tool. They're extremely efficient in simple scenarios, but not as powerful or flexible as server processes.

If you're unsure whether you should use a `GenServer` or an ETS table, it's best to start with a `GenServer`. This will be a simple solution, and in many cases the performance will be sufficient. If you establish that a particular server is a bottleneck, you can see if an ETS table would be a good fit. In many cases, moving to an ETS table will only require changing the implementation. For example, if you compare the `KeyValue` module to `EtsKeyValue`, you'll notice that they have the same public interface. That made it possible to have a generic `Bench` module that can work with both.

You might wonder why `GenServer` is still used in the ETS-based key/value store. The sole purpose of this process is to keep the table alive. Remember, an ETS table is released from memory when the owner process terminates. Therefore, you need to have a distinct, long-running process that creates and owns the table.

10.3.3 Other ETS operations

So far, we've covered only basic insertions and key-based lookups. These are arguably the most important operations you'll need, together with `:ets.delete/2`, which deletes all rows associated with a given key.

Key-based operations are extremely fast, and you should keep this in mind when structuring your tables. Your aim should be to maximize key-based operations, thus making ETS-related code as fast as possible.

Occasionally you may need to perform non-key-based lookups or modifications, retrieving a list of rows based on value criteria. There are a couple of ways you can do this.

The simplest but least efficient approach is to convert the table to a list using `:ets.tab2list/1`. You can then iterate over the list and filter out your results, such as by using functions from the `Enum` and `Stream` modules.

Another option is to use `:ets.first/1` and `:ets.next/2`, which make it possible to traverse the table iteratively. Keep in mind that this traversal isn't isolated. If you want to make sure no one modifies the table while you're traversing it, you should serialize all writes and traversals in the same process. Alternatively, you can call `:ets.safe_fixtable/2`, which provides some weak guarantees about traversal. If you're iterating a fixed table, you can be certain there won't be any errors, and each element will be visited only once. But an iteration through the fixed table may or may not pick up rows that are inserted during the iteration.

Traversals and `:ets.tab2list/1` aren't very performant. Given that data is always copied from the ETS memory space to the process, you end up copying the entire table. If you only need to fetch a couple of rows based on non-key criteria, this is overkill and a waste of resources.

A better alternative is to rely on *match patterns* — features that allow you to describe the data you want to retrieve.

Match patterns

Match patterns are a simple way to match individual rows. For example, let's say you're managing a to-do list in an ETS table:

```
iex(1)> todo_list = :ets.new(:todo_list, [:bag])
iex(2)> :ets.insert(todo_list, {~D[2023-05-24], "Dentist"})
iex(3)> :ets.insert(todo_list, {~D[2023-05-24], "Shopping"})
iex(4)> :ets.insert(todo_list, {~D[2023-05-30], "Dentist"})
```

Here you use a `bag` ETS table because it allows you to store multiple rows with the same key (date).

Most often, you'll want to query a table by key, asking, "What appointments are on the given date?"

```
iex(5)> :ets.lookup(todo_list, ~D[2023-05-24])
[{~D[2023-05-24], "Dentist"}, {~D[2023-05-24], "Shopping"}]
```

Occasionally you may be interested in obtaining all dates for an appointment type. Here's an example of how to do this using match patterns:

```
iex(6)> :ets.match_object(todo_list, {:_, "Dentist"})
[{~D[2023-05-24], "Dentist"}, {~D[2023-05-30], "Dentist"}]
```

The function `:ets.match_object/2` accepts a match pattern — a tuple that describes the shape of the row. The atom `:_` indicates that you accept any value, so the pattern `{:_, "Dentist"}` essentially matches all rows where the second element is `"Dentist"`.

Notice that this isn't classical pattern matching. Instead, this tuple is passed to `:ets.match_object/2`, which iterates through all rows and returns the matching ones. Therefore, when you don't care about a tuple element, you must pass an atom (`:_`) instead of a typical match-all anonymous variable (_). It's also worth mentioning the `:ets.match_delete/2` function, which can be used to delete multiple objects with a single statement.

In addition to being a bit shorter, match patterns have an important performance benefit over simple traversal. Recall that data is always copied from the ETS table to the selected process. If you used `:ets.tab2list/1` or plain traversal, you'd have to copy every single row into your own process. In contrast, `:ets.match_object/2` performs filtering in the ETS memory

space, which is more efficient.

Going beyond match patterns, it's possible to perform even richer queries, specifying more complex filters and even choosing individual fields you want to return. This is done by writing a full-blown *match specification* that consists of the following parts:

- *Head* — A match pattern describing the rows you want to select
- *Guard* — Additional filters
- *Result* — The shape of the returned data

Such specifications can be passed to the `:ets.select/2` function, which produces the corresponding result.

Match specifications can become complicated quickly, as you can see by looking at the documentation for `:ets.select/2` (erlang.org/doc/man/ets.html#select-2). To make that task simpler, take a look at the third-party library called ex2ms (github.com/ericmj/ex2ms).

Other use cases for ETS

Managing server-wide shared state is arguably the most common use case for ETS tables. In addition, ETS tables can be used to allow processes to persist their data. Remember from chapters 8 and 9 that processes lose their state on termination. If you want to preserve state across process restarts, the simplest way is to use a public ETS table as a means of providing in-memory state persistence. This should work reasonably quickly and allow you to recover from process crashes.

But be careful about taking this road. As mentioned in chapters 8 and 9, it's generally better to recover from a crash with clean state. You should also consider whether you can restore state based on data from other processes. Persisting state in the ETS table (or anywhere else, for that matter) should be used mostly for critical processes that are part of your error kernel.

It's also possible to use ETS tables as a faster alternative to immutable data structures, such as maps. Because changes to ETS tables are destructive and data is immediately released, there won't be any garbage-collection penalty

involved, so you can expect more predictable latency, with fewer deviations.

There is a caveat, though. Remember that data is copied between an ETS table and a client process. Consequently, if your row data is complex and large, ETS tables may yield worse performance than pure, immutable data structures. Another important downside of ETS tables is that unlike plain data, they can't be sent over the network to another BEAM instance. That means relying on ETS makes it harder to take advantage of distribution facilities (described in chapter 12).

In general, you should avoid using ETS and instead favor immutable structures as much as possible. Resort to ETS only in cases where you can obtain significant performance gains.

Beyond ETS

Erlang ships with two facilities that are closely related to ETS and that provide a simple way of implementing an embedded database that runs in the BEAM OS process. I won't discuss these features in detail in this book, but they deserve a brief mention so you can be aware that they exist and research them more deeply on your own.

The first feature, *disk-based* ETS (DETS, erlang.org/doc/man/dets.html) is disk-based term storage. Just like ETS, DETS relies on the concept of tables, and each table is managed in a single file. The interface of the corresponding `:dets` module is somewhat similar to ETS, but more limited in features. DETS provides a simple way of persisting data to disk. Basic isolation is supported — concurrent writes are allowed, even when you're storing to the same row.

Erlang also ships with a database called Mnesia (erlang.org/doc/apps/mnesia/users_guide.html), built on top of ETS and DETS, that has many interesting features:

- Mnesia is an embedded database — it runs in the same BEAM instance as the rest of your Elixir/Erlang code.
- Data consists of Erlang terms.

- Tables can be in-memory (powered by ETS) or disk-based (powered by DETS).
- Some typical database features are provided, such as complex transactions, dirty operations, and fast searches via secondary indices.
- Sharding and replication are supported.

These features make Mnesia a compelling option for storing data. You initialize the database from your startup Elixir/Erlang code, and you're good to go. This has the huge benefit of allowing you to run the entire system in a single OS process.

On the downside, Mnesia is a somewhat esoteric database and isn't used much outside the Elixir/Erlang community. This means there's less community and tooling support compared to popular DBMS solutions. It also takes some trickery to make Mnesia work on a larger scale. For example, one problem is that disk-based tables can't exceed 2 GB (this is a limitation of the underlying DETS storage), which means you have to fragment larger tables.

10.3.4 Exercise: process registry

Now is a good time to practice a bit. A textbook example of ETS in practice is a process registry. The `Registry` module uses a smart combination of `GenServer` and ETS to obtain maximum efficiency. In this exercise, you'll implement a basic version of a `:unique` registry.

Here's an example of using such a registry:

```
iex(1)> SimpleRegistry.start_link()
{:ok, #PID<0.89.0>}

iex(2)> SimpleRegistry.register(:some_name)     #1
:ok                                             #1

iex(3)> SimpleRegistry.register(:some_name)     #2
:error                                          #2

iex(4)> SimpleRegistry.whereis(:some_name)      #3
#PID<0.87.0>                                    #3

iex(5)> SimpleRegistry.whereis(:unregistered_name)   #4
nil
```

The interface of `SimpleRegistry` is very basic. The server process is started and registered locally. Then, any process can register itself by invoking `SimpleRegistry.register/1`, passing an arbitrary term for the process key. The function returns `:ok` on successful registration, or `:error` if the name is occupied. Lookup is done by invoking `SimpleRegistry.whereis/1`, which returns the pid for the given key, or `nil` if no process is registered under a given name.

In addition, the registry process can detect a termination of each registered process and remove all the registration entries for that process.

`SimpleRegistry` has no other fancy features of the `Registry` module, such as support for via tuples, duplicate registrations, or multiple registry instances.

Here's how you could build such a registry:

1. Implement the first version of `SimpleRegistry` as a `GenServer`. Both `register` and `whereis` will be implemented as calls.
2. The state of the `GenServer` should be a map, where keys are registered names and values are pids.
3. While handling the `:register` call, the registry process should link to the caller, so it can detect the process termination and deregister it. Therefore, the registry server also has to trap exits (by invoking `Process.flag(:trap_exit, true)` in `init/1`). Links are chosen over monitors to ensure that registered processes are informed if the registry process goes down.
4. The registry process should handle `{:EXIT, pid, reason}` in its `handle_info`, and remove all entries for the given process from the map.

Once you have this registry in place, you can consider moving some parts out of the server process. In particular, by using ETS tables, it's possible to perform both registration and lookup in client processes. Here's how:

1. During `init/1`, the registry process should create a named ETS table with public access. This table will map names to pids so the registry process doesn't need to maintain any state.
2. Registration can be done via `:ets.insert_new/2`

(erlang.org/doc/man/ets.html#insert_new-2). This function will only insert the new entry if there's no entry under the given key. Therefore, you can safely call this function simultaneously from separate processes. The function returns a boolean to indicate whether the entry has been inserted or not.

3. Prior to invoking `:ets.insert_new/2`, the caller process should link to the server process by invoking `Process.flag(:trap_exit, true)`. You can wrap this linking in `SimpleRegistry.register/1`.
4. The implementation of `whereis/1` boils down to invoking `:ets.lookup/2` and matching the result.
5. Finally, the server process still needs to handle `:EXIT` messages and remove the entries for the terminated processes. This can easily be done with the help of `:ets.match_delete/2`.

This exercise is a bit more involved than previous ones, but it's a nice synthesis of some techniques you've seen in the past few chapters. If you get stuck, take a look at the solution in the process_registry folder. There, you'll find both versions — the basic one implemented completely in a `GenServer`, and a more performant one that uses an ETS table to store registrations.

10.4 Summary

- Tasks can be used to run OTP-compliant concurrent job processes.
- Agents can be used to simplify the implementation of processes that manage some state but don't need to handle any plain messages.
- ETS tables can be used to improve performance in some cases, such as shared key/value memory structures.

11 Working with components

This chapter covers

- Creating OTP applications
- Working with dependencies
- Building a web server
- Configuring applications

It's time to turn our attention toward producing releasable systems that can be deployed and started. To reach that goal, you need to learn about OTP *applications,* which let you organize your system into reusable components. Applications are a standard way of building Elixir/Erlang production systems and libraries, and relying on them brings various benefits, such as dependency management, simplified system starting, and the ability to build standalone, deployable releases.

In this chapter, you'll learn how to create applications and work with dependencies. In the process, you'll turn your to-do system into a proper OTP application and use some third-party libraries from the Erlang and Elixir ecosystem to expose an HTTP interface for your existing system. There's a lot of work ahead, so let's get started with OTP applications.

11.1 OTP applications

An OTP application is a component that consists of multiple modules and that can depend on other applications. This makes it possible to start the entire system with a single function call. As you're about to see, it's reasonably easy to turn the code into an application. Your current version of the to-do system is already an OTP application, but there are some minor details you can improve. You'll see this in action shortly; first let's look at what OTP applications consist of.

11.1.1 Creating applications with the mix tool

An application is an OTP-specific construct. The resource that defines an application is called an *application resource file* — a plain-text file written in Erlang terms that describes the application (don't worry, you won't need to write this directly; you'll instead rely on the `mix` tool to do this for you). This file contains several pieces of information such as:

- The application name and version, and a description
- A list of application modules
- A list of application dependencies (which must be applications themselves)
- An optional application-callback module

Relying on the `mix` tool simplifies and automates some of the work of generating application resource files. For example, the application resource file must contain a list of all application modules. When you use `mix`, this list is generated for you automatically, based on the modules defined in the source code.

Some things, such as the application name, version, and description, must of course be provided by you. The `mix` tool can then use this data and your source code to generate the corresponding resource file while compiling the project.

Let's see this in practice. Go to a temporary folder, and run `mix new hello_world --sup`. This command creates the hello_world folder with the minimum `mix` project skeleton. The parameter `--sup` makes the `mix` tool generate the application callback module and start the empty (childless) supervisor from it.

You can now change to the hello_world folder and start the system with the familiar `iex -S mix`. On the surface, nothing spectacular happens. But `mix` automatically starts your application, which you can verify by calling `Application.started_applications/0`:

```
iex(1)> Application.started_applications()
[
  {:iex, ~c"iex", ~c"1.15.0"},
  {:hello_world, ~c"hello_world", ~c"0.1.0"},   #1
  {:logger, ~c"logger", ~c"1.15.0"},
```

```
  {:mix, ~c"mix", ~c"1.15.0"},
  {:elixir, ~c"elixir", ~c"1.15.0"},
  {:compiler, ~c"ERTS  CXC 138 10", ~c"8.3"},
  {:stdlib, ~c"ERTS  CXC 138 10", ~c"5.0"},
  {:kernel, ~c"ERTS  CXC 138 10", ~c"9.0"}
]
```

As you can see, the `hello_world` application is running, together with some additional applications such as Elixir's `mix`, `iex`, and `elixir`, and Erlang's `stdlib` and `kernel`.

You'll see the benefits of this shortly, but first let's look at how the application is described. The main place where you specify an application is in the mix.exs file. Here are the full contents of the generated file (comments are stripped out):

```
defmodule HelloWorld.MixProject do
  use Mix.Project

  def project do              #1
    [                         #1
      app: :hello_world,      #1
      version: "0.1.0",       #1
      elixir: "~> 1.15",      #1
      start_permanent: Mix.env() == :prod,    #1
      deps: deps()            #1
    ]                         #1
  end                         #1

  def application do          #2
    [                         #2
      extra_applications: [:logger],       #2
      mod: {HelloWorld.Application, []}  #2
    ]                         #2
  end                         #2

  defp deps do        #3
    []                #3
  end                 #3
end
```

The first interesting thing happens in the `project/0` function, where you describe the `mix` project. The `app: :hello_world` gives a name to your application. Only an atom is allowed as an application name, and you can use

this atom to start and stop the application at runtime.

The application is described in the `application/0` function. Here you specify some options that will eventually make it to the application resource file. In this case, the description includes the list of other Erlang and Elixir applications you depend on, together with the module that will be used to start the application. By default, Elixir's `:logger` application is listed (https://hexdocs.pm/logger/Logger.html).

Finally, the `deps` function returns the list of third-party dependencies — other libraries you want to use in your project. By default, this list is empty. You'll see how dependencies are used a bit later in this chapter.

11.1.2 The application behavior

The critical part of the application description is `mod:` `{HelloWorld.Application, []}`, provided in mix.exs by `application/0`. This part specifies the module that will be used to start the application. When the application is started, the function `HelloWorld.Application.start/2` is called.

Obviously, you need to implement the `HelloWorld.Application` module. This is done for you by the `mix` tool, so let's see what it looks like:

```
defmodule HelloWorld.Application do
  use Application                   #1

  def start(_type, _args) do      #2
    children = []                  #3
    opts = [strategy: :one_for_one, name: HelloWorld.Supervisor]
    Supervisor.start_link(children, opts)            #3
  end
end
```

An application is an OTP behaviour, powered by the `Application` module (https://hexdocs.pm/elixir/Application.html), which is a wrapper around Erlang's `:application` module (https://www.erlang.org/doc/man/application.html). To be able to work with `Application`, you must implement your own callback module and define some callback functions.

At minimum, your callback module must contain the `start/2` function. The arguments passed are the application start type (which you'll usually ignore) and an arbitrary argument (a term specified in mix.exs under the `mod` key). See the official documentation (https://hexdocs.pm/elixir/Application.html#c:start/2) for details.

The task of the `start/2` callback is to start the top-level process of your system, which should usually be a supervisor. The function returns its result in the form `{:ok, pid}`, or `{:error, reason}` if something went wrong.

11.1.3 Starting the application

To start the application in the running BEAM instance, you can call `Application.start/1`. This function first looks for the application resource file (which is generated by `mix`) and interprets its contents. Then it verifies whether all the applications you're depending on are started. Finally, the application is started by calling the callback module's `start/2` function. The `Application.ensure_all_started/2` function is also available, which recursively starts all dependencies that aren't yet started.

Usually you won't need to invoke these functions because `mix` automatically starts the application implemented by the project. Calling `iex -S mix` automatically starts the application together with its dependencies.

It should be noted that you can't start multiple instances of a single application. Trying to start an already-running application will return an error.

```
$ iex -S mix

iex(1)> Application.start(:hello_world)
{:error, {:already_started, :hello_world}}
```

You can stop the application using `Application.stop/1`:

```
iex(2)> Application.stop(:hello_world)
[info] Application hello_world exited: :stopped
```

Stopping the application terminates its top-level process. If that process is a

supervisor, it will stop its own children before it stops itself. This is why it's important to organize the processes in the supervision tree, as described in chapter 9. Doing so ensures that the application is stopped in the controlled manner, leaving no dangling processes behind.

`Application.stop/1` stops only the specified application, leaving dependencies (other applications) running. To stop the entire system in a controlled way, you can invoke `System.stop/0`. This function will take down all the OTP applications, and then the BEAM instance itself. Both `Application.stop/1` and `System.stop/0` work in a polite way. Every process in the supervision tree can perform some final cleanup in its `terminate/1` callback, as explained in section 9.1.6.

11.1.4 Library applications

You don't need to provide the `mod:` … option from the `application/0` function in mix.exs:

```
defmodule HelloWorld.Application do
  ...

  def application do
    []
  end

  ...
end
```

In this case, there's no application callback module, which in turn means there's no top-level process to be started. This is still a proper OTP application. You can even start it and stop it.

What's the purpose of such applications? This technique is used for *library applications*: components that don't need to create their own supervision tree. As the name indicates, these are usually libraries, such as a JSON or a CSV parser. Erlang's own `stdlib` application (https://erlang.org/doc/apps/stdlib/index.html) is a pure library application because it exposes various utility modules but doesn't need to manage its own supervision tree.

Library applications are useful because you can list them as runtime dependencies. This plays an important role when you start to assemble the deployable release, as you'll see in chapter 13.

11.1.5 Implementing the application callback

Equipped with this knowledge, you can turn your to-do system into a proper application. As mentioned earlier, it's already an OTP application, albeit a library one that doesn't implement a callback module for the application behavior. In fact, there is usually a 1:1 relationship between a mix project and an OTP application. A mix project implements exactly one OTP application. An exception is a so called umbrella project which can contain multiple OTP applications.

Given that the to-do system runs a set of its own processes under a supervision tree, it makes sense to implement the application callback module. Once you do that, the system can be automatically started as soon as you run `iex -S mix`.

The first thing you need is edit the mix.exs file. Recall that you created your initial project back in chapter 7 using the `mix` tool. Therefore, you already have this file in place and only need to add some information. The complete code for mix.exs is provided in the following listing.

Listing 11.1. Specifying application parameters (todo_app/mix.exs)

```
defmodule Todo.MixProject do
  use Mix.Project

  def project do
    [
      app: :todo,
      version: "0.1.0",
      elixir: "~> 1.15",
      start_permanent: Mix.env() == :prod,
      deps: deps()
    ]
  end

  def application do
    [
```

```
      extra_applications: [:logger],
      mod: {Todo.Application, []}        #1
    ]
  end

  defp deps do
    []
  end
end
```

The only change to mix.exs is in the `application/0` function, where the callback module is specified.

Next, you need to implement the callback module. The code is shown next.

Listing 11.2. Implementing the application module (todo_app/lib/todo/application.ex)

```
defmodule Todo.Application do
  use Application

  def start(_, _) do
    Todo.System.start_link()
  end
end
```

As already mentioned, starting the application is as simple as starting the top-level supervisor. Given that you've already structured your system to reside under the `Todo.System` supervisor, this is all it takes to turn your system into a full-blown OTP application.

Let's see it in action:

```
$ iex -S mix

Starting database server             #1
Starting database worker             #1
Starting database worker             #1
Starting database worker             #1
Starting to-do cache                 #1
```

By implementing the OTP application callback, you've made it possible to start the system automatically.

It's worth noting that you get the same benefit when running tests. When you invoke `mix test`, all of the essential processes in your system are started. Back in chapter 7, when you added a couple of tests, you had to manually start the cache process:

```
defmodule Todo.CacheTest do
  use ExUnit.Case

  test "server_process" do
    {:ok, cache} = Todo.Cache.start()      #1
    bob_pid = Todo.Cache.server_process(cache, "bob")

    assert bob_pid != Todo.Cache.server_process("alice")
    assert bob_pid == Todo.Cache.server_process("bob")
  end

  ...
end
```

The code has since seen a bunch of transformations, but in every incarnation there was some version of the preceding pattern. Inside `Todo.CacheTest` you needed to manually start the supporting processes, such as cache. With the system turned into a proper OTP application, this is not the case anymore, as you can see in the following listing.

Listing 11.3. Testing server_process (todo_app/test/todo_cache_test.exs)

```
defmodule Todo.CacheTest do
  use ExUnit.Case

  test "server_process" do
    bob_pid = Todo.Cache.server_process("bob")         #1

    assert bob_pid != Todo.Cache.server_process("alice")
    assert bob_pid == Todo.Cache.server_process("bob")
  end

  ...
end
```

11.1.6 The application folder structure

Let's briefly discuss your compiled application's folder structure. Owing to

the `mix` tool, you usually won't need to worry about this, but it can sometimes be useful to understand the folder structure of a compiled system.

Mix environments

Before we look at the application structure, you should know a bit about *mix environment.* This is a compile-time option that can be used to affect the shape of the compiled code.

Mix projects use three environments: dev, test, and prod. These three environments produce slight variations in the compiled code. For example, in a version compiled for development (the dev environment), you'll likely want to run some extra debug logging, whereas in a version compiled for production (the prod environment) you don't want to include such logging. In a version compiled for tests (the test environment), you want to further reduce the amount of logging, and use a different database to prevent the tests from polluting your development database.

You can introduce your own mix environment if you want to, but this is rarely needed, if ever. The three environment mentioned here should be sufficient to cover all possible scenarios.

For most `mix` tasks, the default environment is dev, indicating that you're dealing with development. One exception to that rule is the test task. When you invoke `mix test`, the mix environment is automatically set to test.

You can specify a mix environment by setting the `MIX_ENV` OS environment variable. By convention, when building for production, you should use the prod environment. To compile the code for prod, you can invoke `MIX_ENV=prod mix compile`. To compile and start a prod version you can invoke `MIX_ENV=prod iex -S mix`.

The prod environment is a frequent source of confusion. Many teams conflate prod with the production machine. A prod version indeed runs on the production machine, but it will also run on other machines, such as staging. Most importantly, it should be possible and simple to start a version compiled for prod on a local developer machine. You won't need to do this frequently,

but it's occasionally useful to analyze the behavior of the system which is as close to the deployed production version. This is exactly what the prod-compiled version is. Therefore, you're advised to support running prod-compiled version locally. Commit to this goal from the start of the project, because doing this later, after the codebase grows significantly, is typicallt much more difficult.

The compiled code structure

Once you compile your project, compiled binaries reside in the _build/project_env folder, where project_env is the mix environment that was in effect during compilation.

Because dev is the default environment, if you run `mix compile` or `iex -S mix`, you get binaries in the _build/dev folder. The OTP itself recommends the following folder convention:

```
lib/
  App1/
    ebin/
    priv/

  App2/
    ebin/
    priv/
  ...
```

Here, App1 and App2 represent application names (such as `todo`). The ebin folder contains compiled binaries (.beam files and application resource files), whereas the priv folder contains application-specific private files (images, compiled C binaries, and so on). This isn't a mandatory structure, but it's a convention used in most Elixir/Erlang projects, including Elixir and Erlang. Some tools may rely on this structure, so it's best to follow this convention.

Luckily, you don't need to maintain this structure yourself because `mix` does so automatically. The final folder structure of a compiled `mix` project has the following shape:

```
YourProjectFolder
_build
```

```
dev
  lib
    App1
      ebin
      priv

    App2
    ...
```

In addition to your application, the lib folder also contains the application dependencies, except for the standard applications included in Elixir and Erlang, which reside in the folder where Elixir and Erlang are installed, and which are accessible via the load path.

As mentioned, the application resource file resides in lib/YourApp/ebin and is named YourApp.app. For the to-do system, the file resides in _build/dev/lib/todo/ebin/ (relative to the root project folder). When you attempt to start the application, the generic application behavior looks for the resource file in the load paths (the same paths that are searched for compiled binaries).

This concludes our discussion of application basics. Now that you have a grasp of some theory, let's look at how to work with dependencies.

Deployable systems

Applications play an important role in building a deployable system. You'll learn about this in chapter 13, where we'll discuss OTP releases. In short, the general idea is to assemble a minimal self-contained system that includes only required applications and the Erlang runtime. For this to work, you must turn your code into an OTP application, because only then can you specify dependencies to other applications. This is ultimately used to bundle all required applications into a single deployable release.

Therefore, anything reusable in Elixir/Erlang should reside in an application. This holds even for Elixir/Erlang: examples include the `elixir` application, which bundles the Elixir standard library, as well as `iex` and `mix`, which are implemented as separate applications. The same thing happens in Erlang, which is divided into many applications

(https://www.erlang.org/doc/applications.html), such as `kernel` and `stdlib`.

11.2 Working with dependencies

Depending on third-party libraries is an important feature. As soon as you start developing more complicated projects, you'll probably want to use various libraries, such as web frameworks, the JSON parser, and database drivers, to name a few examples. Such libraries can simplify various side concerns, prevent you from going down too many rabbit holes, and allow you to focus on the core challenges of your system.

For example, in the current implementation of the to-do system, you maintain a small pool of database workers, which allows you to have controlled parallelism of database operations. You implemented this pool completely from scratch back in chapter 7. But as it happens, managing a pool of processes is a frequent pattern in Elixir and Erlang, so a couple of third-party libraries offer a solution to this challenge. In this chapter, you'll replace your naive implementation of the process pool with a proven and battle-tested process pool library.

11.2.1 Adding a dependency

In this section, you'll add a dependency to the Poolboy library (https://github.com/devinus/poolboy). This library provides a mature implementation of a process pool.

The dependency to an external library must be specified in the mix.exs file, as illustrated in the next listing.

Listing 11.4. Adding an external dependency (todo_poolboy/mix.exs)

```
defmodule Todo.MixProject do
  ...

  defp deps do
    [
      {:poolboy, "~> 1.5"}     #1
    ]
  end
```

```
end
```

An external dependency is specified as a tuple. The first element is always an atom, which should correspond to the application name of the dependency. The second element in the tuple, `"~> 1.5"`, is the version requirement. Here you indicate that you want version 1.5 or any later 1.x version. For more information about the version syntax, take a look at the official documentation at https://hexdocs.pm/elixir/Version.html#module-requirements.

At this point, you've specified that your project depends on an external library, so you need to fetch your dependencies. This can be done by running `mix deps.get` from the command line. Dependencies are fetched from Elixir's external package manager, which is called Hex (https://hex.pm). Other possible dependency sources include the GitHub repository, a Git repository, or a local folder. For more information, take a look at the official documentation at https://hexdocs.pm/mix/Mix.Tasks.Deps.html.

Running `mix deps.get` fetches all dependencies (recursively) and stores the reference to the exact version of each dependency in the mix.lock file, unless mix.lock already exists on the disk, in which case this file is consulted to fetch the proper versions of dependencies. This ensures reproducible builds across different machines, so make sure you include mix.lock in the source control where your project resides.

Now that you've fetched all of your dependencies, you can build the entire system by running `mix compile`, which will compile all the dependencies and the project. It's worth mentioning that Poolboy is an Erlang library, but `mix` will still know how to compile it.

11.2.2 Adapting the pool

With these preparations in place, you can start adapting the pool implementation.

Using Poolboy requires starting a process called pool manager that manages a pool of workers. While starting the pool manager, you pass the desired pool size (the number of worker processes), and the module that powers each

worker. During its startup, the pool manager starts the worker processes as its children.

Other processes can ask the pool manager to give them the pid of one worker. This operation is called *checkout*. Once a process gets the worker's pid, it can issue requests to that worker. When the client process doesn't need the worker process anymore, it notifies the pool manager. This operation is called *checkin*.

This workflow is a bit more sophisticated than your own simple pool. The checkout and checkin operations allow the pool manager to keep track of which worker processes are being used. If some workers are available, a client can immediately get a worker. Otherwise, if all the workers are checked out, the client will have to wait. As soon as the first worker is returned to the pool, a waiting client will check out that worker.

Poolboy also relies on monitors and links to detect the termination of a client. If a client checks out a worker and then crashes, the pool manager process will detect it and return the worker to the pool. Likewise, if a worker process crashes, a new one will be started.

Equipped with this knowledge, you're going to adapt the pool. First, you need to start the pool manager somewhere in your supervision tree. You could do that by invoking `:poolboy.start_link`, but there's a slightly more elegant way. You can invoke `:poolboy.child_spec/3`, which describes how Poolboy is supposed to be started. Therefore, to switch the database to the Poolboy-powered pool, you only need to change the implementation of `Todo.Database.child_spec/1`. The code is provided in the following listing.

Listing 11.5. Starting a Poolboy-powered pool (todo_poolboy/lib/todo/database.ex)

```
defmodule Todo.Database do
  @db_folder "./persist"

  def child_spec(_) do
    File.mkdir_p!(@db_folder)

    :poolboy.child_spec(
      __MODULE__,            #1
      [                          #2
```

```
      name: {:local, __MODULE__},
      worker_module: Todo.DatabaseWorker,
      size: 3
    ],

    [@db_folder]      #3
  )
end

...
end
```

The first argument passed to `:poolboy.child_spec/3` is the ID of the child. This information is needed by the parent supervisor. Here, you use the module name (`Todo.Database`) as the ID.

The second argument is the pool options. The `:name` option states that the pool manager process should be locally registered, so you can interact with it without needing to know its pid. The `:worker_module` option specifies the module that will power each worker process, whereas `:size` specifies the pool size.

The final argument to `:poolboy.child_spec/3` is a list of arguments passed to the `start_link` of each worker when they're being started.

These particular arguments state that you want three worker processes, with each worker powered by the `Todo.DatabaseWorker` module. The pool manager will start each worker by invoking `Todo.DatabaseWorker.start_link(@db_folder)`.

It's worth noting that with this change, you don't need `Todo.Database.start_link` anymore because the new specification states that the database should be started by invoking `:poolboy.start_link`.

Next you need to adapt the `store` and `get` functions of `Todo.Database`. Previously these functions selected a worker ID, and then forwarded this ID to the corresponding function in the `Todo.DatabaseWorker` module. Now these functions need to check out the worker from the pool, make a request to the worker, and return the worker to the pool. All of this can be easily done with the function `:poolboy.transaction/2`. The code is provided in the

following listing.

Listing 11.6. Adapted operation functions (todo_poolboy/lib/todo/database.ex)

```
defmodule Todo.Database do
  ...

  def store(key, data) do
    :poolboy.transaction(        #1
      __MODULE__,
      fn worker_pid ->           #2
        Todo.DatabaseWorker.store(worker_pid, key, data)
      end
    )
  end

  def get(key) do
    :poolboy.transaction(
      __MODULE__,
      fn worker_pid ->
        Todo.DatabaseWorker.get(worker_pid, key)
      end
    )
  end
end
```

Here you invoke `:poolboy.transaction/2`, passing the registered name of the pool manager. This will issue a checkout request to fetch a single worker. Once a worker is available, the provided lambda is invoked. When the lambda is finished, `:poolboy.transaction/2` will return the worker to the pool.

In the provided lambda, you get the pid of the checked-out worker and issue a `Todo.DatabaseWorker` request. This means that you need to slightly change the implementation of `Todo.DatabaseWorker`. Previously the functions from this module accepted a worker ID and then did a registry lookup to find the worker process. In this version, the discovery of the pid is performed by the pool, so you don't need to register the workers anymore or perform any lookup to discover the process. The new implementation is provided in the following listing.

Listing 11.7. Adapted worker interface functions (todo_poolboy/lib/todo/database_worker.ex)

```
defmodule Todo.DatabaseWorker do
  use GenServer

  def start_link(db_folder) do
    GenServer.start_link(__MODULE__, db_folder)
  end

  def store(pid, key, data) do
    GenServer.cast(pid, {:store, key, data})
  end

  def get(pid, key) do
    GenServer.call(pid, {:get, key})
  end

  ...
end
```

That's all it took to change the pool implementation. The code in the client modules, `Todo.System`, and `Todo.Server`, as well as the testing code remains exactly the same. This is thanks to the fact that you hid the implementation details behind the interface functions.

11.2.3 Visualizing the system

Once you have a full-blown OTP application, you can visualize it with the help of a tool called `observer`, which is part of the standard Erlang/OTP distribution.

Let's see this in action. First, start the system, and then create two to-do servers:

```
$ iex -S mix

iex(1)> Todo.Cache.server_process("Alice")
iex(2)> Todo.Cache.server_process("Bob")
```

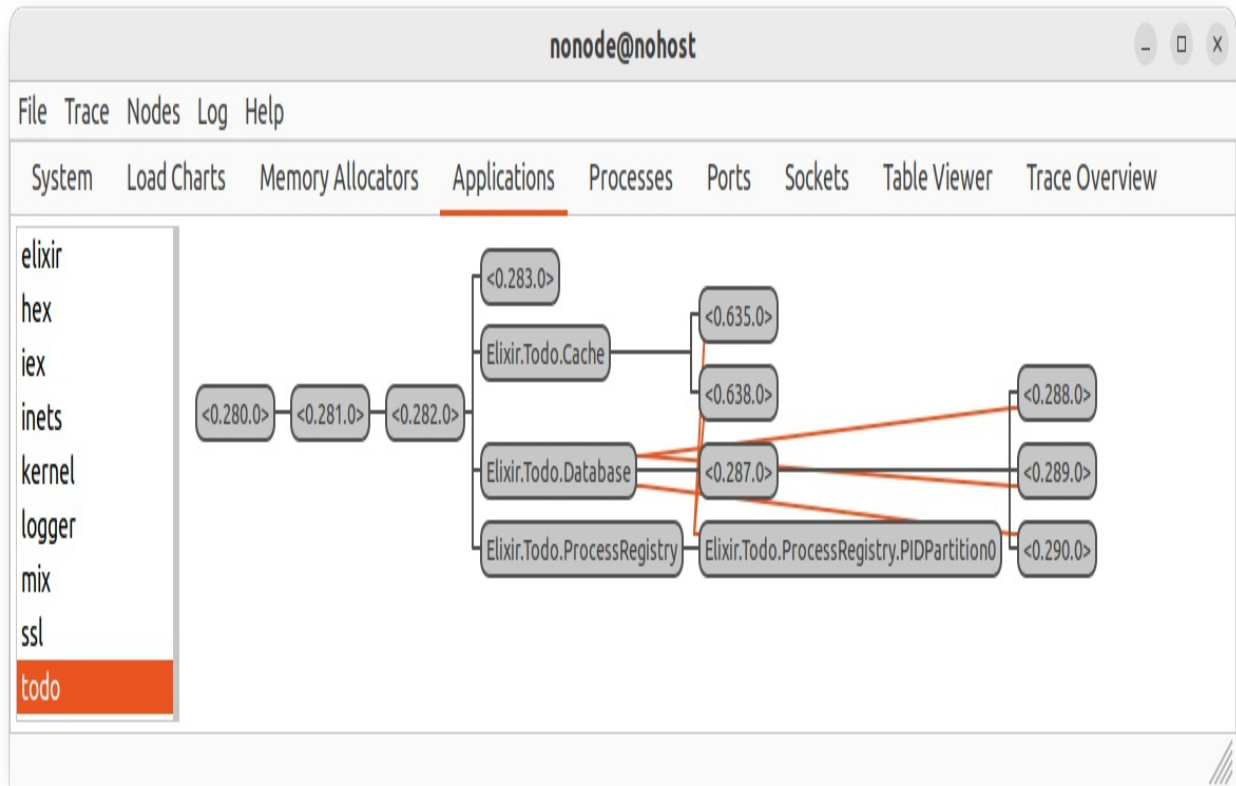Now you can start the `observer` tool:

```
iex(3)> :observer.start()
```

A GUI window should appear, which presents some basic information about the system. If you click on the Applications tab, you can see the supervision

tree of your application, as illustrated in figure 11.1.

Figure 11.1. Observing the application

The two top processes, with pids <0.280.0> and <0.281.0> are the processes used by OTP to manage the application. The third process, with the pid <0.282.0>, is the top-level `Todo.System` supervisor. Then you can see the four children: the process with the pid <0.283.0>, `Todo.Cache`, `Todo.Database`, and `Todo.ProcessRegistry`.

If available, the observer uses the registered name of a process. Otherwise, the process pid is used to display the process. In this example, the first process with the pid <0.283.0> is the metrics reporter powered by the `Todo.Metrics` module.

Under the cache, you can see two children, which are to-do servers that you just started from the shell. You can easily verify this by double-clicking on the process box and clicking the State button in the new window. In the last row, you'll see the process state, which will contain the name of the to-do

list.

The `observer` tool can be useful for visualizing the behavior of the running system. For example, in the Processes tab you can see a list of all processes running in the system, and easily figure out which processes are very busy or use a lot of memory. This tab is frequently used to find bottlenecks in the system. You can even use the `observer` to visualize the system running in production. This will be explained in chapter 13.

11.3 Building a web server

The time has finally come to introduce an HTTP interface into your to-do system. You'll implement a rudimentary server, wrapping only `entries` and `add_entry` requests.

Our focus won't be so much on the details and finesse of web servers. Instead, the aim is to demonstrate how you can work with OTP applications and to show you how everything connects in a simplistic simulation of a real-world system.

11.3.1 Choosing dependencies

You could, of course, implement the entire server from scratch, but that would be too much work. Instead you'll reach for a couple of existing libraries to make your life easier. Note that it's not the purpose of this section to provide detailed descriptions of those libraries; they're just a simple means to an end (a basic working HTTP server). You're encouraged to research those libraries in more detail on your own.

Various web server frameworks and libraries are available for both Elixir and Erlang. If you plan on doing any serious production development, you should definitely take a look at the Phoenix Framework (https://phoenixframework.org/). Phoenix is versatile and highly modular, so it's a great choice for powering all kinds of web servers. At first glance, Phoenix might seem a bit daunting, so we'll keep things simple here and reach for two lower-level libraries.

The first one is an Erlang library called Cowboy (https://github.com/extend/cowboy); it's a fairly lightweight and efficient HTTP server library that's popular choice for HTTP servers in Erlang and Elixir ecosystems.

But you won't use Cowboy directly. Instead, you'll interface with it via the Plug library (https://github.com/elixir-lang/plug), which provides a unified API that hides the details of the actual HTTP server implementation.

You can add both libraries to your project with a single dependency, as shown in the following listing.

Listing 11.8. External dependencies for the web server (todo_web/mix.exs)

```
defmodule Todo.Mixfile do
  ...

  defp deps do
    [
      {:poolboy, "~> 1.5"},
      {:plug_cowboy, "~> 2.6"} #1
    ]
  end
end
```

We're adding a new dependency called `plug_cowboy`. This library provides the glue code between plug and cowboy. Internally, this library depends on plug and cowboy libraries, which are therefore transitive dependencies of our project. We don't have to explicitly list them. The `mix` tool can resolve this out automatically.

11.3.2 Starting the server

With the dependencies configured, you can now run `mix deps.get` and start implementing the HTTP interface. As mentioned, your primary interface for working with Cowboy is Plug. Plug is a reasonably complex library, and I won't provide an in-depth treatment here. Our focus is on getting a basic version running and understanding how all the pieces work together.

To start a server powered by Plug and Cowboy, you can reach for

`Plug.Cowboy.child_spec/1`. This function returns a child specification describing how to start the processes responsible for the HTTP server part of the system. As usual, it's best to wrap this in a dedicated module. The code is provided in the following listing.

Listing 11.9. HTTP server specification (todo_web/lib/todo/web.ex)

```elixir
defmodule Todo.Web do
  ...

  def child_spec(_arg) do
    Plug.Cowboy.child_spec(
      scheme: :http,
      options: [port: 5454],
      plug: __MODULE__
    )
  end

  ...
end
```

The argument passed to `Plug.Cowboy.child_spec/1` provides the server options. Here you specify that you want to serve HTTP traffic on port 5454. The final option, `:plug`, indicates that some function from this module will be invoked to handle the request.

When you include `Todo.Web` in a supervisor, a bunch of new processes will be started. There will be at least one process that listens on a given port and accepts requests. Then, each distinct TCP connection will be handled in a separate process, and your callbacks (which you have to implement) will be invoked in those request-specific processes. But you don't need to care about any of these details. Owing to `child_spec/1`, the Plug library keeps these internals to itself.

Note

Remember that applications are singletons — you can start only one instance of a distinct application in a running BEAM instance. But this doesn't mean you can run only one HTTP server in your system. The Plug and Cowboy

applications can be considered factories of HTTP servers. When you start these applications, no HTTP server is started yet. You use `child_spec/1` to inject an HTTP server somewhere in your supervision tree. You can, of course, run multiple HTTP servers in your system. For example, you could add another HTTP server for administration purposes.

With `child_spec/1` in place, you can inject the HTTP server into the supervision tree. This is done in the `Todo.System` supervisor.

Listing 11.10. Starting the HTTP server (todo_web/lib/todo/system.ex)

```
defmodule Todo.System do
  def start_link do
    Supervisor.start_link(
      [
        Todo.Metrics,
        Todo.ProcessRegistry,
        Todo.Database,
        Todo.Cache,
        Todo.Web        #1
      ],
      strategy: :one_for_one
    )
  end
end
```

Thanks to wrapping and proper naming, the `Todo.System` module clearly describes what the system is made of: metrics, a process registry, a database, a cache, and a web server.

11.3.3 Handling requests

Now you can start handling some requests. Let's introduce support for `add_entry`. This will be a `POST` request. To keep the implementation simple, you'll transfer all parameters via URL. The example request looks like this:

```
http://localhost:5454/add_entry?list=bob&date=2023-12-19&title=De
```

Begin by setting up a route for this request. The skeleton is provided next.

Listing 11.11. Setting up a route for add_entry (todo_web/lib/todo/web.ex)

```
defmodule Todo.Web do
  use Plug.Router          #1

  plug :match              #2
  plug :dispatch           #2

  ...

  post "/add_entry" do     #3
    ...
  end

  ...
end
```

There are some strange constructs here, which are part of how Plug is used. I'll provide a simple explanation, but you don't need to understand all this because it isn't the aim of this exercise.

Calling use Plug.Router adds some functions to your module. This is similar to how you used the Elixir behavior helper (by calling use GenServer). For the most part, the functions imported will be used internally by Plug.

Expressions such as plug :match and plug :dispatch deserve special mention. These calls will perform some additional compile-time work that will allow you to match different HTTP requests. These expressions are examples of Elixir macro invocations, and they'll be resolved at compilation time. As a result, you'll get some additional functions in your module (which are again used only by Plug).

Finally, you call a post macro to define request-handling code. The post macro works similarly to the test macro you saw in chapter 7. Under the hood, this macro will generate a function that's used by all the other generated boilerplate you got by calling the plug macro and use Plug.Router. The generated function will look roughly like the following:

```
defp do_match(conn, "POST", ["add_entry"], _) do
  ...
end
```

The code of the generated do_match function will contain the code which was

passed to the `plug` macro.

The generated `do_match` function is invoked by some other generated code that exists in your module courtesy of the `use Plug.Router` and `plug :match` expressions.

If you issue multiple calls to the `post` macro, you'll have multiple clauses for this `do_match` function, each clause corresponding to the single route you're handling. You'll end up with something like this:

```
defp do_match(conn, "POST", ["add_entry"], _), do: ...
defp do_match(conn, "POST", ["delete_entry"], _), do: ...
...
defp do_match(conn, _, _, _), do: ...    #1
```

You can check this yourself with the help of the tool called decompile (https://github.com/michalmuskala/decompile). Install the tool per README instructions, change to the todo project root folder, compile the project, and then run `mix decompile Todo.Web --to expanded`. This will generate the file Elixir.Todo.Web.ex, which will contain the Elixir code after all the macros have been expanded.

Long story short, the code you provide to `post "/add_entry"` is invoked when an HTTP `POST` request with an `/add_entry` path arrives at your server.

Let's now look at the implementation of the request handler.

Listing 11.12. Implementing the add_entry request (todo_web/lib/todo/web.ex)

```
defmodule Todo.Web do
  ...

  post "/add_entry" do
    conn = Plug.Conn.fetch_query_params(conn)         #1
    list_name = Map.fetch!(conn.params, "list")       #1
    title = Map.fetch!(conn.params, "title")          #1
    date = Date.from_iso8601!(Map.fetch!(conn.params, "date"))

    list_name              #2
    |> Todo.Cache.server_process()         #2
    |> Todo.Server.add_entry(%{title: title, date: date})   #2
```

```
    conn            #3
    |> Plug.Conn.put_resp_content_type("text/plain")    #3
    |> Plug.Conn.send_resp(200, "OK")      #3
  end

  ...
end
```

Notice that in the request handler you use the `conn` variable, which doesn't exist anywhere. This variable is brought to you by the `post` macro, which generates this variable and binds it to the proper value.

As the name implies, the `conn` variable holds your connection. This is an instance of a `Plug.Conn` struct that holds a TCP socket together with information about the state of the request you're processing. From your handler code, you must return the modified connection, which will hold response information, such as its status and body.

The implementation of the request handler consists of three parts. You first decode the input parameters. Then you invoke the code that performs some action in your system, in this case adding the new entry. Finally, you respond to the client.

To decode input parameters, you need to invoke `Plug.Conn.fetch_query_params/1`. This function returns a new version of the connection structure with the `params` field containing request parameters (in the form of a map). This is essentially a caching technique. Plug caches the result of `fetch_query_params` in the connection struct, so repeated calls to `fetch_query_params` won't result in excessive parsing.

Once you have the parameters ready, you can add the entry. Finally, you set the response content type, status, and body.

The same approach is used to handle the `entries` request. The code is provided in the following listing.

Listing 11.13. Implementing the entries request (todo_web/lib/todo/web.ex)

```
defmodule Todo.Web do
  ...
```

```
get "/entries" do
  conn = Plug.Conn.fetch_query_params(conn)
  list_name = Map.fetch!(conn.params, "list")
  date = Date.from_iso8601!(Map.fetch!(conn.params, "date"))

  entries =
    list_name
    |> Todo.Cache.server_process()
    |> Todo.Server.entries(date)

  formatted_entries =
    entries
    |> Enum.map(&"#{&1.date} #{&1.title}")
    |> Enum.join("\n")

  conn
  |> Plug.Conn.put_resp_content_type("text/plain")
  |> Plug.Conn.send_resp(200, formatted_entries)
end

...
end
```

This code follows the same approach as in `add_entry`. Because `entries` is a GET request, you use the `get` macro instead of `post`.

In this implementation, you decode the parameters and then fetch the desired entries. You then produce the textual representation sent to the client and return the response.

Note

The presented code organization treats the HTTP server as merely an interface to the core of the system. Notice how in both requests, you convert the input into domain-specific types, and then invoke HTTP-agnostic code, such as `Todo.Server.add_entry` to perform the operation. The core of the system which we implemented in the previous chapters doesn't know or care that it's running in the context of an HTTP request. Such separation of concerns makes the code easier to work with, test, and expand.

At this point, you can start the system with `iex -S mix` and issue requests.

For example, this is what you get using a command-line `curl` tool:

```
$ curl -d "" \
  "http://localhost:5454/add_entry?list=bob&date=2023-12-19&title
OK

$ curl "http://localhost:5454/entries?list=bob&date=2023-12-19"
2023-12-19 Dentist
```

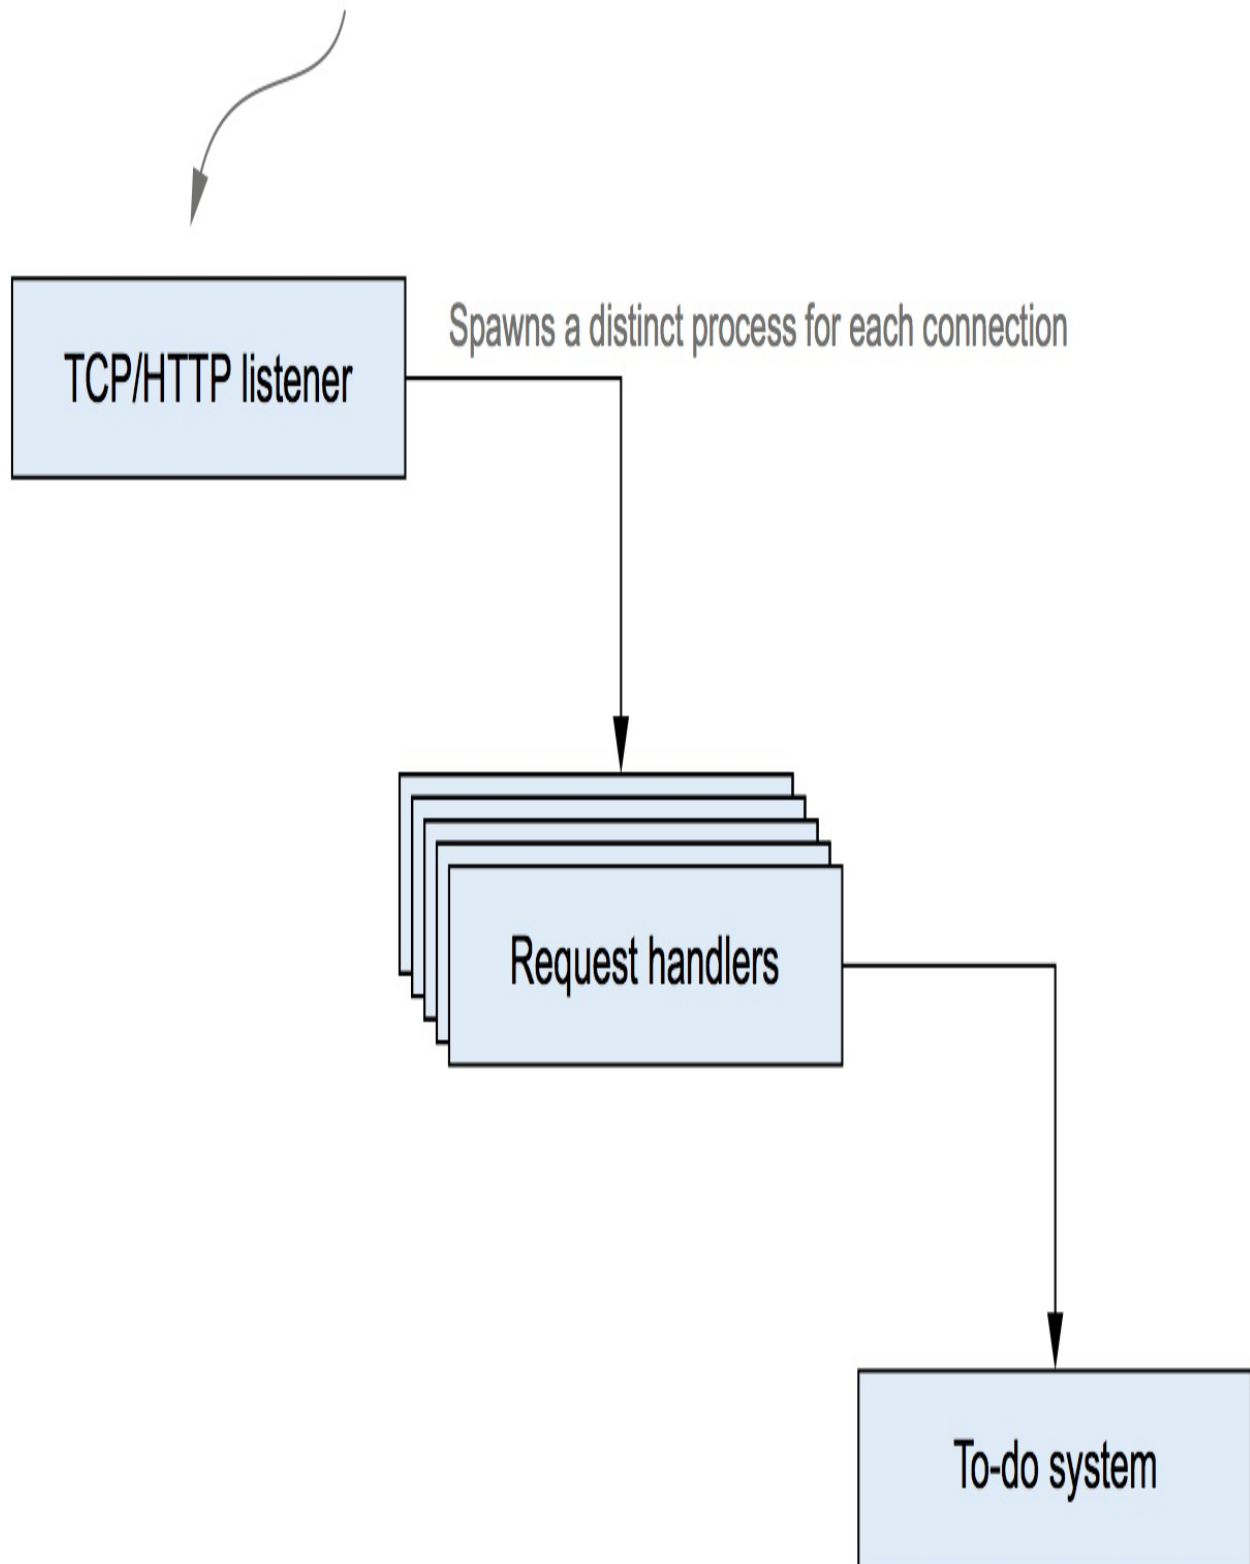This proves that your system is working, but let's see how everything combines.

11.3.4 Reasoning about the system

First, let's look at how the HTTP server works. The simplified idea is illustrated in figure 11.2.

Figure 11.2. Requests are handled in separate processes

Listens for incoming requests on a given port

TCP/HTTP listener

Spawns a distinct process for each connection
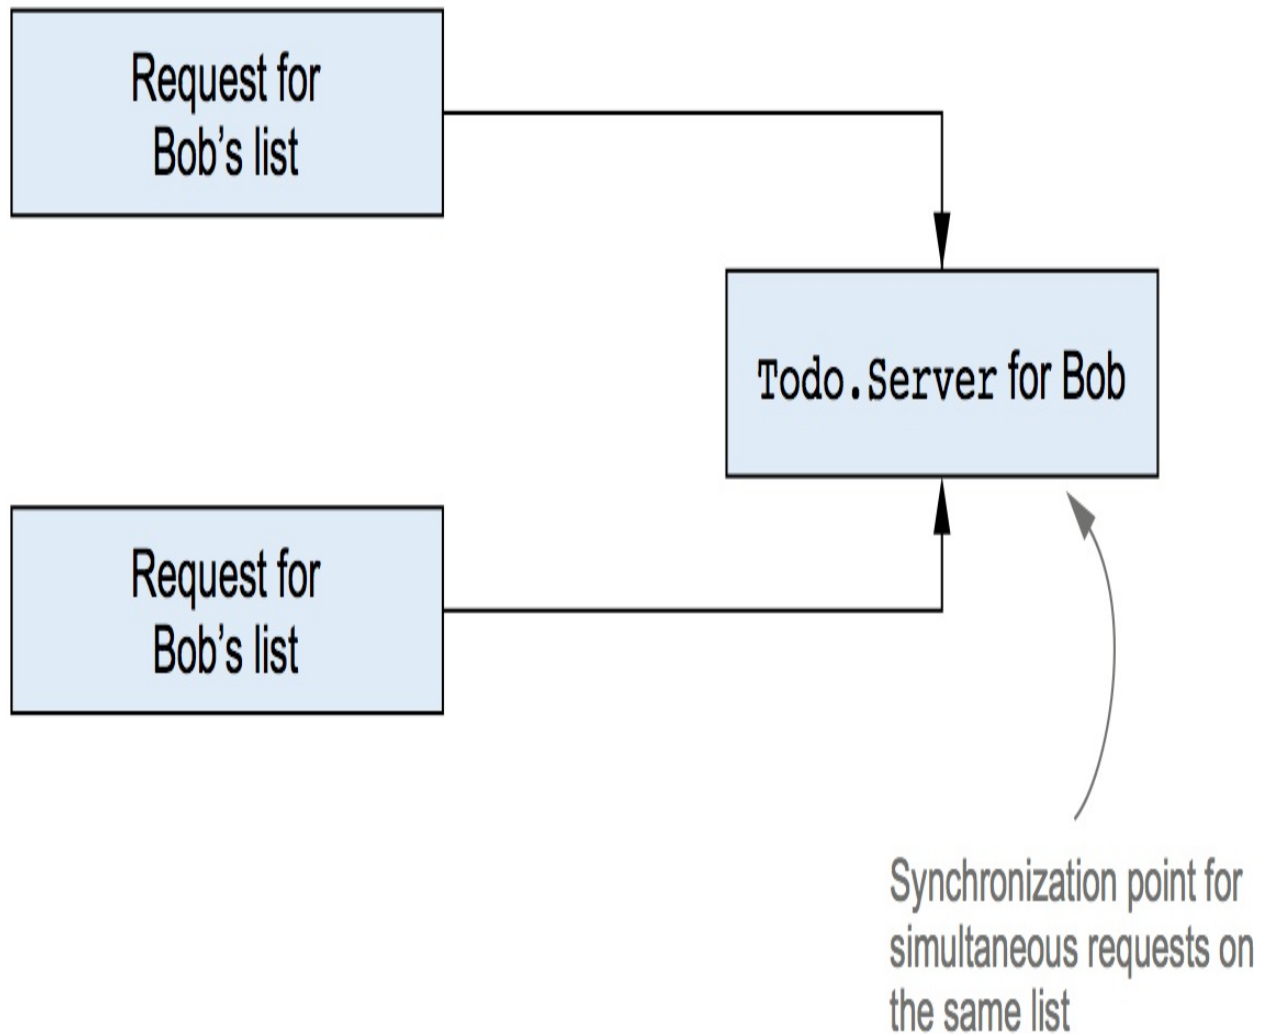
Request handlers

To-do system

The most important thing to notice is that each connection is managed in a distinct process. In practice, this means different requests are handled in different processes. There's no special magic here — this is how the underlying Cowboy web server is implemented. It uses one process to listen on a port, and then it spawns a separate process for each incoming request.

This architecture has all sorts of benefits due to the way BEAM treats processes. Because processes are concurrent, CPU resources are maximally used and the system is scalable. Because processes are lightweight, you can easily manage a large number of simultaneous connections. Moreover, thanks to the BEAM scheduler being preemptive, you can be certain that occasional long-running, CPU-intensive requests won't paralyze the entire system. Finally, due to process isolation, a crash in a single request won't affect the rest of the system.

Processes also make it easy to reason about the system. For example, you can be certain that independent requests won't block each other, whereas multiple requests on the same to-do list are synchronized, as illustrated in figure 11.3.

Figure 11.3. Independent simultaneous requests are handled concurrently, whereas requests on the same to-do list are synchronized.

```
Request for
Alice's list
```
→
```
Todo.Server for Alice
```

```
Request for
Bob's list
```

```
Request for
Bob's list
```
→
```
Todo.Server for Bob
```

Synchronization point for
simultaneous requests on
the same list

This is due to the way the to-do cache works. Whenever you want to

manipulate Bob's list, you first ask the to-do cache to return the pid of the process in charge. All requests for Bob's list go through that same process and are therefore handled one by one.

Performance

After so much development, it's worth measuring the performance of the system. I quickly tested the system using a tool called `wrk` (https://github.com/wg/wrk). A brief 30-second load test on my machine gives a throughput of about 40,000 requests per second, with an average latency of 1 ms, and a 99 percentile latency of 7 ms. While the test was running, all the CPU cores were completely busy, which proves that there are no bottlenecks in the system. The system is highly concurrent and able to use all of the cores at its disposal.

The observed results are quite decent, especially given that we made some naive decisions in the implementation. Here are a couple of possible issues in the implementation:

- On every list modification, you store the entire list.
- Date-based lookups in the to-do list abstraction iterate through the entire list.
- Every to-do server lookup goes through a to-do cache dynamic supervisor.

There's definitely some room for improvement here, but it's comforting to know that the system performs well out of the box, even though you didn't do any particular fine-tuning.

Calls vs. casts

Let's discuss some effects that calls and casts may have on your system. To refresh your memory, *casts* are fire-and-forget requests. A caller sends a message to the server and then immediately moves on to do something else. In contrast, a *call* is a blocking request, where a caller waits for the server to respond.

Remember that you've opted to use casts for all operations except where you need to return a response. This was a somewhat arbitrary decision, made mostly for didactic purposes. In reality, casts have a drawback: you don't know what happened with your request. This, in turn, means you may be giving false responses to end users, as illustrated in figure 11.4.

Figure 11.4. Using casts reduces the certainty of your responses.
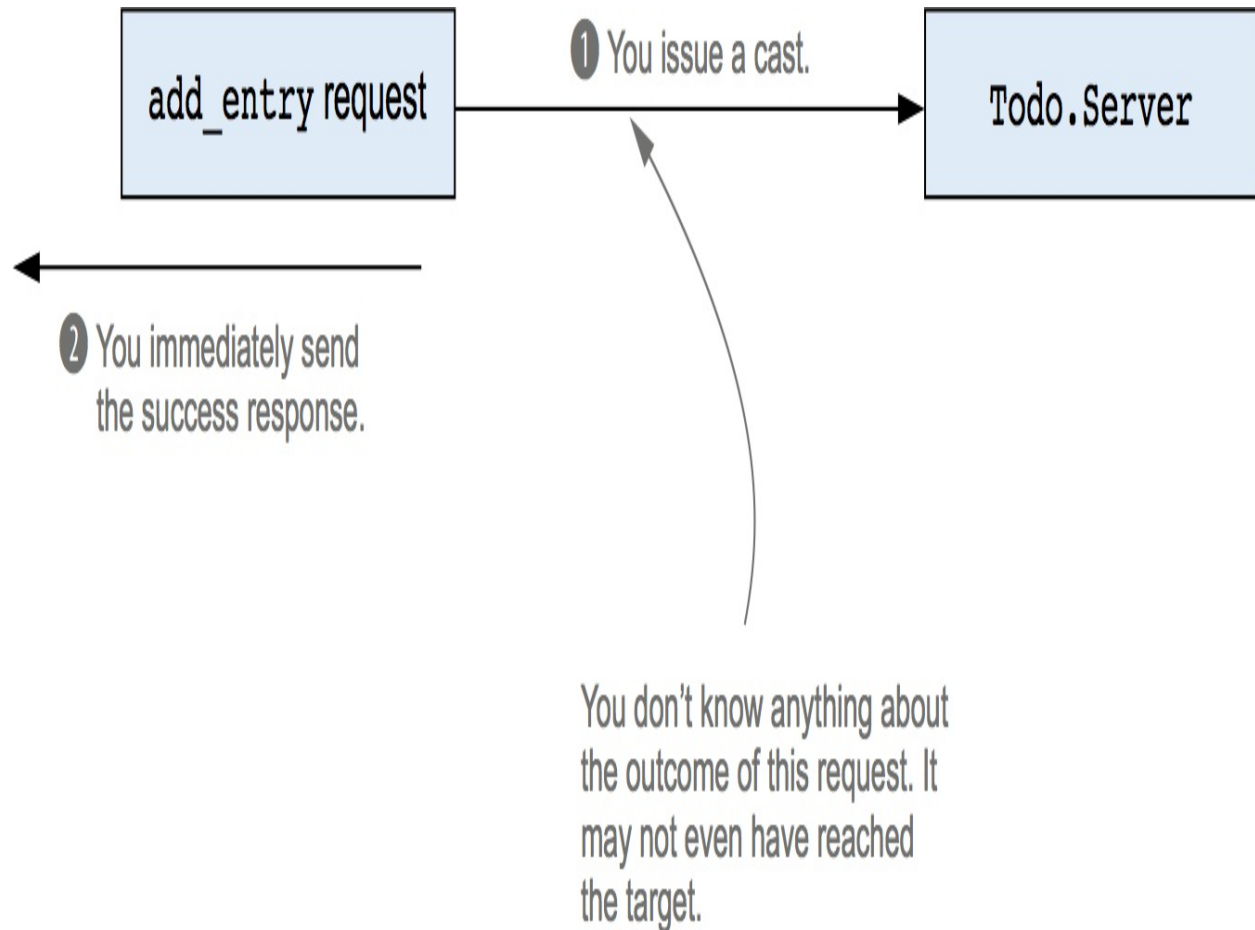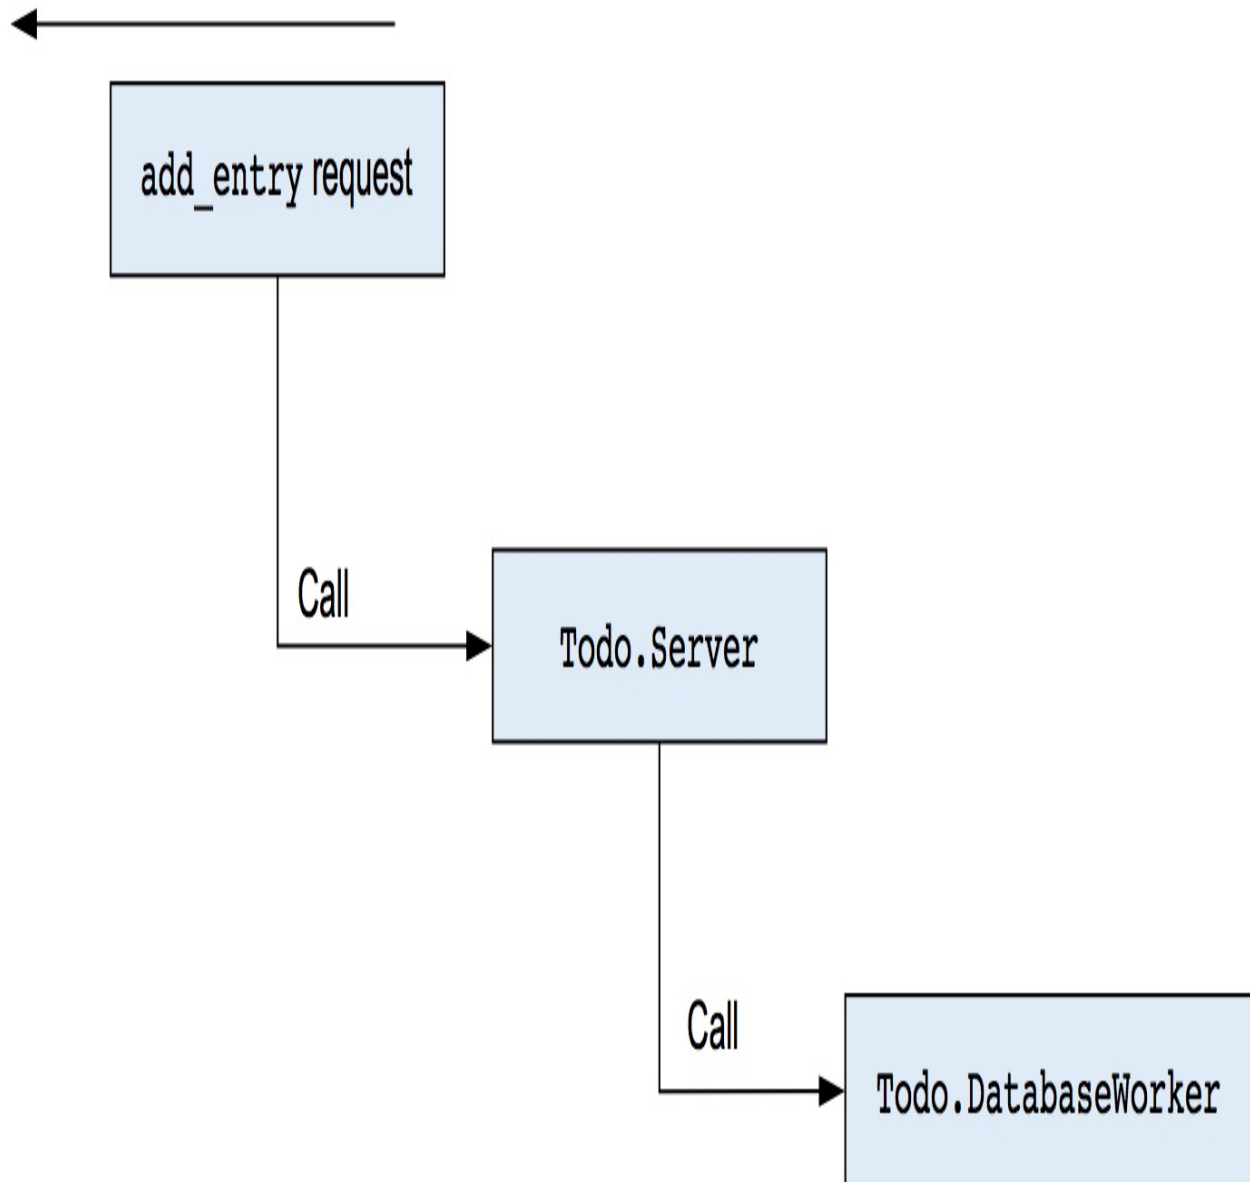

Because you use a cast to add a to-do entry in your system, you have no way of knowing what happened with your request. When you're telling the end user that you succeeded, this is a guess rather than a truthful statement.

Figure 11.5. Using calls promotes consistency but reduces the responsiveness of the system.

The response is sent only
after the entry has been
successfully stored to the
database.

←————————————

```
┌─────────────────────────┐
│                         │
│    add_entry request    │
│                         │
└─────────────────────────┘
          │
          │
       Call
          │
          └──────▶ ┌─────────────────────────┐
                   │                         │
                   │      Todo.Server        │
                   │                         │
                   └─────────────────────────┘
                              │
                              │
                           Call
                              │
                              └──────▶ ┌─────────────────────────────┐
                                       │                             │
                                       │   Todo.DatabaseWorker       │
                                       │                             │
                                       └─────────────────────────────┘
```
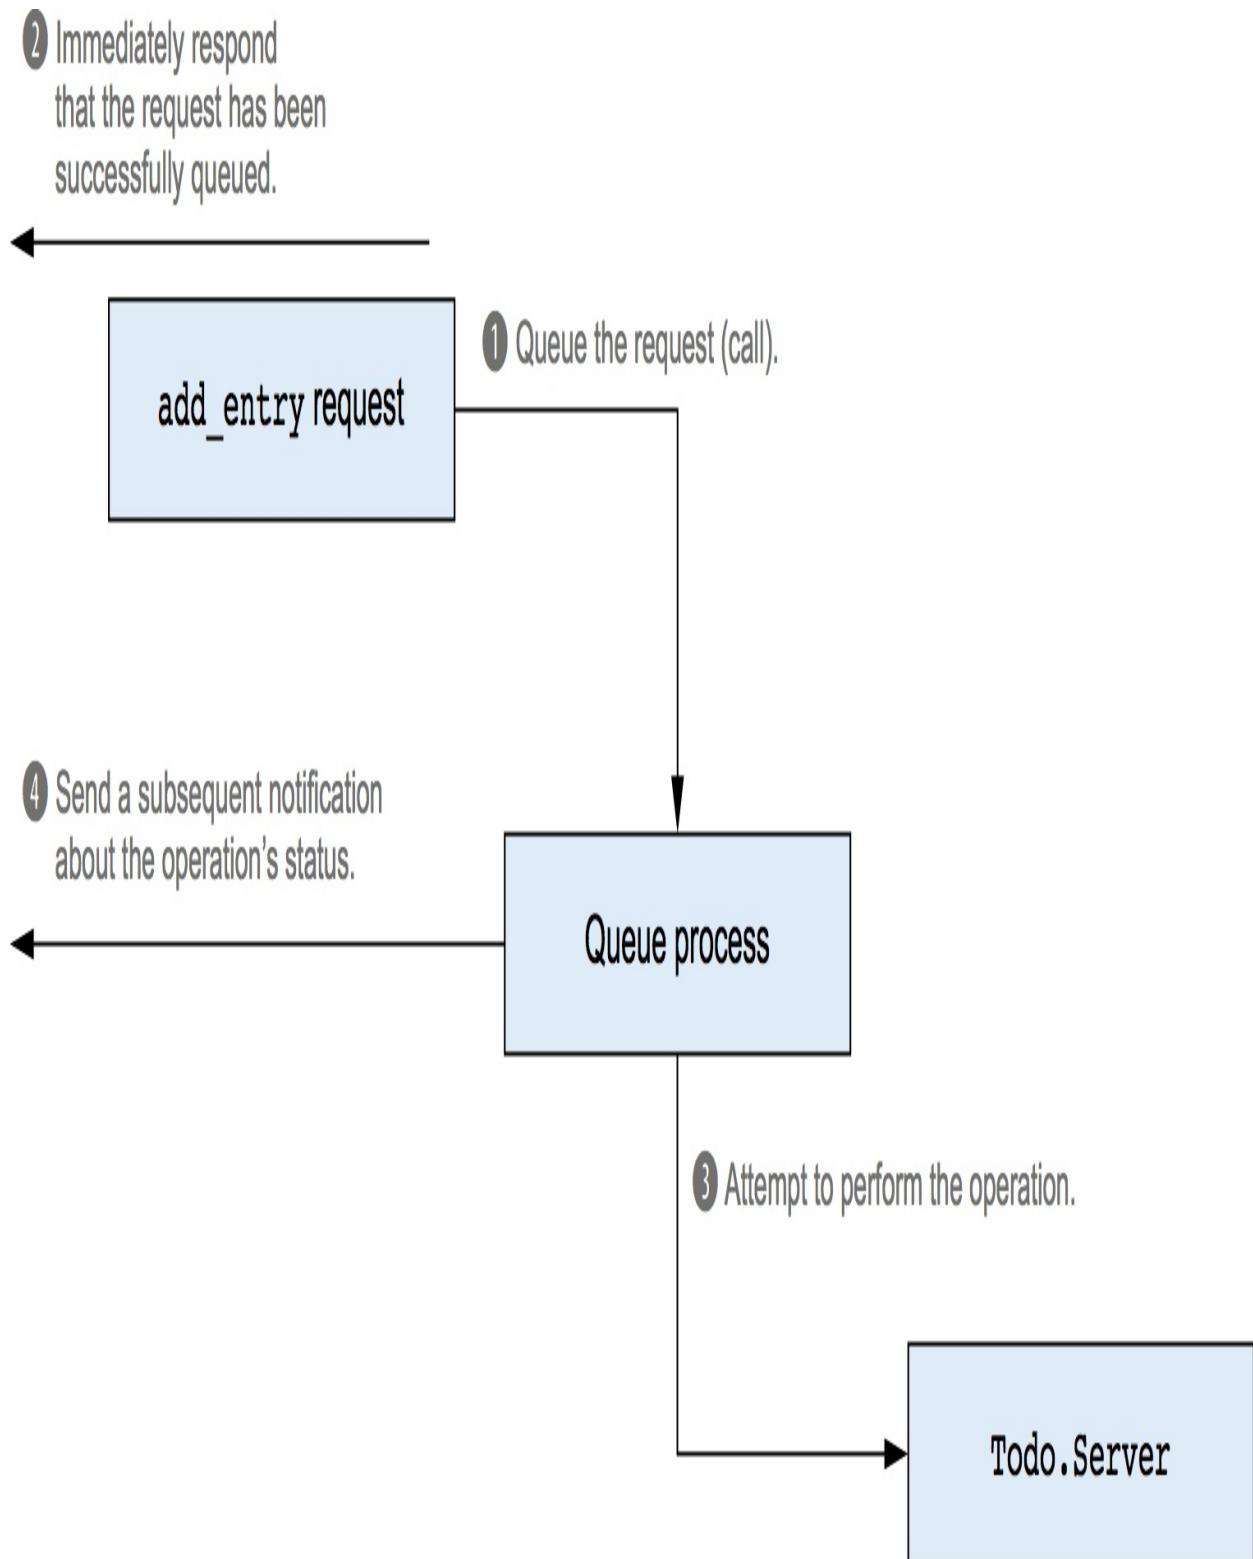
Obviously, the simple way to resolve this is to use calls, which are
synchronous, meaning the client must wait until the response arrives, as
illustrated in figure 11.5.

This approach is more consistent: you return success only when you're

certain that the entry has been stored. But the downside is that the entire system now depends on the throughput of database workers, and as you may recall, you're running only three workers, and you're using a pretty inefficient database.

This can be resolved by introducing an intermediate process. The idea is to provide an immediate response stating that the request has been queued. Then you do your best to process the request, and you send a subsequent notification about the request's status. See figure 11.6.

Figure 11.6. Queuing the request and sending a notification about the status

② Immediately respond
that the request has been
successfully queued.

←

add_entry request

❶ Queue the request (call).

④ Send a subsequent notification
about the operation's status.

←

Queue process

❸ Attempt to perform the operation.

Todo.Server

This scheme is definitely more elaborate and involved, so it's not appropriate for simpler cases where using plain calls is enough. But in cases when the

load is very high and end-to-end operations can take longer, introducing an intermediate process may be beneficial. This process can increase the responsiveness of the system while retaining consistency. Moreover, this process can serve as a proactive agent in dealing with congestion and increased load. If the system becomes overloaded and the queue starts piling up, you can refuse to take more requests into the queue process until you regain some breathing space.

As always, no single approach works for all cases. Using calls can be a reasonable first attempt because it promotes consistency. Later you can easily switch to casts or introduce an intermediate process, depending on the specific situation.

If you need to perform more complex load management, you can introduce an intermediate process. Instead of rolling your own solution, you can consider using the `GenStage` library (https://github.com/elixir-lang/gen_stage), which allows you to build various pipelines of producers and consumers.

At this point, you're finished implementing a basic HTTP server. Next we'll talk about application configuration.

11.4 Configuring applications

An OTP application can be configured using a feature called an *application environment* — a key/value, in-memory store where keys are atoms, and values are Elixir terms. You can provide application environment values through config script files — Elixir scripts that reside in the config folder.

The `mix` tool makes sure that the configuration is loaded into the application environment before the application is started. Finally, you can retrieve an environment value using functions from the `Application` module.

11.4.1 Application environment

Let's look at a simple example. The `todo` HTTP server currently listens on a hard-coded port, 5454. In this example, you'll make the HTTP port

configurable via OS environment variable.

The most typical way of setting an application environment is to use the config/runtime.exs file. This is a script which is evaluated at runtime, just before the application is started. In here you can set the environment configuration of your application, as well as any of its dependencies.

To make the HTTP port configurable, you can read the configuration from the OS environment variable, and store it into the application environment, as shown in the following listing:

Listing 11.14. Configuring the HTTP port (todo_env/config/runtime.exs)

```
import Config                                         #1

http_port = System.get_env("TODO_HTTP_PORT", "5454")  #2
config :todo, http_port: String.to_integer(http_port) #3
```

Here you use `System.get_env/2` to read the OS environment variable `TODO_HTTP_PORT`, using the default value of 5454 if the variable is not set. Then, you use the imported `Config.config/2` function to configure the `:http_port` setting of the `:todo` application.

As soon as you start the system, this setting is available in the application environment, and you can retrieve it; for example, with the function `Application.fetch_env!/2`:

```
$ iex -S mix

iex(1)> Application.fetch_env!(:todo, :http_port)
5454
```

Setting the OS env var before the system is started affects the application environment:

```
$ TODO_HTTP_PORT=1337 iex -S mix

iex(1)> Application.fetch_env!(:todo, :http_port)
1337
```

Now you can adapt the code of `Todo.Web` to read the port from the

application environment. The change is shown next.

Listing 11.15. Fetching the http_port setting (todo_env/lib/todo/web.ex)

```
defmodule Todo.Web do
  ...

  def child_spec(_arg) do
    Plug.Cowboy.child_spec(
      scheme: :http,
      options: [port: Application.fetch_env!(:todo, :http_port)],
      plug: __MODULE__
    )
  end

  ...
end
```

In this particular case, you use `Application.fetch_env!/2`, which will raise an error if the HTTP port isn't configured. To provide a default value if none is set, you can use `Application.get_env/3` (https://hexdocs.pm/elixir/Application.html#get_env/3).

11.4.2 Varying configuration

In some cases, you might want to use different settings in different mix environments. For example, you currently always use the same HTTP port for development and for tests. As a consequence, you can't run tests if the to-do system is started in development. The solution for this problem is to use a different HTTP port in the test mix environment.

ⓘ Note

It's worth noting that the term "environment" is overloaded. An application environment is a key/value store that holds various settings for your OTP application. A `mix` environment determines the compilation target, such as development, test, or production. Finally, the OS environment consists of the OS-level variables present at the current shell session.

To vary the settings in different mix environments, you need to make some changes to config/runtime.exs. The code is shown next.

Listing 11.16. Mix environment-specific settings (todo_env/config/runtime.exs)

```
import Config

http_port =
  if config_env() != :test,                                #1
    do: System.get_env("TODO_HTTP_PORT", "5454"),          #1
    else: System.get_env("TODO_TEST_HTTP_PORT", "5455")    #1

config :todo, http_port: String.to_integer(http_port)
```

Here you rely on the imported `Config.config_env/0` to figure out the mix environment which was used to build the project. Then, in the test environment you'll use a different OS env variable and a different default http port. The app env name is always the same (`:http_port`), so the rest of the code doesn't have to be changed.

Let's quickly verify if this works. Start the application in the default dev environment:

```
$ iex -S mix

iex(1)> Application.fetch_env!(:todo, :http_port)
5454
```

Now, try the same in the test environment:

```
$ MIX_ENV=test iex -S mix

iex(1)> Application.fetch_env!(:todo, :http_port)
5455
```

At this point, you can run tests, even if the system is started in another OS process.

There is one additional inconvenience. The database folder is currently always the same regardless of the mix environment. As a result, running the tests will pollute your current dev data.

To avoid this, you can make the database folder configurable, using a different OS env variable, and a different default value in test env, just like you did with the HTTP port. For the sake of brevity the code is not presented here, but you're advised to try it on your own, or check the implementation in the todo_env folder.

11.4.3 Config script considerations

It's worth mentioning that Elixir also supports build-time config scripts. If you create the config/config.exs file, this script will be evaluated before the project is compiled. As a result, any setting you read from external sources, such as OS environment, is done on the build machine.

If you're building the system on a different machine than the one where you're running it (which is a good practice), the system might end up being improperly configured. Because of this subtle confusion, and due to its special semantics (it's evaluated at compile-time), try to use config.exs as little as possible. In most cases runtime.exs should serve all your needs.

If you're developing a library, try to avoid taking your parameters through the application environment. If a library is accepting its parameters via application environment, its users are forced to provide the values through config scripts. A much more flexible approach is to design the API of your library to accept all the parameters as function arguments. With such an approach, library users have complete freedom in deciding how to provide the parameters. They can hardcode the parameters, configure them through the application environment, or fetch the parameters from external sources at runtime. For more details, refer to the Library Guidelines page in the official documentation (https://hexdocs.pm/elixir/library-guidelines.html#avoid-application-configuration).

11.5 Summary

- An OTP application is a reusable component. The application can run the entire supervision tree or just provide utility modules (as in a library application).
- A non-library application is a callback module that must start the

supervision tree.

- Applications allow you to specify runtime dependencies to other applications. This isn't the same as compile-time dependencies that allow you to fetch external code and compile it from your project.
- Application environments allow you to provide application settings. This can be useful if you want to expose some configuration to the system operators.